

# Lecture 02:

## Ethereum Mechanics

---

Akash Khosla

HOMESTEAD





[ethereum.org](https://ethereum.org)

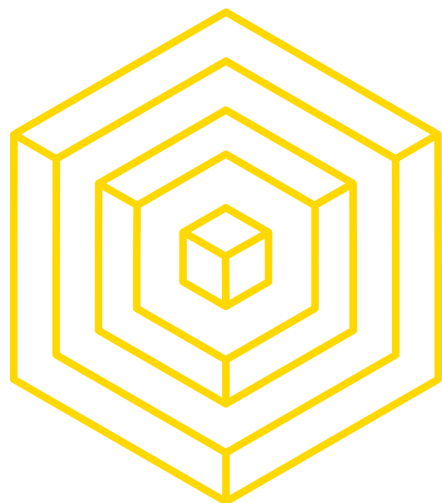


**BLOCKCHAIN**  
AT BERKELEY



# LECTURE OUTLINE

- 1  ETHEREUM PREVIEW
- 2  BITCOIN LIMITATIONS
- 3  ETHEREUM DEEP DIVE
- 4  SMART CONTRACTS



1

# ETHEREUM PREVIEW

# ethereum

~~buzzwords~~

overview

AUTHOR: PHILIP HAYES

BLOCKCHAIN

TRUSTLESS

DECENTRALIZED APPS

SMART CONTRACTS

DAOs


BITCOIN 2.0



# WHAT IS ETHEREUM?

## VITALIK'S CHILD

- Ethereum is a **decentralized** platform designed to run **smart contracts**.
  - Resistant to censorship
  - **Distributed server**
- Ethereum maintains a native asset called **ether** that is the basis of value in the Ethereum ecosystem, allows for aligning incentives
- **Account based**, not **UTXO based**
  - Remember: a UTXO is not a transaction, but simple a mechanism to manipulate and maintain state for transactions
  - **Ethereum has transactions which are used to manipulate the global account state**

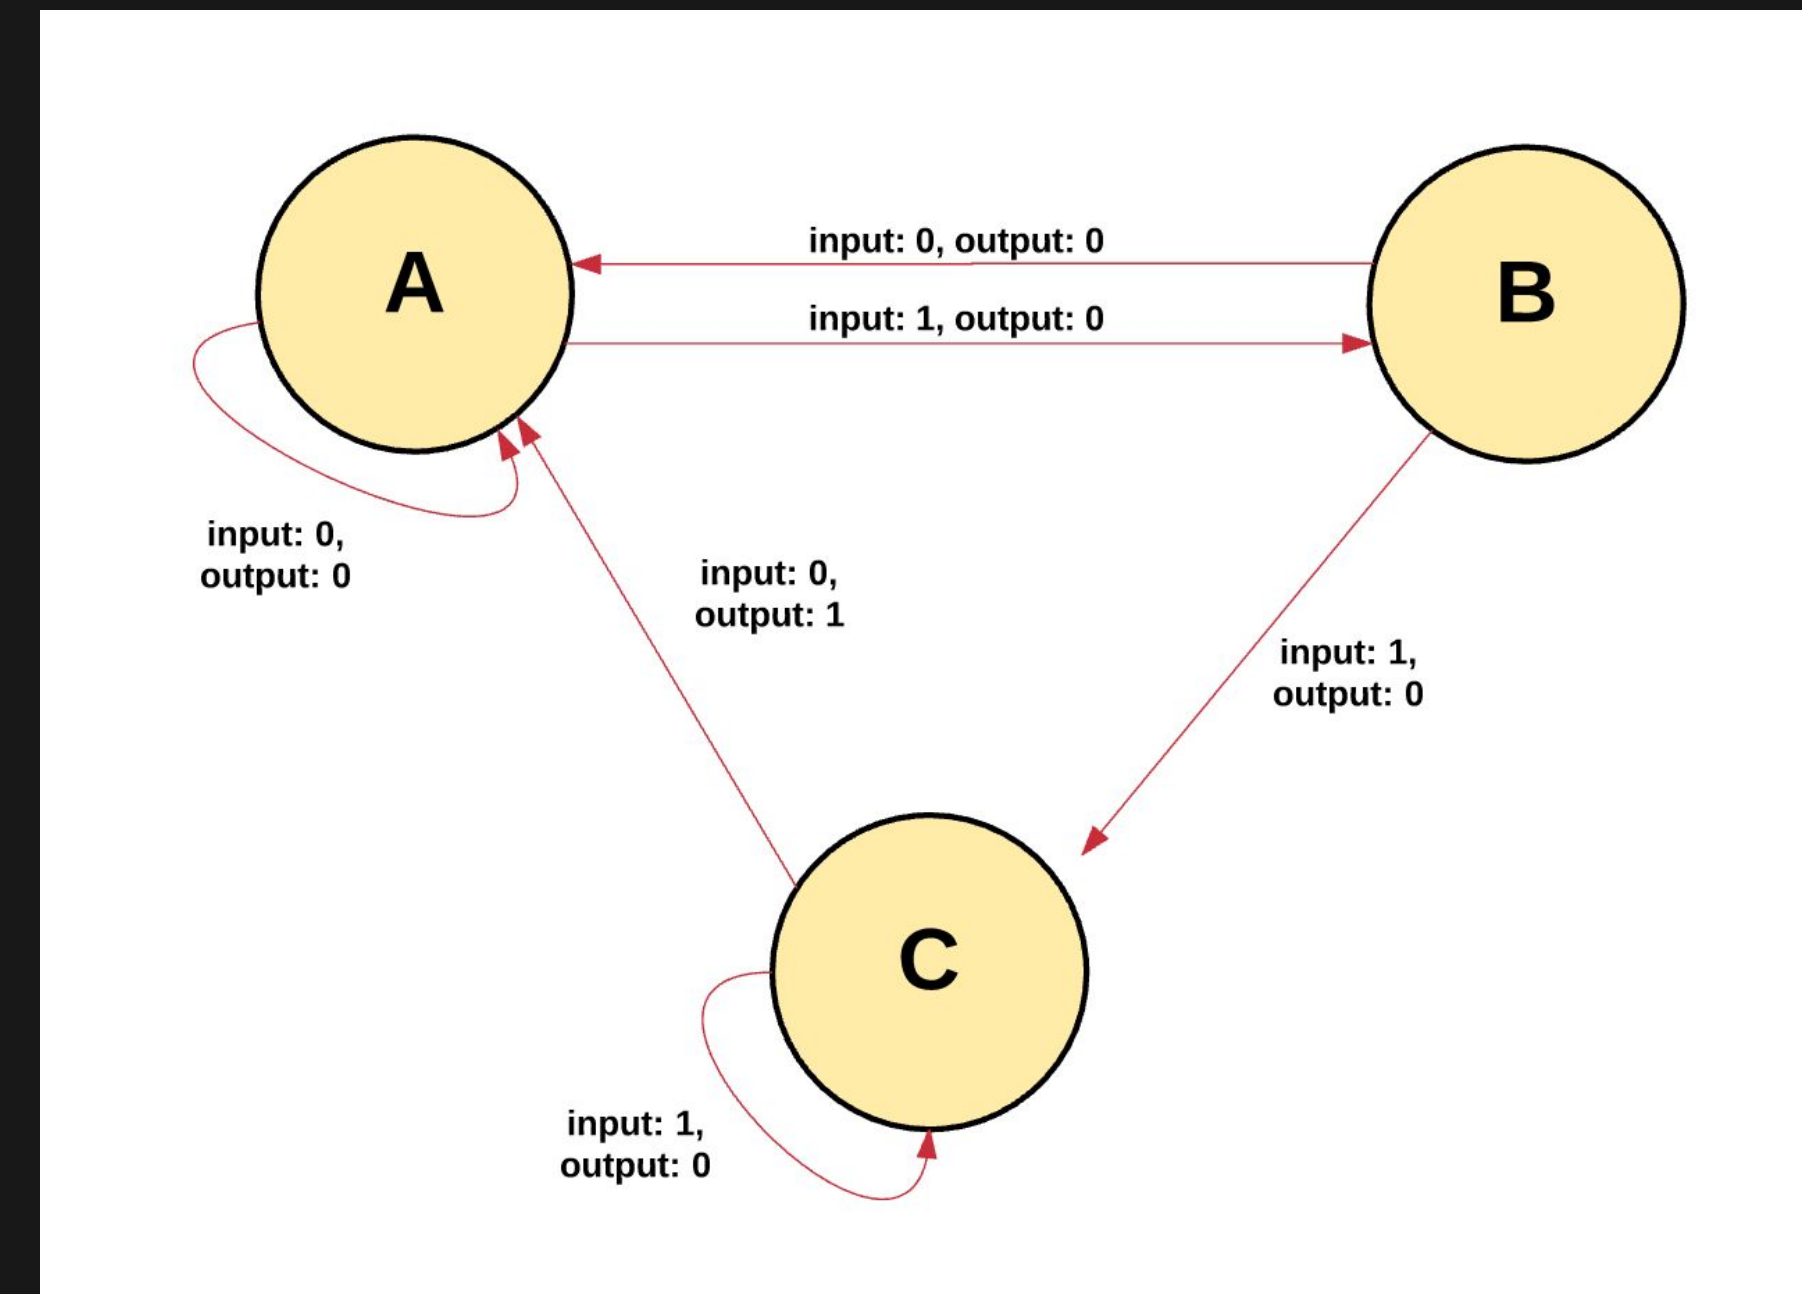




# WHAT IS THE BLOCKCHAIN FOR ETHEREUM

## PUTIN'S FAVORITE CRYPTOCURRENCY

- For Ethereum - a blockchain is a “cryptographically secure transactional singleton machine with shared-state.”
  - **Cryptographically secure** - Can't create fake transactions, erase transactions because of complex mathematical algorithms.
  - **Transaction singleton machine** - single instance of the machine for all the transactions being created in the system (“global truth”)
  - **Shared-state** - state stored on this machine is shared and open to everyone



Source

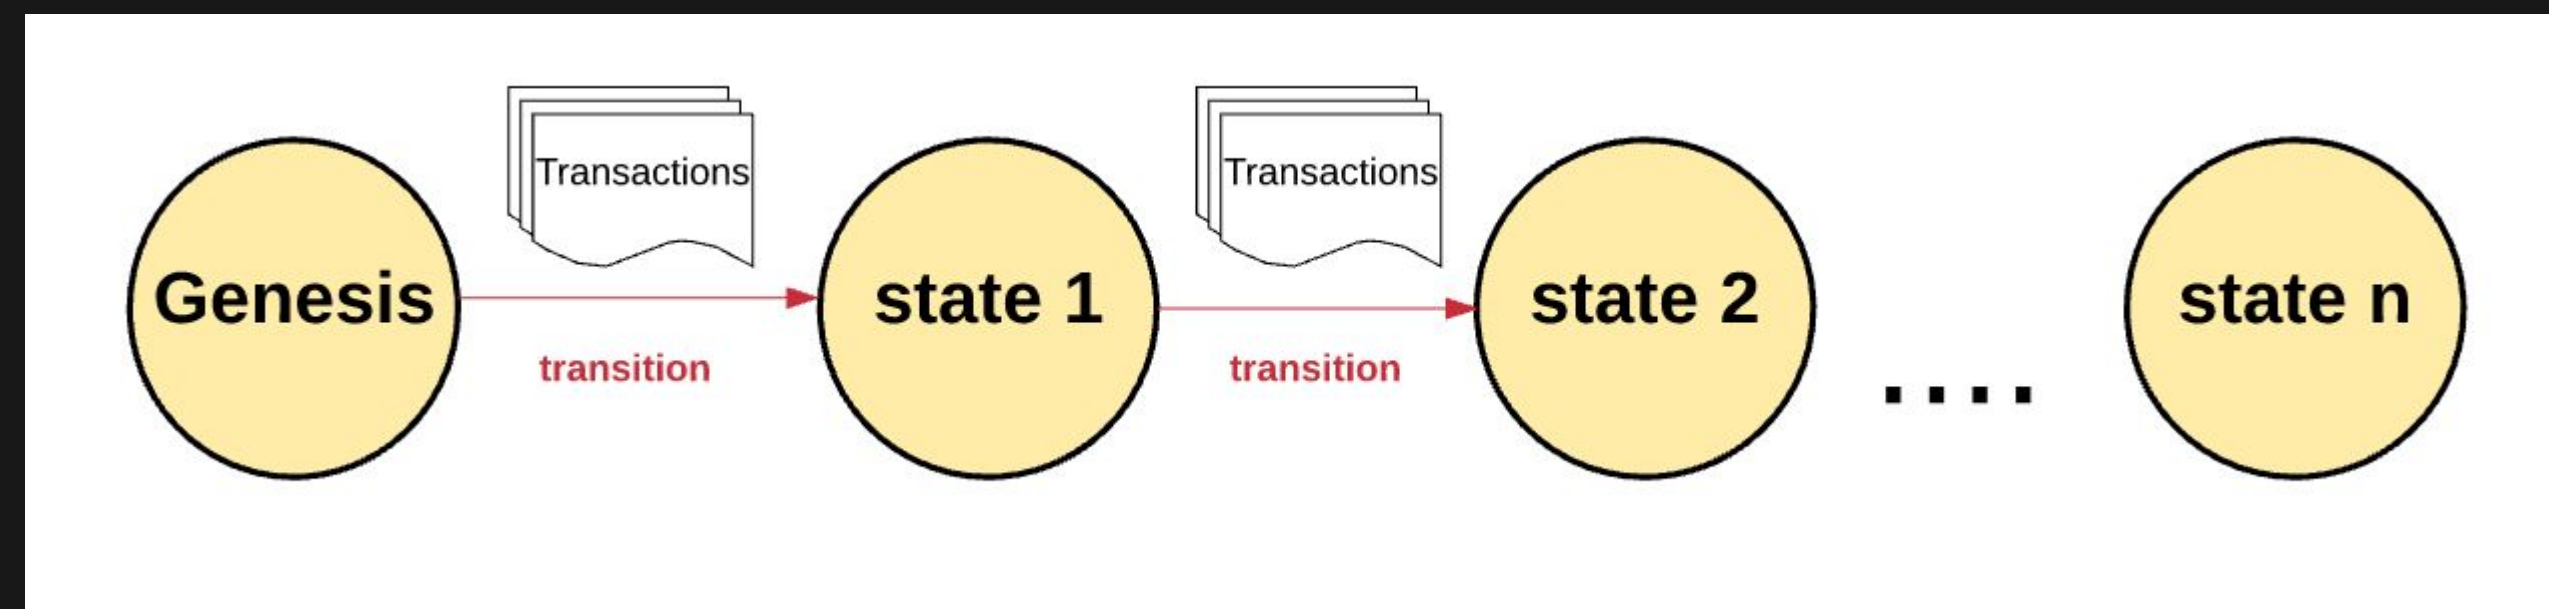





# WHAT IS A STATE MACHINE

## PUTIN'S FAVORITE CRYPTOCURRENCY

- Ethereum blockchain is essentially a **transaction based state machine**
  - A **state machine** refers to something that will read a series of inputs and based on those inputs, will transition to new state
  - This state machine begins with **genesis state** - state at which no transactions have occurred on the network
  - Once transactions are executed, the genesis state transitions to another state, where the **final state** is the current state of the Ethereum network.



[Source](#)



# BITCOIN LIMITATIONS





# HOW BITCOIN WORKS

## STATE TRANSITION PERSPECTIVE

- Bitcoin operates on running client software that agrees on state transitions
- **State:** the ownership status of all existing bitcoins
- **State Transition Function:** takes a state and transaction and outputs a new result
- $\text{APPLY}(S, TX) \rightarrow S'$

```
/* Bitcoin blockchain update algorithm */
```

```
def APPLY(state, TX):
```

```
    For each input in TX: # Input includes UTX0 and signature
```

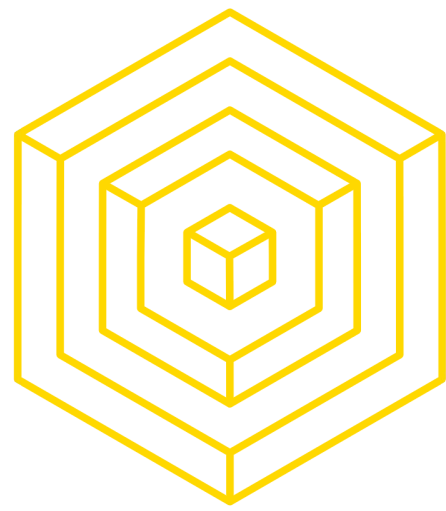
```
        if input.UTX0 not in state: return ERROR # referenced UTX0 not in state
```

```
        # provided signature doesn't match the owner of the UTX0
```

```
        if input.signature != lookup_UTX0(input).signature: return ERROR
```

```
    if sum(input UTX0s) < sum(output UTX0s): return ERROR
```

```
    S.update # Remove input UTX0s and add all output UTX0s
```



# HOW BTC BLOCK VALIDATION WORKS

## STATE TRANSITION PERSPECTIVE

- The previous algorithm works perfect for a centralized currency system
- To make a decentralized currency system, need **consensus** and **state transition** systems
- Produce a package of transactions every 10 minutes called **blocks**.

```
def VALIDATE_BLOCK(block):  
    is_valid(block.prev)  
    assert(block.timestamp > block.prev.timestamp)  
    # less than 2 hours (120 min) into the future  
    assert(block.timestamp < block.prev.timestamp + 120)  
    check_pow(block) # valid Proof of Work, used for consensus  
    TX_list = block.prev.transactions  
    for i in range(len(TX_list)):  
        # S[0] is the state, block.prev.S[0] end of the prev block state  
        block.S[i] = APPLY(block.prev.S[i], TX_list[i]) # return False on error  
    return True
```

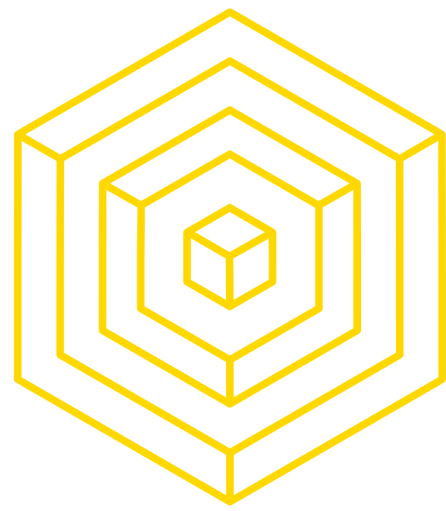


# ALTERNATIVE BLOCKCHAIN APPLICATIONS

## BITCOIN PERSPECTIVE

- **Namecoin** - Decentralized Name Registration
  - Tor, Bitcoin and other protocols often use a pseudorandom hash like `1LW79wp5ZBqaHW1jL5TCiBCrhQYtHagUWy` to identify them
- Why not have human readable names on chain
- Oldest implementation of name reg to use first-to-file paradigm
  - First registration succeeds and the second fails
  - Bitcoin consensus a perfect platform
  - Not easy to write with Bitcoin script

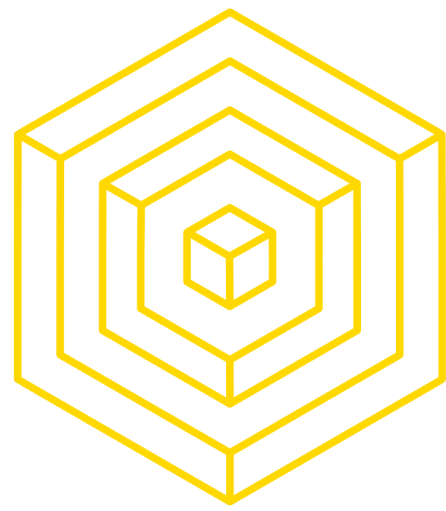




# ALTERNATIVE BLOCKCHAIN APPLICATIONS

## BITCOIN PERSPECTIVE

- **Colored coins** - Allow people to create their own currencies
  - Issue a new currency by publicly assigning color to a specific Bitcoin UTXO
  - Recursively define the color of UTXOs to be the same as the color of the inputs that the transaction creating them spent
  - Therefore can maintain wallet of colored UTXOs and send them around like regular bitcoins
- **Metacoins** - Protocols that live on top of Bitcoin
  - Provide an alternative state transition function
  - Let bitcoin handle the mining and networking infrastructure
  - Financial contracts, name registration, decentralized exchange

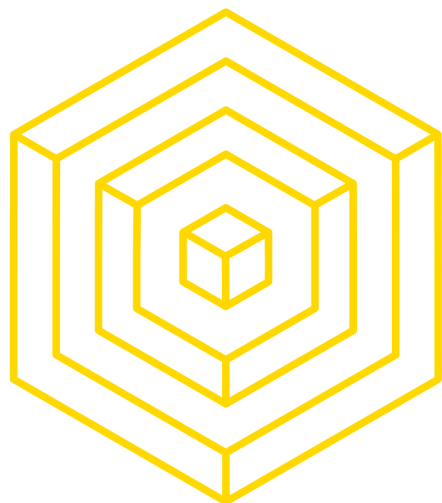


# **LIMITATIONS OF BITCOIN SCRIPT**

## BITCOIN PERSPECTIVE

- **Lack of Turing Completeness**
  - Bitcoin script supports a fair amount of computation but is missing loops
  - Why might that be a good thing?
- **Value-blindness**
  - No way for a UTXO script to provide fine grained control over the amount withdrawn
  - UTXOs are all or nothing, and it's expensive to have varying denominations
- **Lack of State**
  - Only unspent and spent UTXOs, binary state, no “efficient” withdrawal limits
- **Blockchain-blindness**
  - UTXOs are blind to blockchain data like the nonce, timestamp and prev block hash
  - Lost out on potential source of randomness





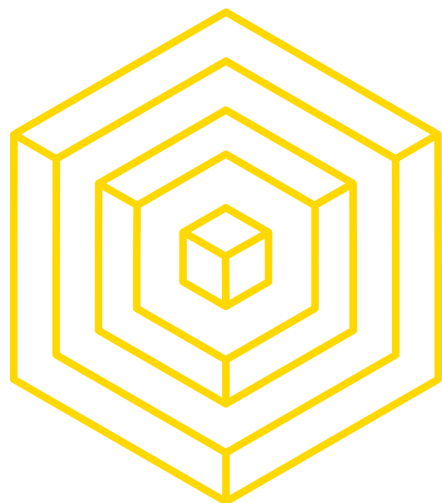
# 3

## ETHEREUM DEEP DIVE



## Ethereum

Ethereum	Bitcoin
<p>Smart Contract Platform</p> <ul style="list-style-type: none"> <li>○ Complex and feature-rich</li> <li>○ <b>Turing complete scripting language</b></li> </ul>	<p>Decentralized Asset</p> <ul style="list-style-type: none"> <li>○ Simple and robust</li> <li>○ Simple stack-based language; not Turing complete</li> </ul>
<b>Account-based for transaction</b>	<b>UTXO-based for transactions</b>
<p><b>Turing complete scripting language</b> that enables smart contracts. Ether asset is not the primary goal. Uses proof of work, plans to do proof of stake.</p>	<p><b>Uses Bitcoin Script, a very simple stack based language.</b> Currency itself the main goal. Uses proof of work.</p>
~12 sec block time - using ethash mining	~10 min block time - using sha256 mining



# 3.1 ACCOUNTS



# UTXO VS. ACCOUNTS

## THE ULTIMATE COMPARISON

- A Bitcoin user's available **balance** is the **sum of unspent transaction outputs** for which they own the private keys to the output addresses
- Instead Ethereum uses a different concept, called **Accounts**, which already keeps track of **balance**

Bitcoin:

Bob owns private keys to set of UTXOs

5 BTC  $\Rightarrow$  Bob

3 BTC  $\Rightarrow$  Bob

2 BTC  $\Rightarrow$  Bob

Ethereum:

Evan owns private keys to an account

address: "0xfa38b..."

balance: 10 ETH


code:  $c := a + b$




# ACCOUNTS

## HOW ETHEREUM MANAGES IDENTITY

- The global shared state of ethereum is comprised of many small objects (**Accounts**) that are able to interact with one another through a message-passing framework
- An **Account** has state and 20 byte/160 bit byte identifier, i.e:
  - `0x91fff4cbd6159a527ca4dcce2e3937431086c662`
- Two Types of Accounts:
  - **Externally owned**, which are controlled private keys and have no code associated with them.
  - **Contract accounts**, which are controlled by their code and have code associated with them.



[Source](#)



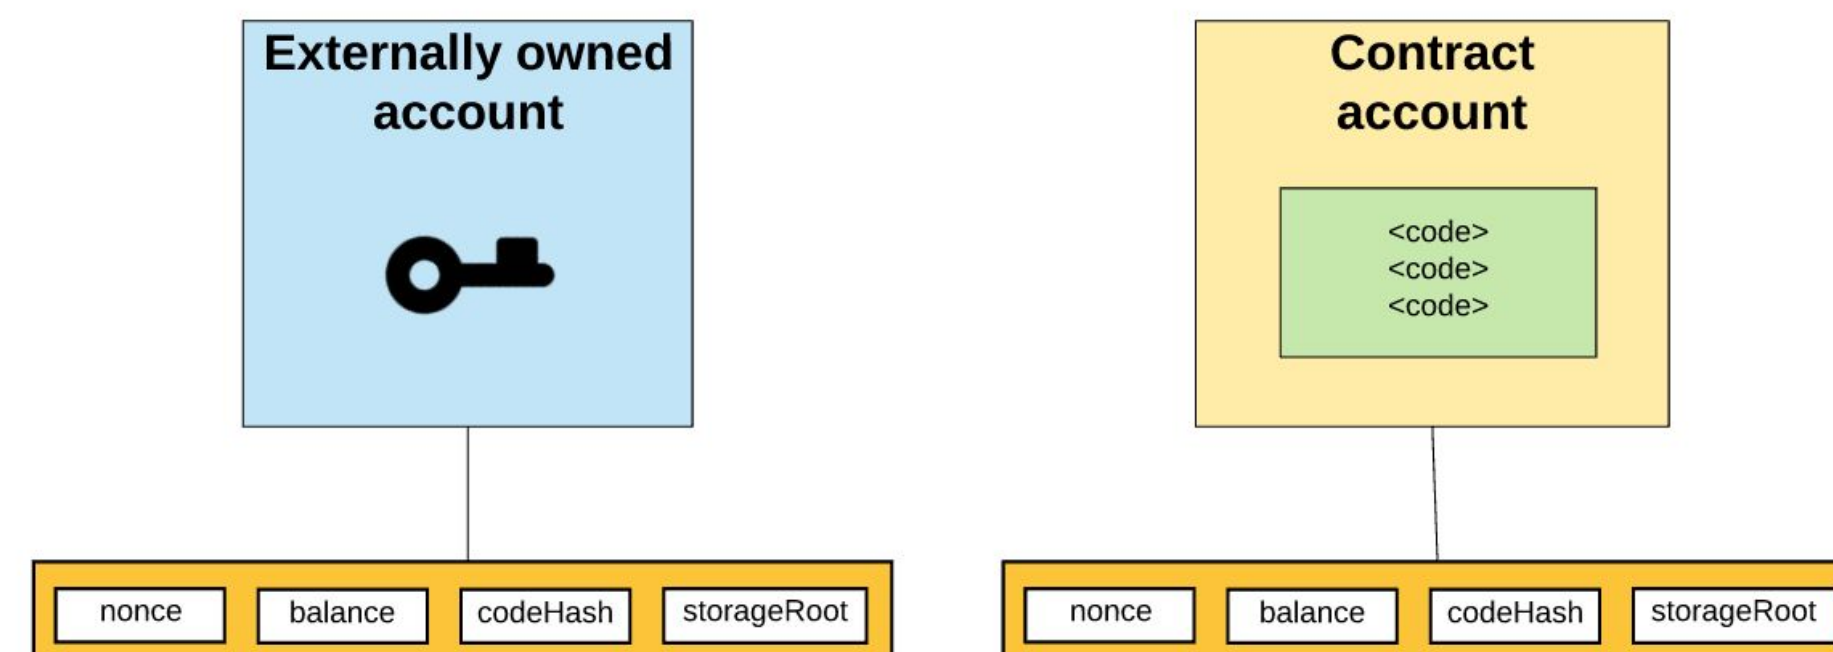




# WHAT IS STORED IN ACCOUNT STATE


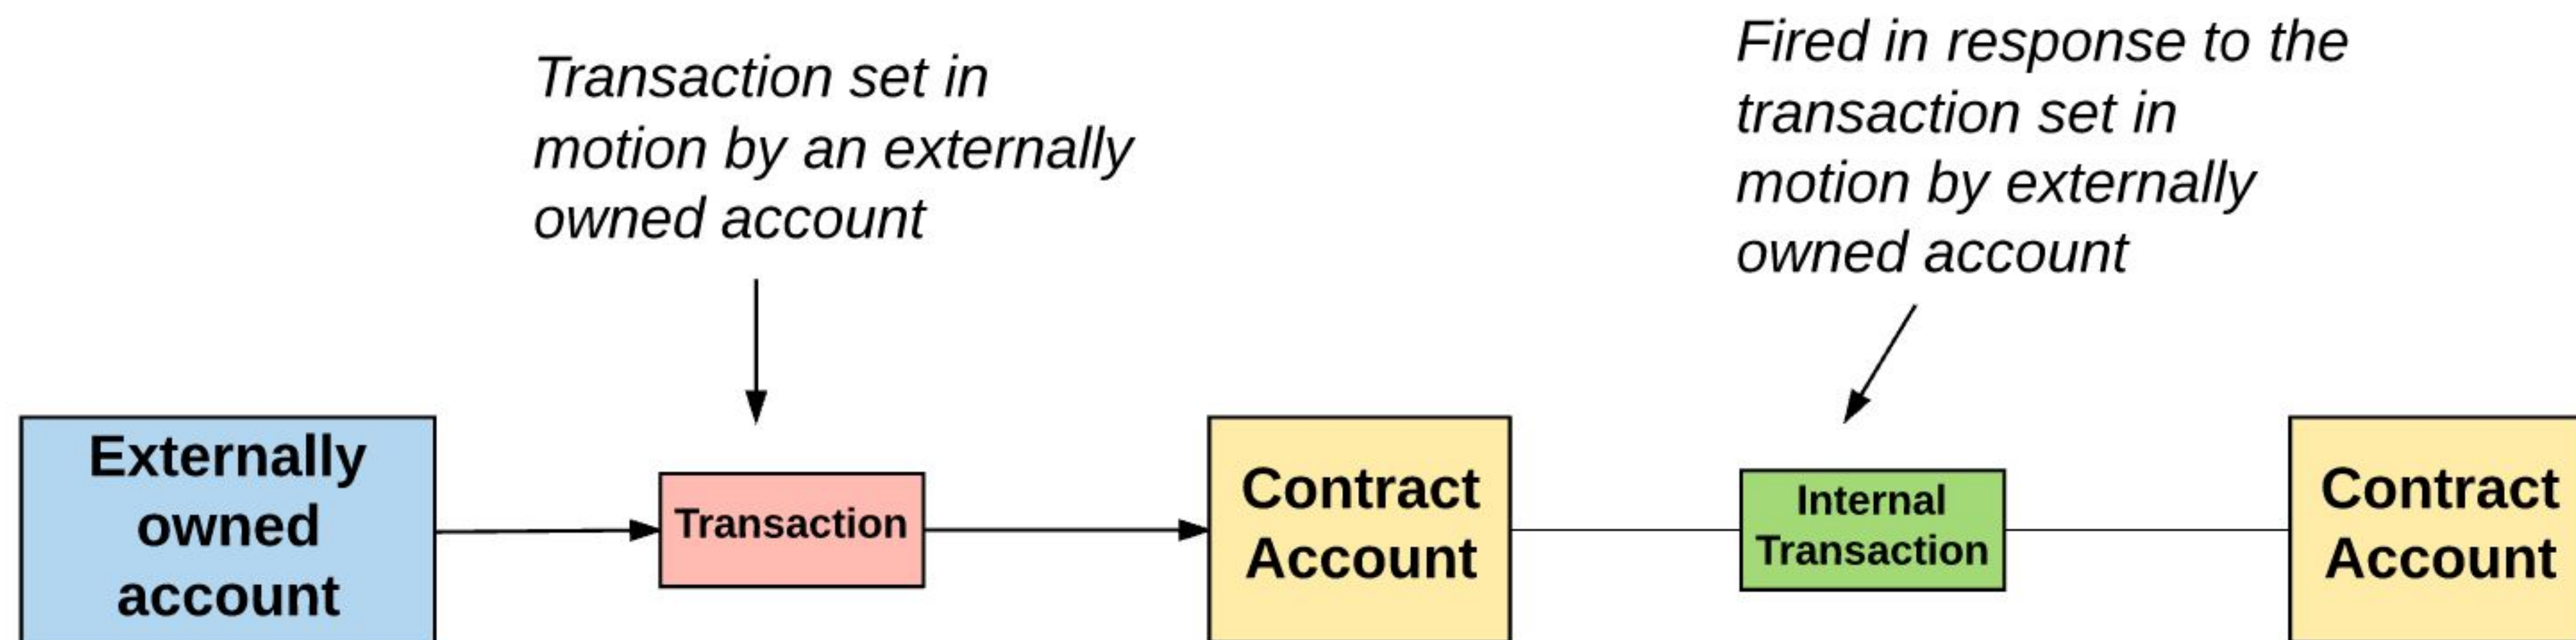
## FOUR COMPONENTS

- **nonce**: number of transactions sent from external account, number of contracts created if contract account
- **balance**: The number of Wei owned by this address. 1e18 Wei per 1 Ether.
- **storageRoot**: A hash of the root node of a Merkle Patricia tree. This tree encodes the hash of the storage contents of this account, and is empty by default.
- **codeHash**: The hash of the EVM (Ethereum Virtual Machine—more on this later) code of contract account. Hash( " " ) for external accounts.






# EXTERNALLY OWNED VS CONTRACT ACCOUNTS



Any action that occurs on the Ethereum blockchain is always set in motion by transactions fired from externally controlled accounts.


[Source](#)



# ALL ACCOUNTS == NETWORK STATE

GLOBAL “TRUTH”

- **State of all accounts is the state of the Ethereum network**
  - Entire Ethereum network agrees on the current balance, storage state, contract code, etc. of every single account
- **Ethereum network state is updated with every block**
  - A block takes the previous state and produces a new network state
    - every node has to agree upon new network state
- **Accounts interact** with network, other accounts, other contracts, and contract state **through transactions**

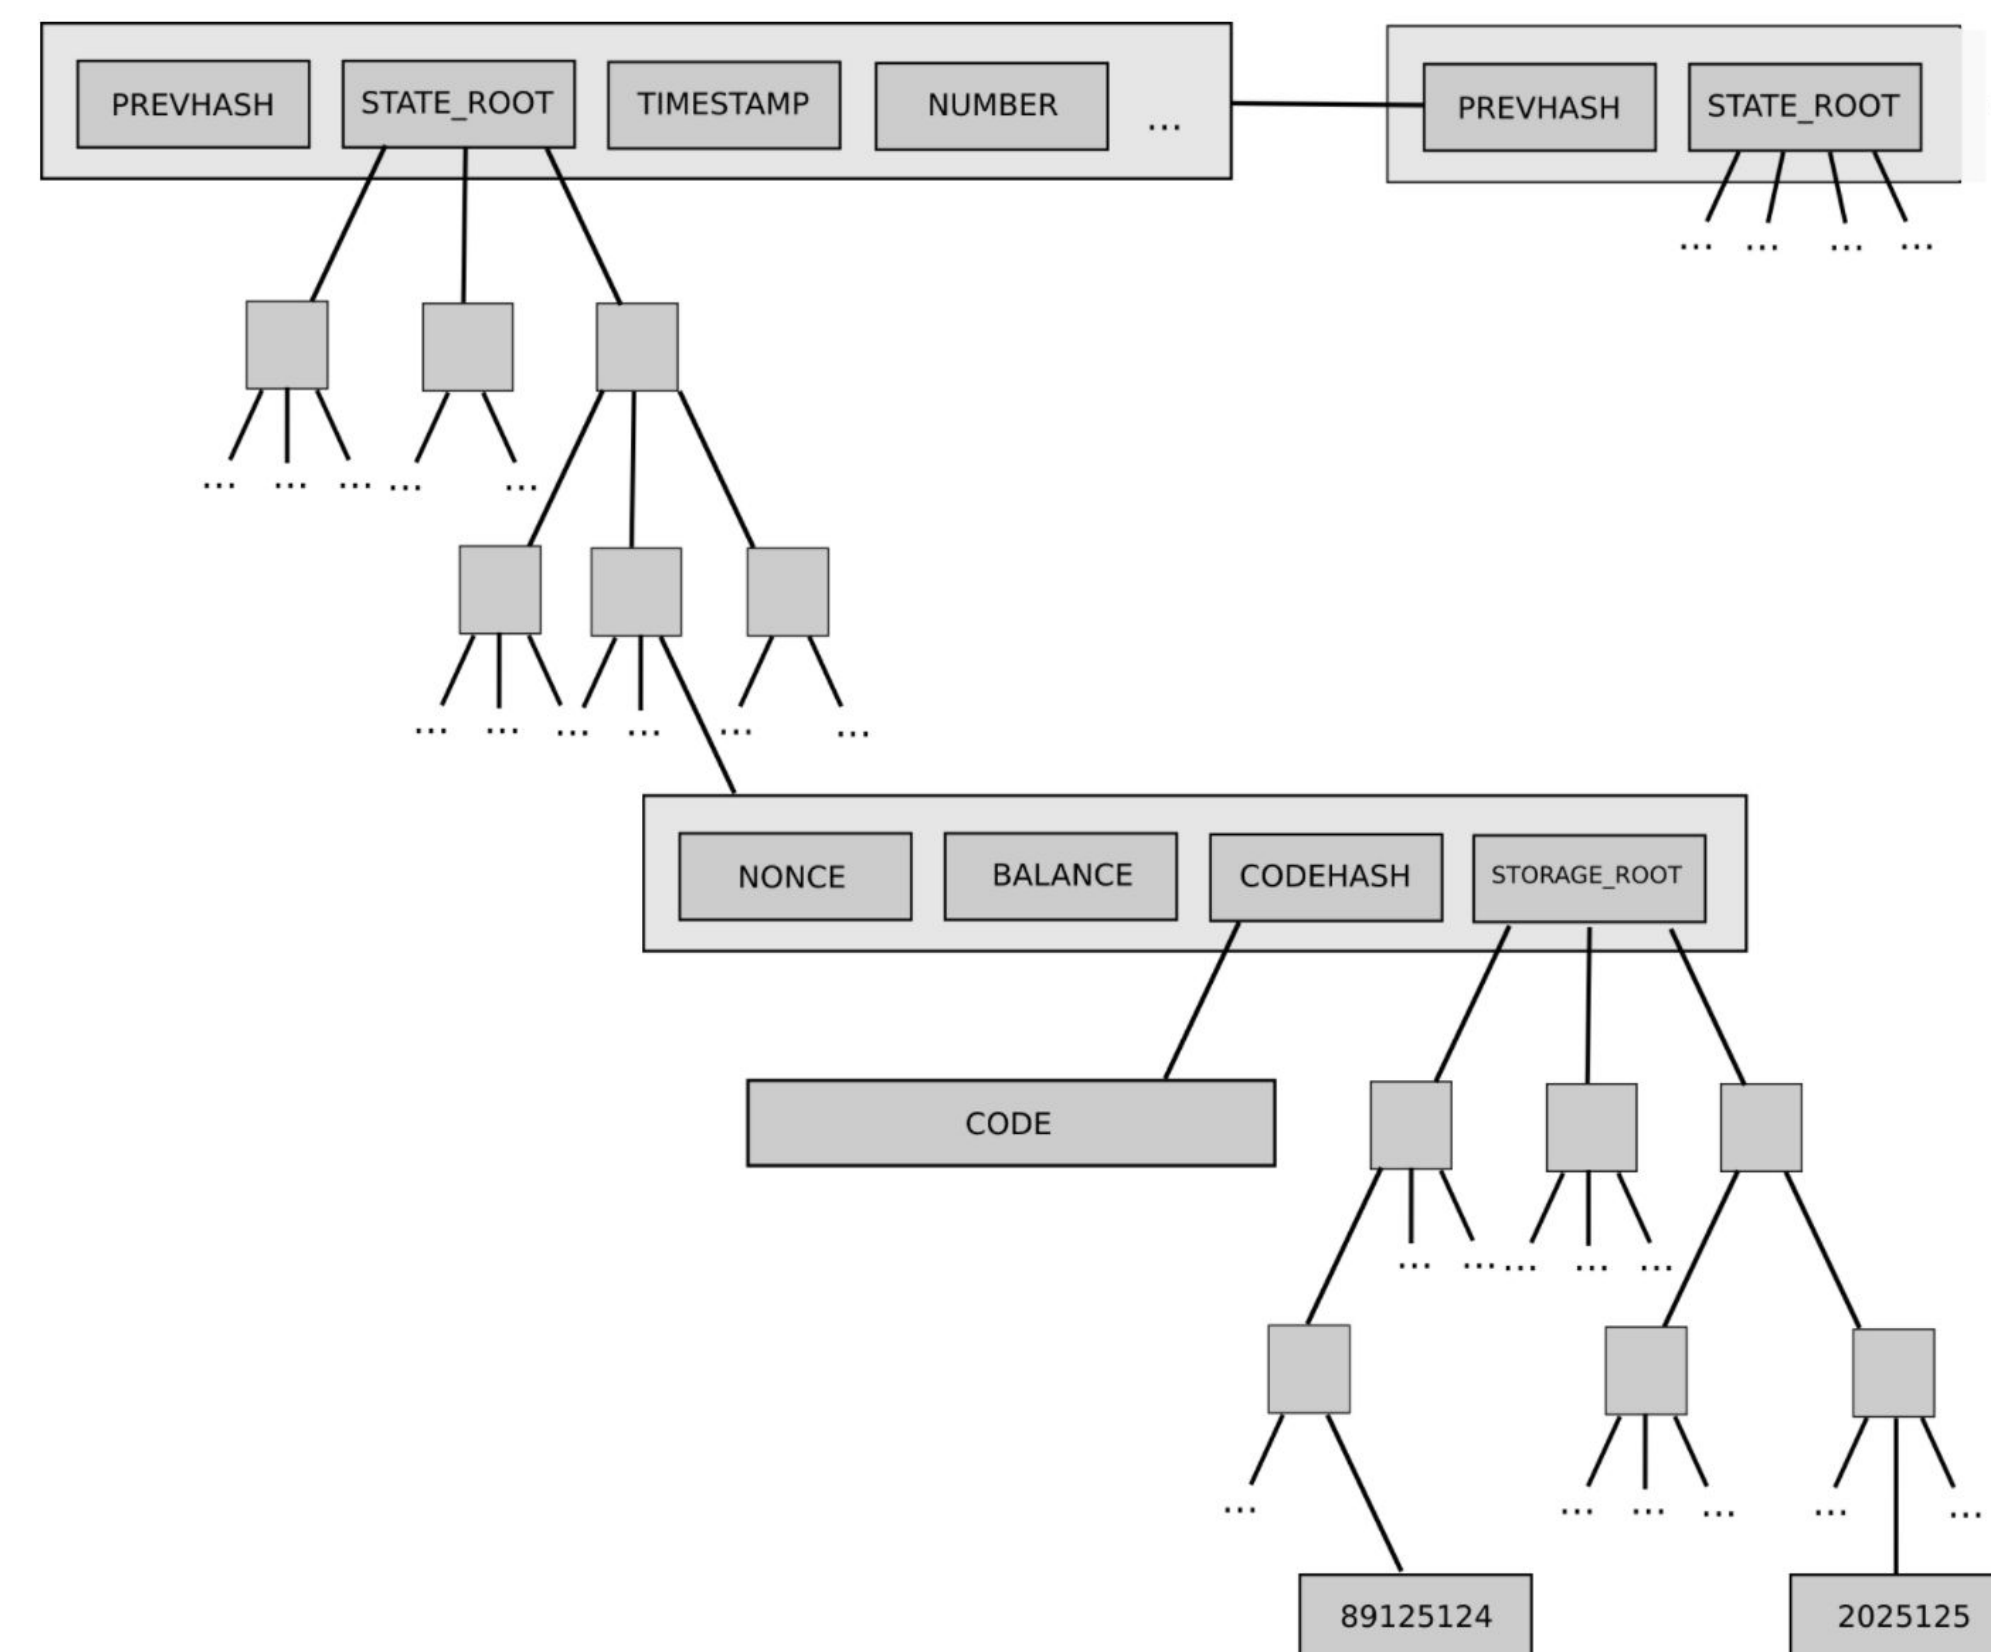




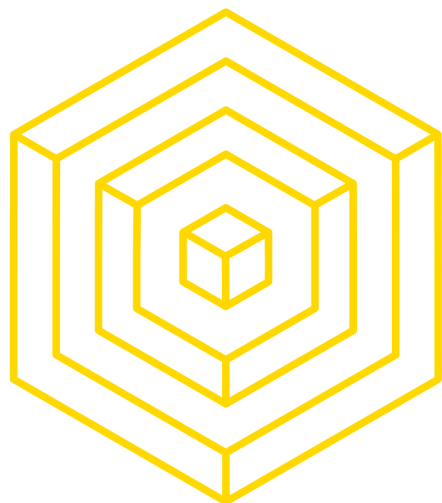
# ACCOUNTS RATIONALE

## WHY USE ACCOUNTS

- **Space Savings**
  - Nodes only need to update each account's balance instead of storing every UTXO
- **More Intuitive**
  - Smart contracts are more easier to program when transferring between accounts with a balance vs. constantly updating a UTXO set to compute user's available balance.
- **Comparable efficiency due to “Merklization”**
  - Use of merkle patricia tries allows for SPV light clients to work



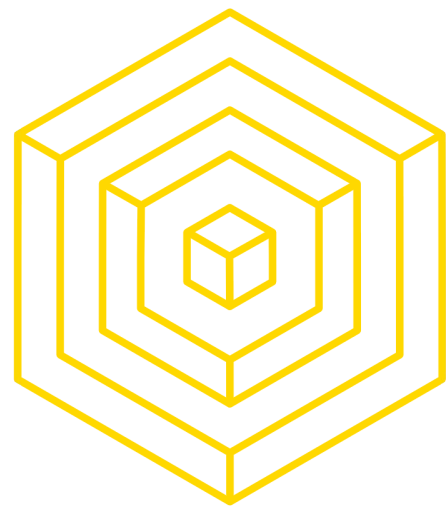




# 3.2

## FEES







# GAS AND PAYMENT

## FEES

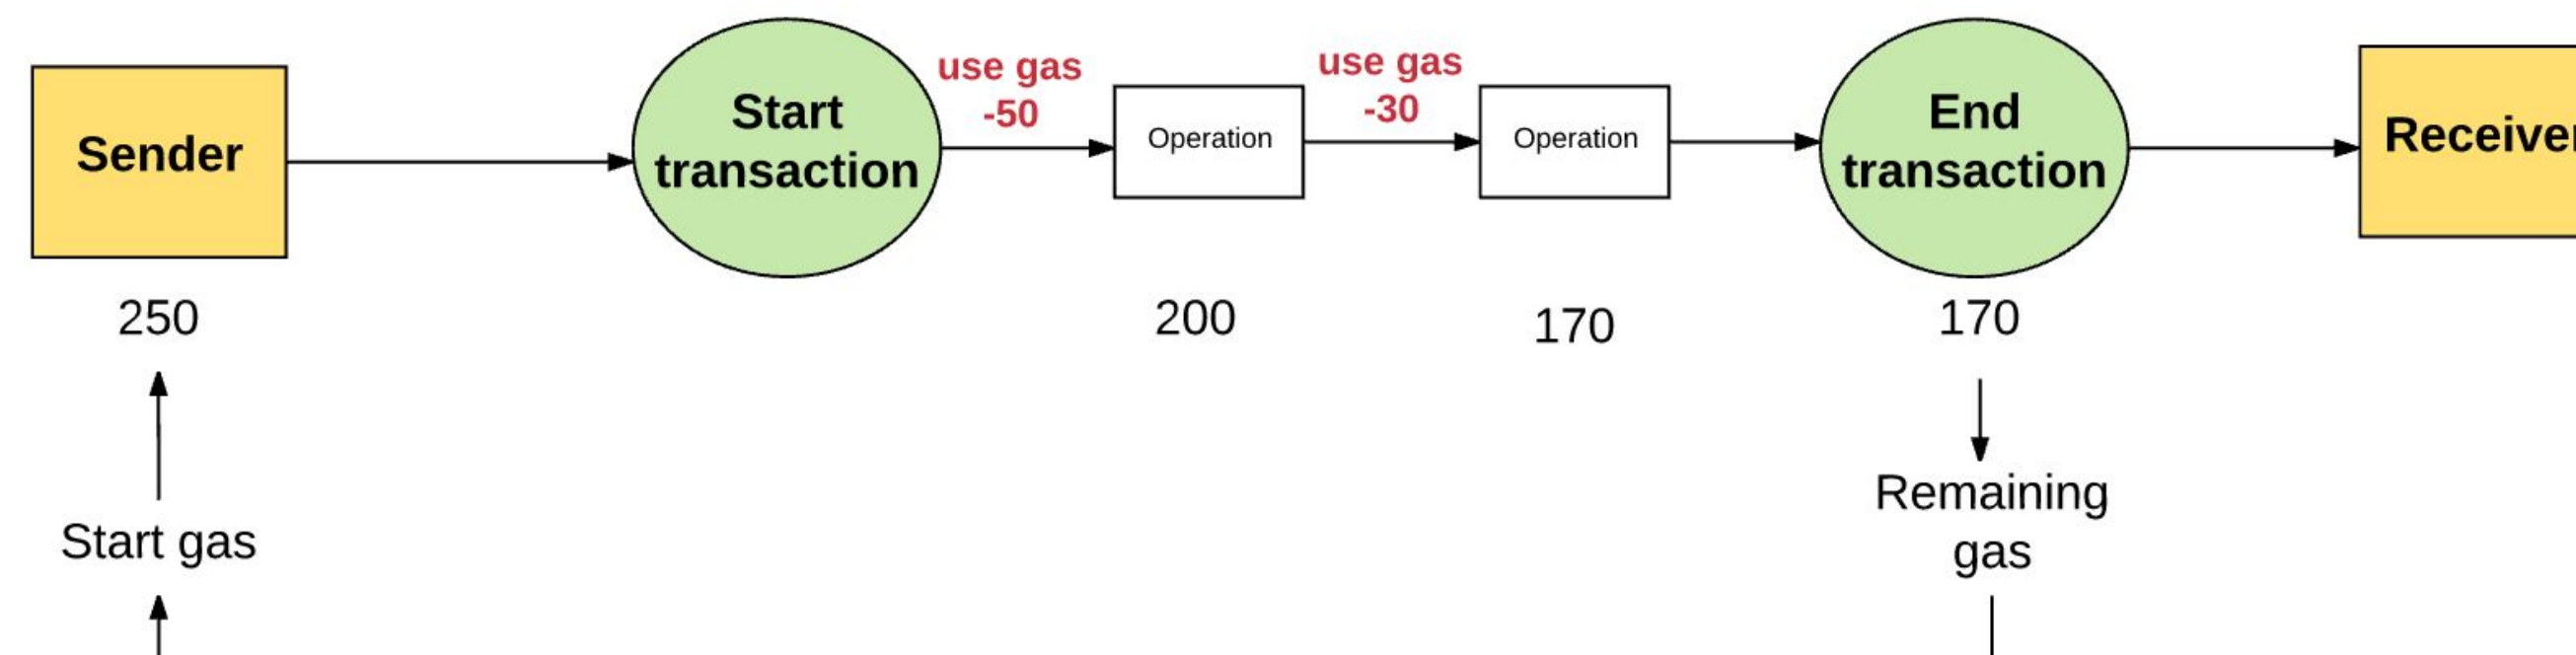
- Every computation that occurs as a result of a transaction on the Ethereum network incurs a fee called **gas**
- **Gas** is the unit used to measure the fees for a particular computation
- **Gas price** is the amount of Ether you are willing to spend on every unit of gas
  - Measured in “gwei” - 1E18 wei = 1 ETH, 1 gwei = 1,000,000,000 wei
- With every transaction, a sender sets a **gas limit** and a **gas price**
  - **gas price** \* **gas limit** = max **amount** of wei **sender is willing to pay** for transaction





# GAS AND PAYMENT FEES

- Example: **Gas Limit:** 50,000, **Gas Price:** 20 gwei
  - Max transaction fee:  $50,000 * 20 \text{ gwei} = 1,000,000,000,000,000,000 = 0.001 \text{ Ether}$
- Since the **gas limit** is the maximum gas the sender is willing to spend money on, if they have enough Ether in their account balance to cover this maximum and the **gas limit** is enough to execute the transaction, transaction will execute
  - `startGas == gasLimit == gasUsed`
- Unused gas is refunded to the sender



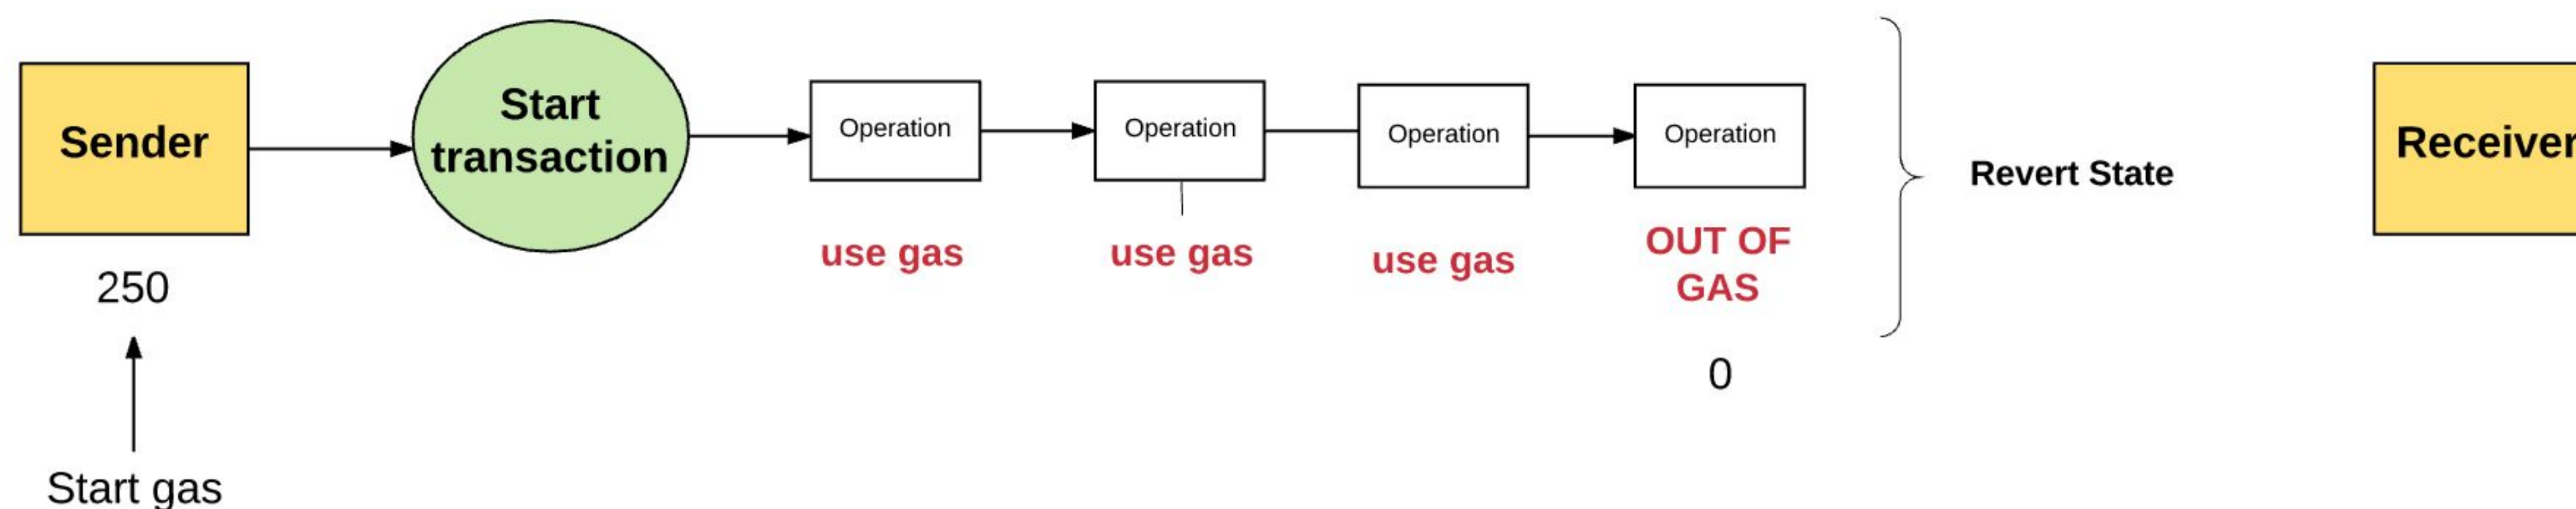
[Source](#)



# GAS AND PAYMENT

## WHAT IF NOT ENOUGH IS SENT?

- Not enough gas to execute the transaction?
  - Transaction runs out of gas and therefore is considered invalid
  - State changes are reversed, failing transaction recorded
  - Since computation was already expended by the network, **none of the gas is refunded**



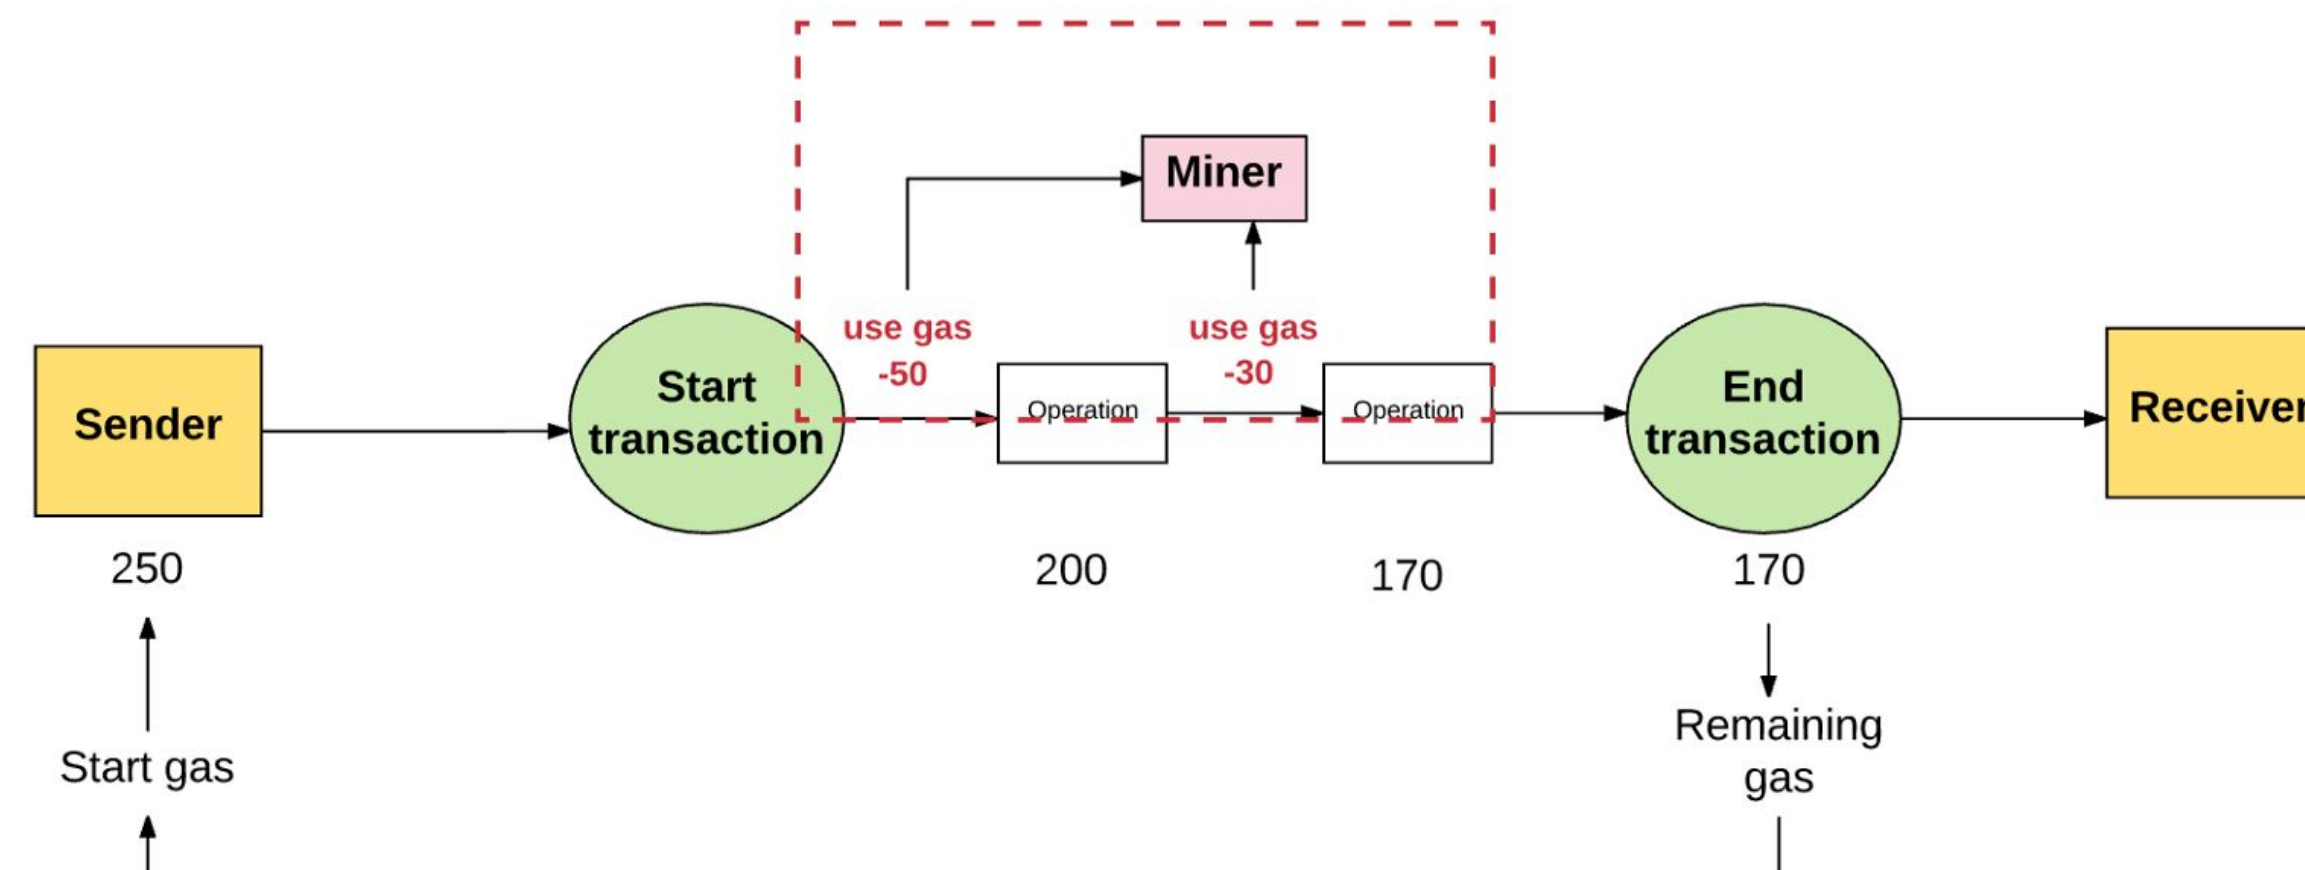
[Source](#)



# GAS AND PAYMENT

## MINER INCENTIVES

- Where does the fee go?
  - **The miner's address**, since they are expanding the effort to run computations and validate transactions - it's the **reward**!
- The higher the gas price the sender is willing to pay, the greater the value the miner derives
  - Therefore more likely to select in block, as miners can choose which transactions to validate and ignore (gas price advertising also possible)




[Source](#)



# STORAGE FEES

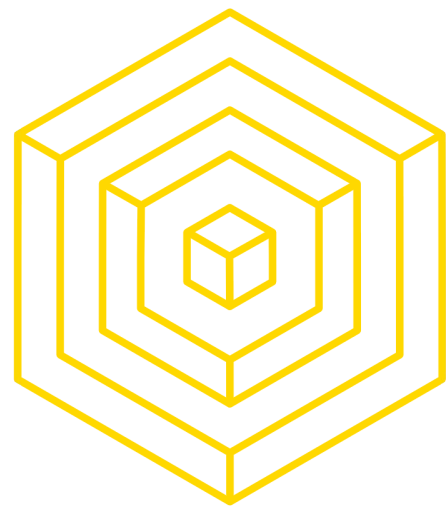
## NOT JUST FOR COMPUTATION

- **Gas** is also used to pay for **storage**, proportion to the smallest multiple of **32 bytes** used
- Increased storage however increases the size of the Ethereum state on all nodes
  - **Therefore incentives to keep data small**
- So if a transaction has a step that clears an entry in storage, the fee is waved and a **refund** is given for free storage space



BLOCKCHAIN FOR DEVELOPERS






# WHY DO WE NEED FEES?

## PARALLEL, DISTRIBUTED COMPUTATION

- Every single operation executed by the network is simultaneously affected by every full node
  - Computation steps on the EVM are therefore very expensive
- **Smart contracts are therefore best used for simple tasks**
  - Business logic or verifying signatures rather than machine learning or file storage
  - Redundantly parallel, no asynchronous or performant parallel execution
- **Turing completeness** allows for loops and susceptibility to the **halting problem**
  - **Halting problem:** inability to determine whether or not a program will run indefinitely
  - No fees means DDOS through infinite loops

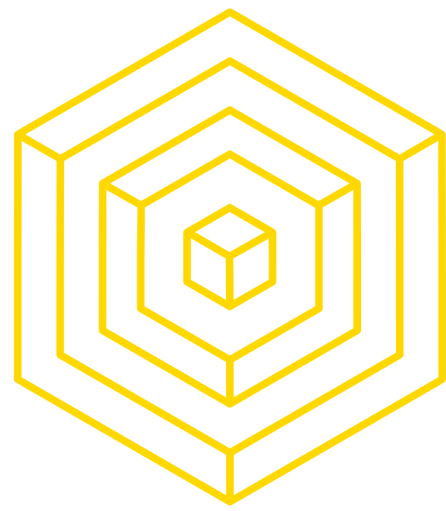


SIGN IN:  
**[tinyurl.com/developers2](https://tinyurl.com/developers2)**

CODE: FEES



# 3.2 TRANSACTIONS AND MESSAGES

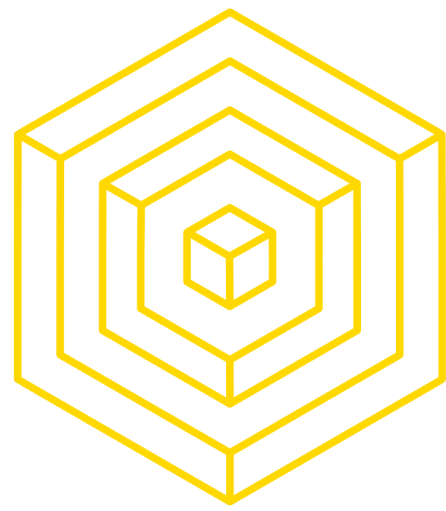


# WHAT IS A TRANSACTION

## STATE CHANGERS

- Transactions move the state of an account within the global state - one state to the next
- **Formal Definition:** A transaction is a cryptographically signed piece of instruction that is generated by an externally owned account, serialized, and then submitted to a blockchain
- Two types: **Message calls and contract creations**



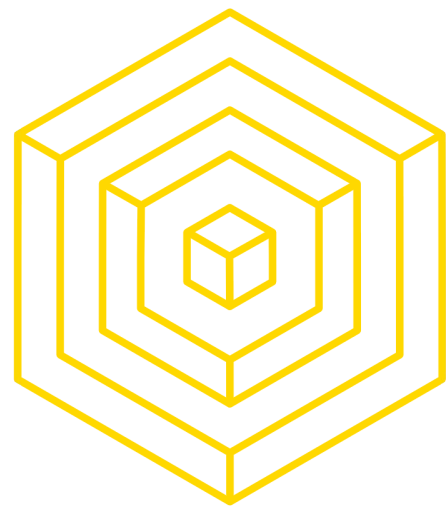


# TRANSACTION COMPONENTS

## STATE CHANGERS

- **nonce**: number of transactions sent by sender
  - **gasPrice**: amount of Wei sender is willing to pay per unit of gas required to execute the transaction
  - **gasLimit**: max amount of gas the sender is willing to pay for executing this transaction, set before any computation is done
  - **to**: address of the recipient
  - **value**: the amount of Wei to be transferred from the sender to the recipient
  - **v, r, s**: used to generate the signature that identifies the sender of the transaction.
- init** (only exists for contract-creating transactions): An EVM code fragment that is used to initialize the new contract account
- **data** (optional field that only exists for message calls): the input data (i.e. parameters) of the message call






# MESSAGE CALLS

## TRANSACTIONS

- Both **message calls** and **contract creating** transactions are always initiated by **externally owned accounts**
  - Transactions bridge the external world to the internal state of Ethereum
- Contracts that exist within the global scope of Ethereum can talk to other contracts using **messages** (internal transactions) to other contracts
- We can think of messages as being similar to transactions, except they are not generated by externally owned accounts, only by contracts
  - Virtual objects within the Ethereum execution environment

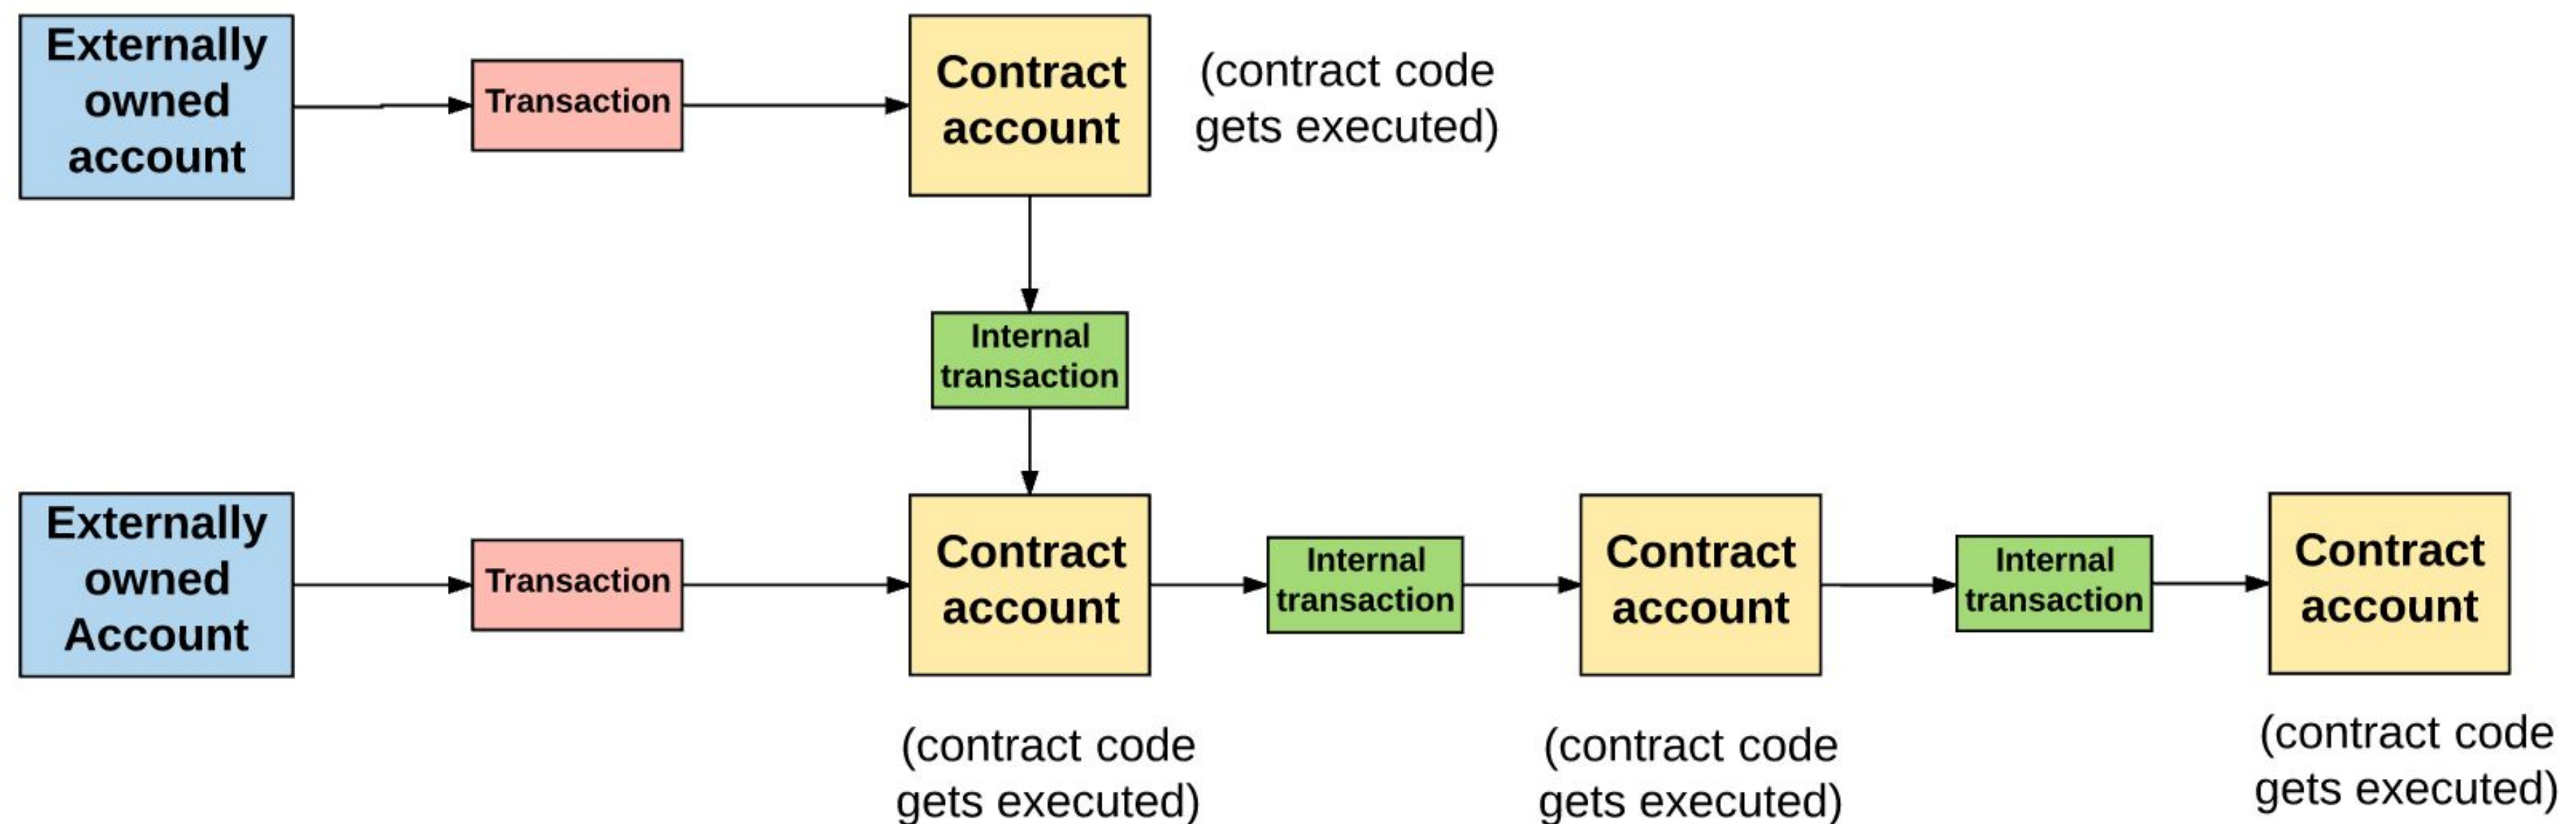




# NESTED MESSAGE CALLS

## TRANSACTIONS

When one contract sends an internal transaction to another contract, the associated code that exists on the recipient contract account is executed.




[Source](#)



# NESTED MESSAGES

## INTERNAL TRANSACTIONS

- Messages do not contain a **gasLimit**
- **gasLimit** determined by the external creator of the original transaction
  - Therefore the **gasLimit** that the externally owned account sets must be high enough to carry out the transaction and any other sub executions that occur as a result of that transaction
  - i.e. Contract-to-Contract messages
- **What if we run out of gas within a parent execution?**
  - Current and subsequent message executions will revert, however the parent execution need not revert





## 4

# SMART CONTRACTS




# WHAT IS A SMART CONTRACT?

## AUTONOMOUS AGENTS

In Ethereum, smart contracts are executed by the network itself!

- **Network consensus** removes need for Trusted Third Party
- Violation of contracts requires **subverting the entire network**
- Allows for secure Peer-to-Peer agreements that live on the blockchain forever







# WHAT IS A SMART CONTRACT?

## CODE YOU CAN POKE

Contracts in Ethereum are like autonomous agents that live inside of Ethereum network

- React to external world when "poked" by transactions (which call functions)
- Have direct control over:
  - internal ether balance
  - internal contract state
  - permanent storage





# SMART CONTRACTS

## SERVERLESS EXECUTION PROTOCOL

- Ethereum Contracts generally serve four purposes:
  - **Store and maintain data**
    - representing something useful to users or other contracts
    - ex: a token currency or organization's membership.
  - **Manage contract or relationship between untrusting users**
    - ex: financial contracts, escrow, insurance.
  - **Provide functions to other contracts**, serving as a software library.
  - **Complex Authentication**
    - ex: M-of-N multisignature access





# APPLICATIONS

## WHAT'S COOL?

- Token Systems
- Financial derivatives and stable value currencies
- Identity and Reputation Systems
- Decentralized File Storage
- Decentralized Autonomous Organizations
- Savings Wallet
- Decentralized Data Feed
- Multisignature Escrow
- Peer-to-Peer Gambling
- Prediction Markets





# TO BE DISCUSSED LATER...

## PROTOCOL SEGMENT

- Ethereum Block Structure
  - Logs
- Omners and Uncles in Ethereum
- GHOST/SPECTRE Protocols
- Merkle Patricia Tries
- Block Difficulty
- Ethereum Execution Model via the EVM
- Formal State Transitions to Finalize a Block
- Ethereum Proof of Work