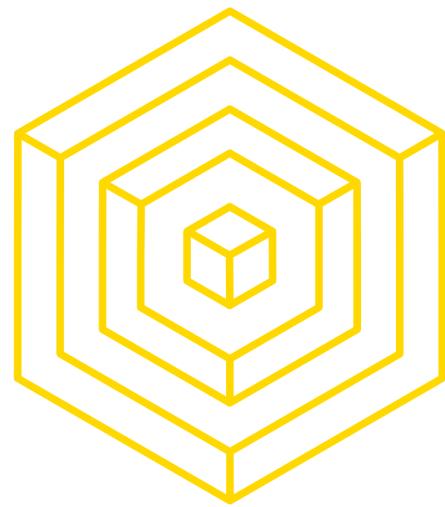


ALTERNATIVE CONSENSUS MECHANISMS

Nadir Akhtar
Brian Ho

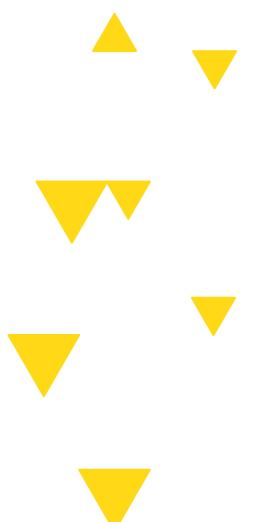


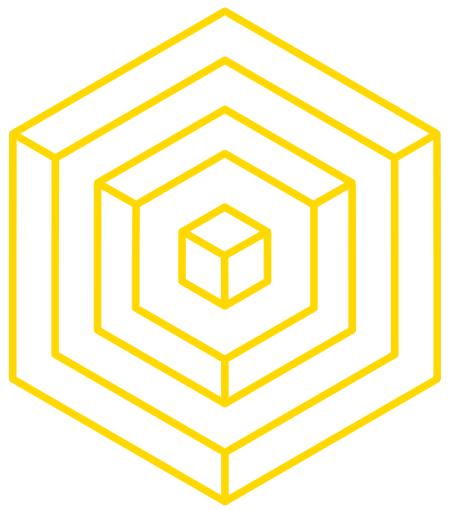
BLOCKCHAIN
AT BERKELEY



LECTURE OVERVIEW

- 1 ► DISTRIBUTED SYSTEMS
- 2 ► PROOF-OF-SOMETHING-ELSE
- 3 ► PROOF-OF-STAKE
- 4 ► VOTING-BASED
CONSENSUS ALGORITHMS
- 5 ► FEDERATED
CONSENSUS

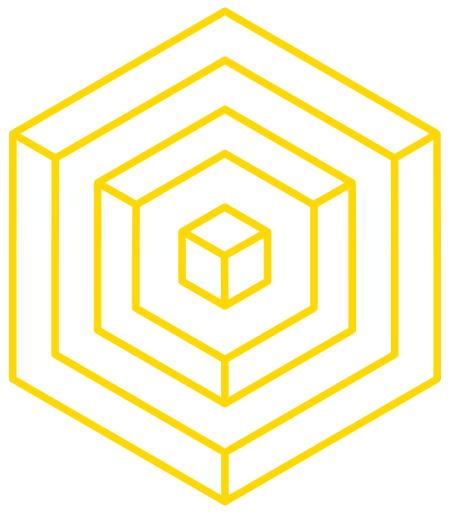




1

DISTRIBUTED SYSTEMS

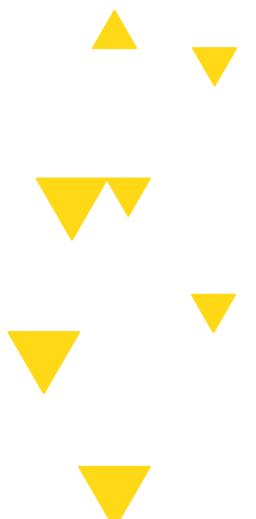
BLOCKCHAIN FUNDAMENTALS LECTURE 8

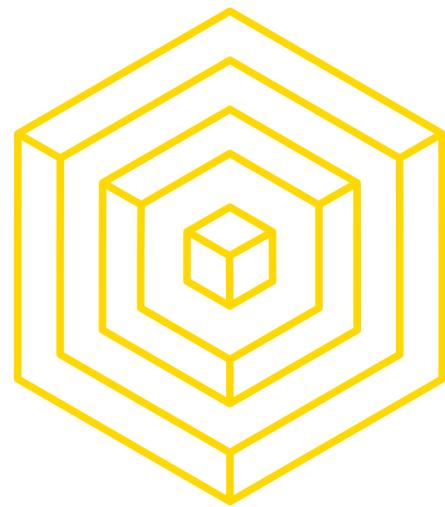


1.1

DIST. SYSTEMS FUNDAMENTALS

BLOCKCHAIN FUNDAMENTALS LECTURE 8

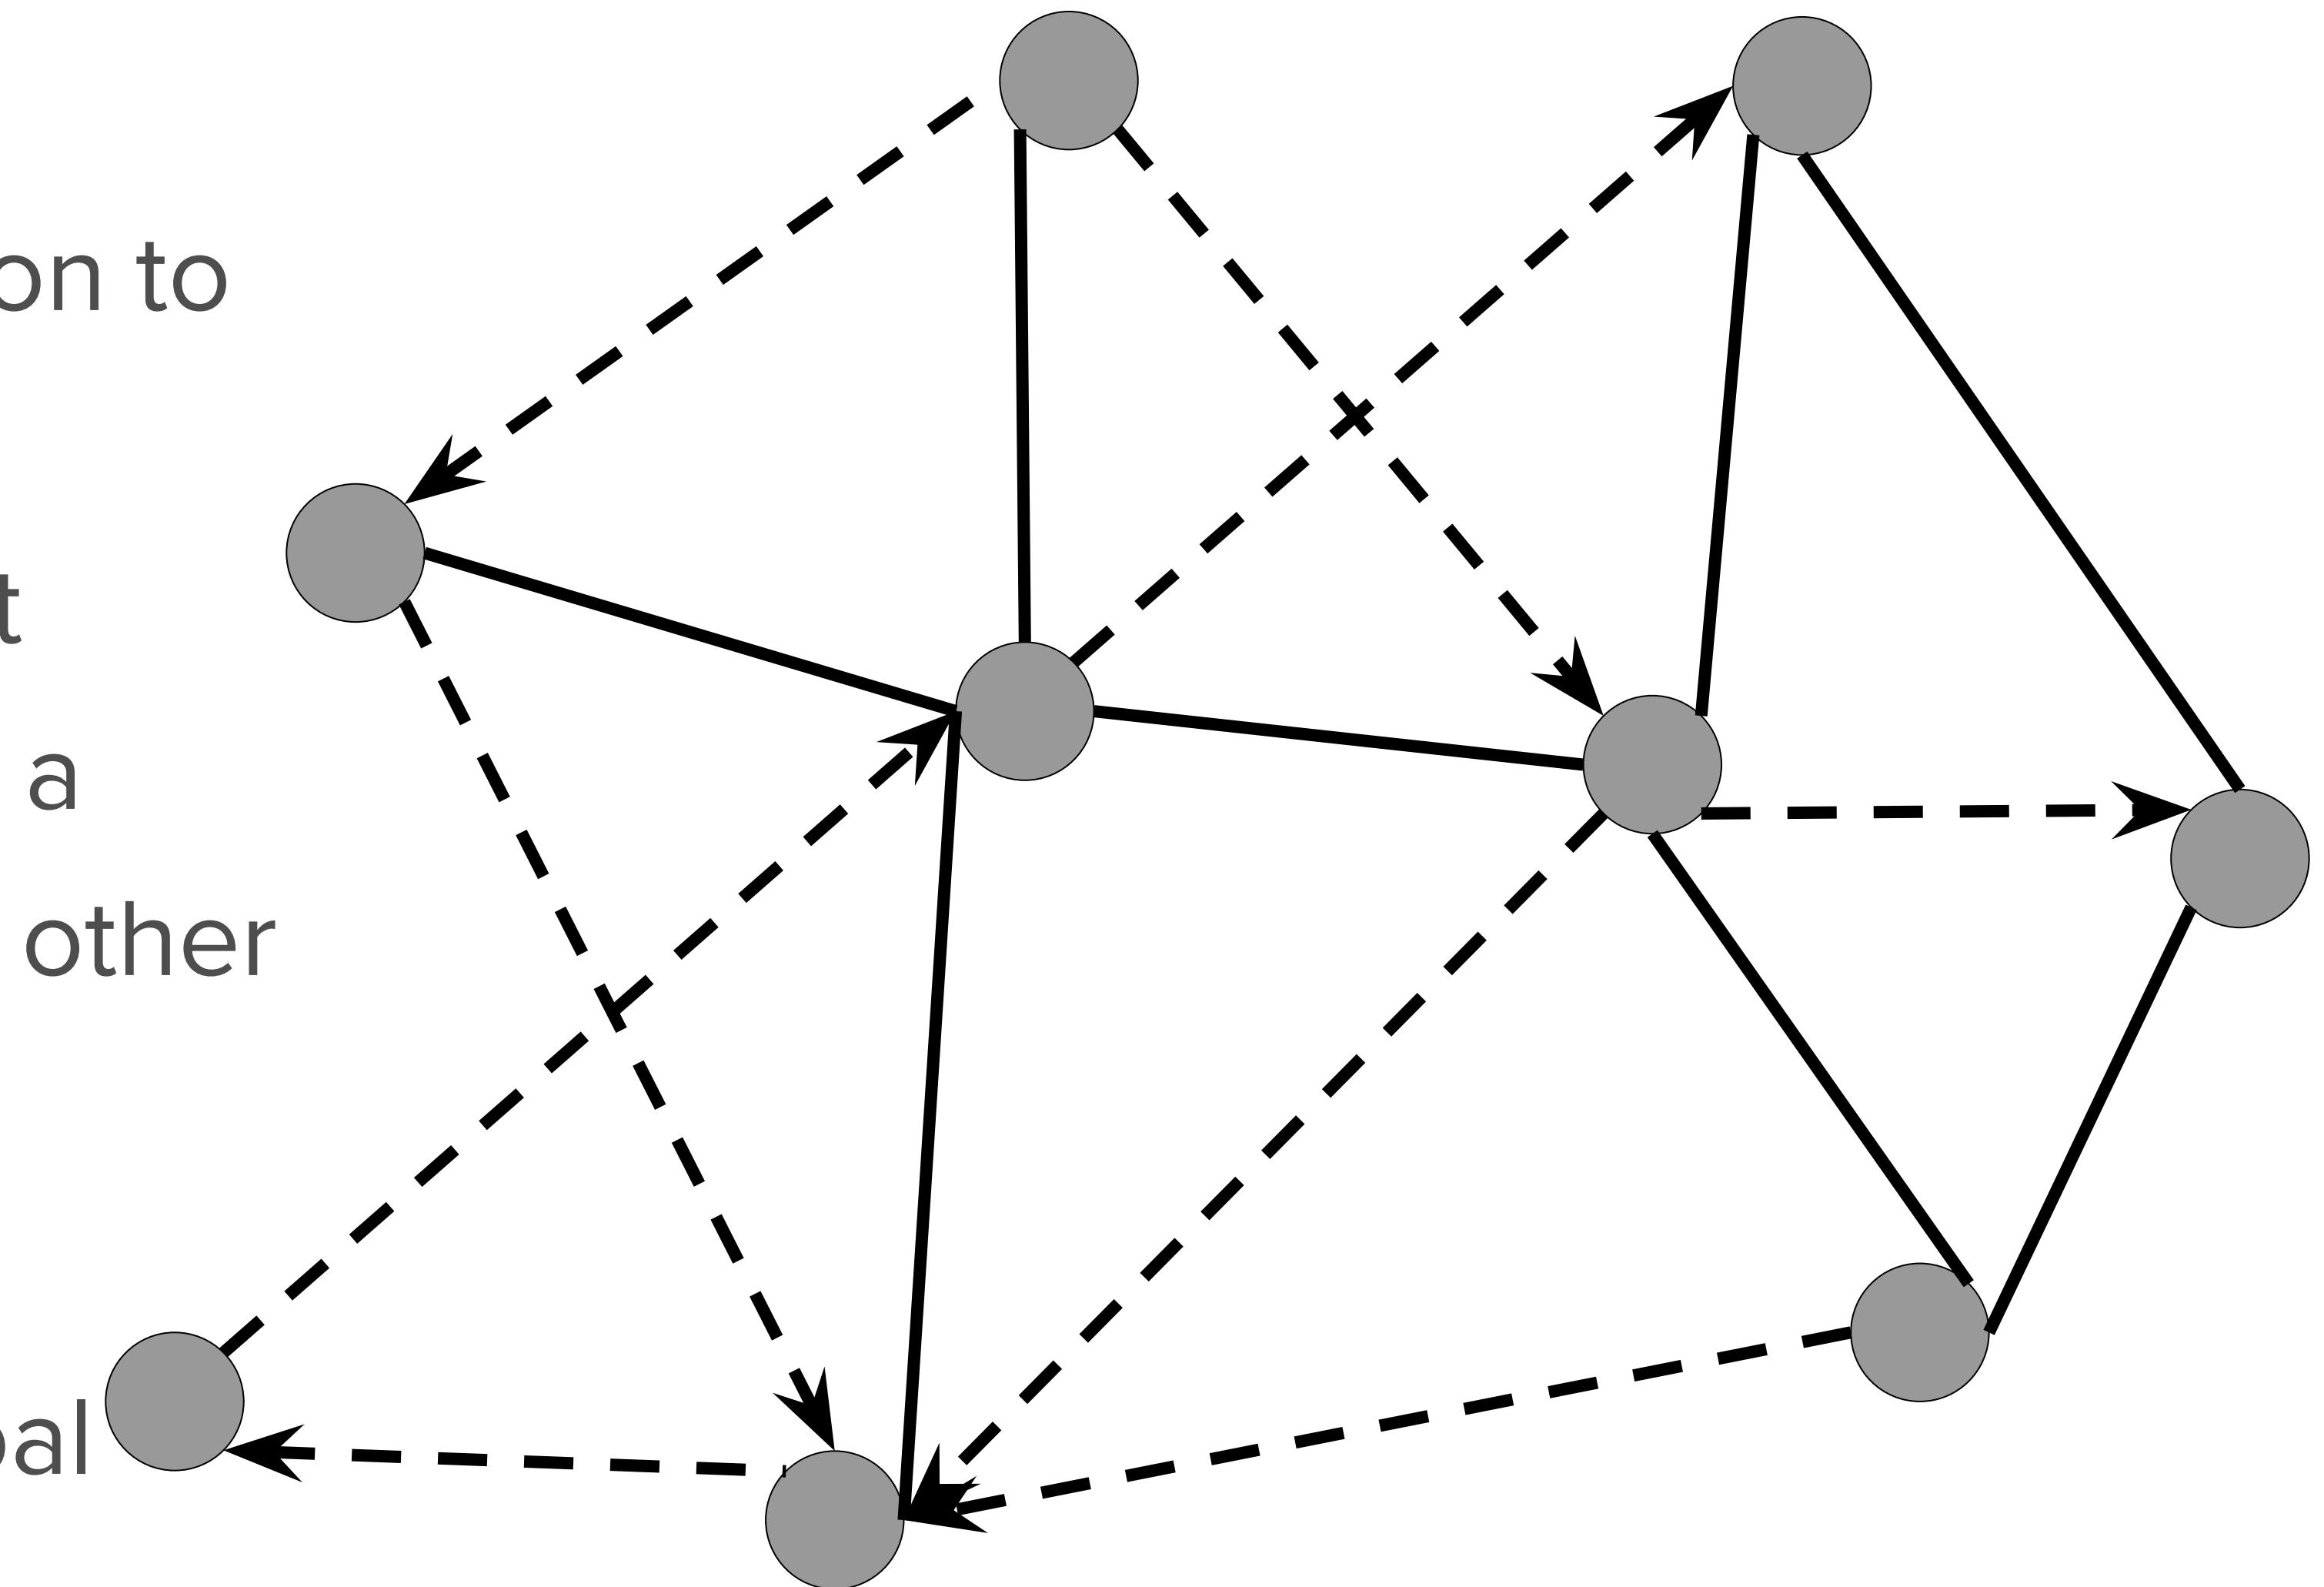


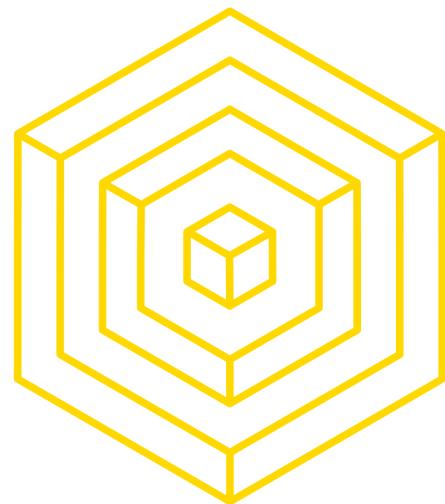


DISTRIBUTED SYSTEMS

INTRODUCTION

- Definition (varies from person to person):
 - A network of independent nodes, each representing a “process,” talking to each other via messages
- Purpose:
 - ▲ ○ Accomplish a common goal

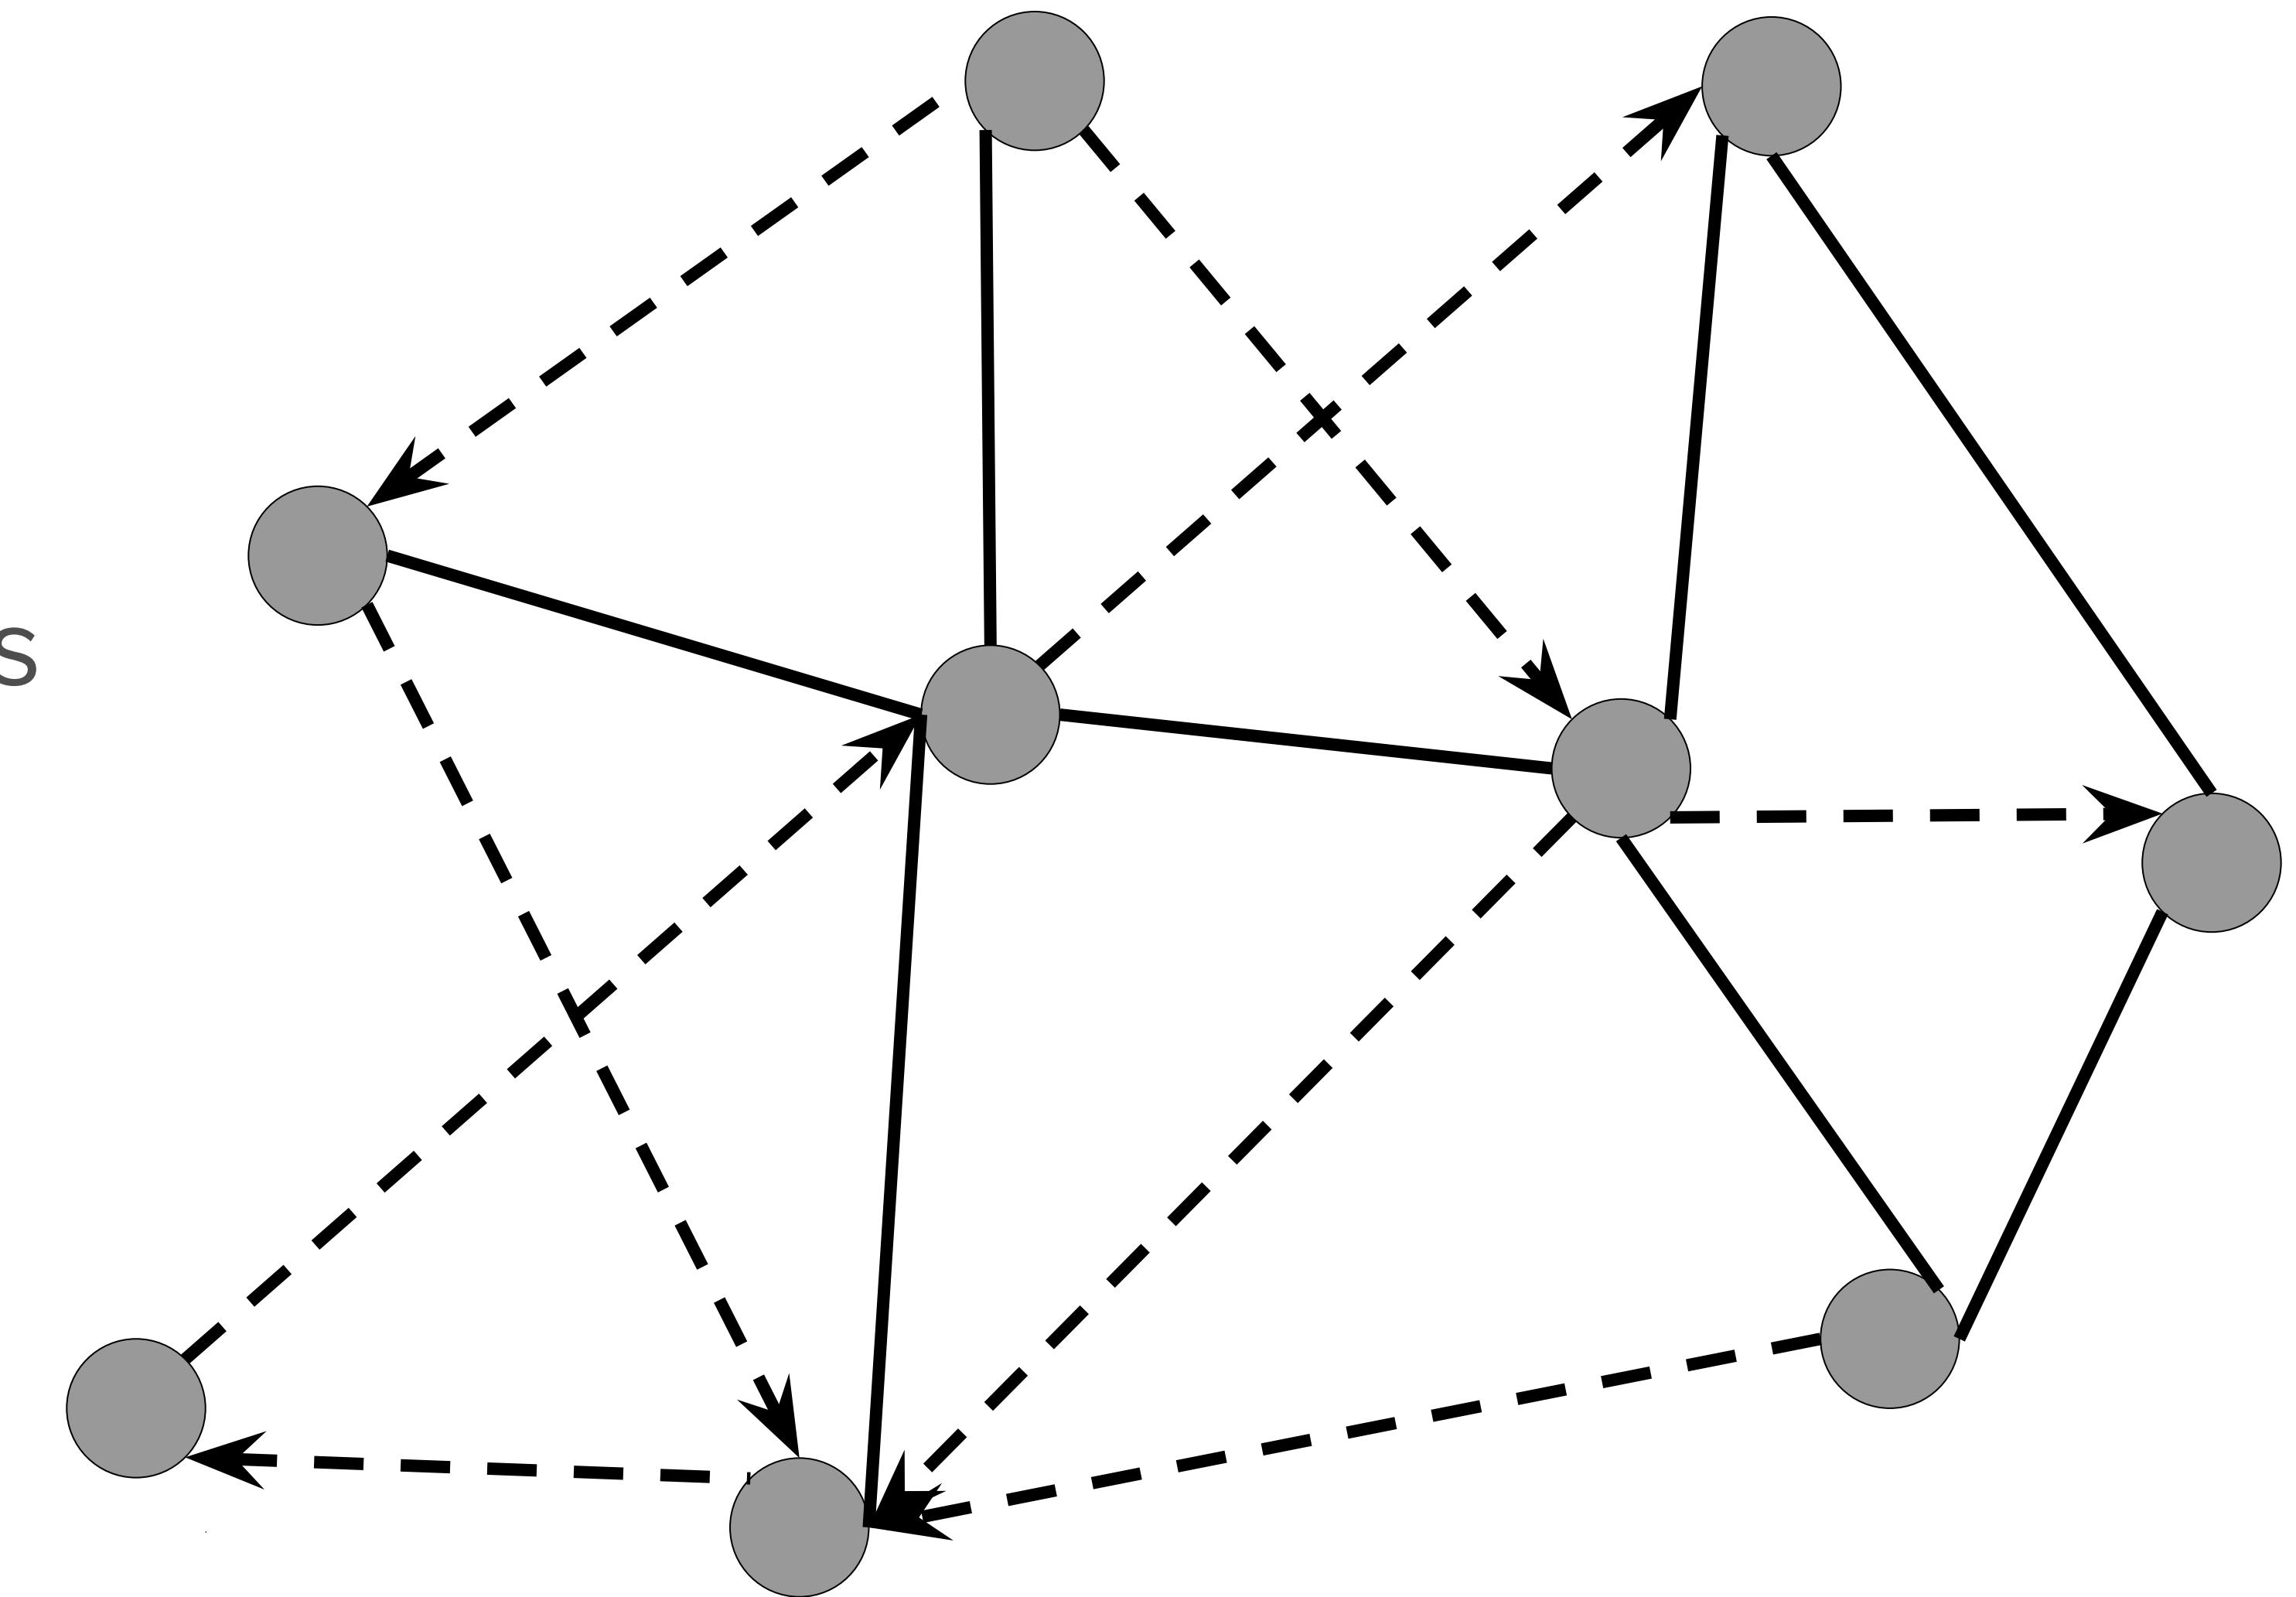


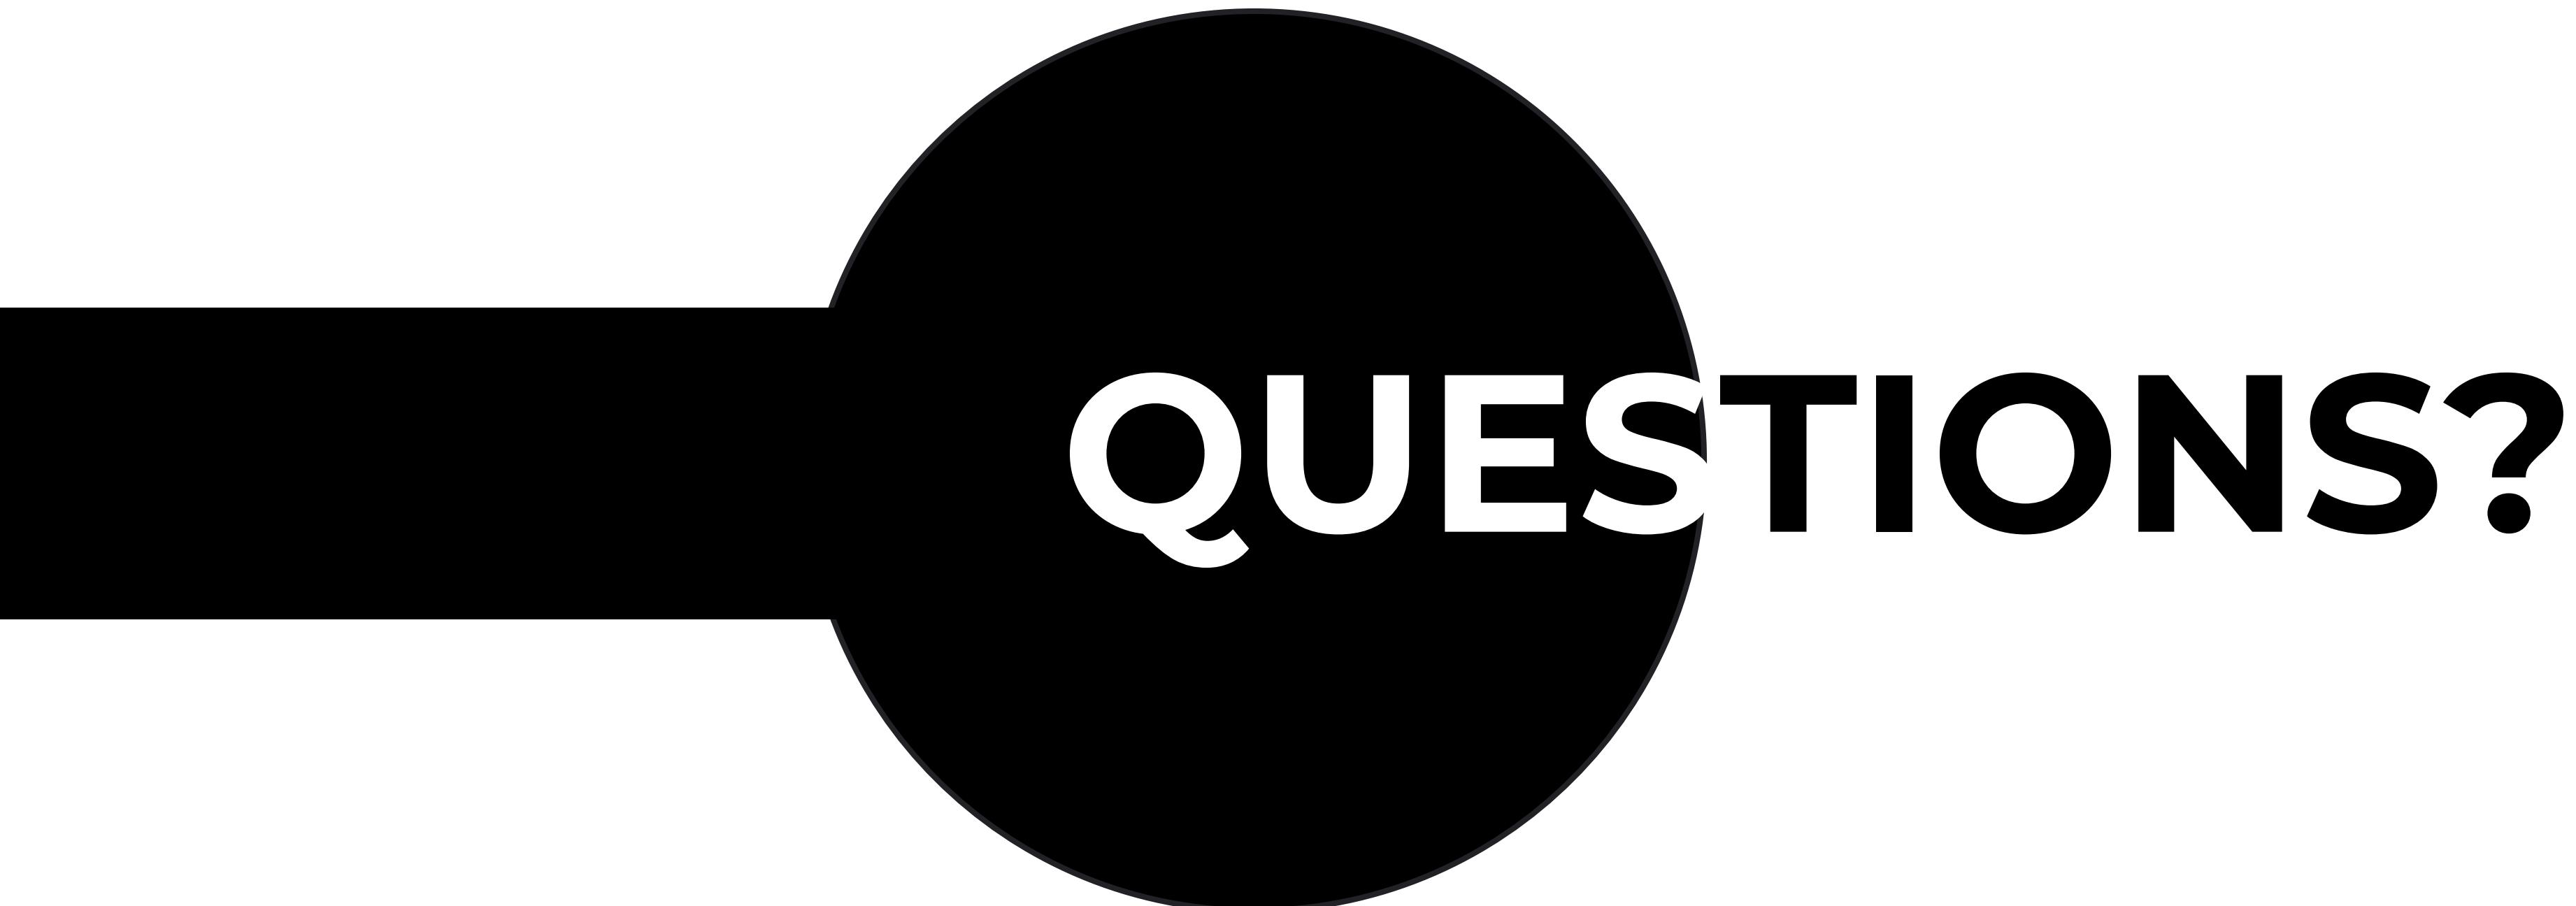


DISTRIBUTED SYSTEMS

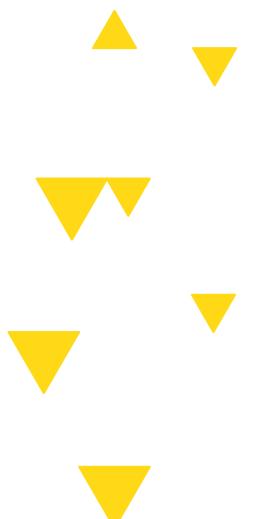
INTRODUCTION

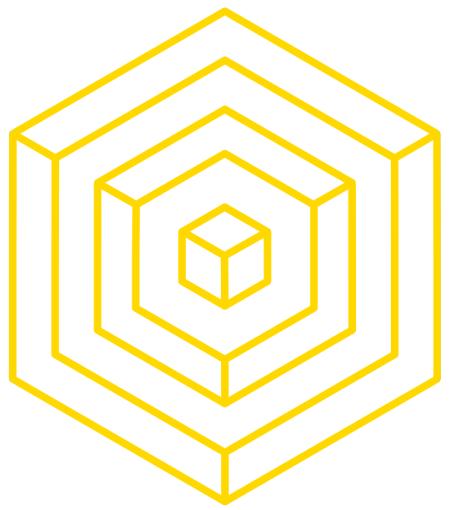
- Properties of distributed systems:
 - Concurrent components
 - Message sharing
 - No global clock
 - Potential failure of individual components





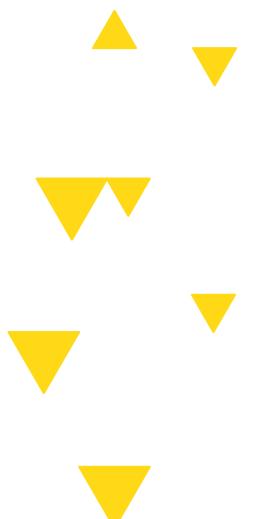
QUESTIONS?

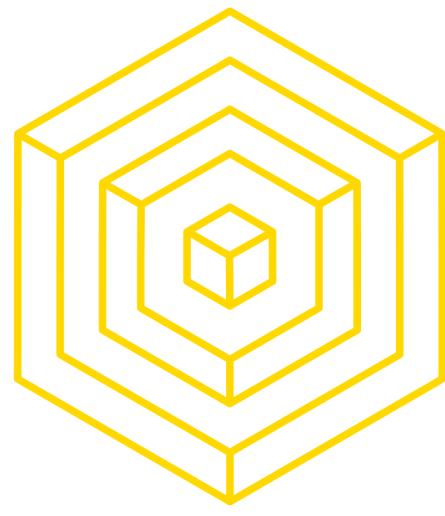




1.2

DIST. SYSTEMS PROPERTIES





DISTRIBUTED SYSTEMS

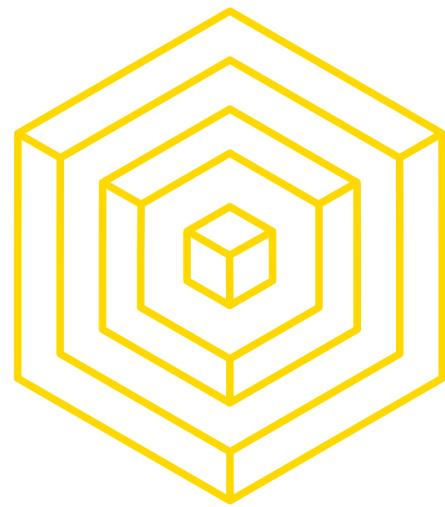
PROPERTIES OF MULTIPROCESS PROGRAMS

Category of Property	Safety	Liveness
Definition	This will <i>not</i> happen	This <i>must</i> happen

Source:
<https://www.microsoft.com/en-us/research/uploads/prod/2016/12/Proving-the-Correctness-of-Multiprocess-Programs.pdf>

AUTHOR: NADIR AKHTAR

BLOCKCHAIN FUNDAMENTALS LECTURE 8



DISTRIBUTED CONSENSUS

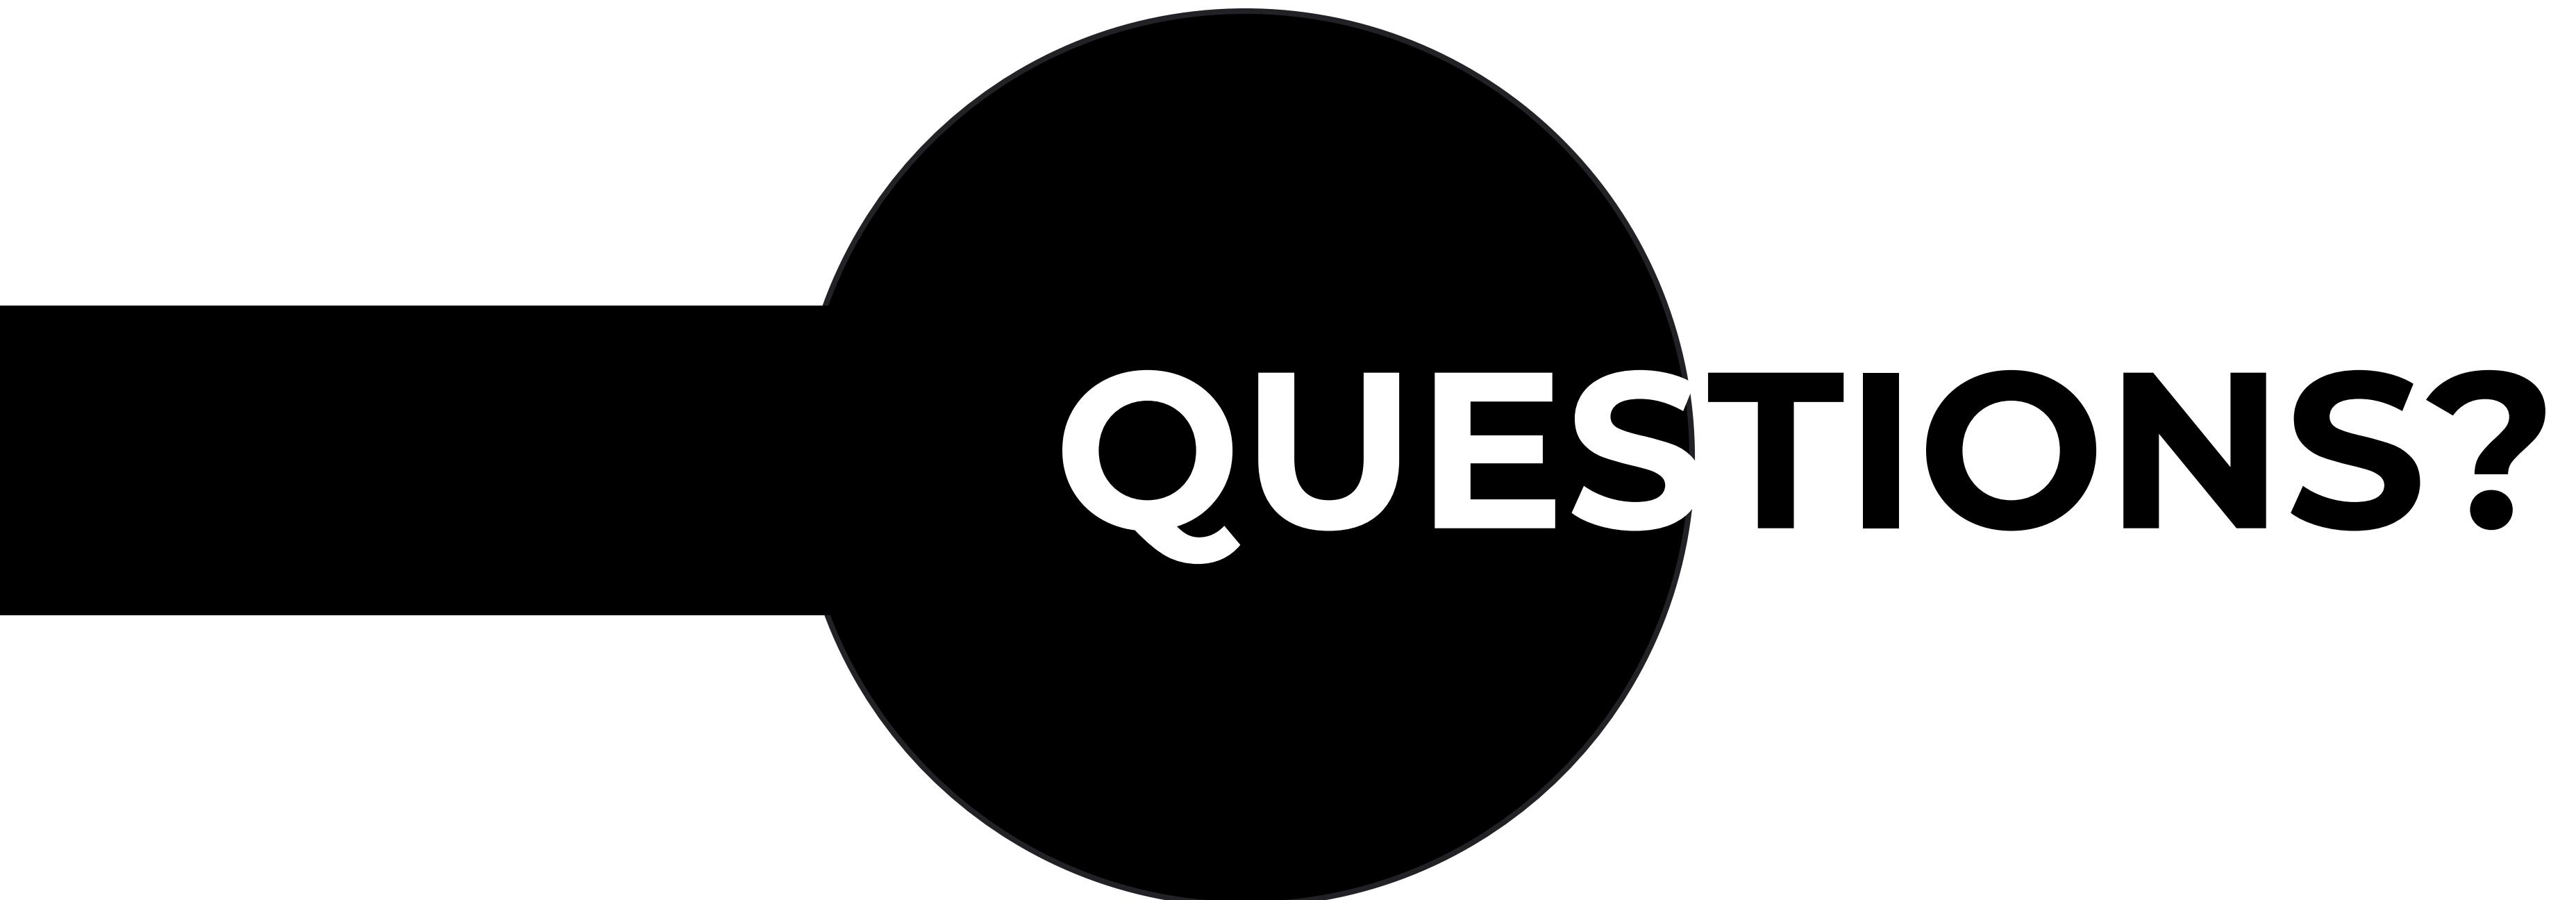
CORRECTNESS

Correctness of a distributed system: achieving its intended goal

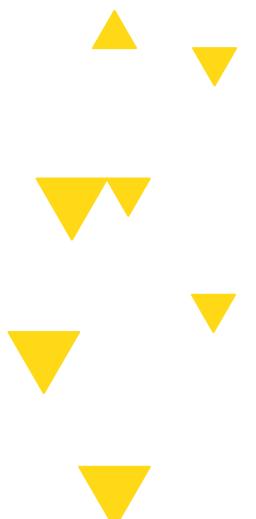
To ensure correctness, one uses a **consensus algorithm** achieving the following:

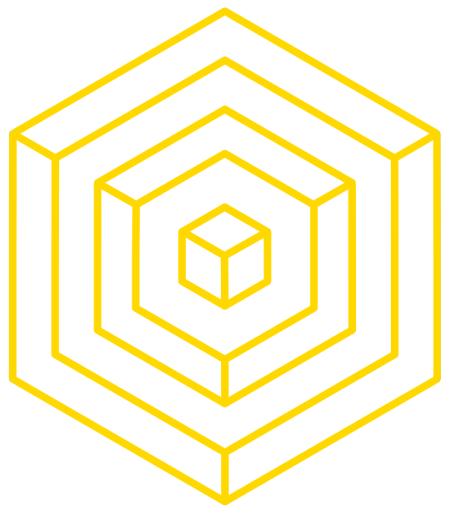
- **Validity:** any value decided upon must be proposed by one of the processes
- **Agreement:** all non-faulty processes must agree on the same value
 - **Agreement and Validity are safety properties:** Honest nodes will never decide on trivial, random, or different values
- **Termination:** all non-faulty nodes eventually decide.
 - **Termination is a liveness property:** All nodes eventually decide on a value





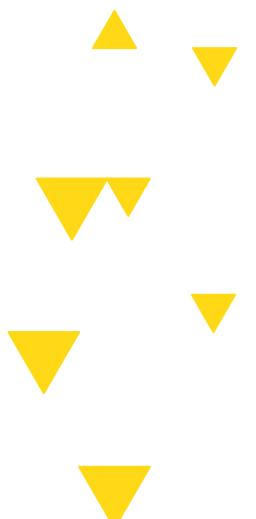
QUESTIONS?

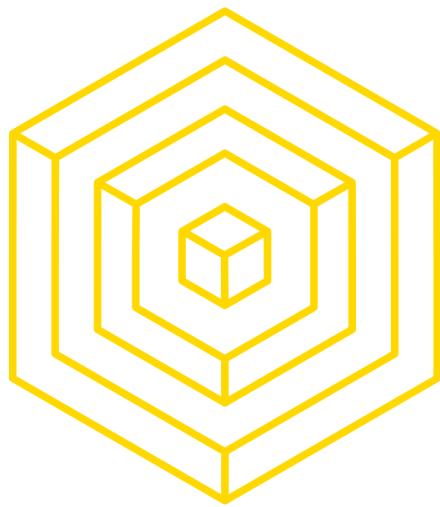




1.3

CAP THEOREM



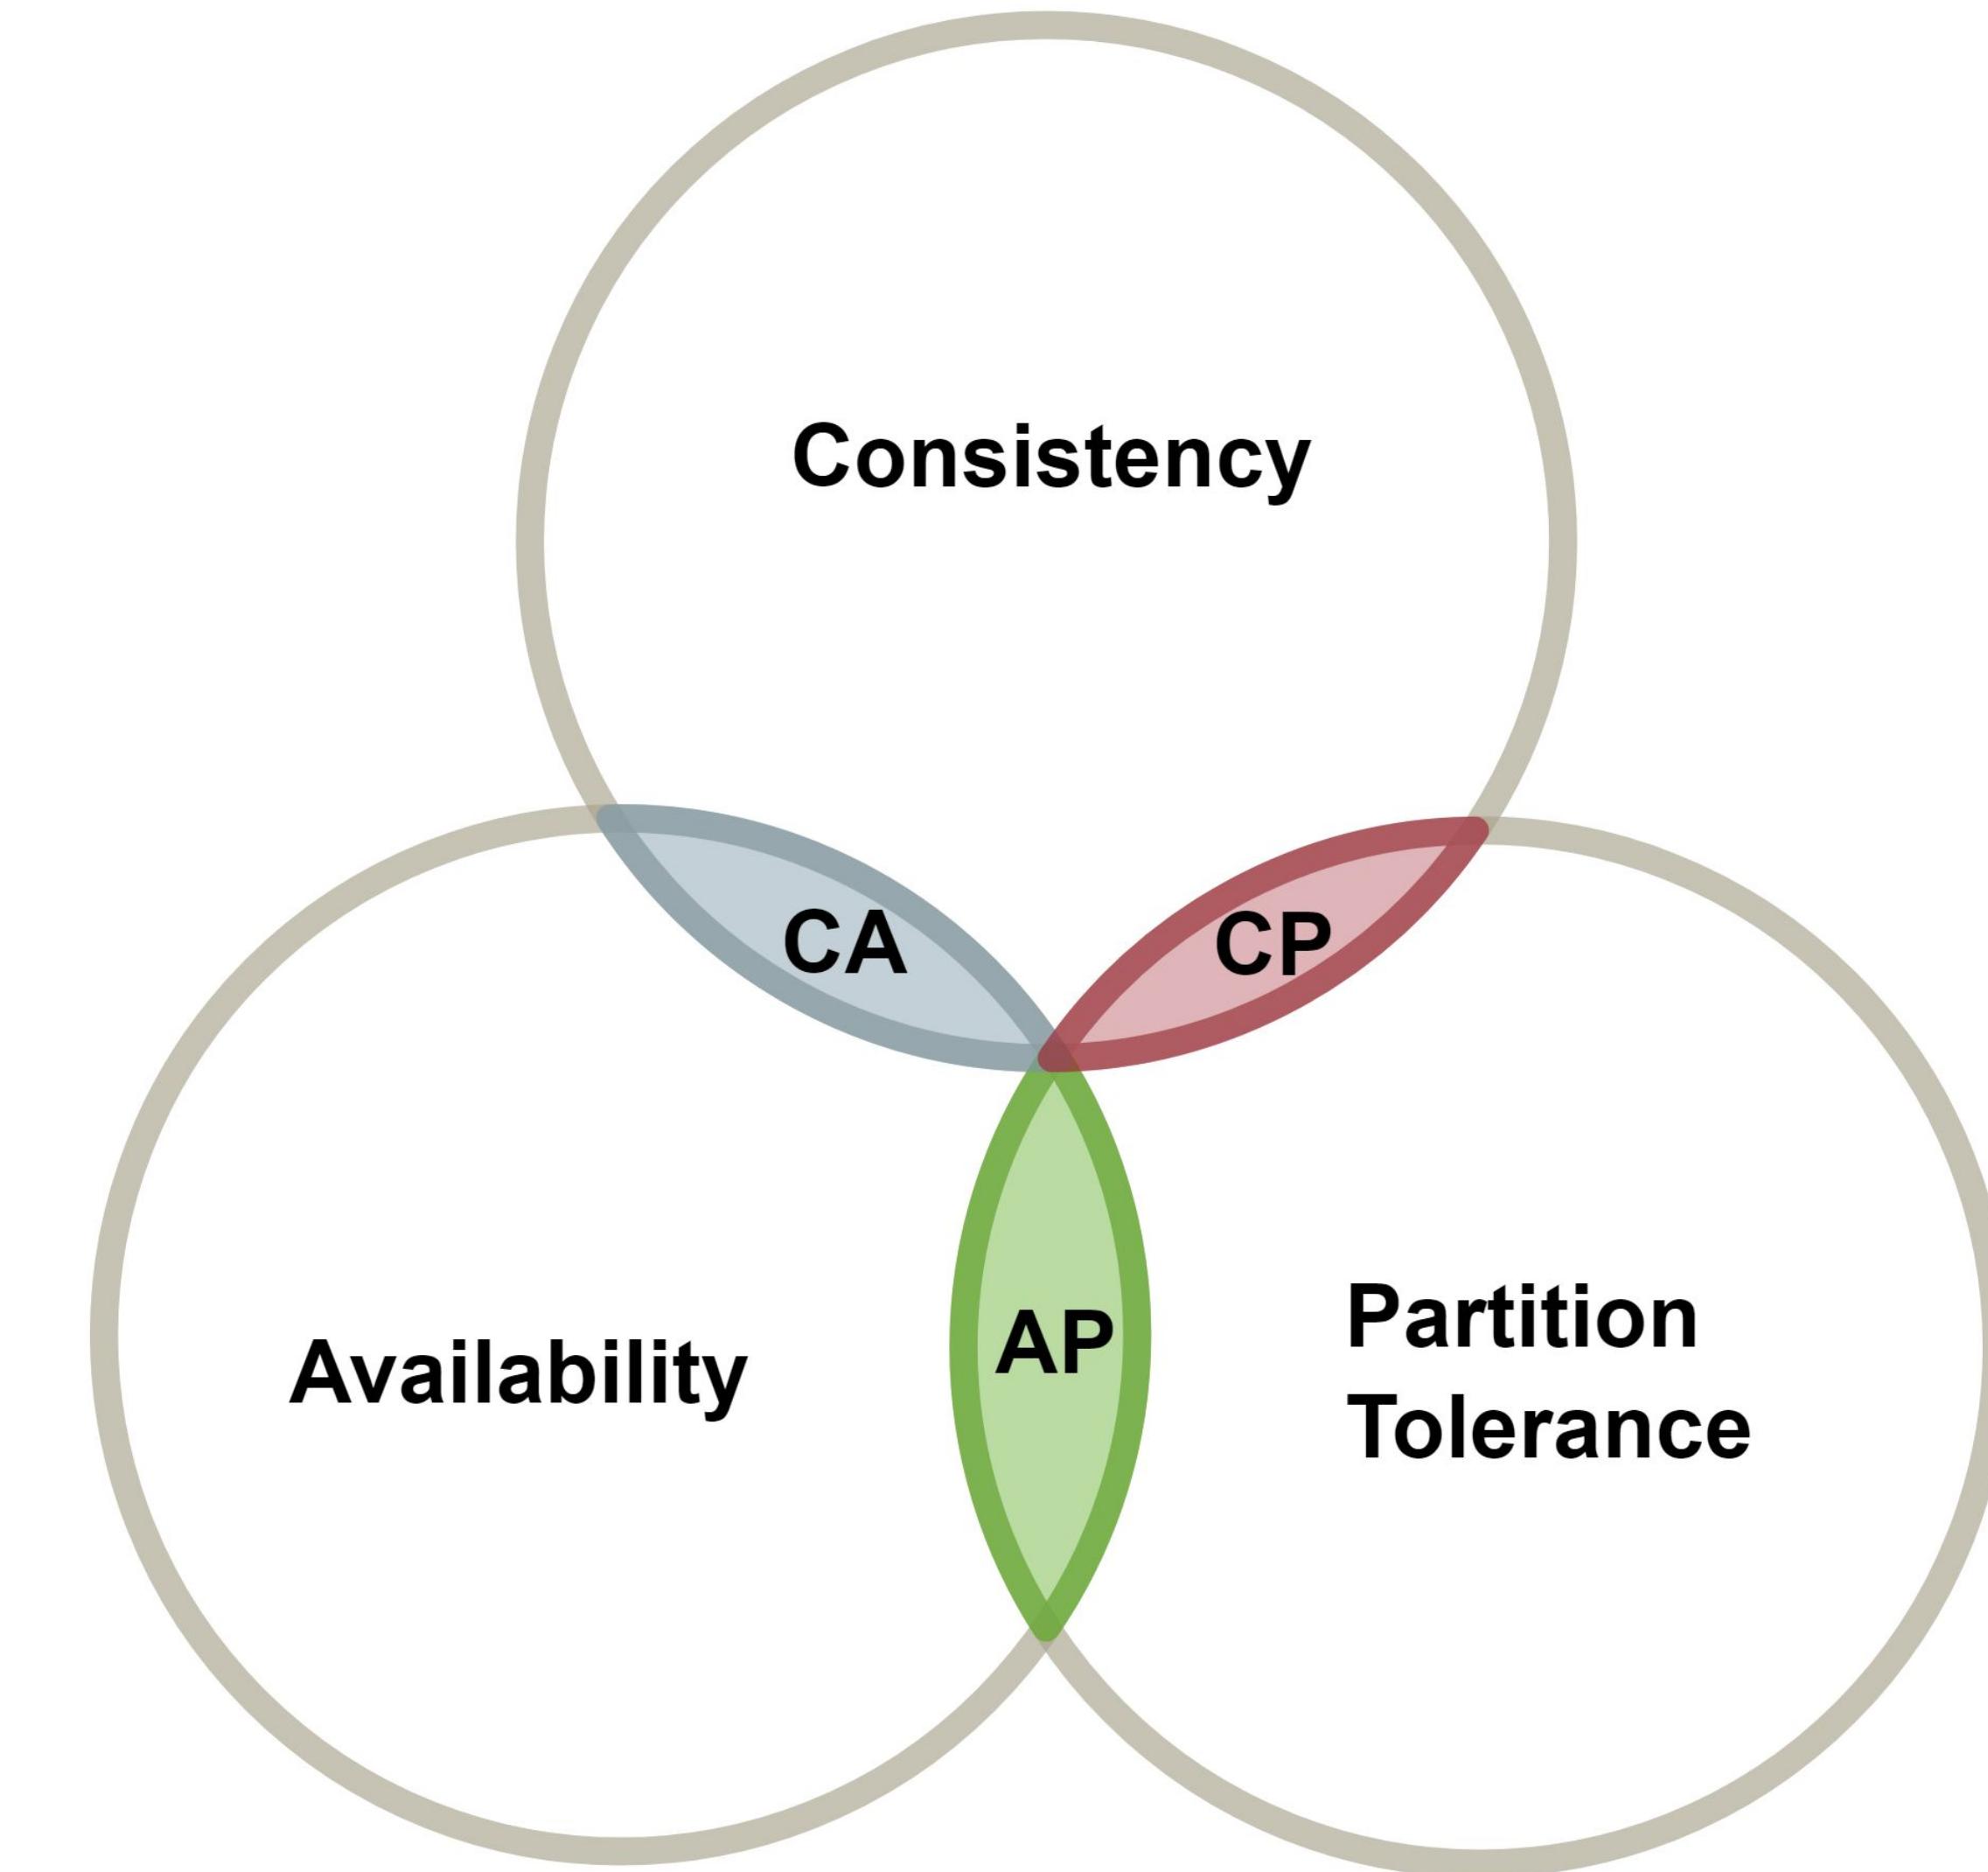


CAP THEOREM

INTRODUCTION

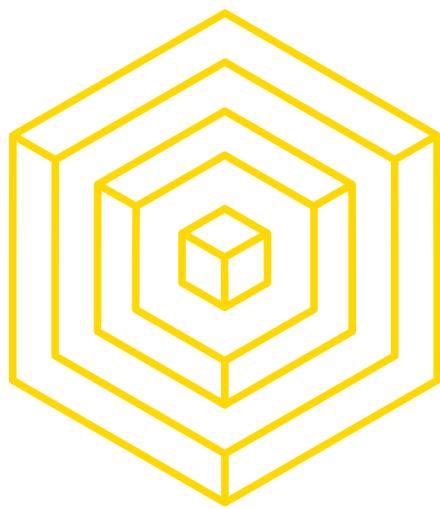
CAP Theorem:
Fundamental theorem for any distributed system pertaining to achievable properties

Image source:
<http://berb.github.io/diploma-thesis/original/resources/cap.svg>



Source:
<https://www.youtube.com/watch?v=Jw1iFr4v58M>

AUTHOR: NADIR AKHTAR

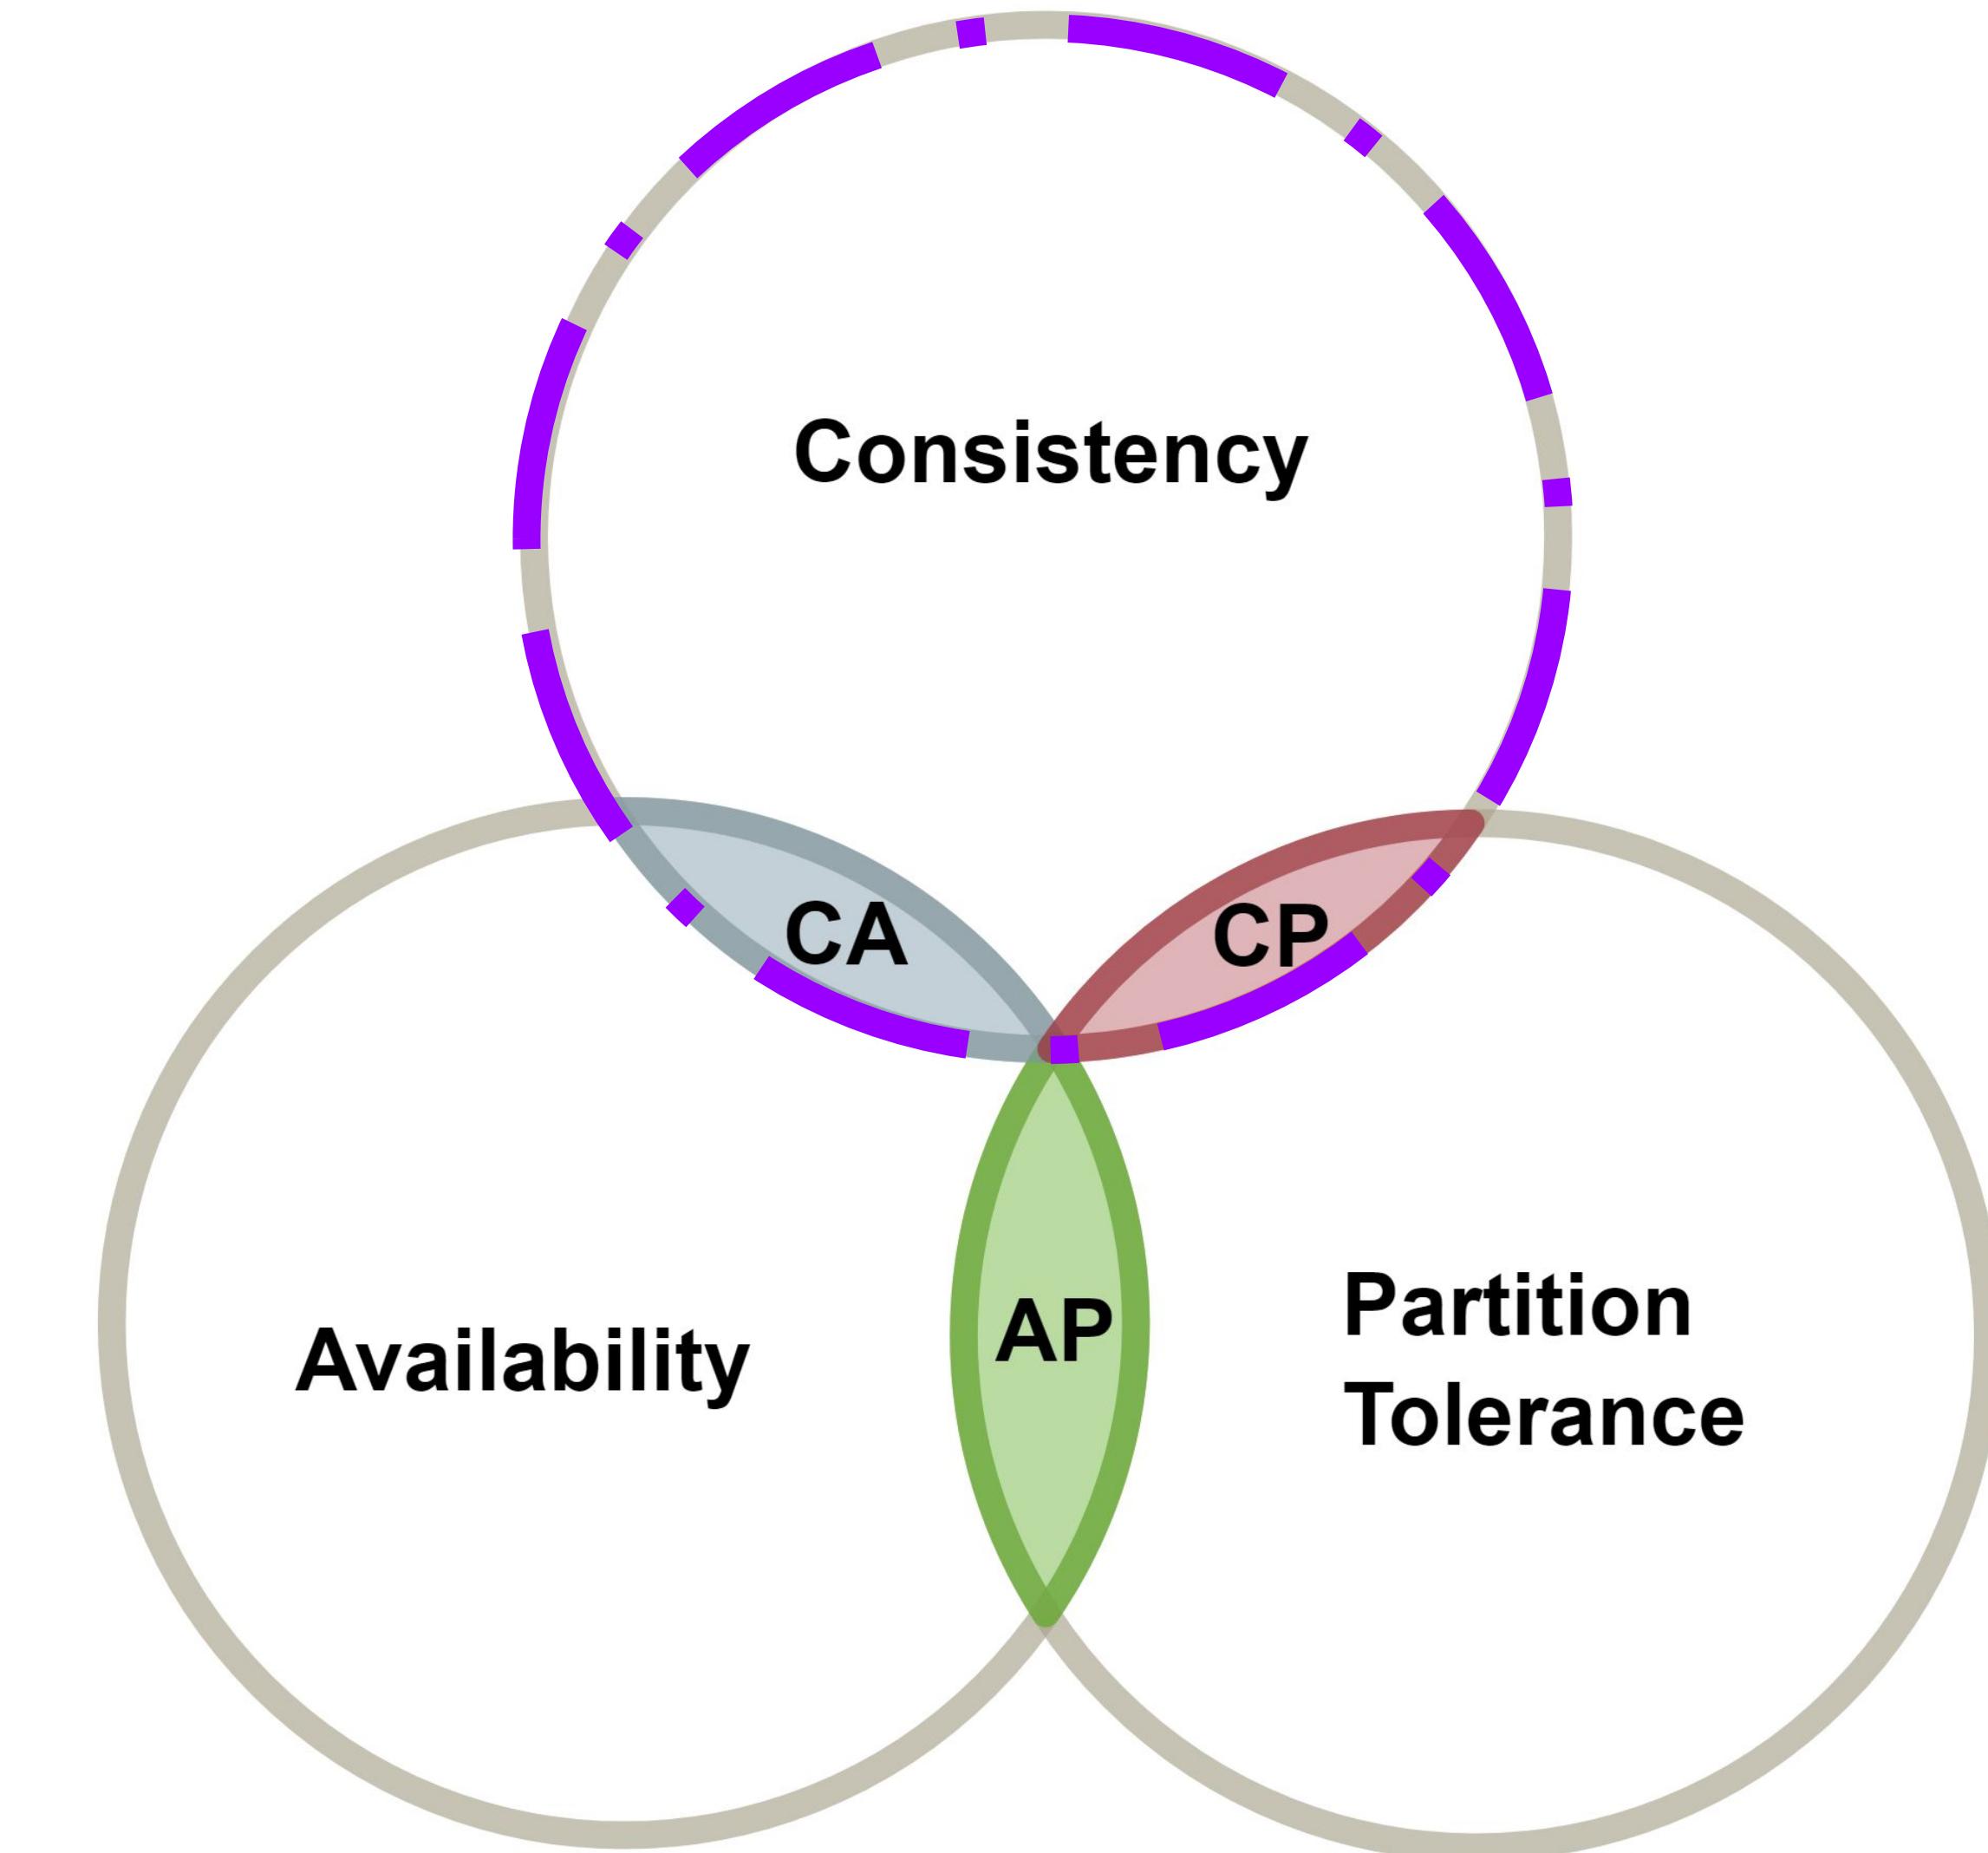


CAP THEOREM

CONSISTENCY

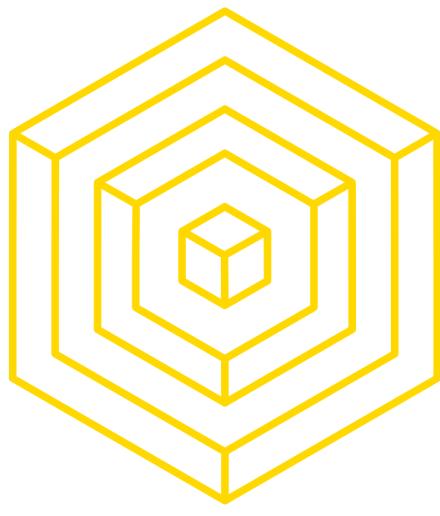
Consistency:
Every node provides
the most recent state

Image source:
<http://berb.github.io/diploma-thesis/original/resources/cap.svg>



Source:
<https://www.youtube.com/watch?v=Jw1iFr4v58M>

AUTHOR: NADIR AKHTAR



CAP THEOREM

AVAILABILITY

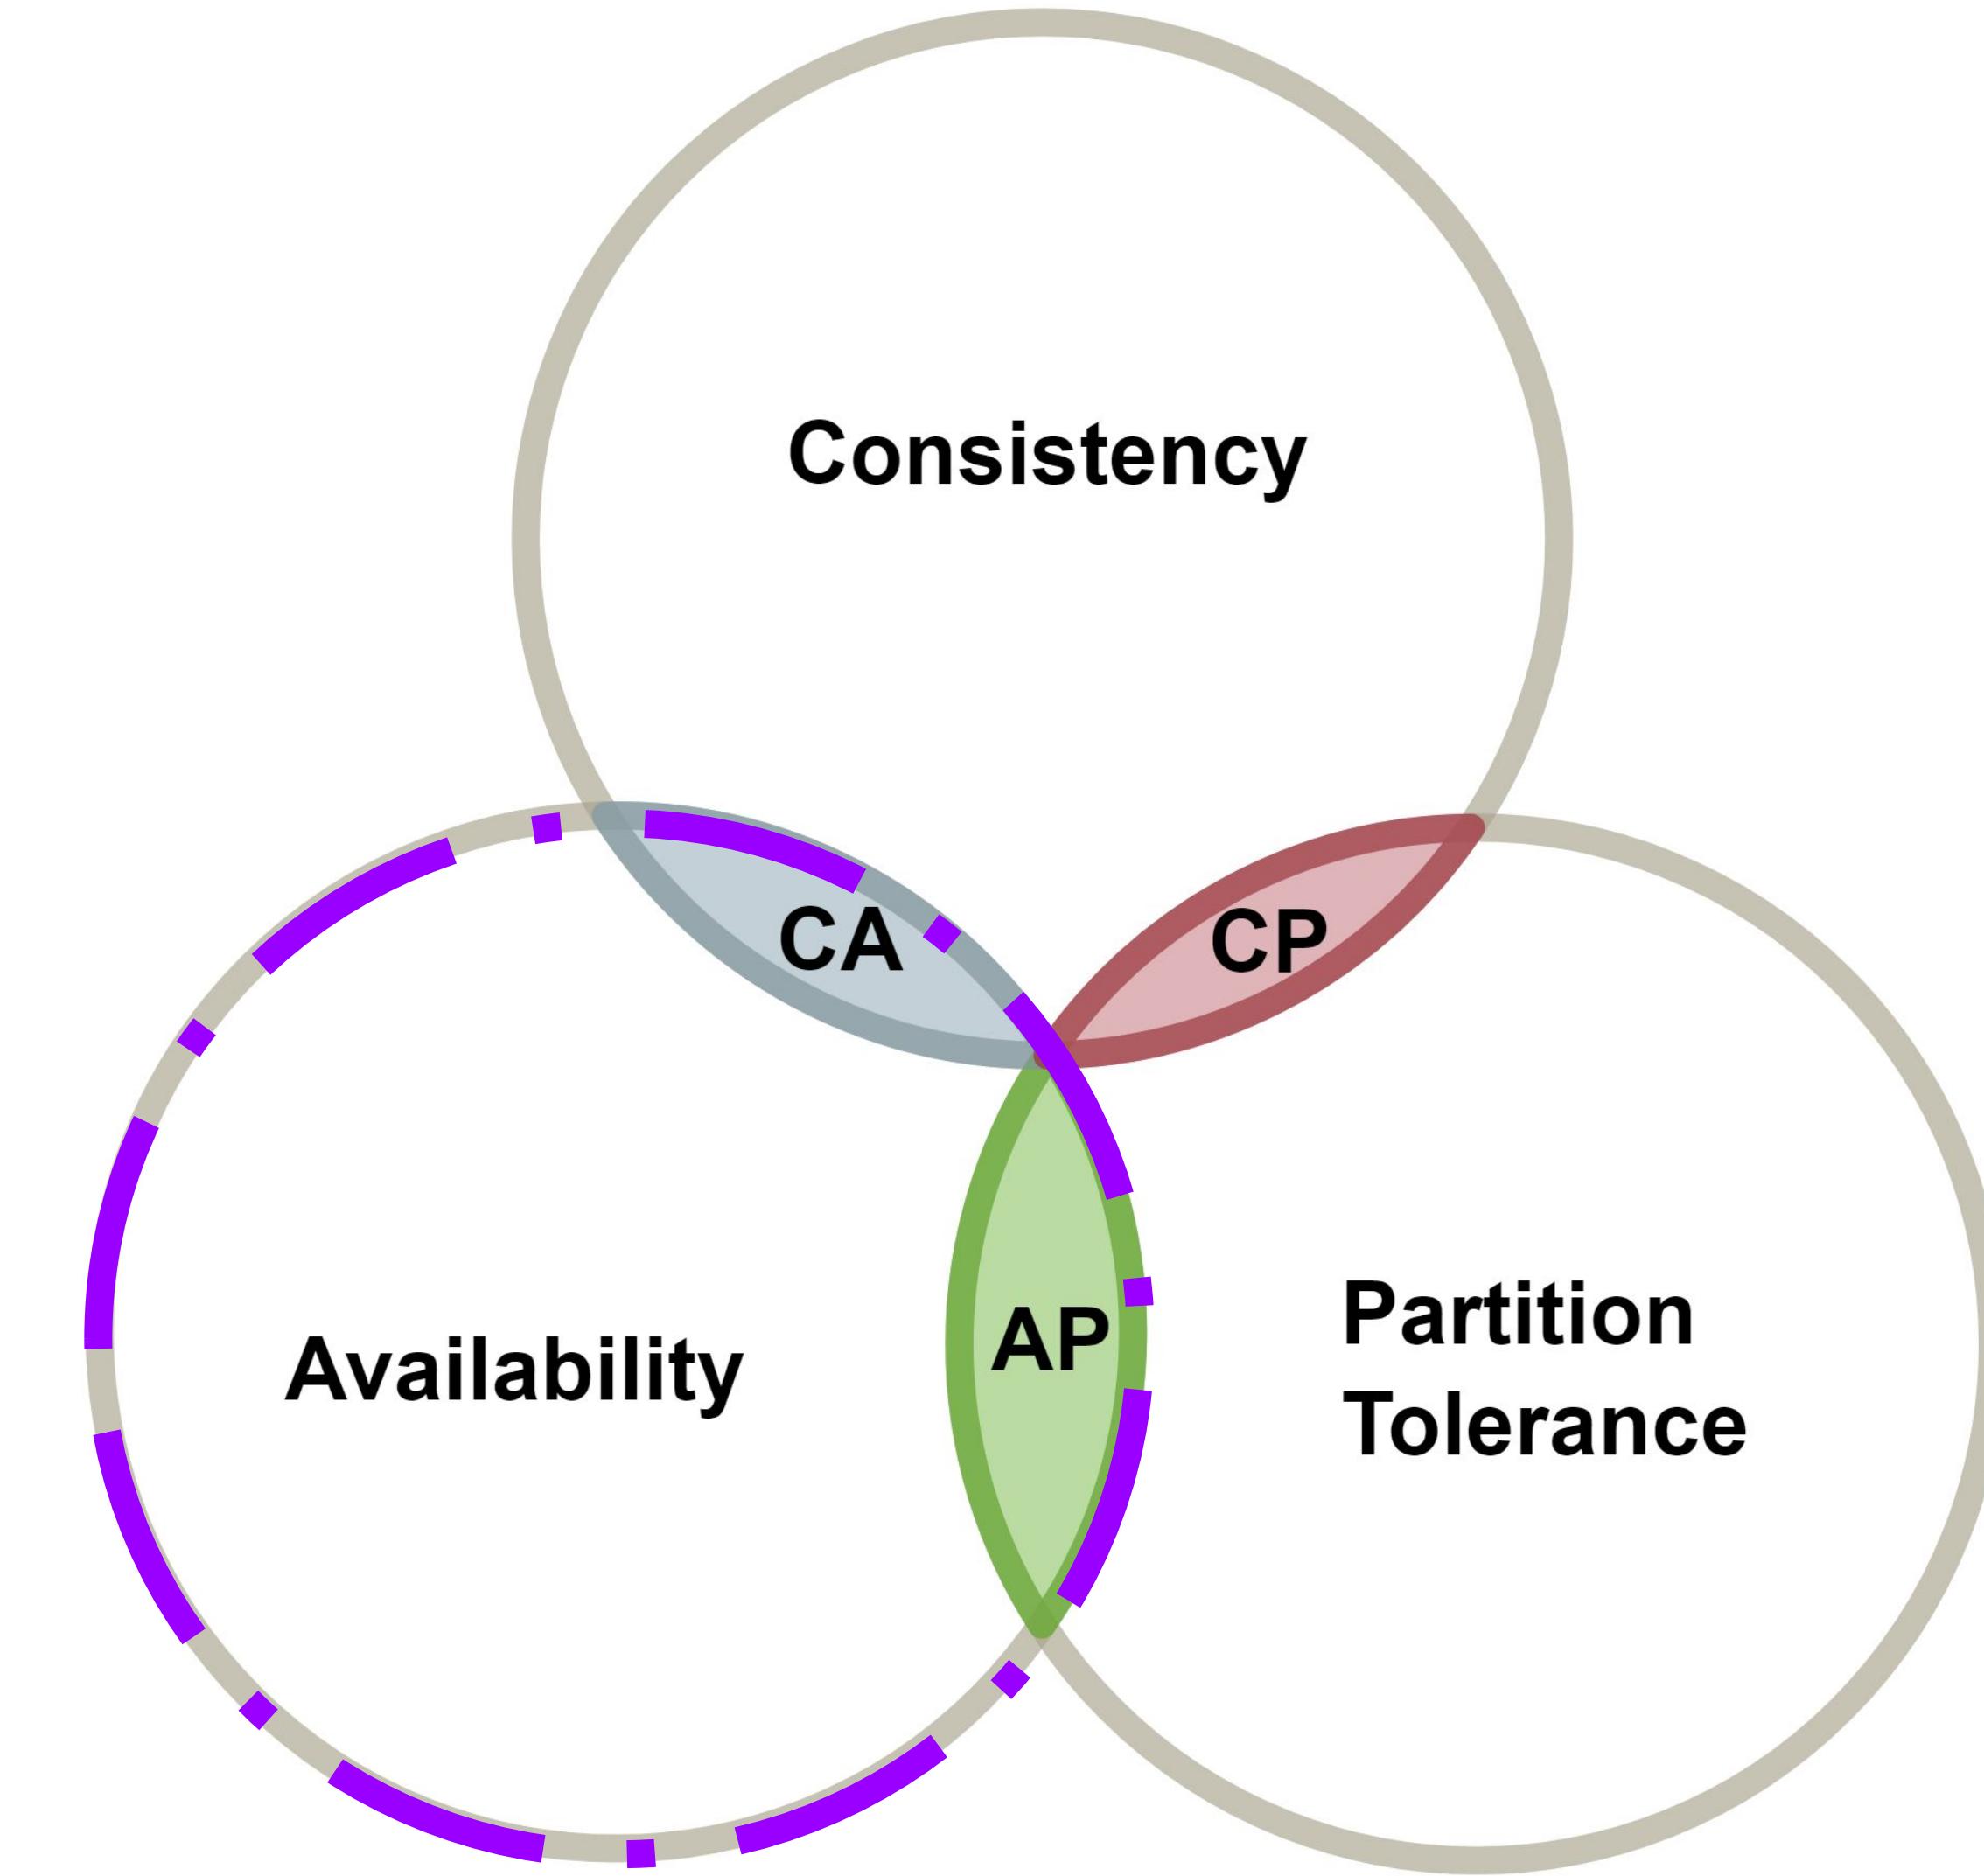
Availability:

Every node has
consistent read and
write access

Image source:

<http://berb.github.io/diploma-thesis/original/resources/cap.svg>

15

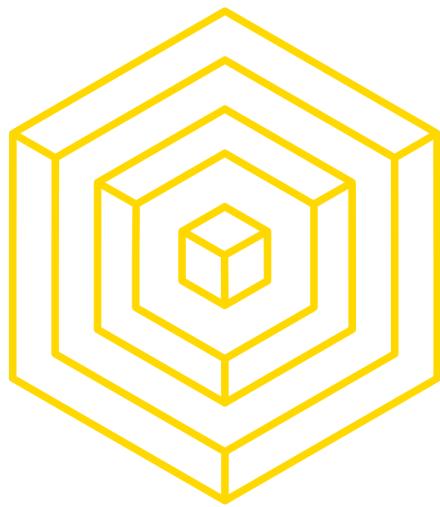


Source:

<https://www.youtube.com/watch?v=Jw1iFr4v58M>

AUTHOR: NADIR AKHTAR

BLOCKCHAIN FUNDAMENTALS LECTURE 8



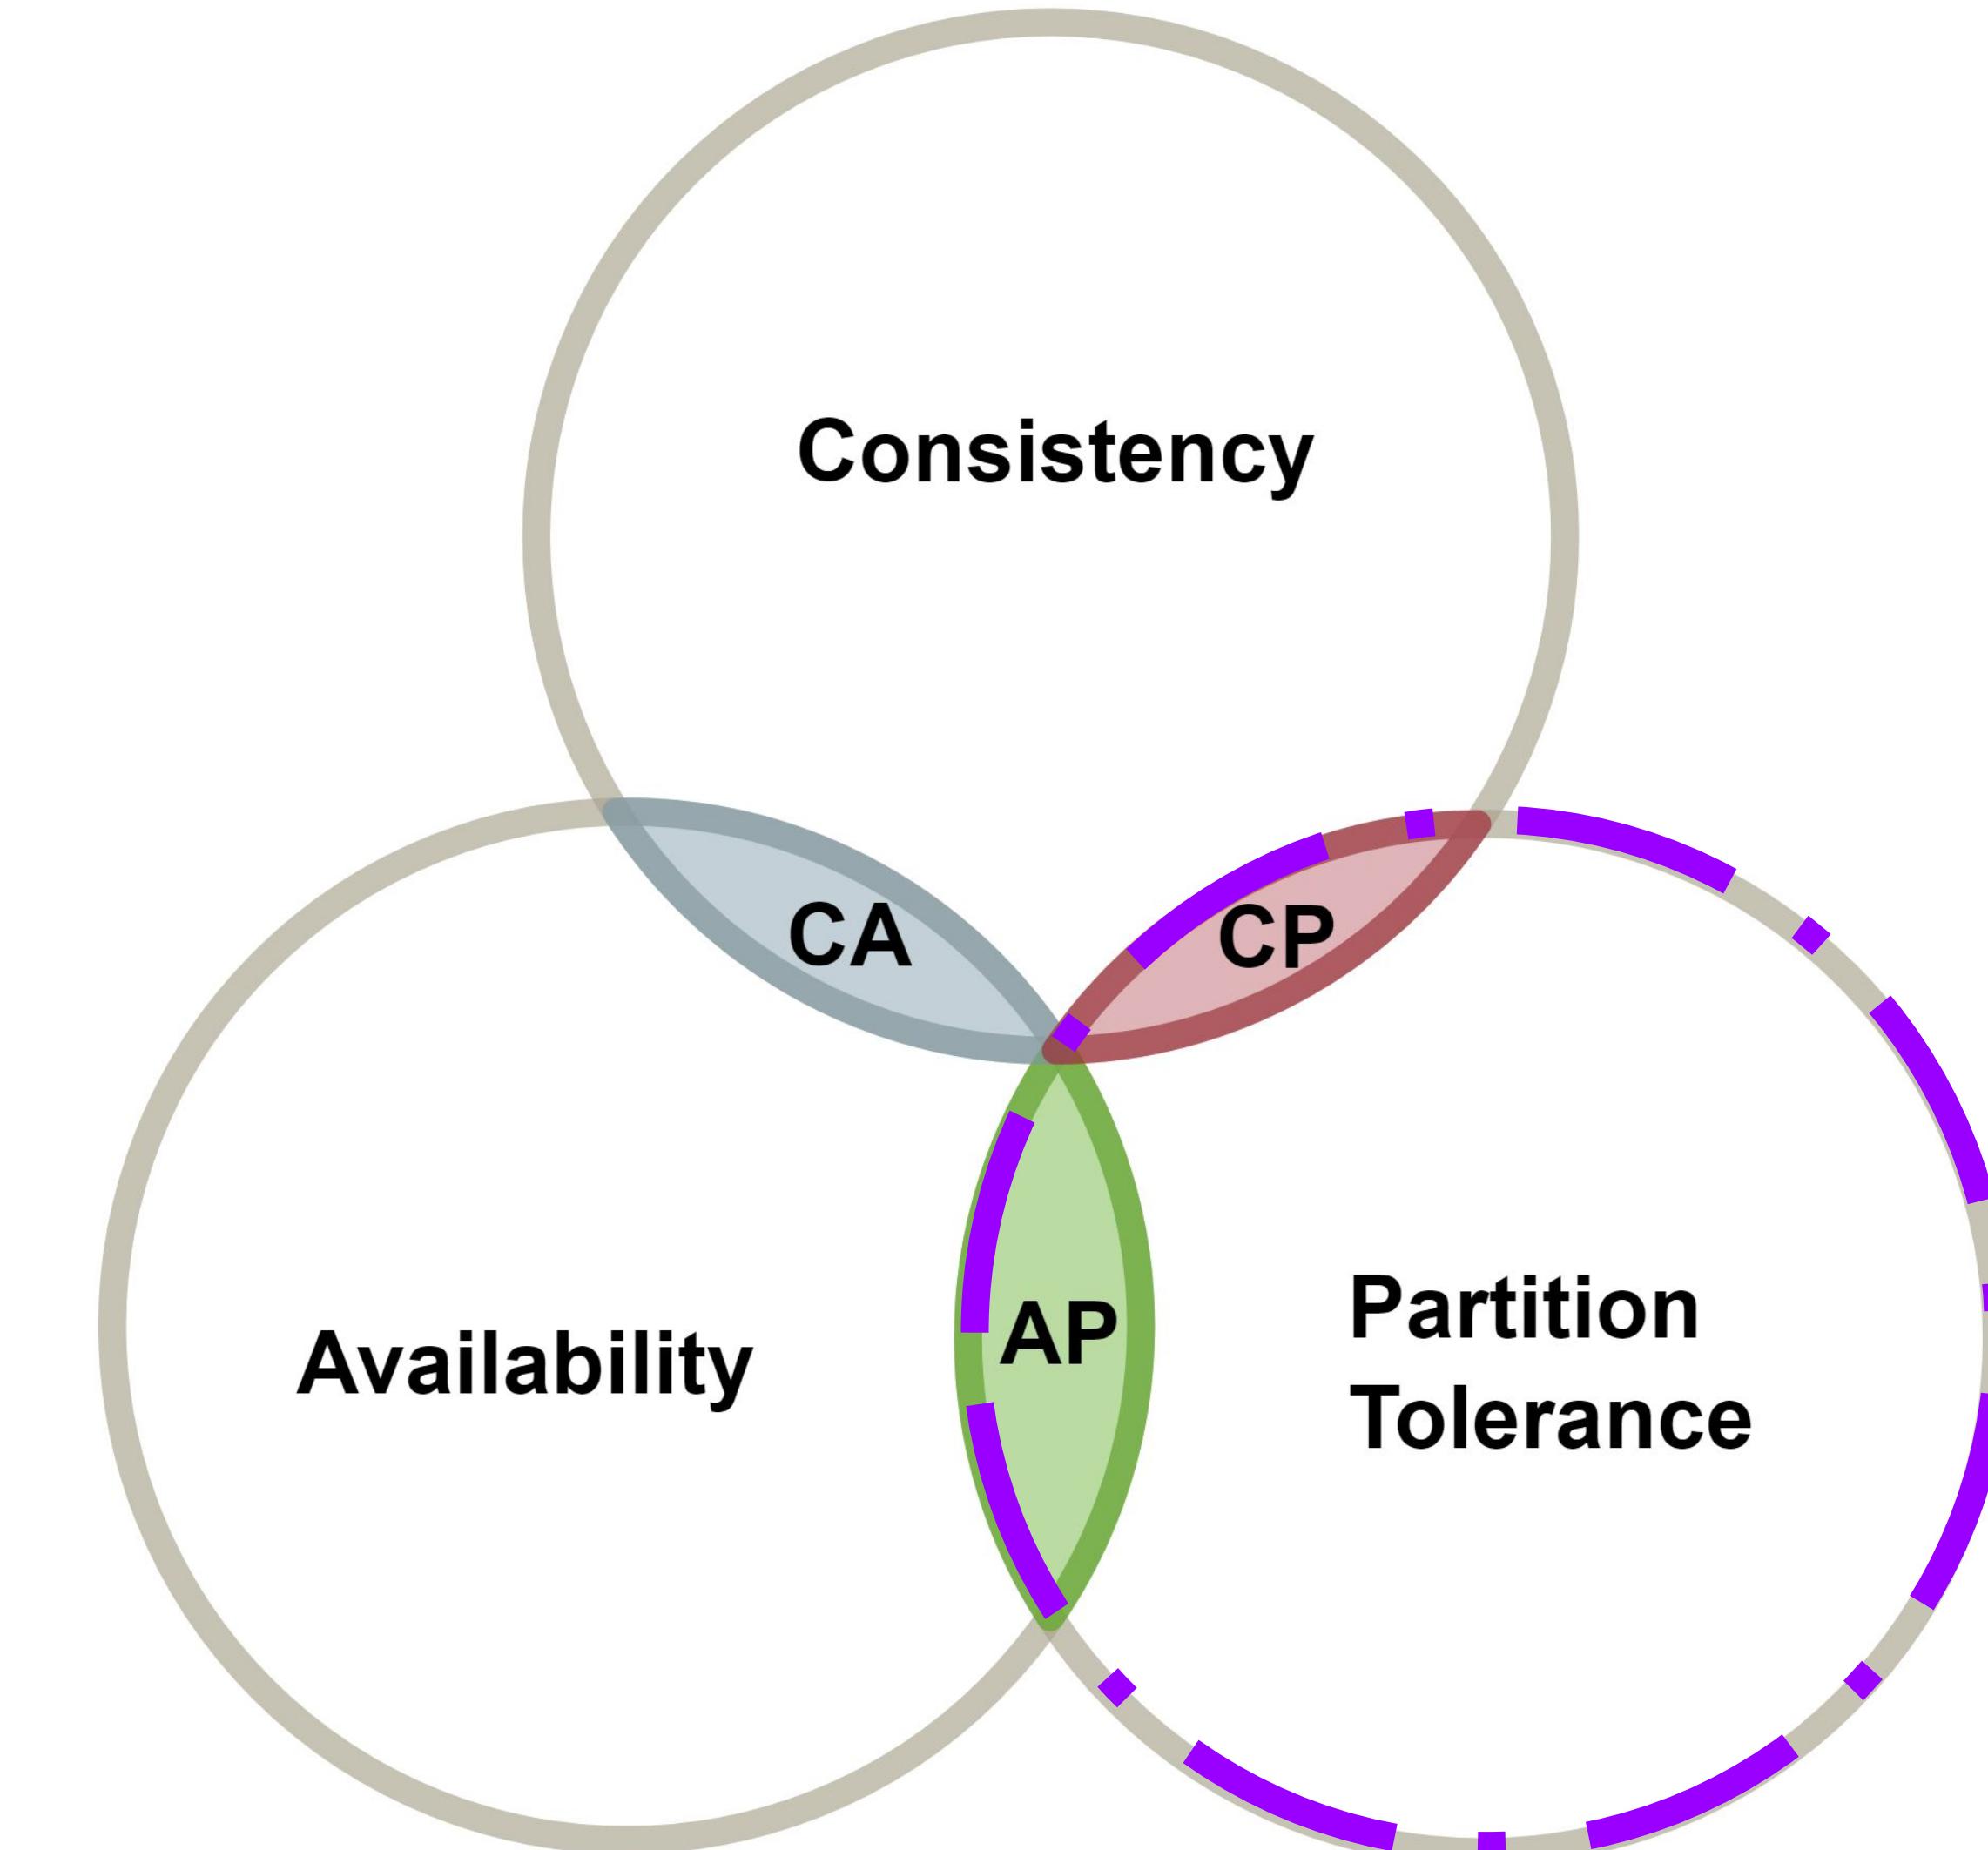
CAP THEOREM

PARTITION TOLERANCE

Image source:

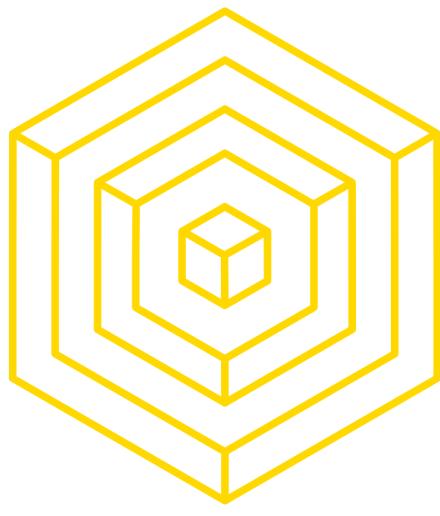
<http://berb.github.io/diploma-thesis/original/resources/cap.svg>

Partition Tolerance:
The system works
despite partitions in
the network



Source:
<https://www.youtube.com/watch?v=Jw1iFr4v58M>

AUTHOR: NADIR AKHTAR

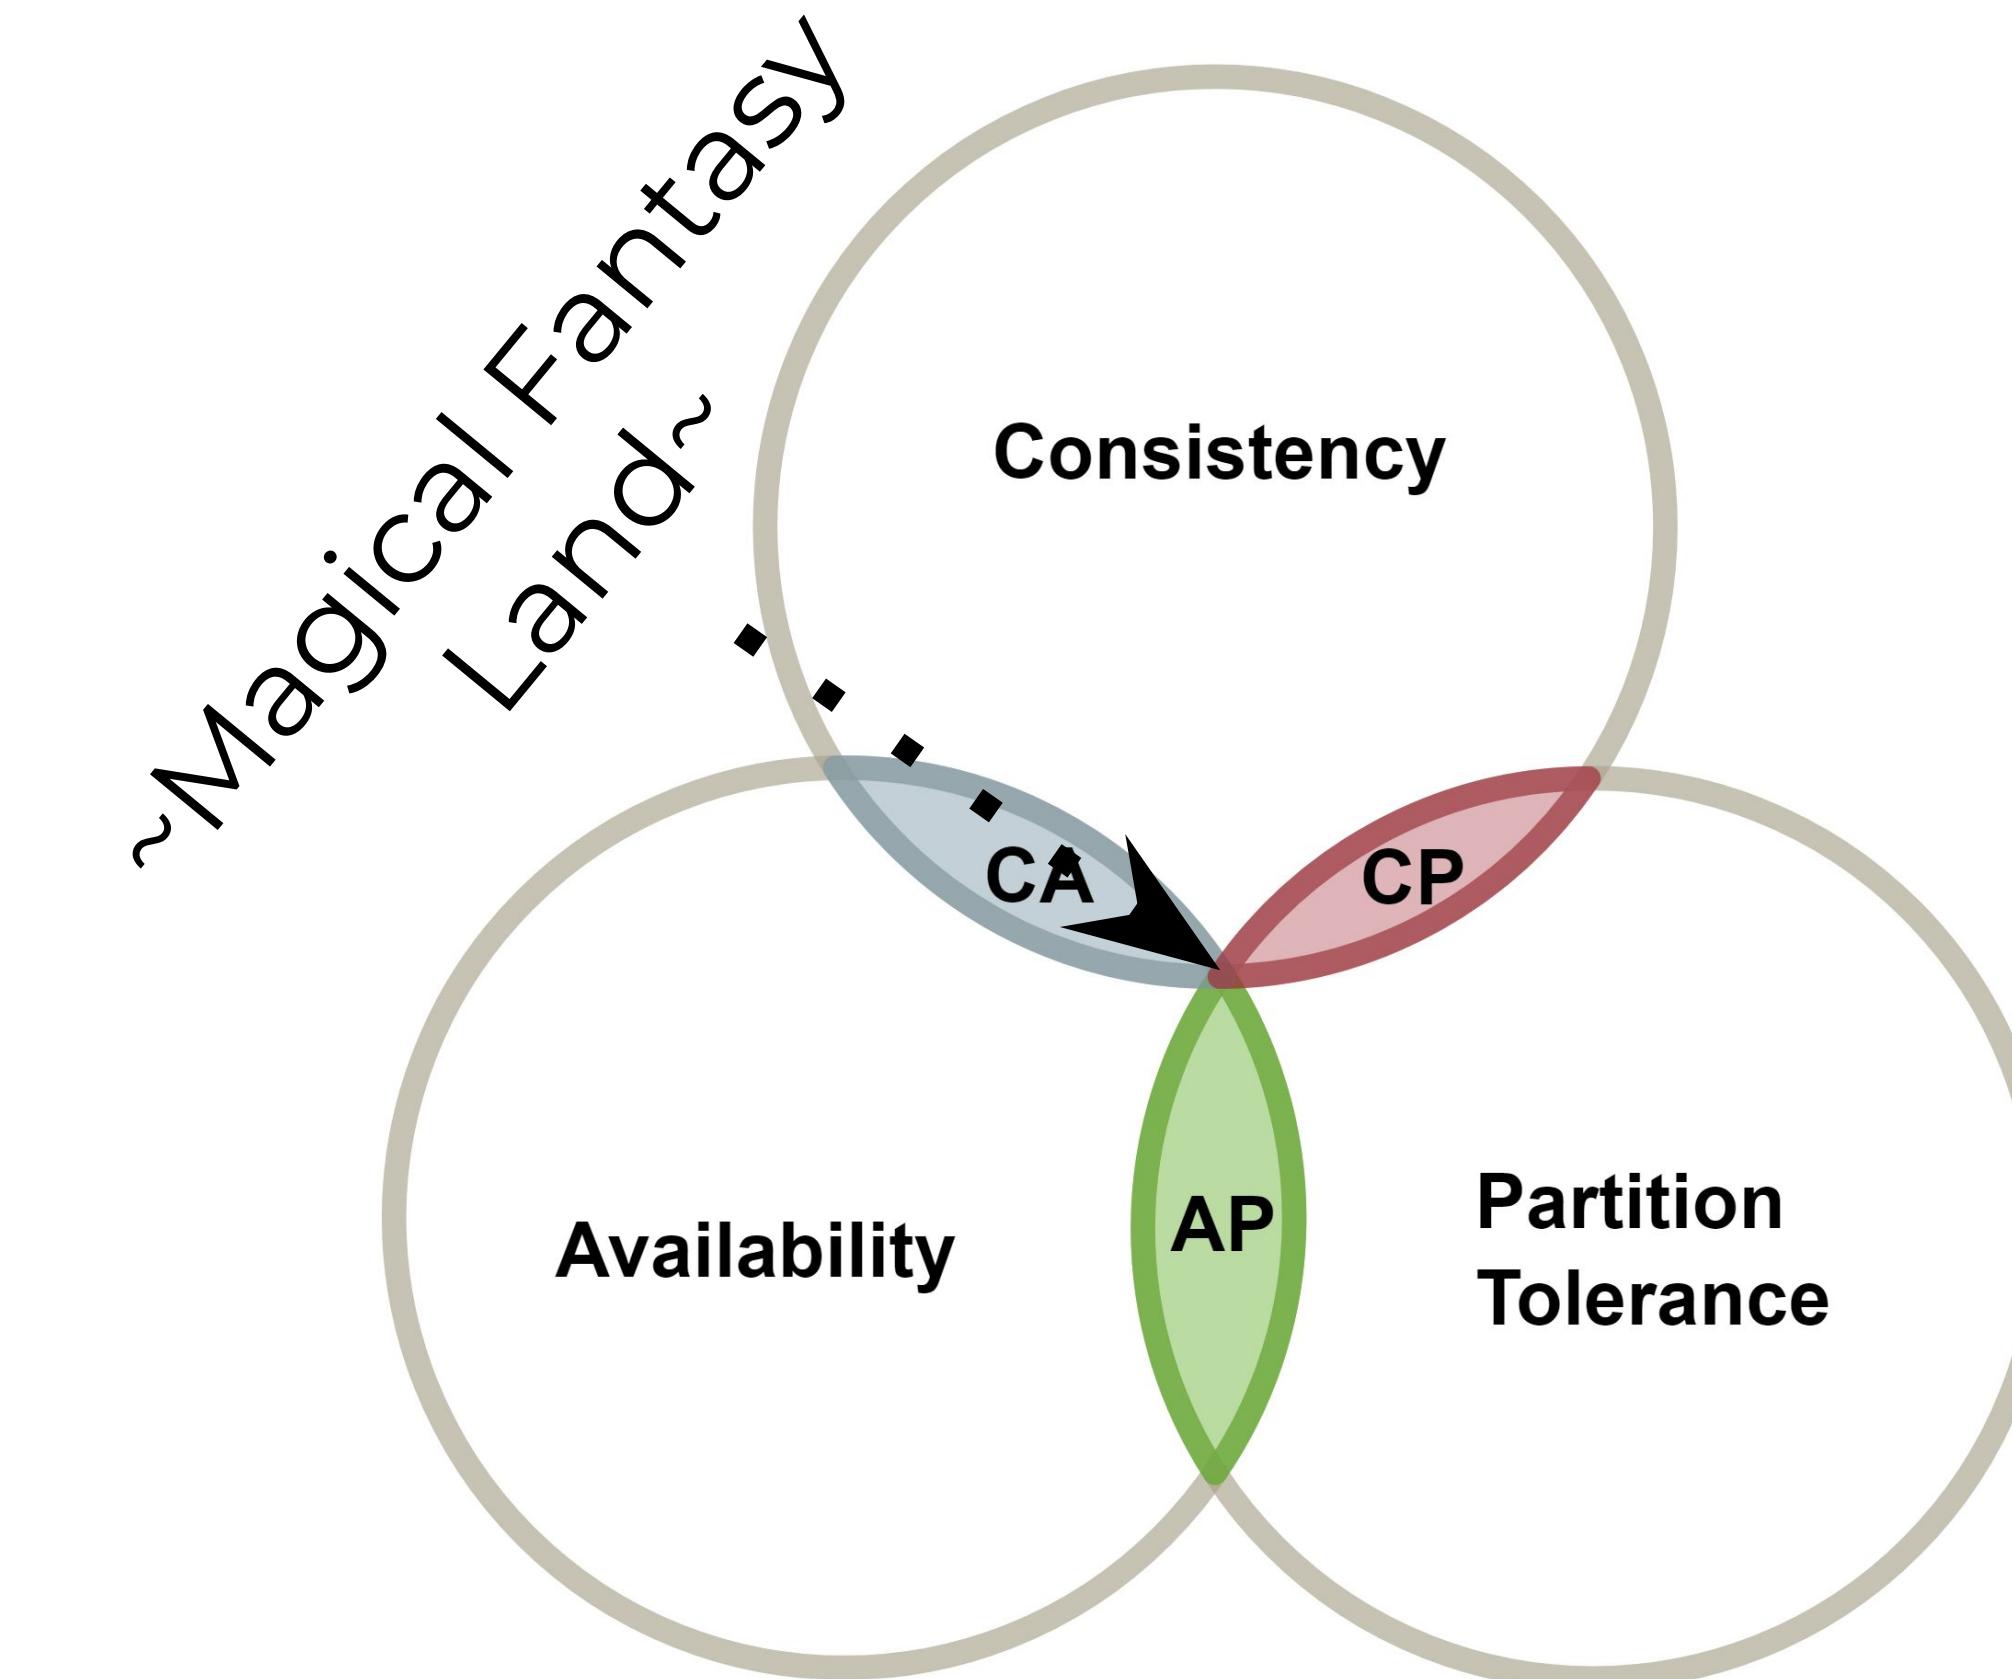


CAP THEOREM

LIMITATIONS

Can only have **two of three**

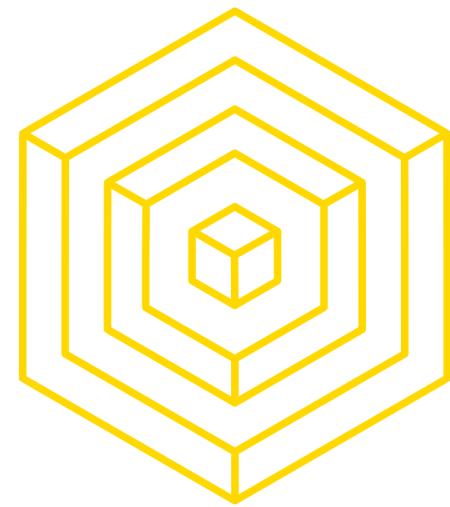
Image source:
<http://berb.github.io/diploma-thesis/original/resources/cap.svg>



Source:

<https://www.youtube.com/watch?v=Jw1iFr4v58M>

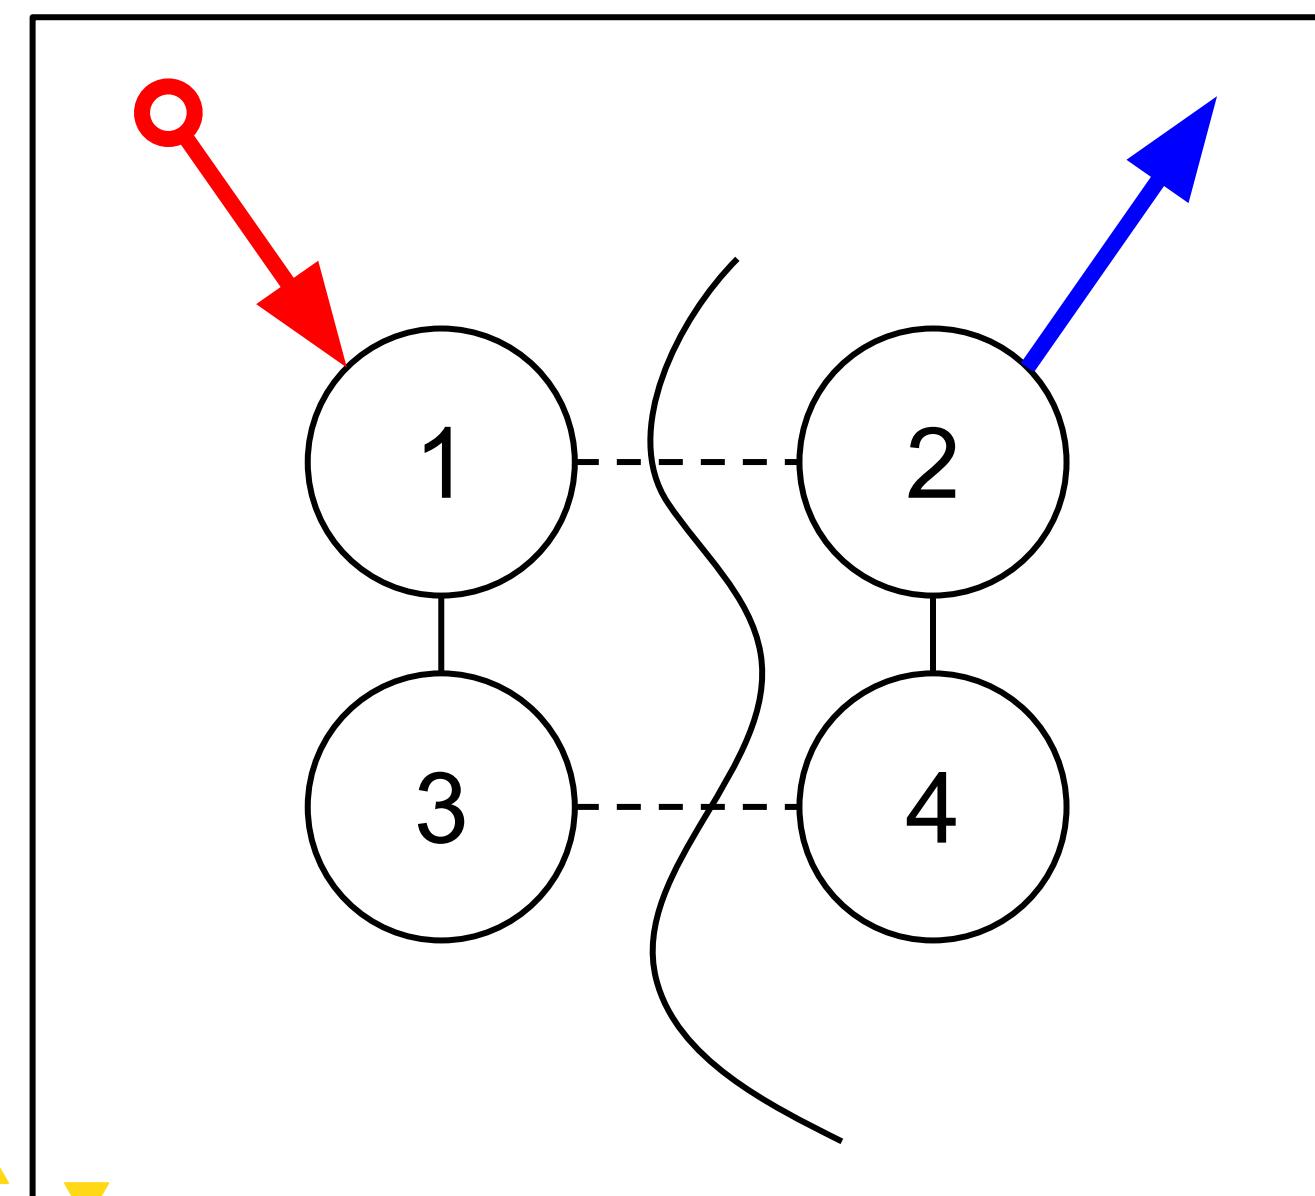
AUTHOR: NADIR AKHTAR



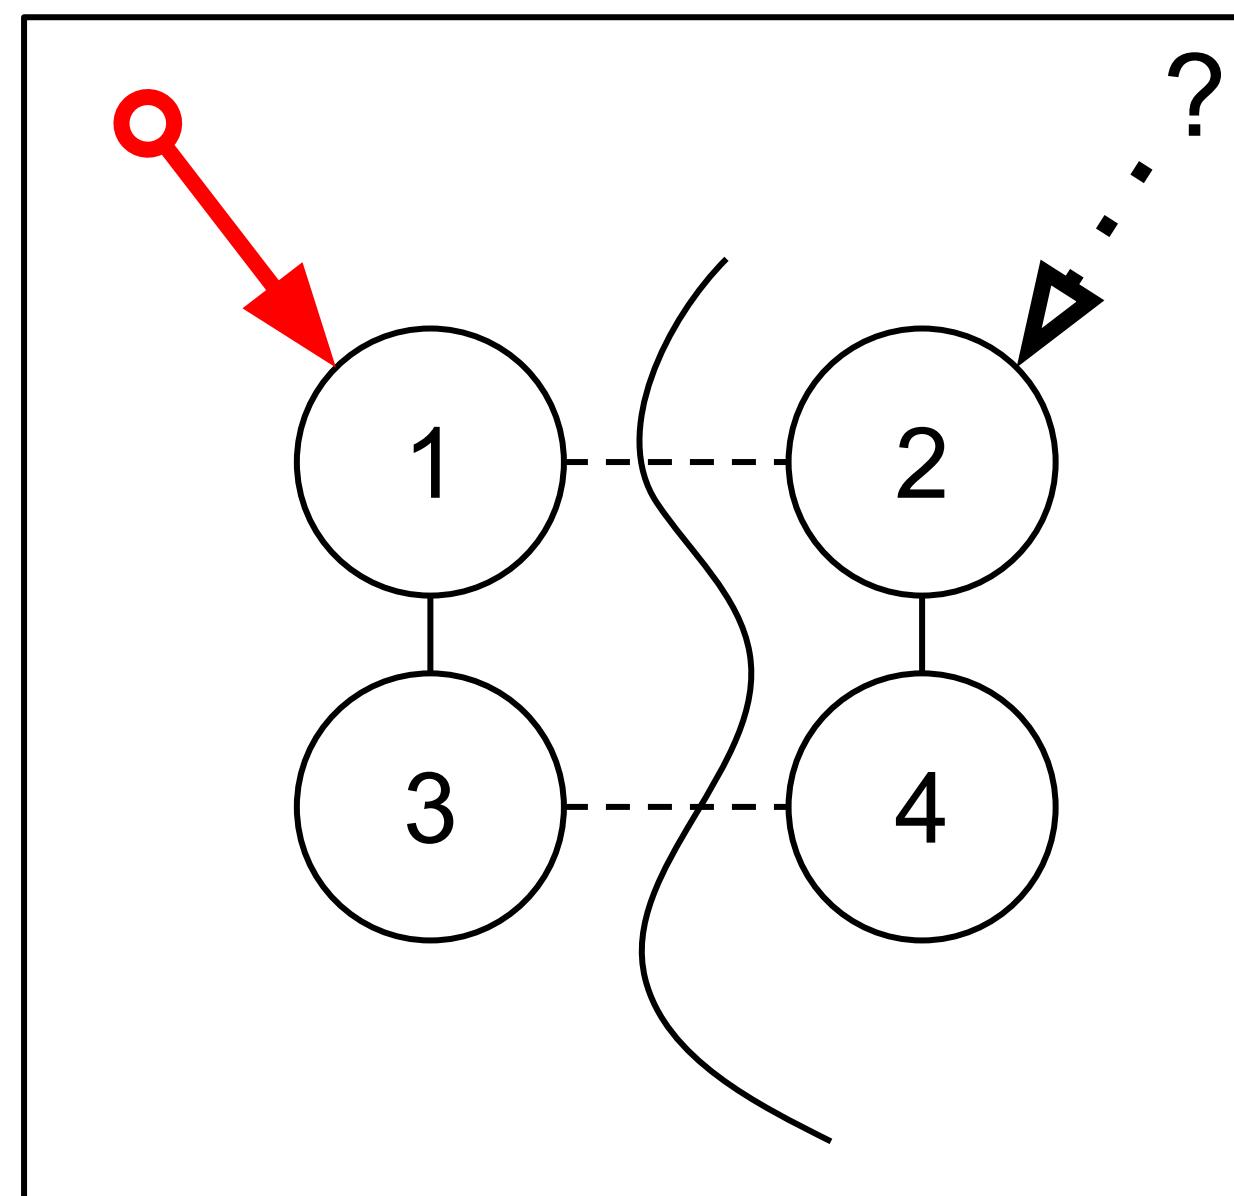
CAP THEOREM

PROOF

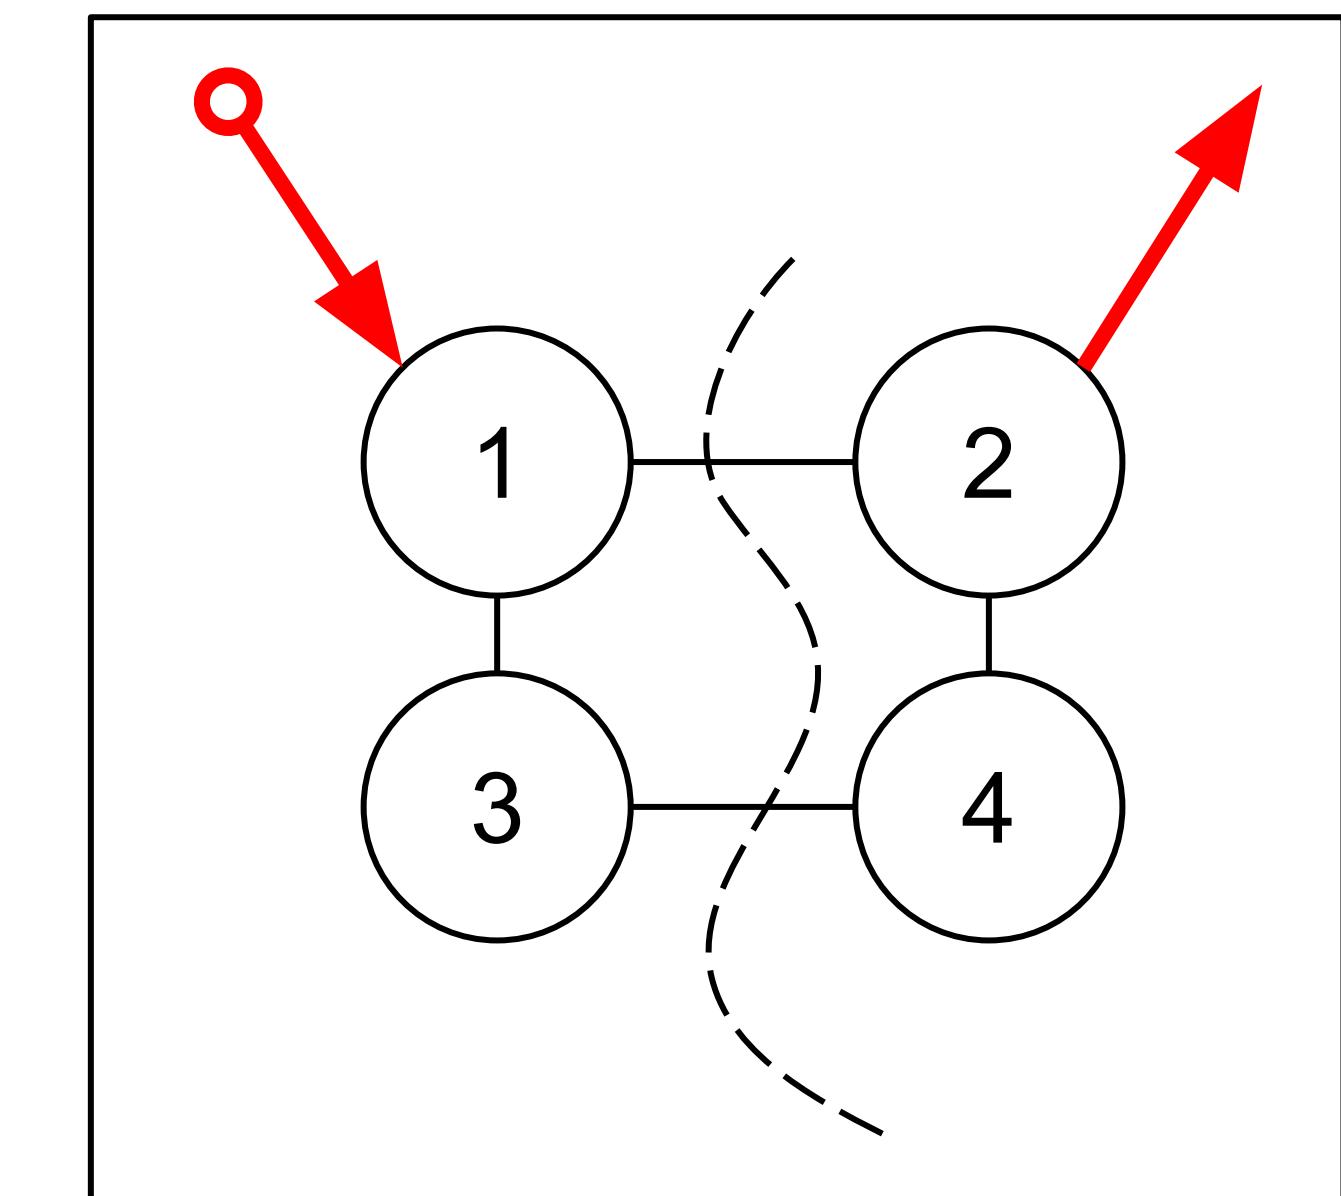
Partition Tolerant +
Available = **Not Consistent**



Partition Tolerant +
Consistent = **Not Available**

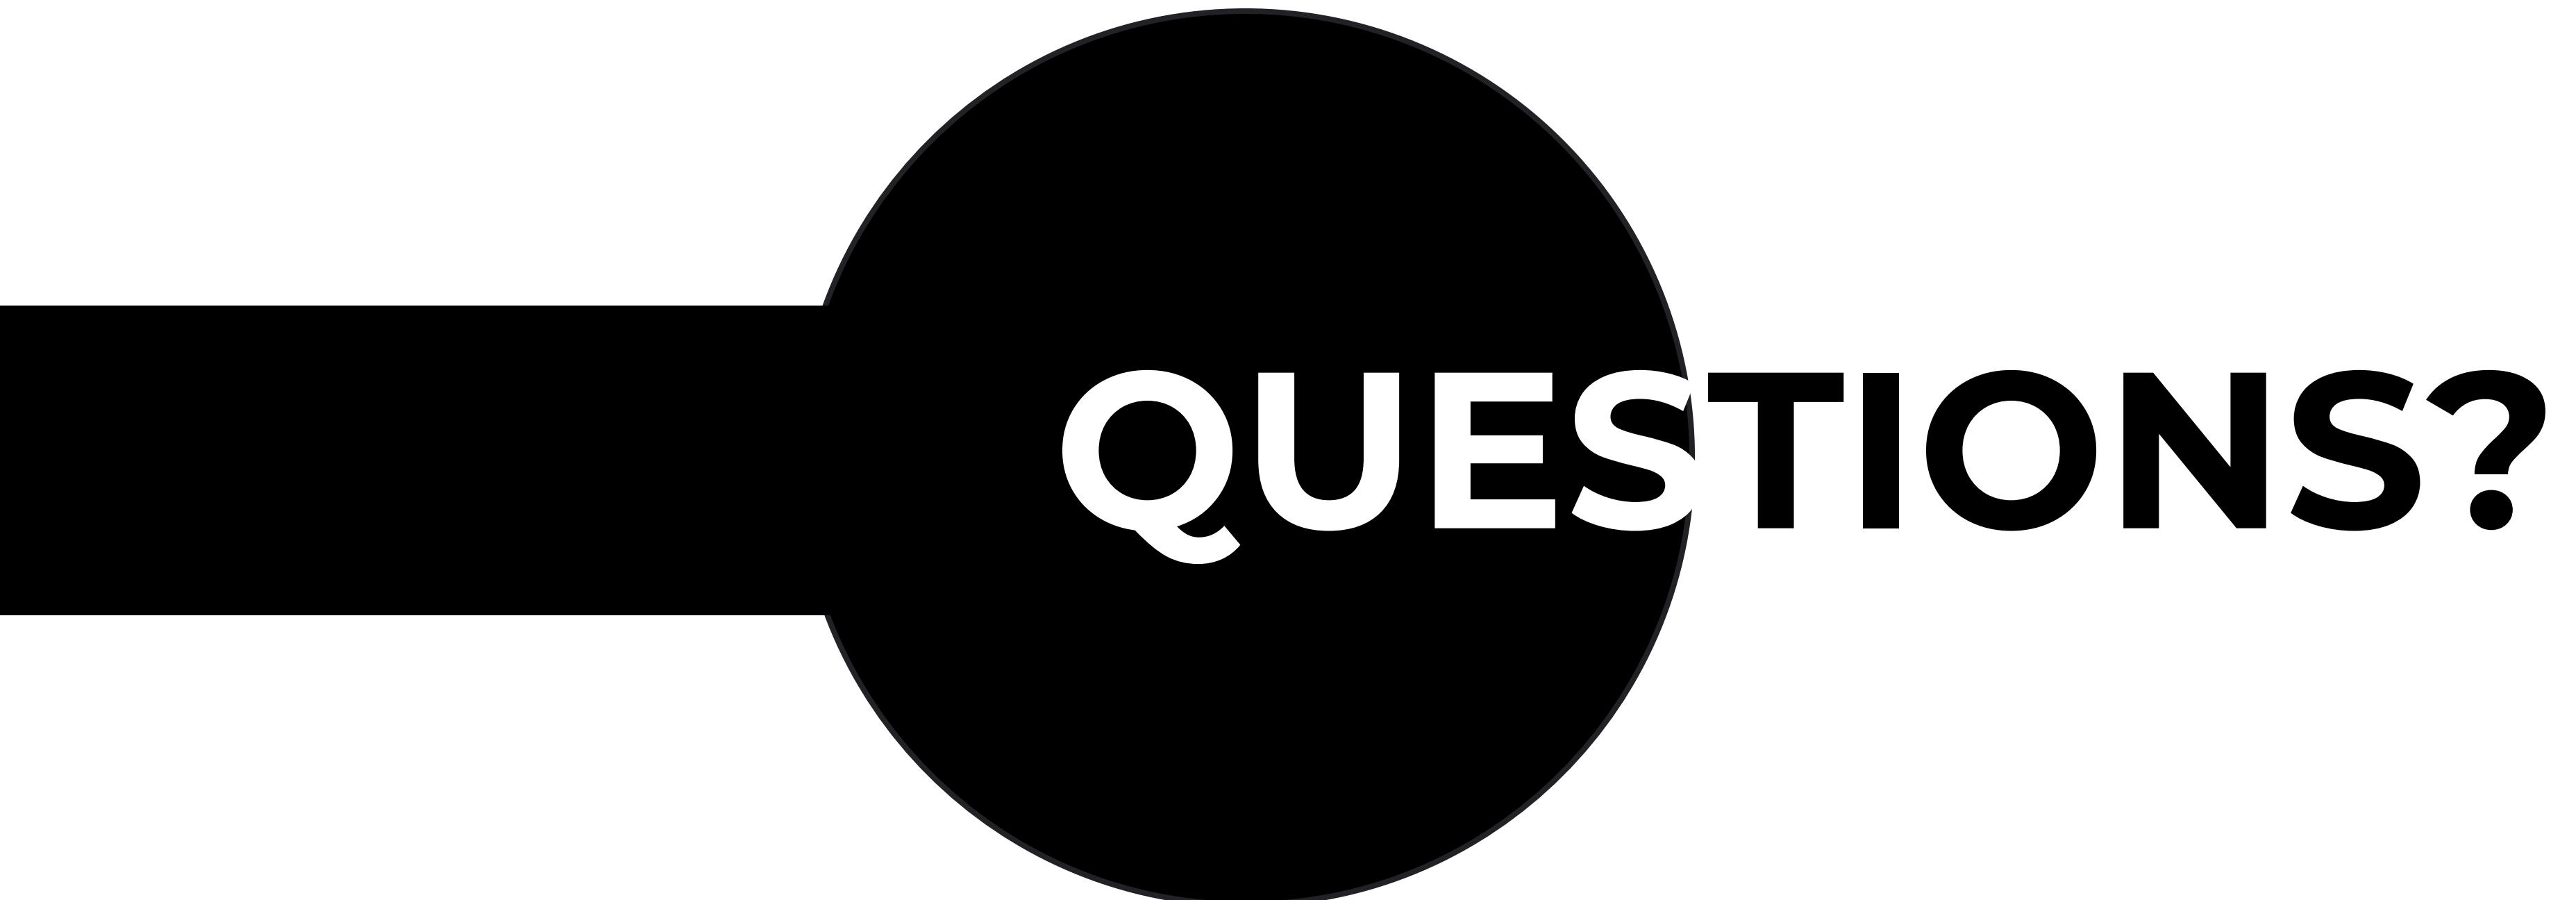


Consistent + Available =
Not Partition Tolerant

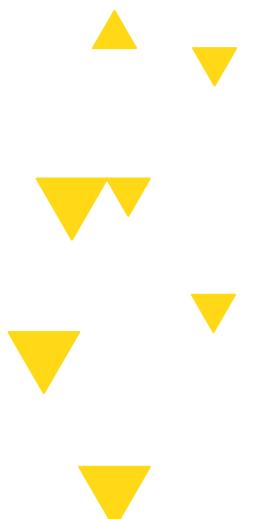


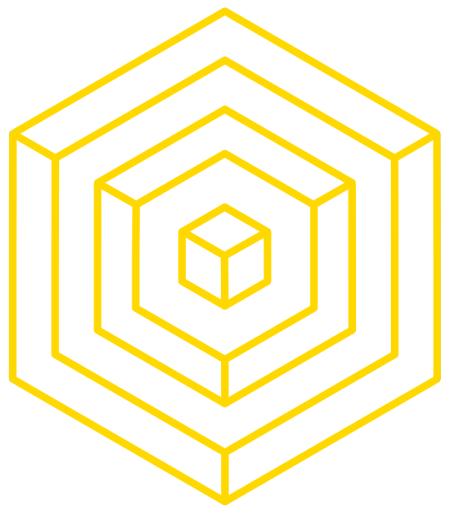
AUTHOR: NADIR AKHTAR

BLOCKCHAIN FUNDAMENTALS LECTURE 8



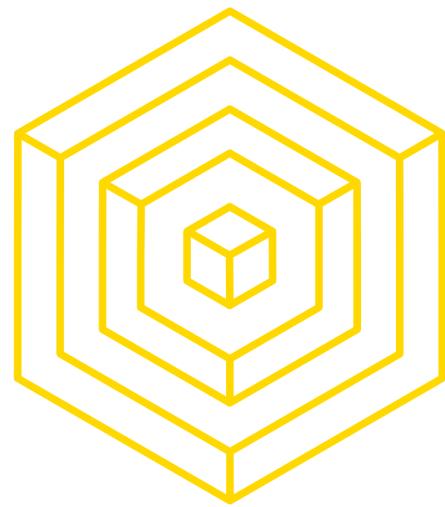
QUESTIONS?





1.4

BYZANTINE GENERALS' PROBLEM



BYZANTINE FAULT TOLERANCE

BYZANTINE GENERALS' PROBLEM

Byzantine Generals' Problem:

- How do we ensure that all honest generals make the same decision?
 - Spoiler alert: No solution for $\frac{1}{3}$ or greater corruption
- Solution to all remaining cases:
Practical Byzantine Fault Tolerance

Source:

[http://www-inst.eecs.berkeley.edu/~cs162
/fa12/hand-outs/Original_Byzantine.pdf](http://www-inst.eecs.berkeley.edu/~cs162/fa12/hand-outs/Original_Byzantine.pdf)



AUTHOR: NADIR AKHTAR

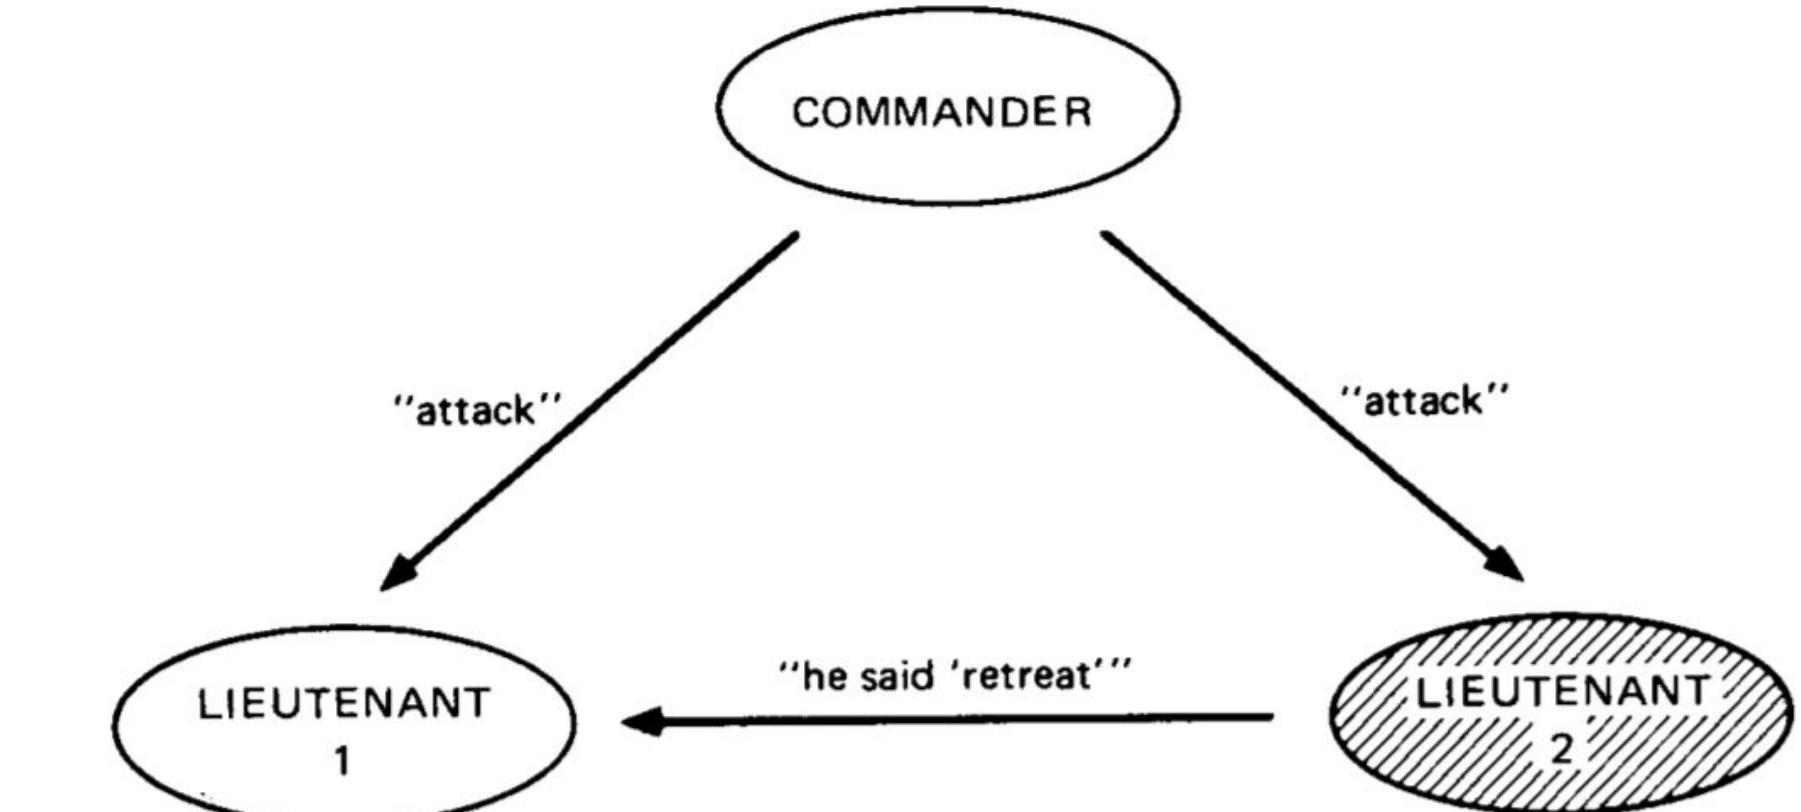


Fig. 1. Lieutenant 2 a traitor.

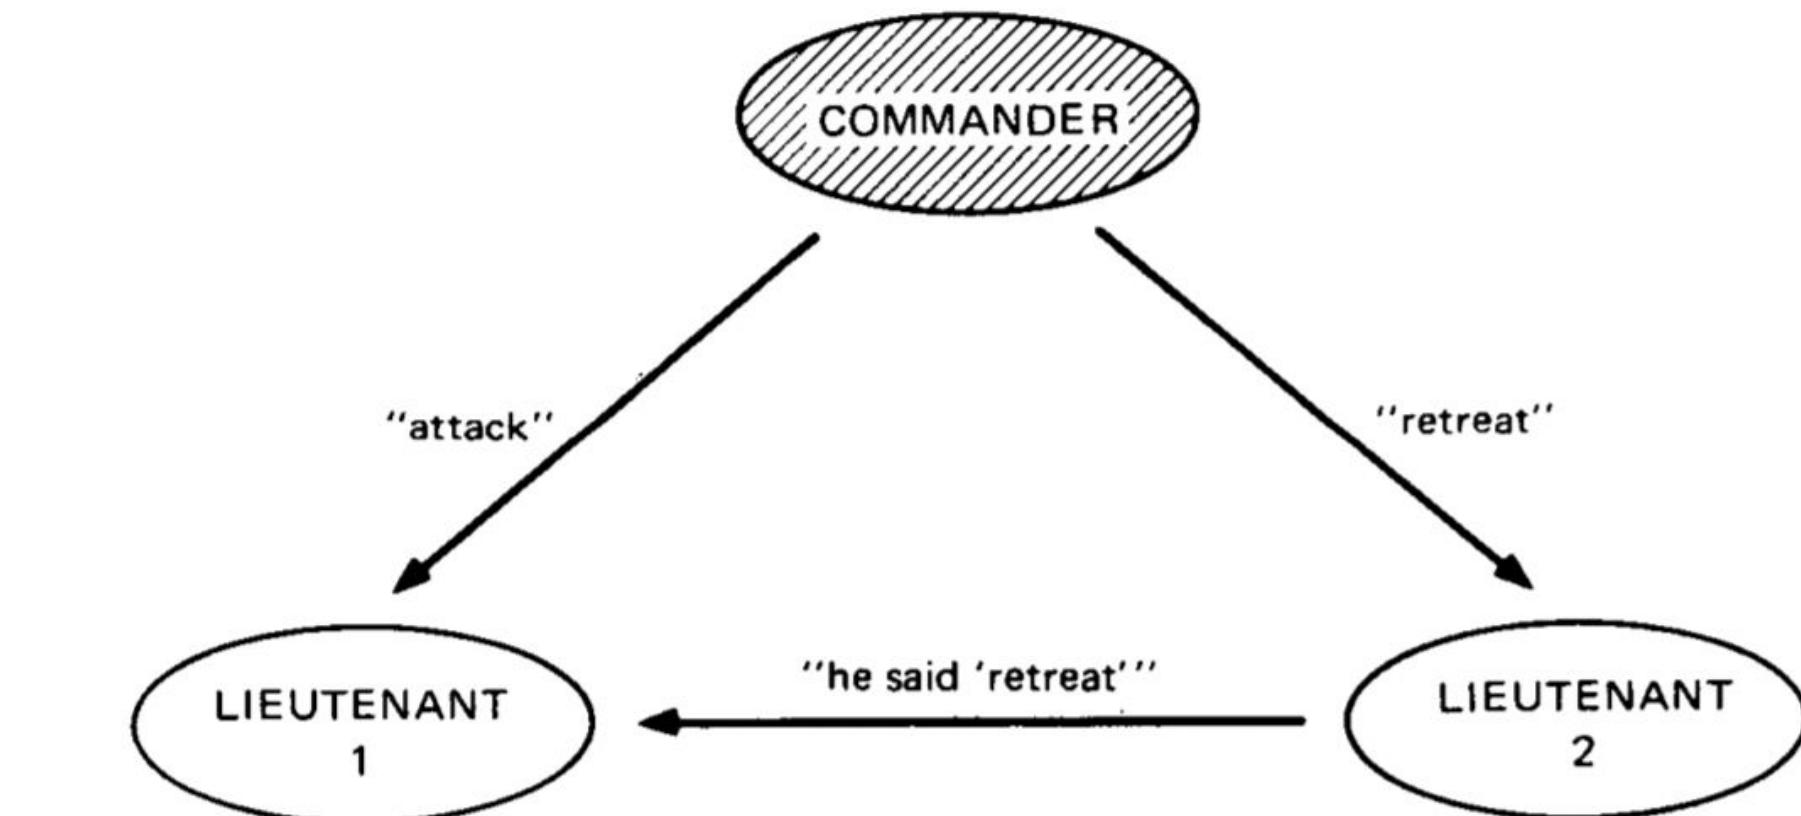
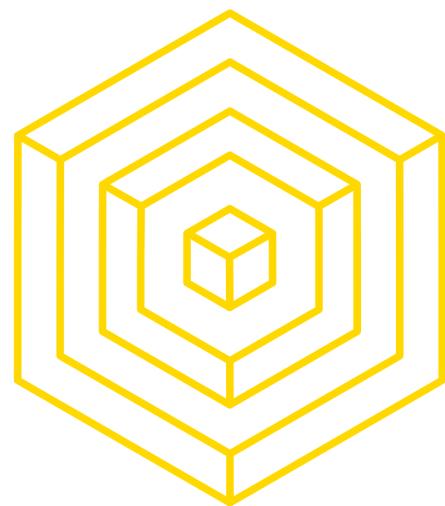


Fig. 2. The commander a traitor.



BYZANTINE FAULT TOLERANCE

DISTINGUISHING FAULT CATEGORIES

Assumptions of fault tolerant systems vs Byzantine fault tolerant systems:

- **Fail-stop fault:** Nodes can crash, not return values, crash detectable by other nodes
- **Byzantine fault:** Nodes can do all of the above *and* send incorrect/corrupted values, corruption or manipulation harder to detect

Source:

[http://www-inst.eecs.berkeley.edu/~cs162
/fa12/hand-outs/Original_Byzantine.pdf](http://www-inst.eecs.berkeley.edu/~cs162/fa12/hand-outs/Original_Byzantine.pdf)

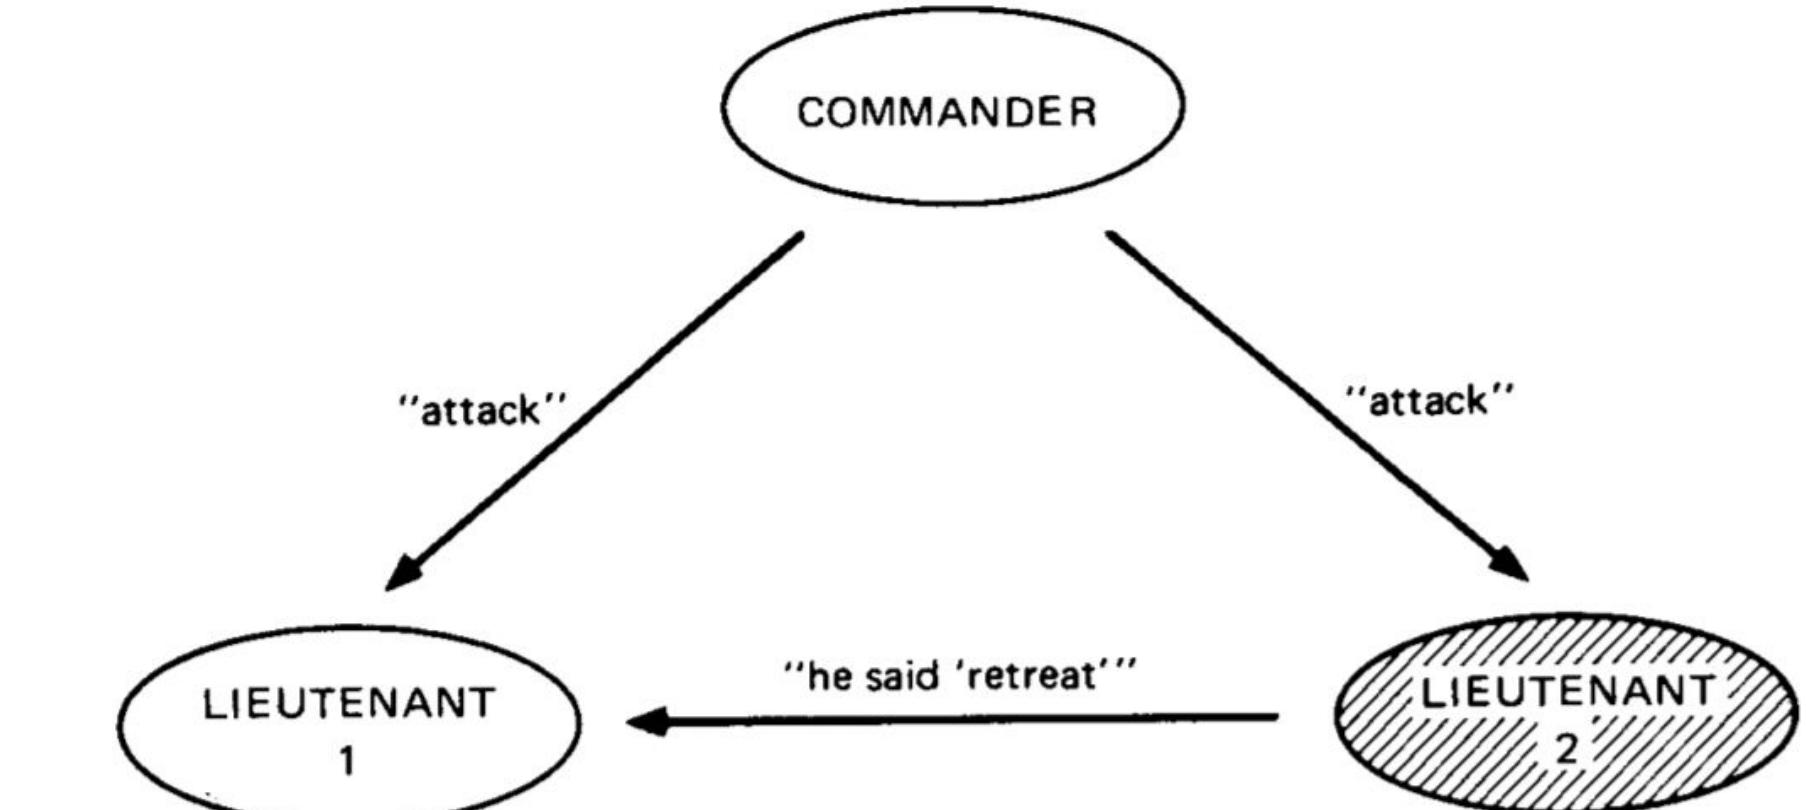


Fig. 1. Lieutenant 2 a traitor.

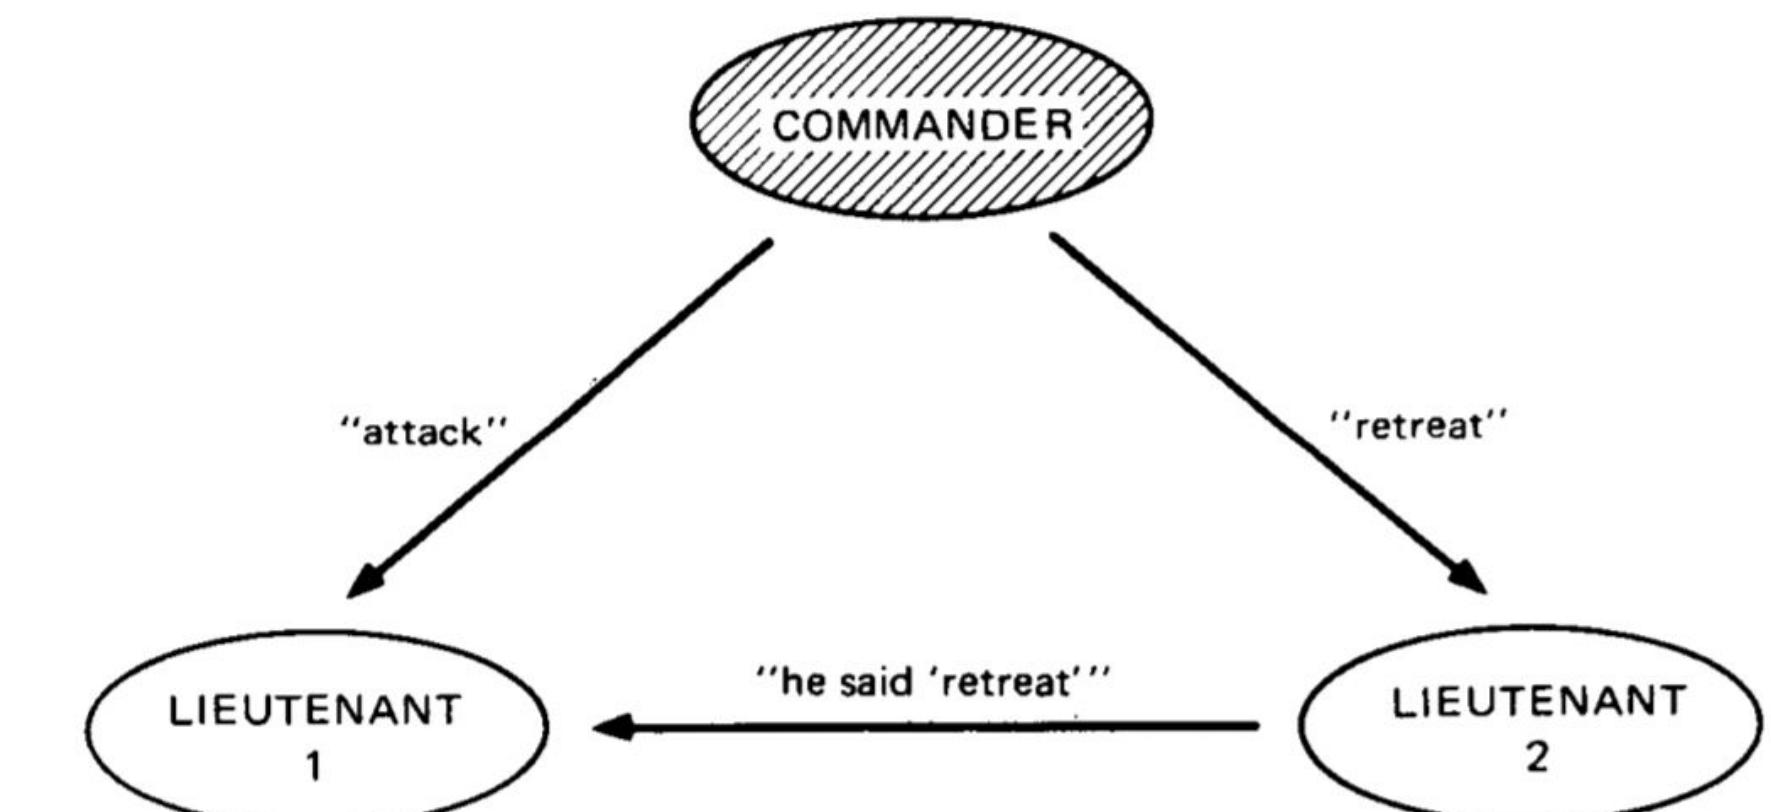
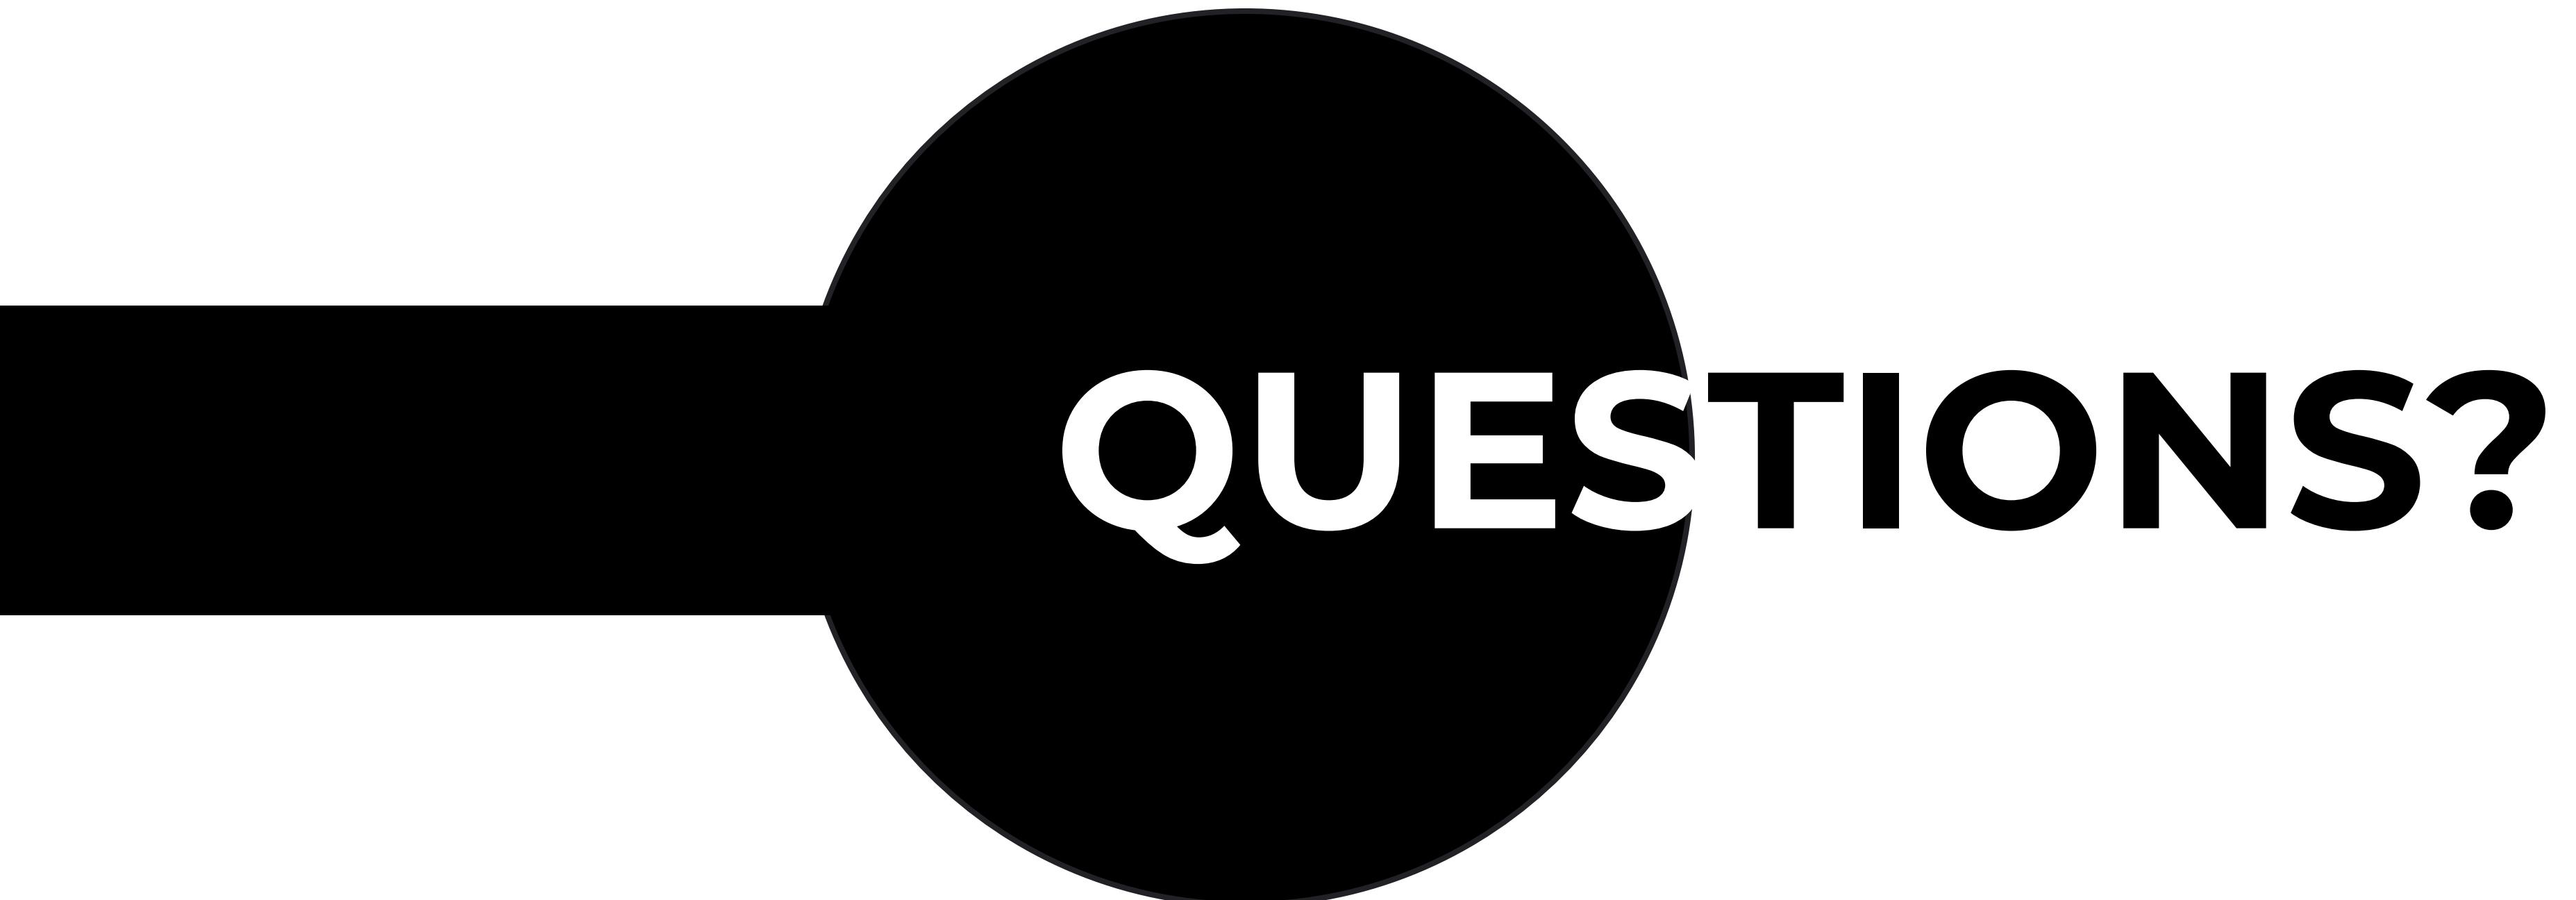
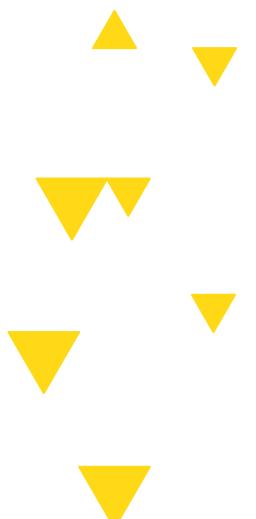
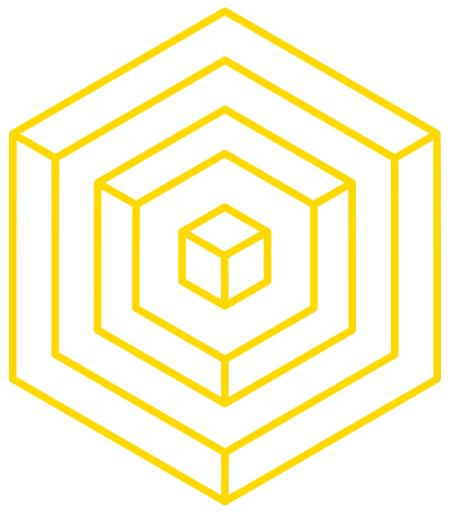


Fig. 2. The commander a traitor.



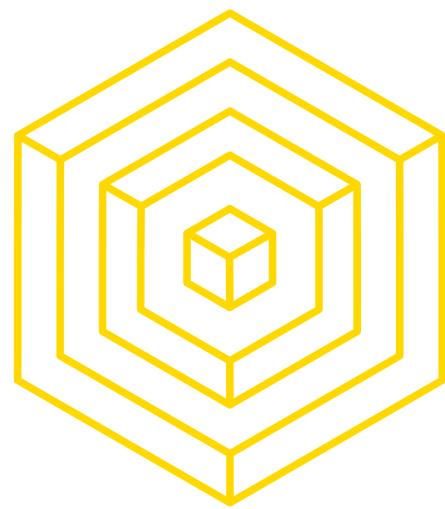
QUESTIONS?





2

PROOF-OF-SOMETHING-ELSE



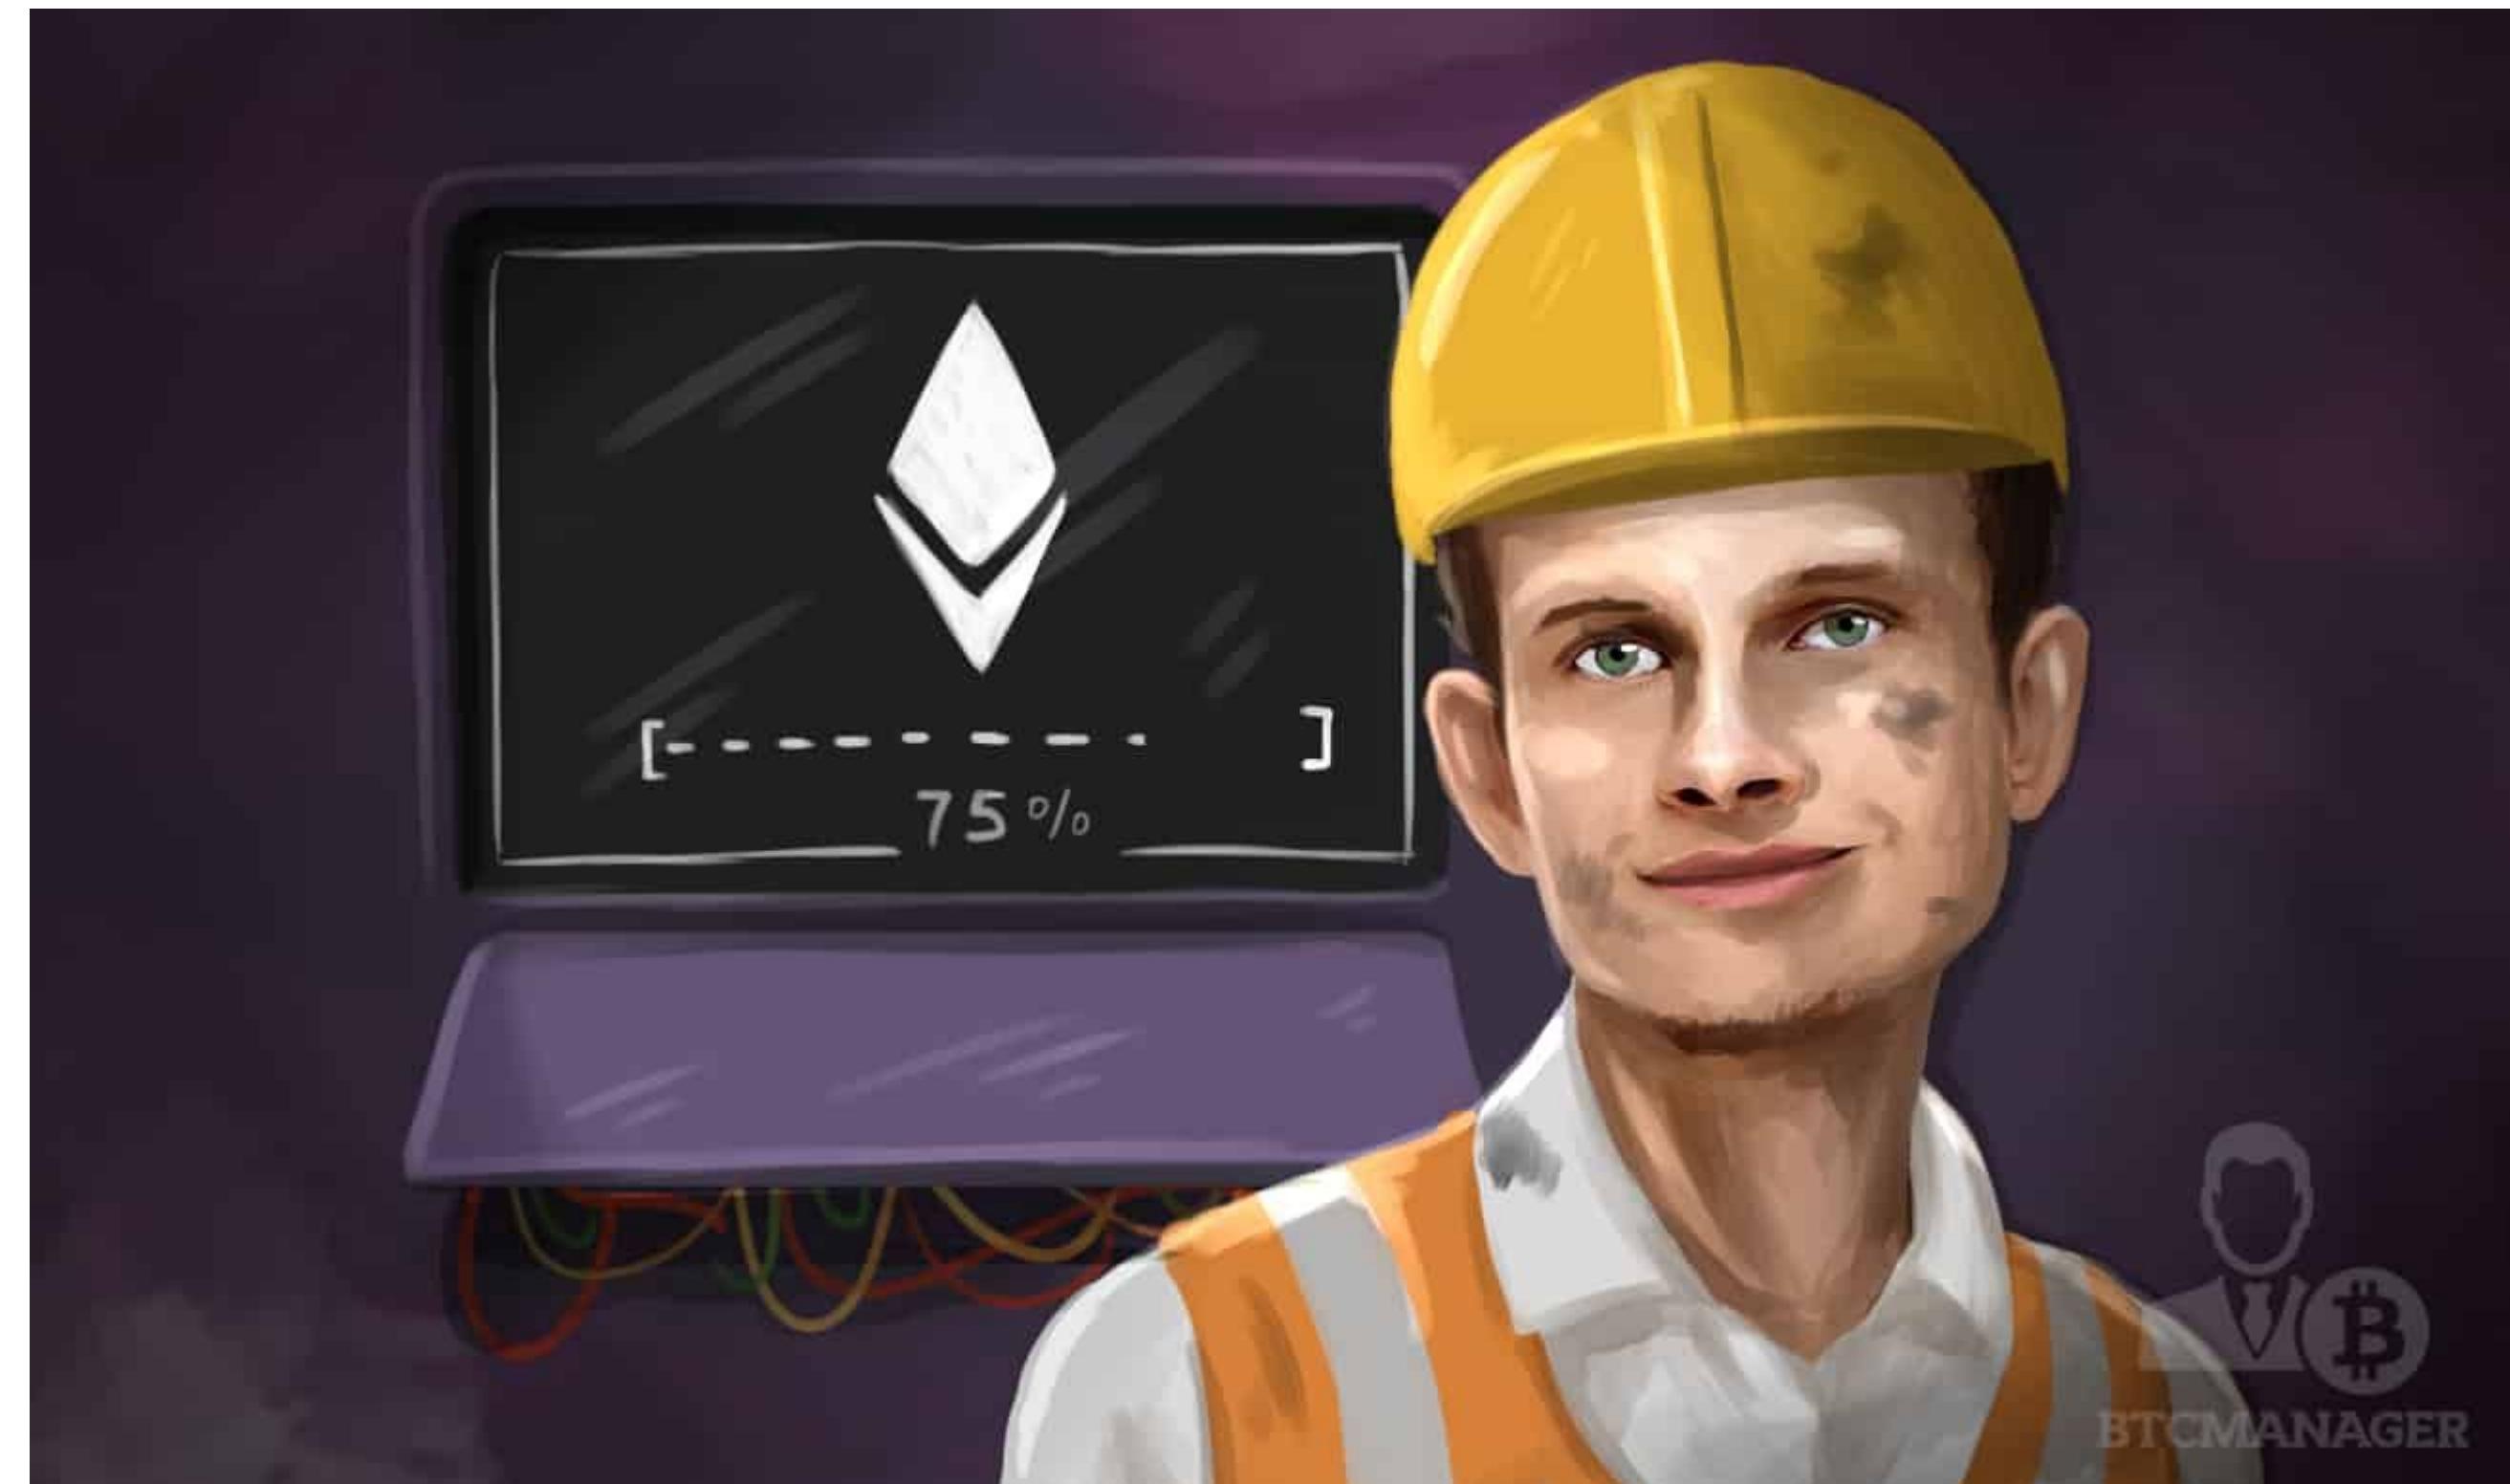
PROOF-OF-SOMETHING-ELSE

PROOF-OF-STAKE

Quick review:

- Validators instead of “miners”
- Locking up “stake”

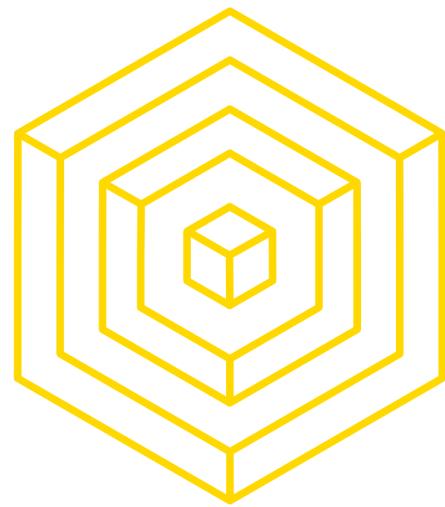
Resource(s) consumed: native currency



Source:
<https://btcmanger.com/buterin-confirms-ethereums-proof-of-stake-75-percent-complete/>

AUTHOR: NADIR AKHTAR

BLOCKCHAIN FUNDAMENTALS LECTURE 8



PROOF-OF-SOMETHING-ELSE

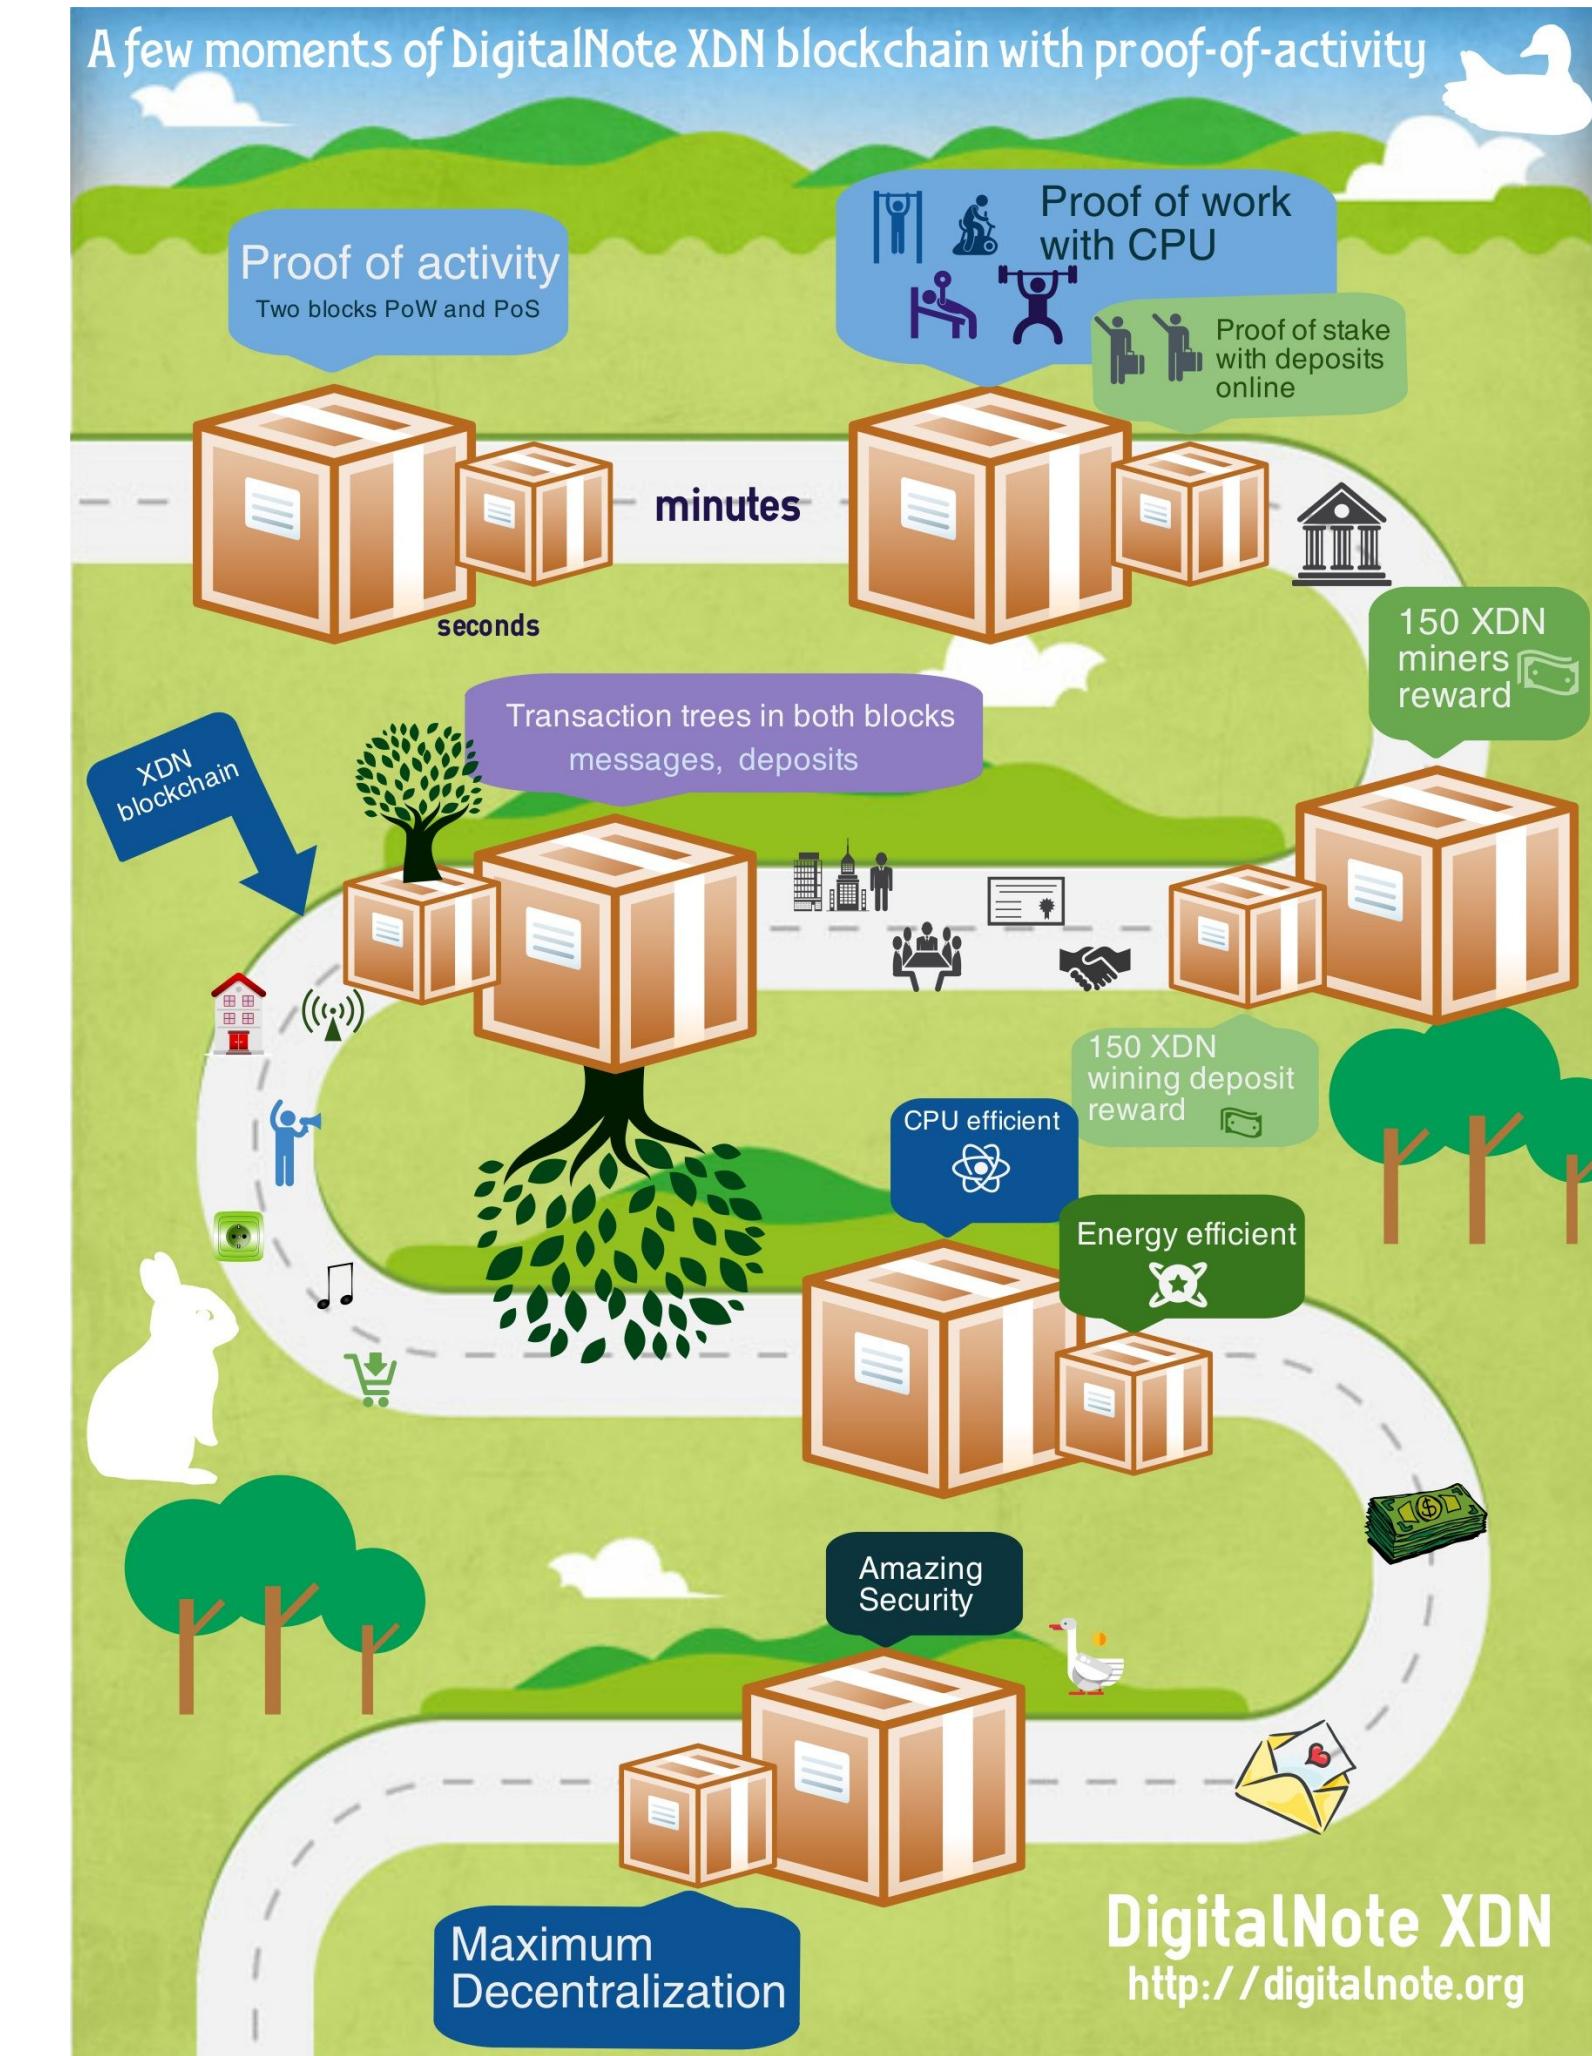
PROOF-OF-ACTIVITY

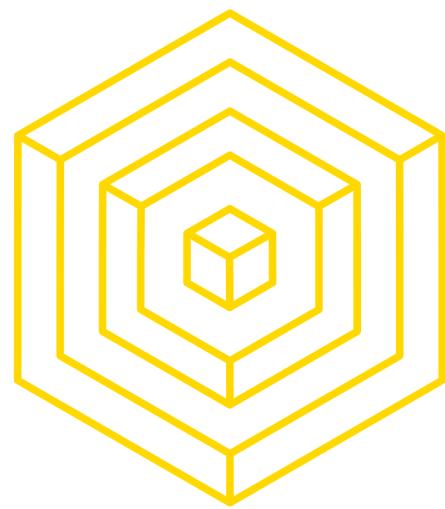
PoW + PoS hybrid:

- PoW: Miners solve and submit block headers to the system (with or without transactions, depending on implementation)
- PoS: Validators vote on valid blocks, receiving awards for signing

Resource(s) consumed: PoW and PoS resources

Source:
<http://digitalnotetalk.org/sites/digitalnotetalk.org/files/media/1459869843.jpg>





PROOF-OF-SOMETHING-ELSE

PROOF-OF-BURN

Overview:

- Send coins to irretrievable address
 - More coins burned, higher likelihood of voting
- Like Proof-of-Stake, but edgier

Bootstrapping mechanism:

- Tie a coin's value to some other coin (e.g. Bitcoin) by demonstrating that users willingly burn Bitcoin to receive this other coin

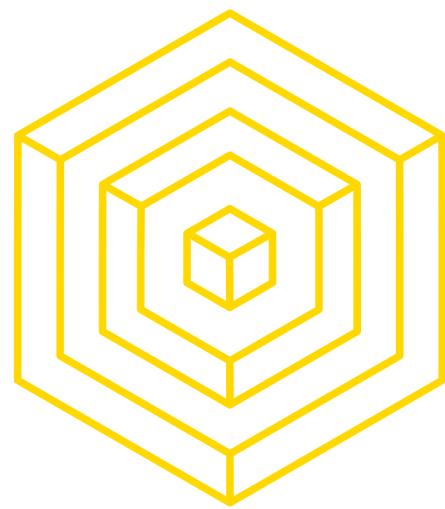
Resource(s) consumed: currency (potentially not native)



AUTHOR: NADIR AKHTAR



Source: https://www.reddit.com/r/dogecoin/comments/2bhqih/with_all_this_talk_about_proof_of_burn_i_thought/



PROOF-OF-SOMETHING-ELSE

PROOF-OF-SPACE

Overview:

- Use disk space to solve challenge
- Nowadays used for file storage
- Variations on implementations:
 - Let users receive coins after voting on blocks
(Also known as Proof-of-Capacity)
 - Give rewards directly for storing information
(Filecoin, Storj, Sia)



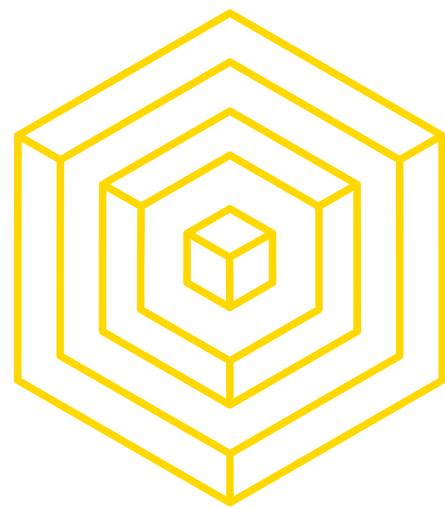
Source:
<https://hackernoon.com/crypto-review-siacoin-sc-b1d0f0a5c78f>



Source:
<https://medium.com/@tokenlot/filecoin-ico-open-to-accredited-investors-only-b7937f24de44>

Resource(s) consumed: storage space

Source:
<https://coinist.com/storj-new-decentralized-storage-solution/>



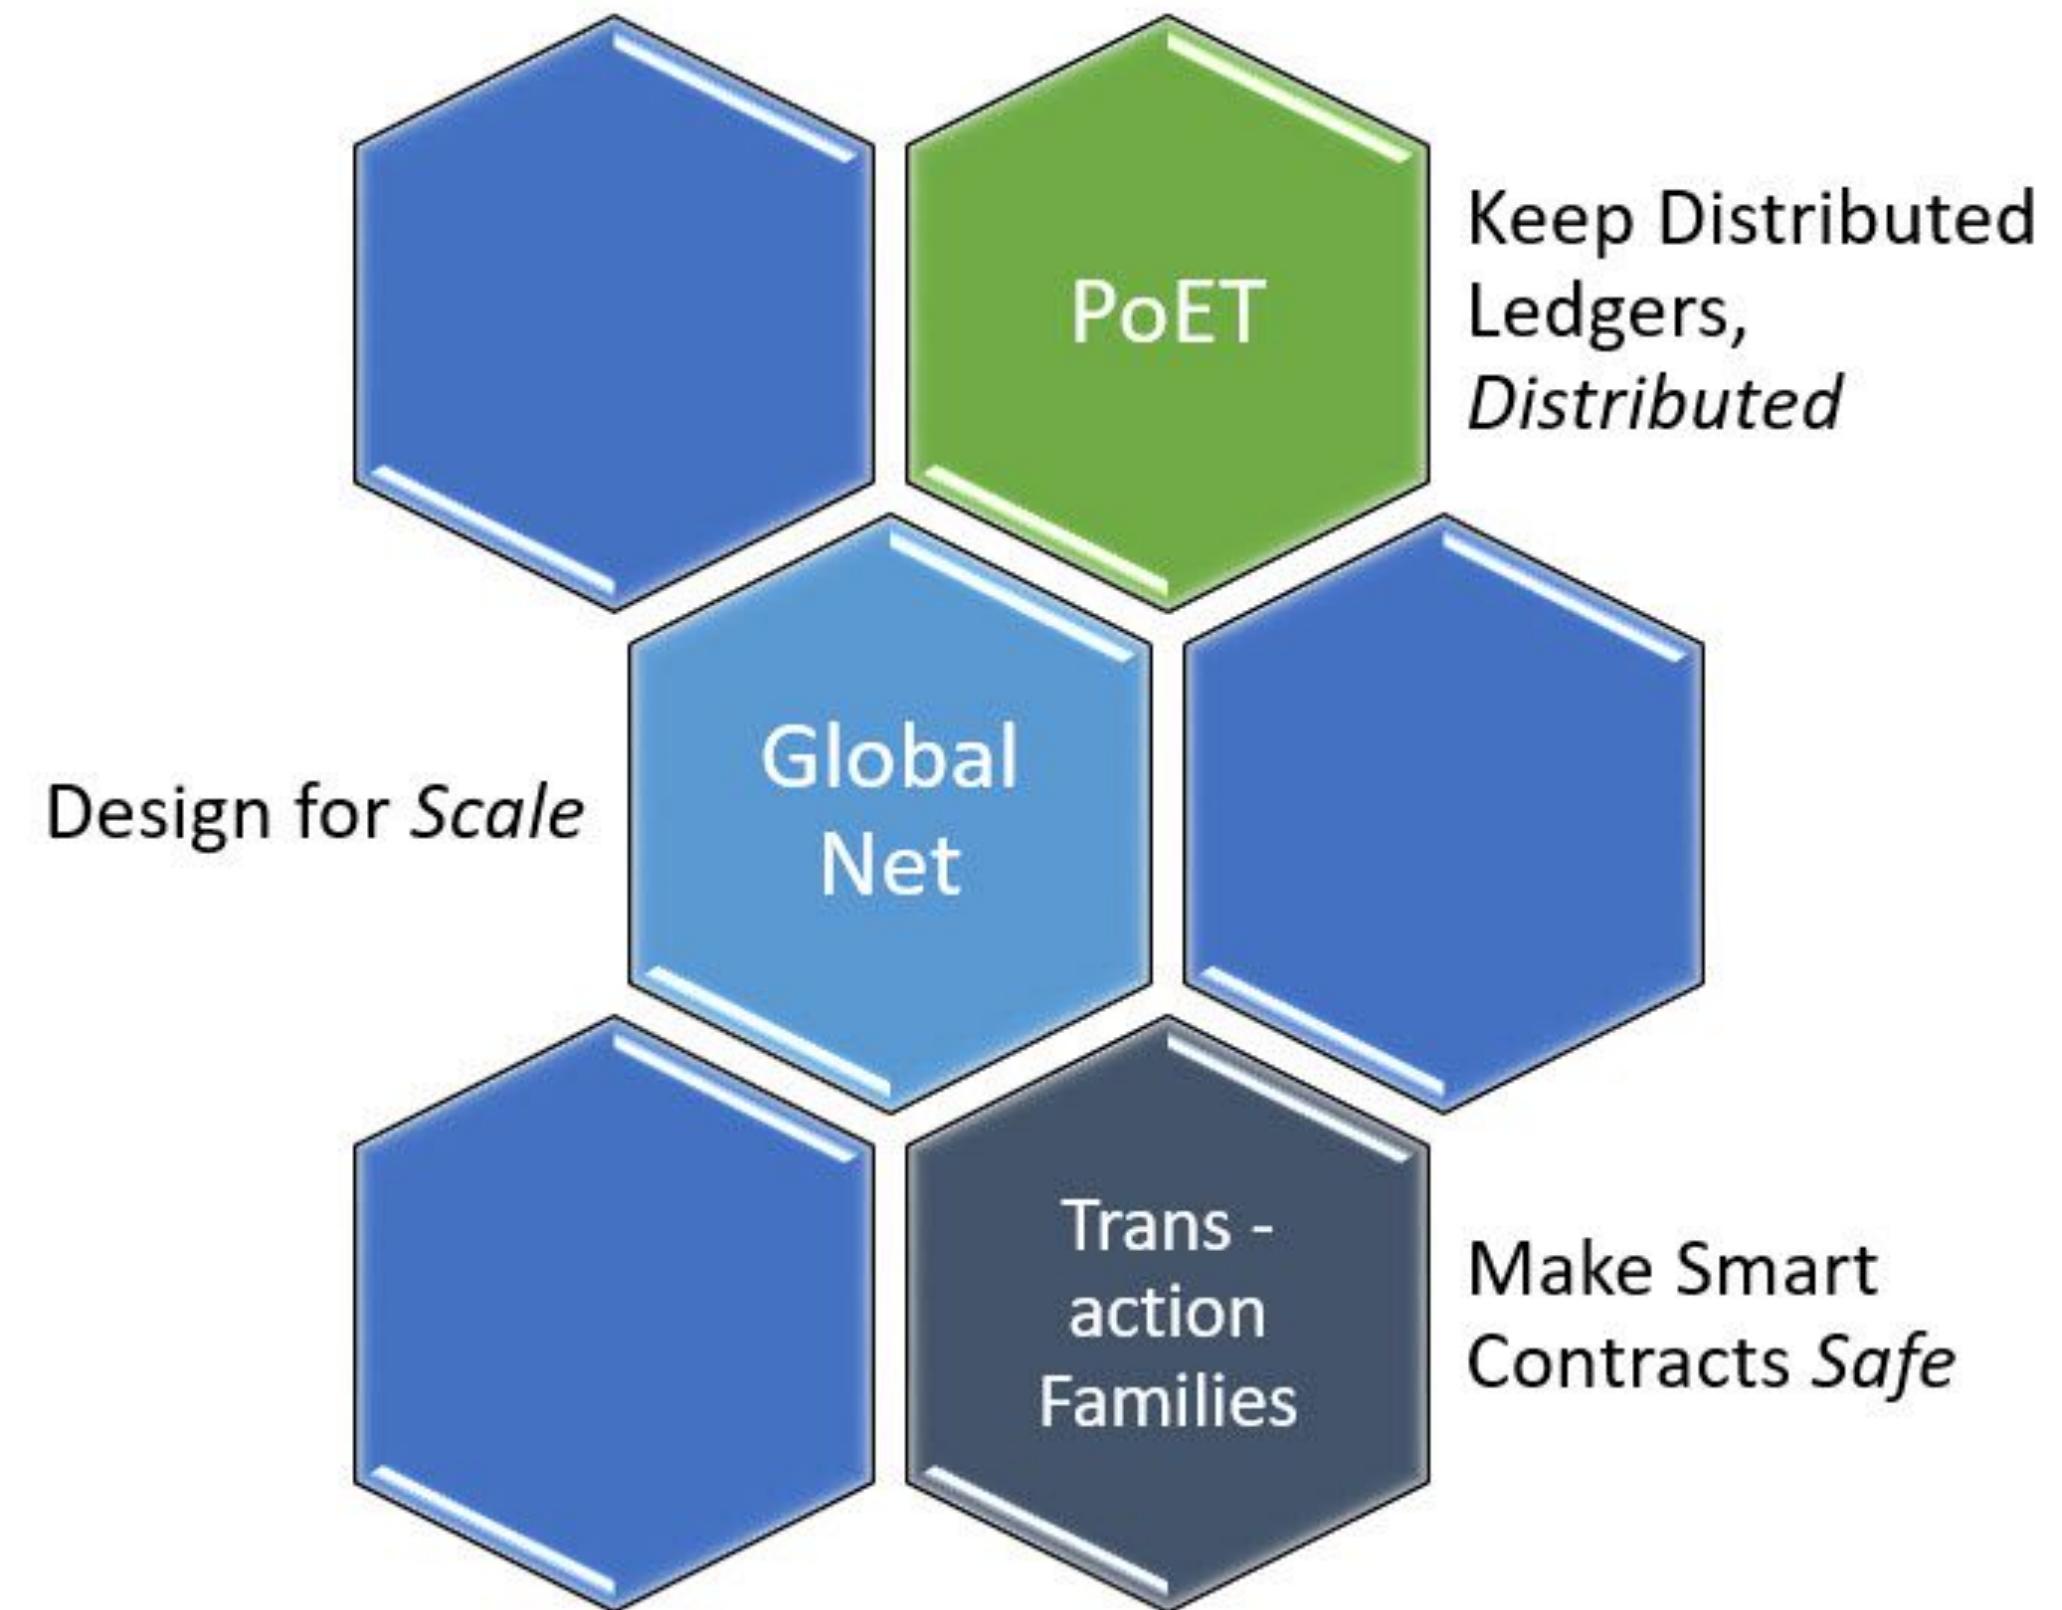
PROOF-OF-SOMETHING-ELSE

PROOF-OF-ELAPSED-TIME

Overview:

- Spend time instead of mining power
 - Uses Trusted Execution Environments (TEEs), particularly Intel SGXs
- Ask your SGX to wait some amount of time decided by the machine, wait for an attribution that it has indeed waited that much time
 - Let the first person to complete this random wait win
- Based on assumptions of randomness and trust in manufacturer (Intel)

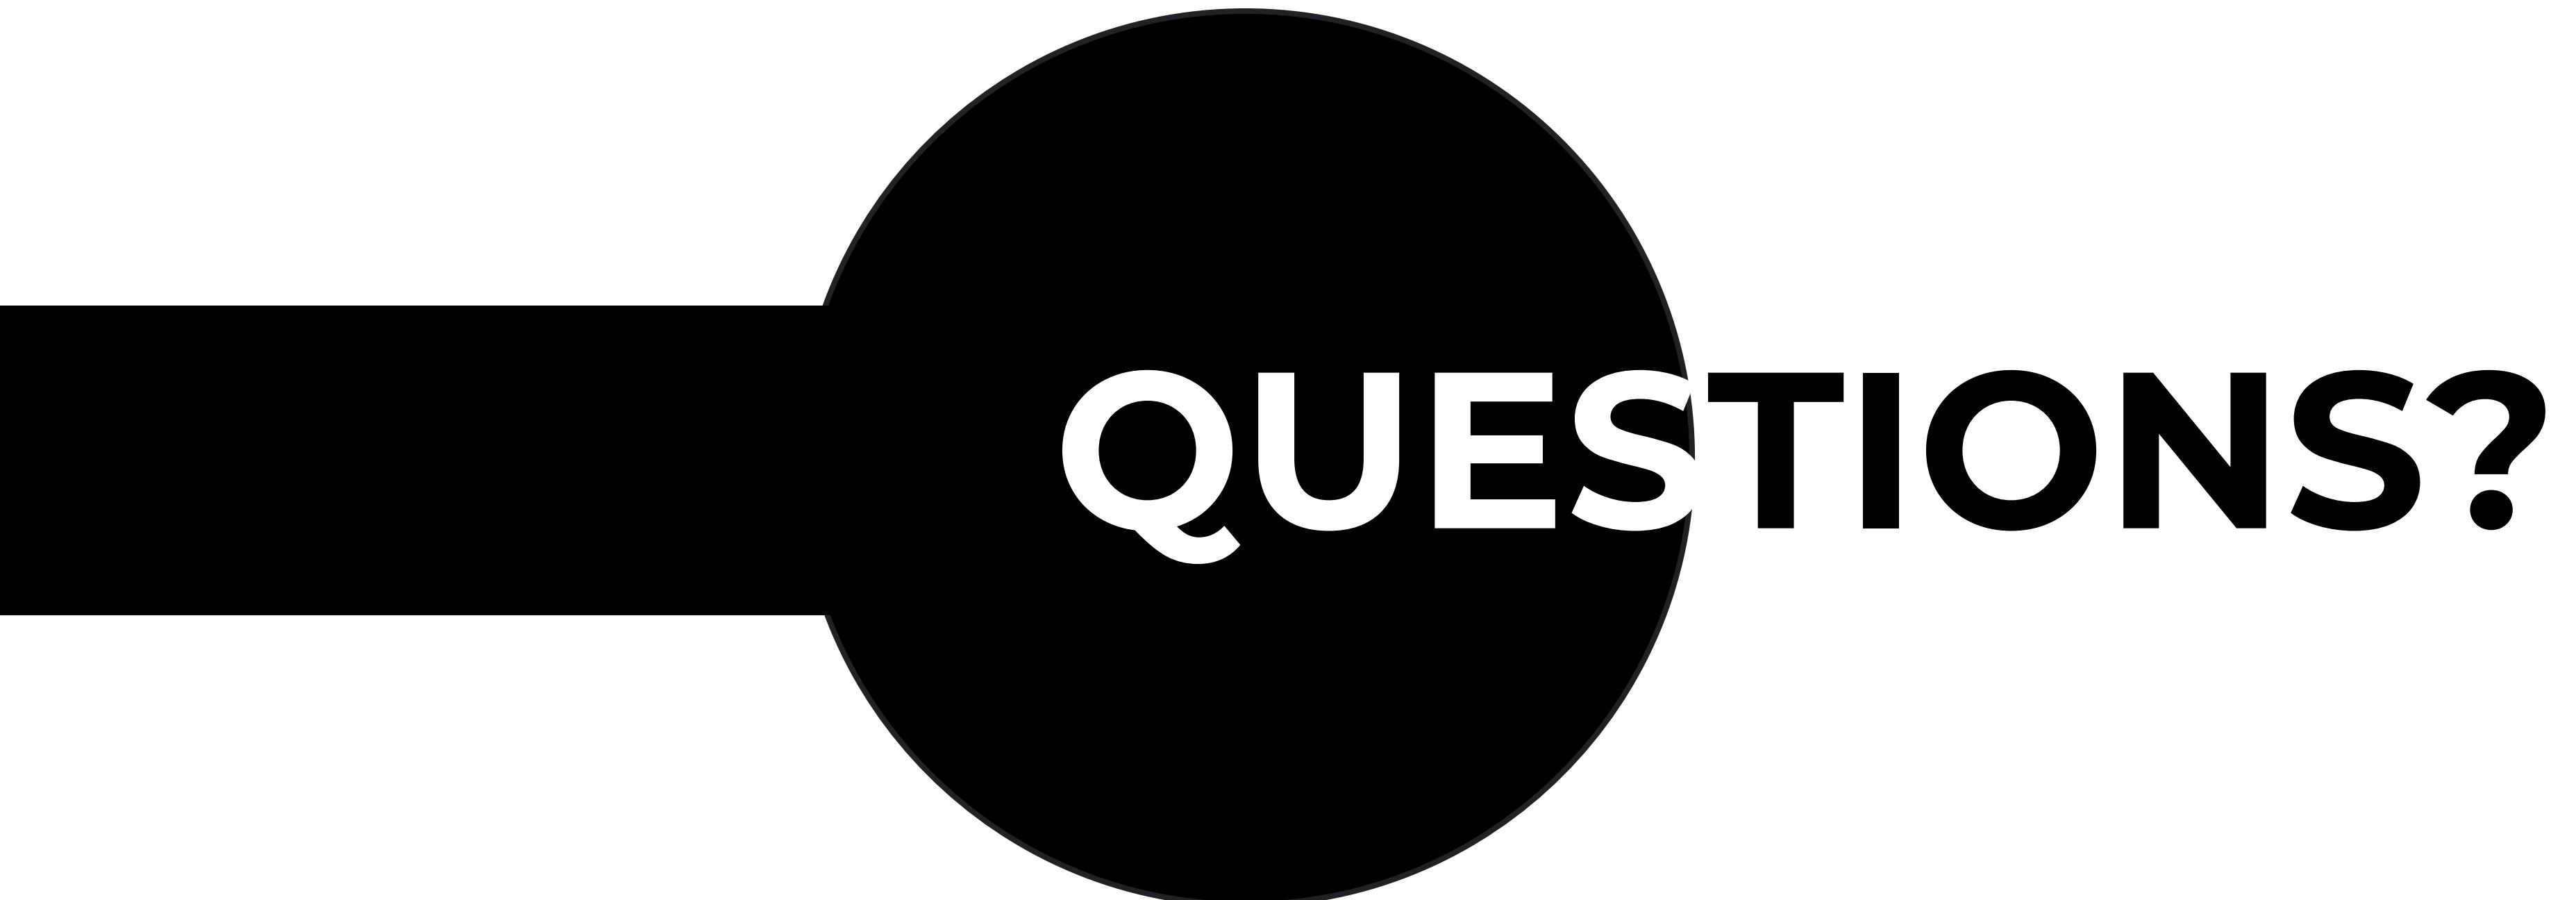
Resource(s) consumed: time



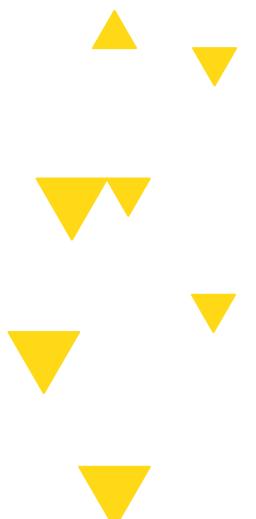
Source: <https://www.altoros.com/blog/wp-content/uploads/2017/02/Hyperledger-Intel-Sawtooth-Lake-difference.jpg>

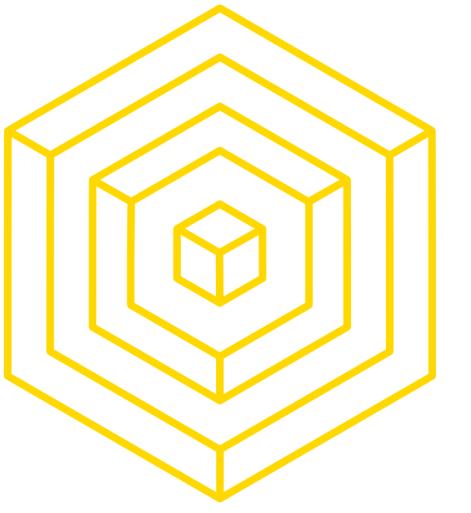


AUTHOR: NADIR AKHTAR



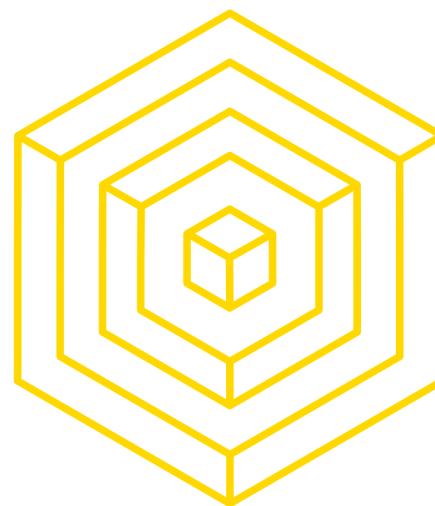
QUESTIONS?





3

PROOF OF STAKE

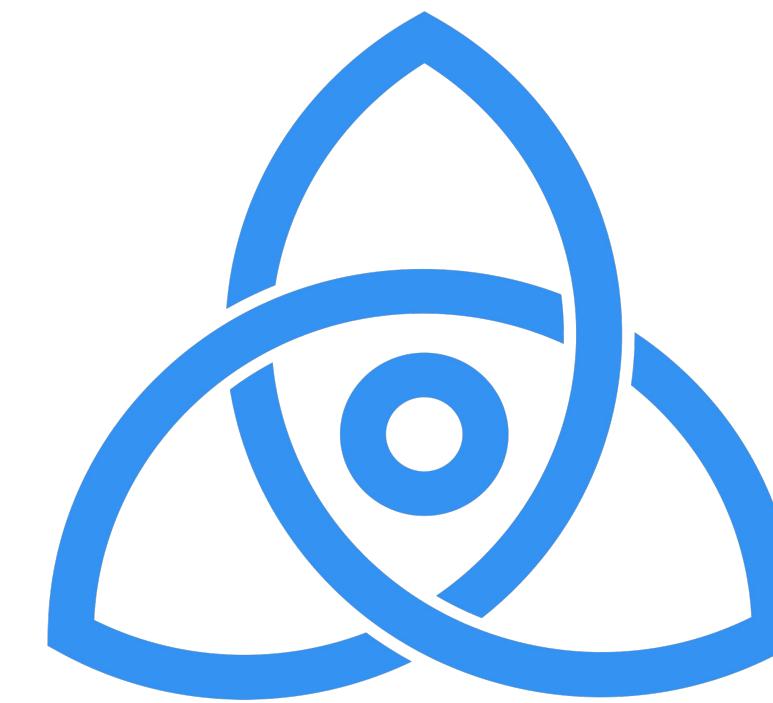


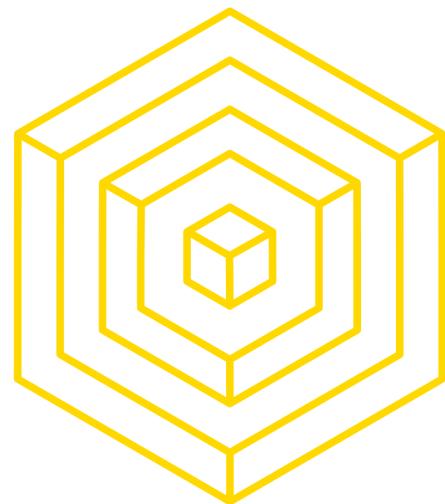
PROOF-OF-STAKE

IMPLEMENTATIONS

- **Tendermint**: mostly asynchronous, BFT consensus protocol with voting power denominated in validator stake; two-stage voting process; favors safety (consistency). Used in blockchain network Cosmos.
- **Casper**: Two versions:
 - *Casper the Friendly GHOST*: Correct-by-Construction (Vlad), which designs the protocol from scratch
 - *Casper the Friendly Finality Gadget* (Vitalik), a PoW/PoS hybrid that uses checkpoints

CØSMOS
INTERNET OF BLOCKCHAINS

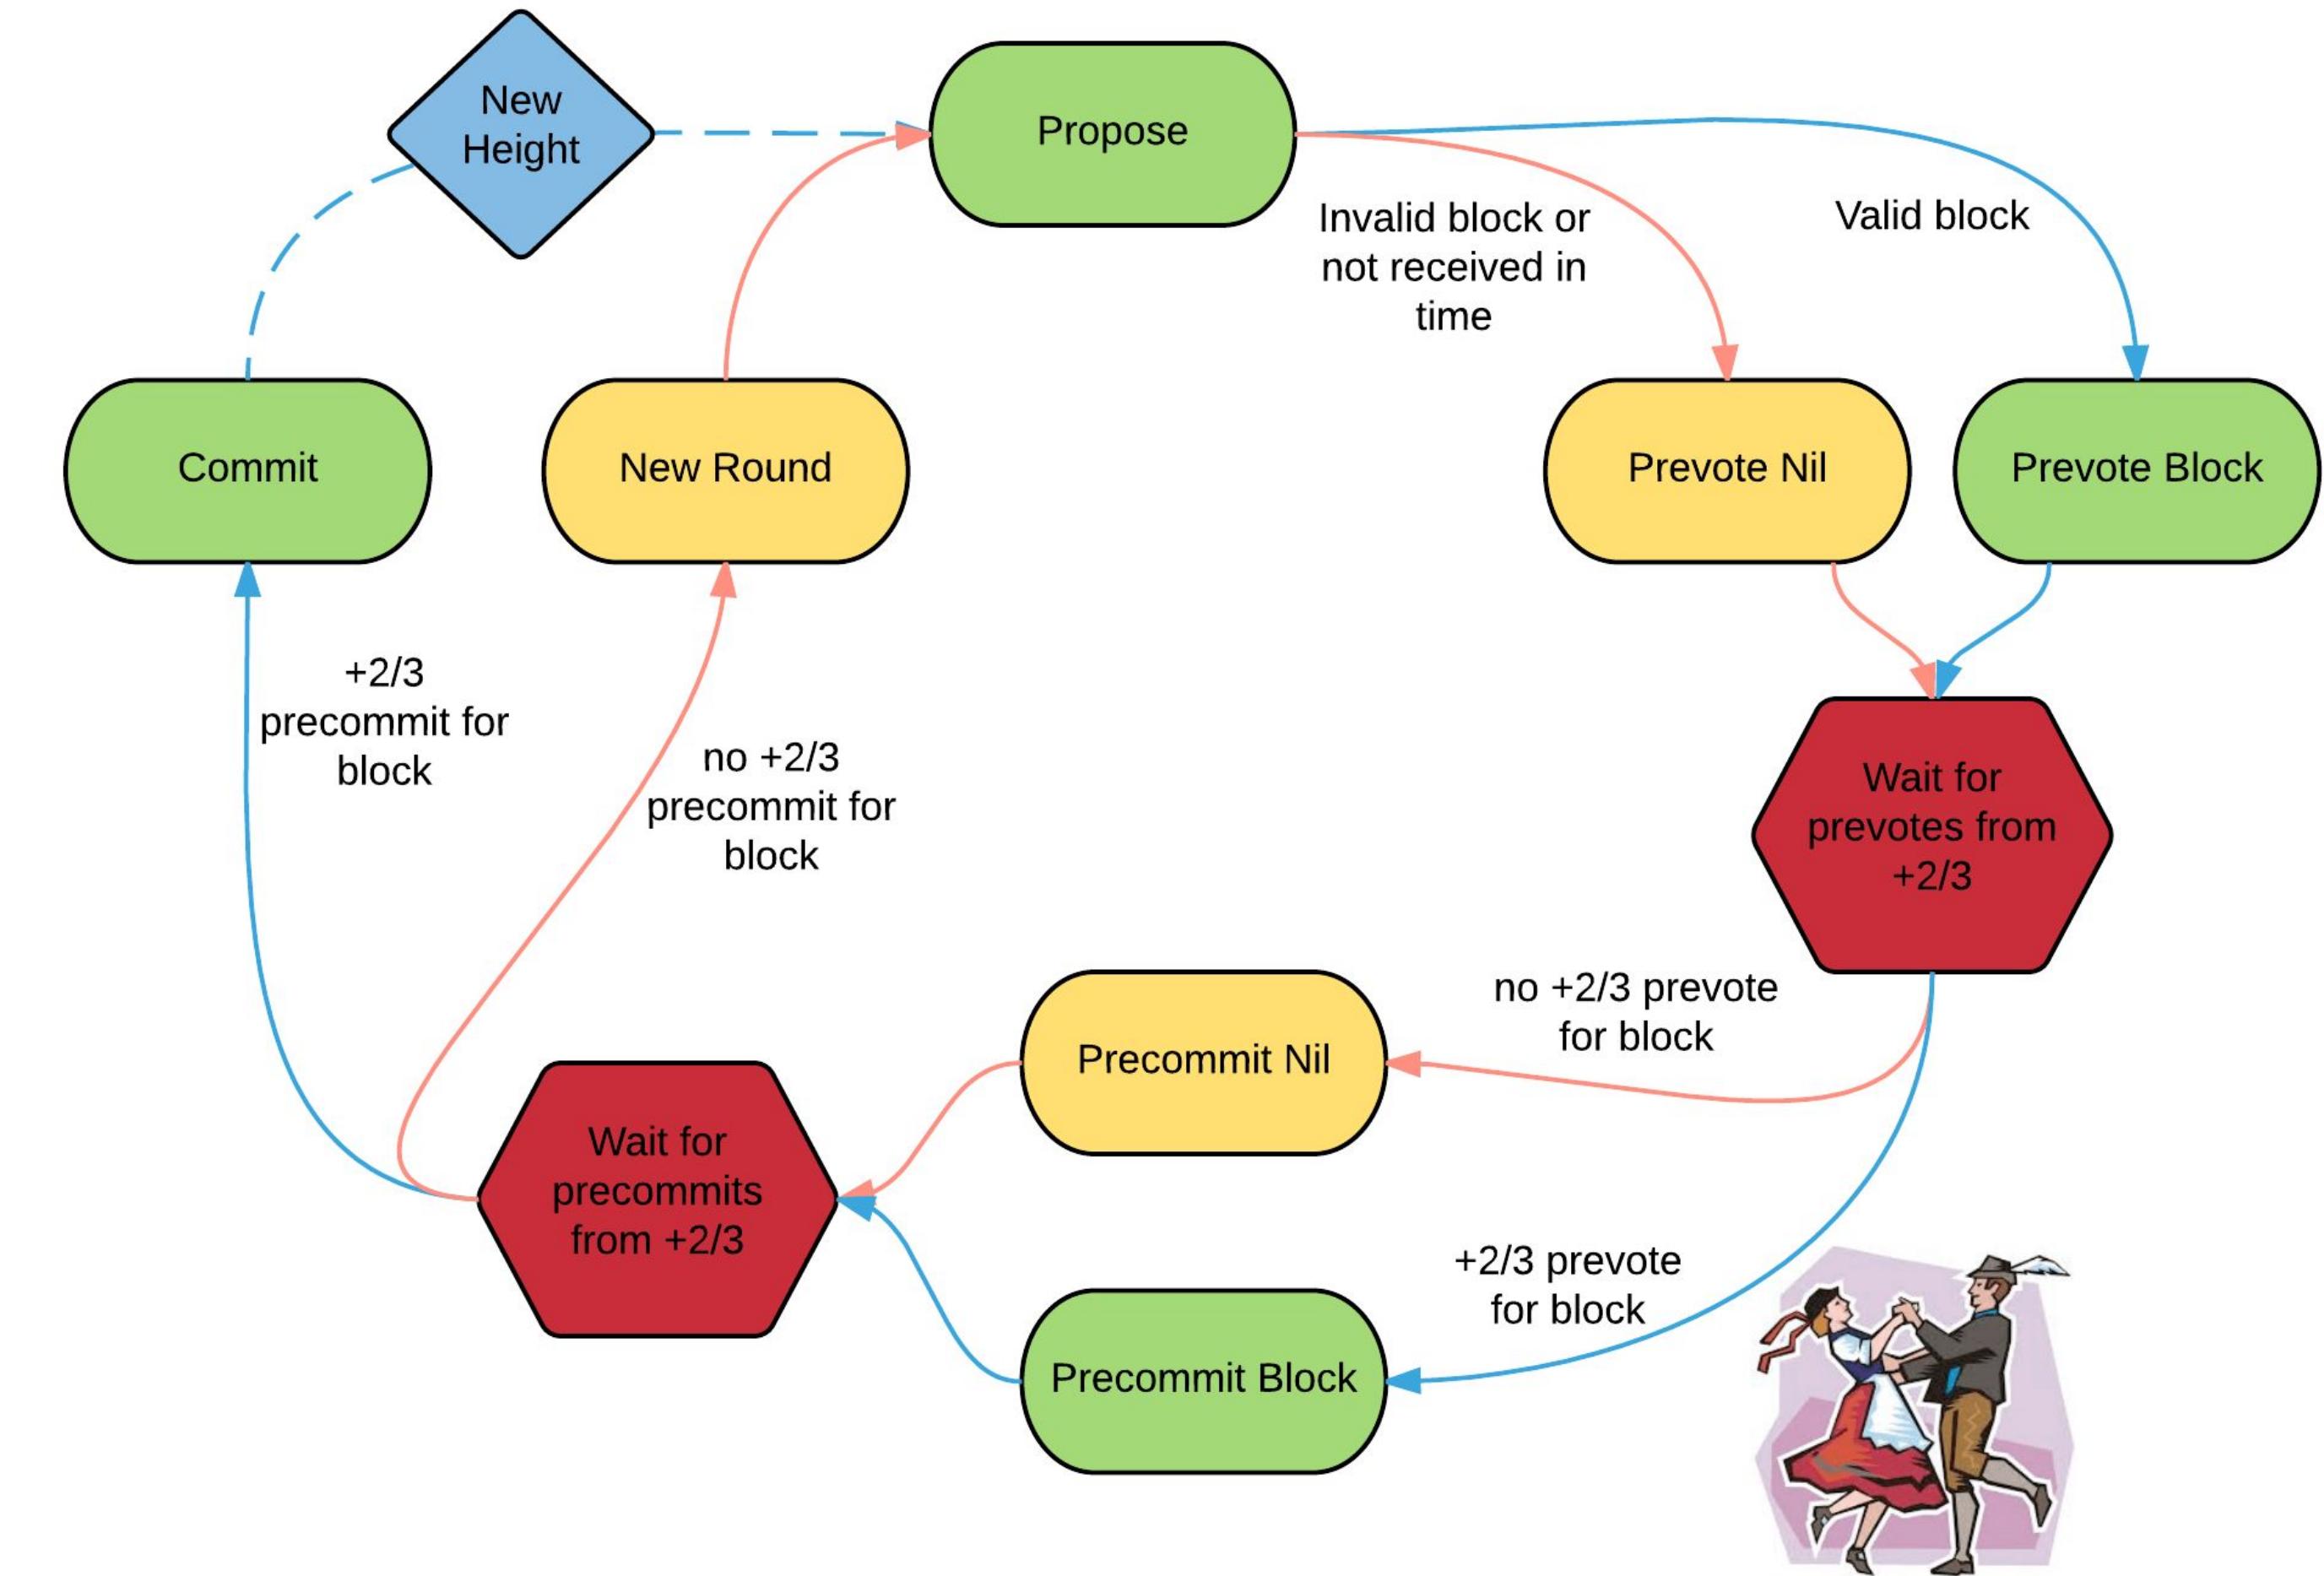




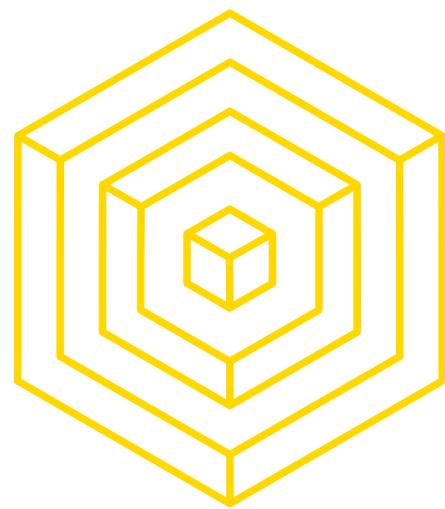
PROOF-OF-STAKE

TENDERMINT

- BFT-based PoS:
 - A chosen set of validators of a specific size at any given point with bonded stake
 - Weakly asynchronous
 - Favors consistency over availability
 - ~~Requires dancing experience~~
 - Energy efficient & fast
 - Excellent for enterprise blockchains



Source:
<http://tendermint.readthedocs.io/projects/tools/en/master/introduction.html>



PROOF-OF-STAKE

CASPER

Two versions, one by Vitalik and one by Vlad Zamfir

- Casper the Friendly Finality Gadget is a chain/BFT-style hybrid
 - Anyone can bond tokens
 - Initially PoW/PoS hybrid, a slow transition to pure PoS
- Casper the Friendly Ghost is a pure PoS protocol
 - Focuses on safety
 - Maintains validator set

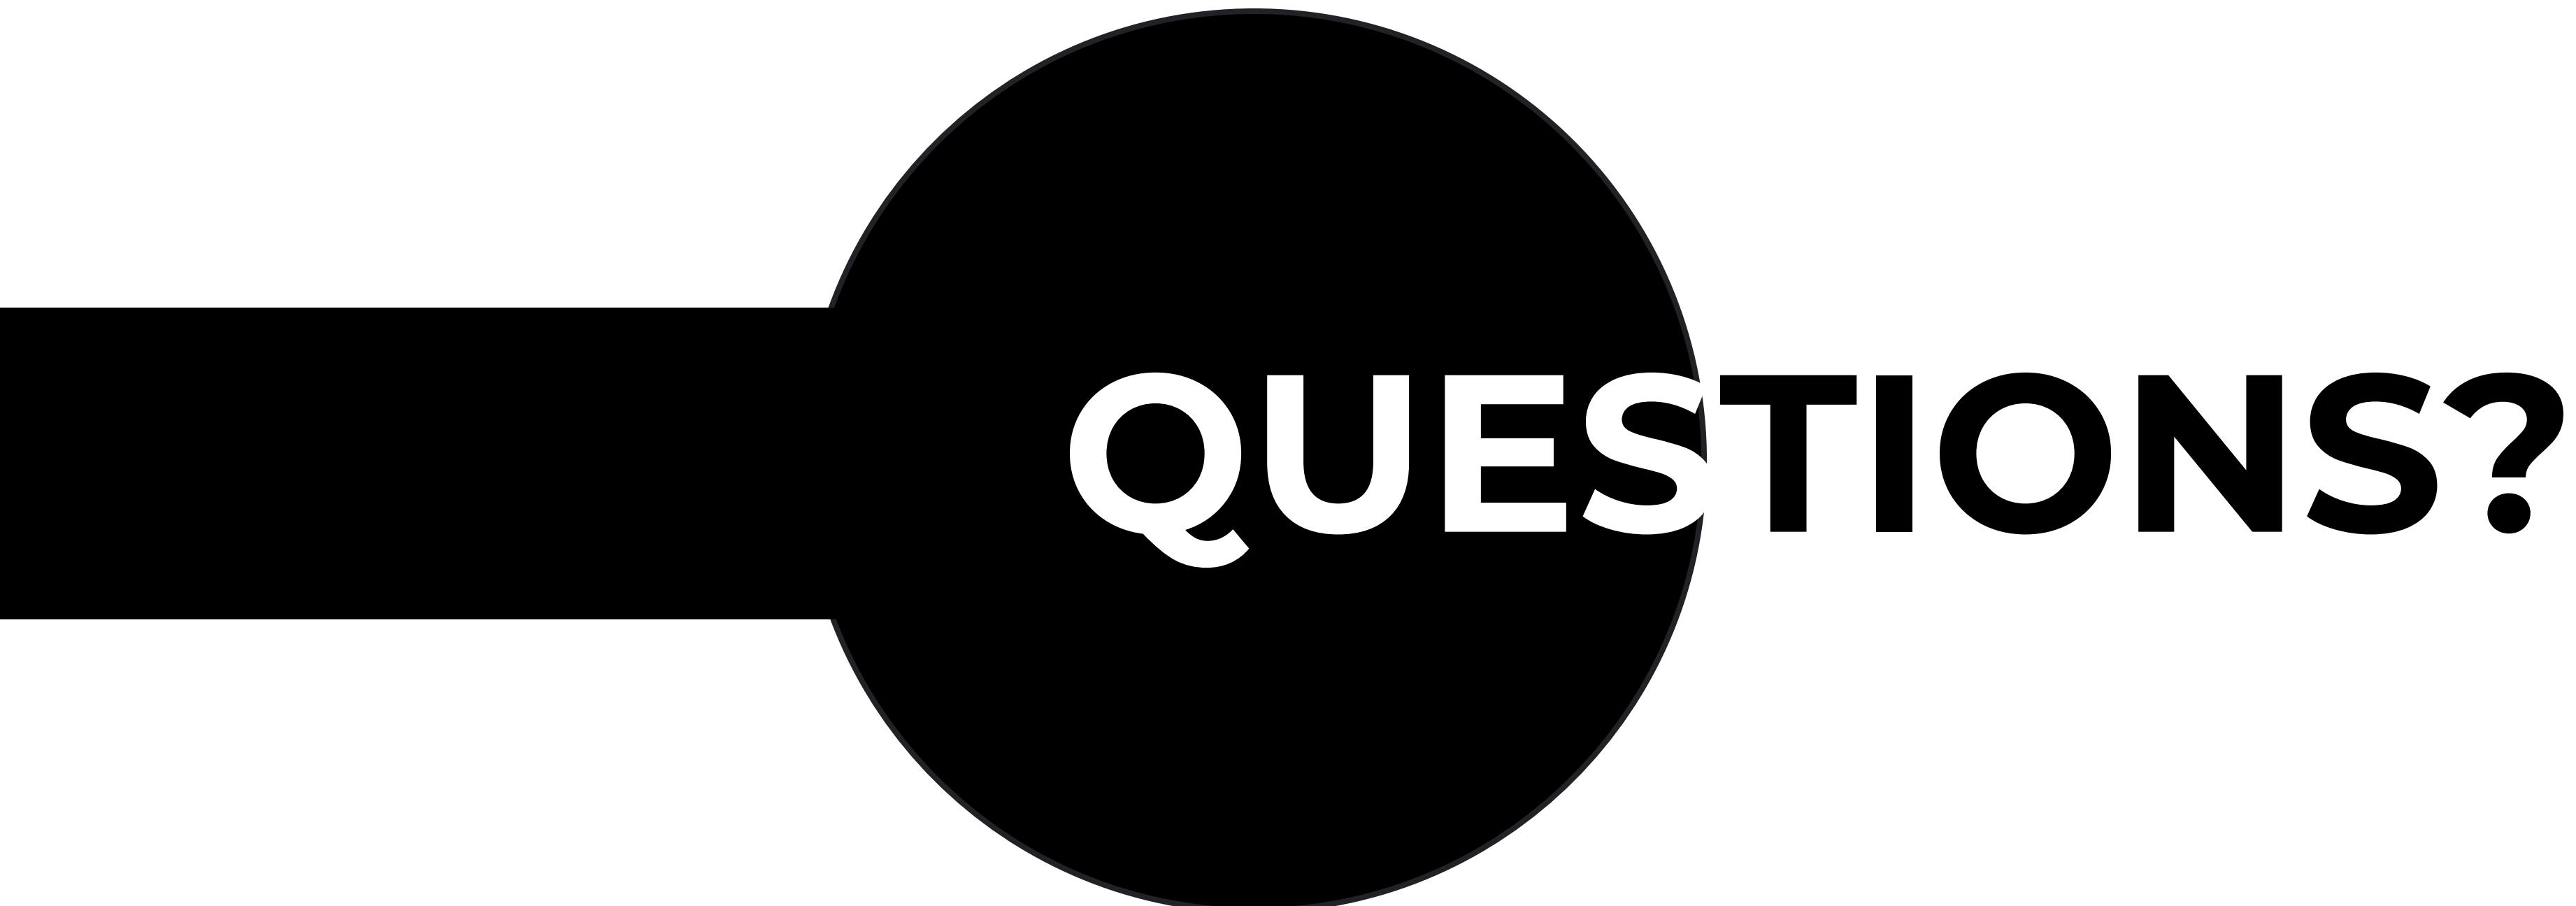


c a s p e r

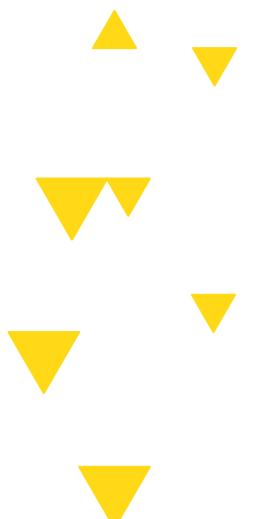
Source:
<https://media.consensys.net/casper-smart-contract-consensus-7be6cfa6f7ec>



AUTHOR: NADIR AKHTAR



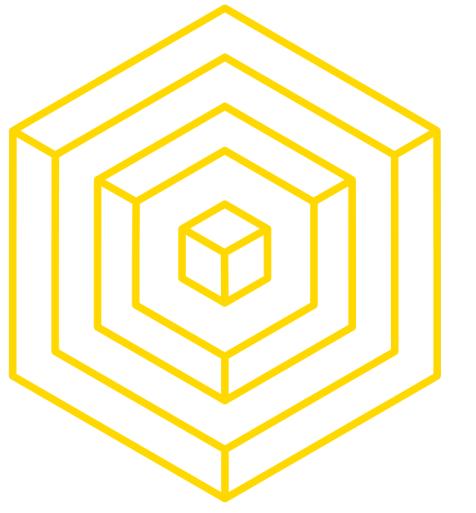
QUESTIONS?





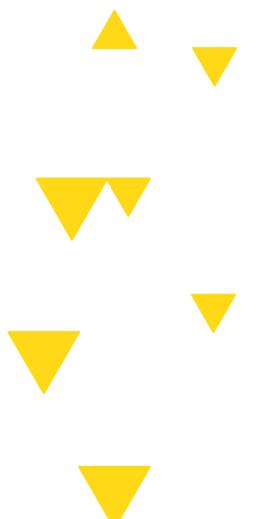
BLOCKCHAIN
AT BERKELEY

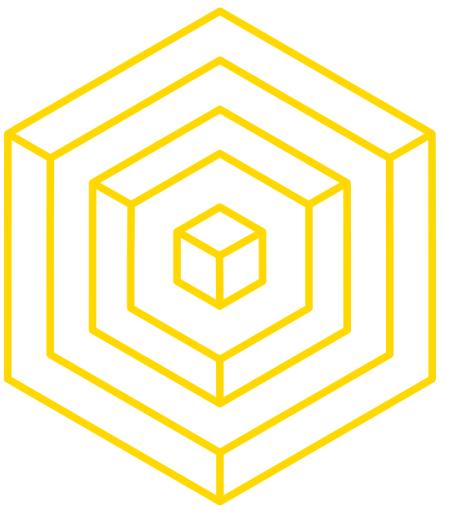
**BREAK
SECTION**



4

VOTING-BASED CONSENSUS ALGORITHMS

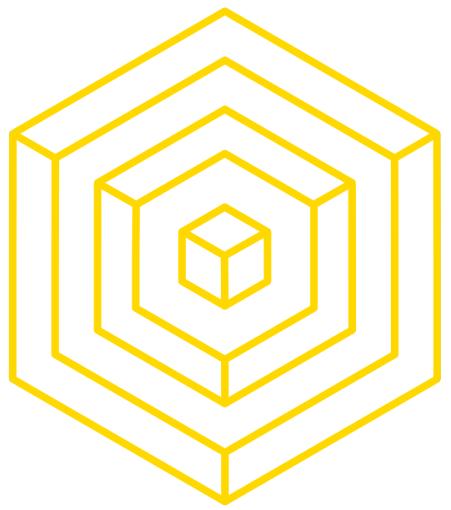




Source: <http://greece10best.com/wp-content/uploads/2017/05/Paxos.jpg>

AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



PAXOS

OVERVIEW

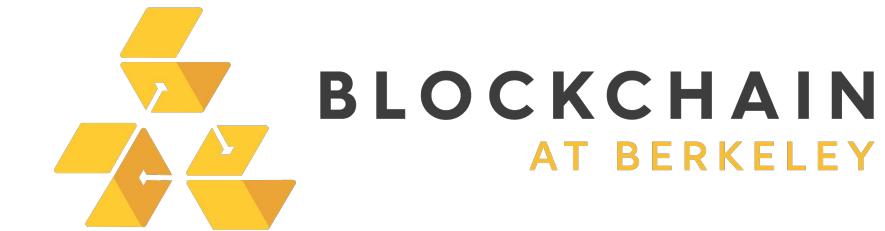
39

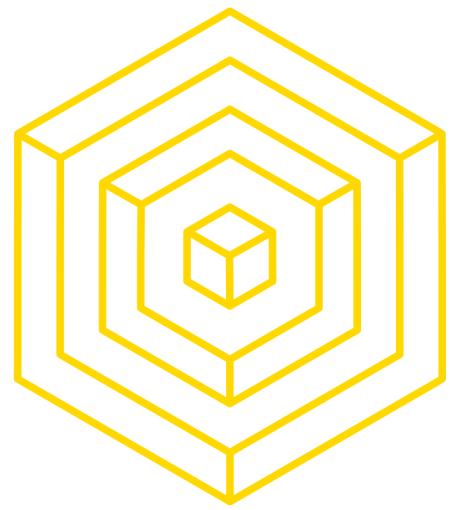


Source: <http://paxos-map.com/paxos-island-general-information/>

AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8





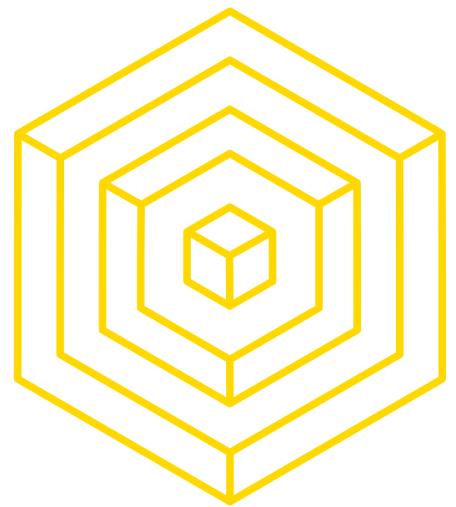
PAXOS

OVERVIEW



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



PAXOS

OVERVIEW

Paxos is a Greek island

but also . . .

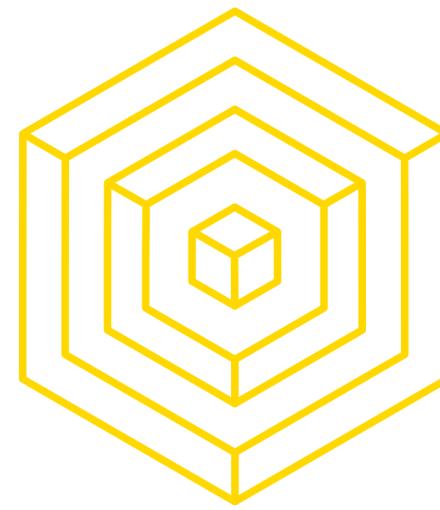
A CONSENSUS
ALGORITHM



Source: <http://www.hikingbikingadventures.com/biking-adventures/bicycling-europe-2/bicycling-greece/bicycling-paxos-island/>



AUTHOR: RUSTIE LIN



PAXOS

TERMINOLOGY

Proposer: advocates a client request, moves protocol forward

Acceptor: a voter

Learner: remembers result, send message to the client

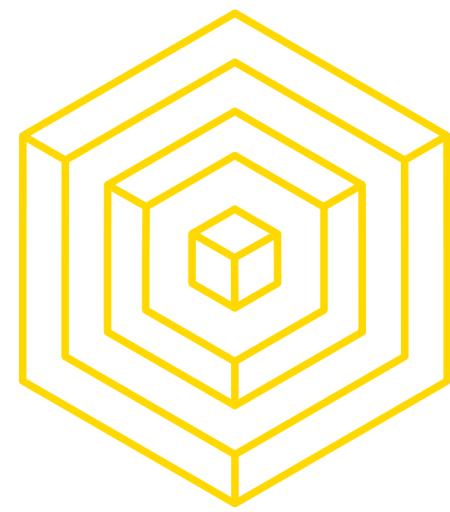
Quorum:

- any majority of Acceptors
- subset of Acceptors, such that any two Quorums share at least one member (they must overlap)



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



PAXOS

MAIN IDEA

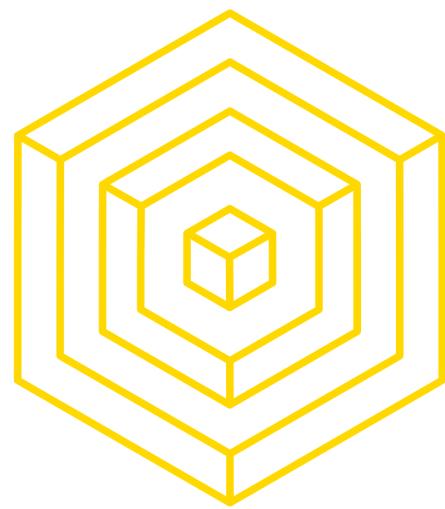
Protocol proceeds over several rounds
where each successful round has 2 phases

1. Client talks to Paxos
2. Within Paxos, Proposer, Acceptor, and Learner discuss
 - a. Pass “ballots”
 - b. Ballot has number and value
3. Learner talks to Client



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



PAXOS IN PRACTICE

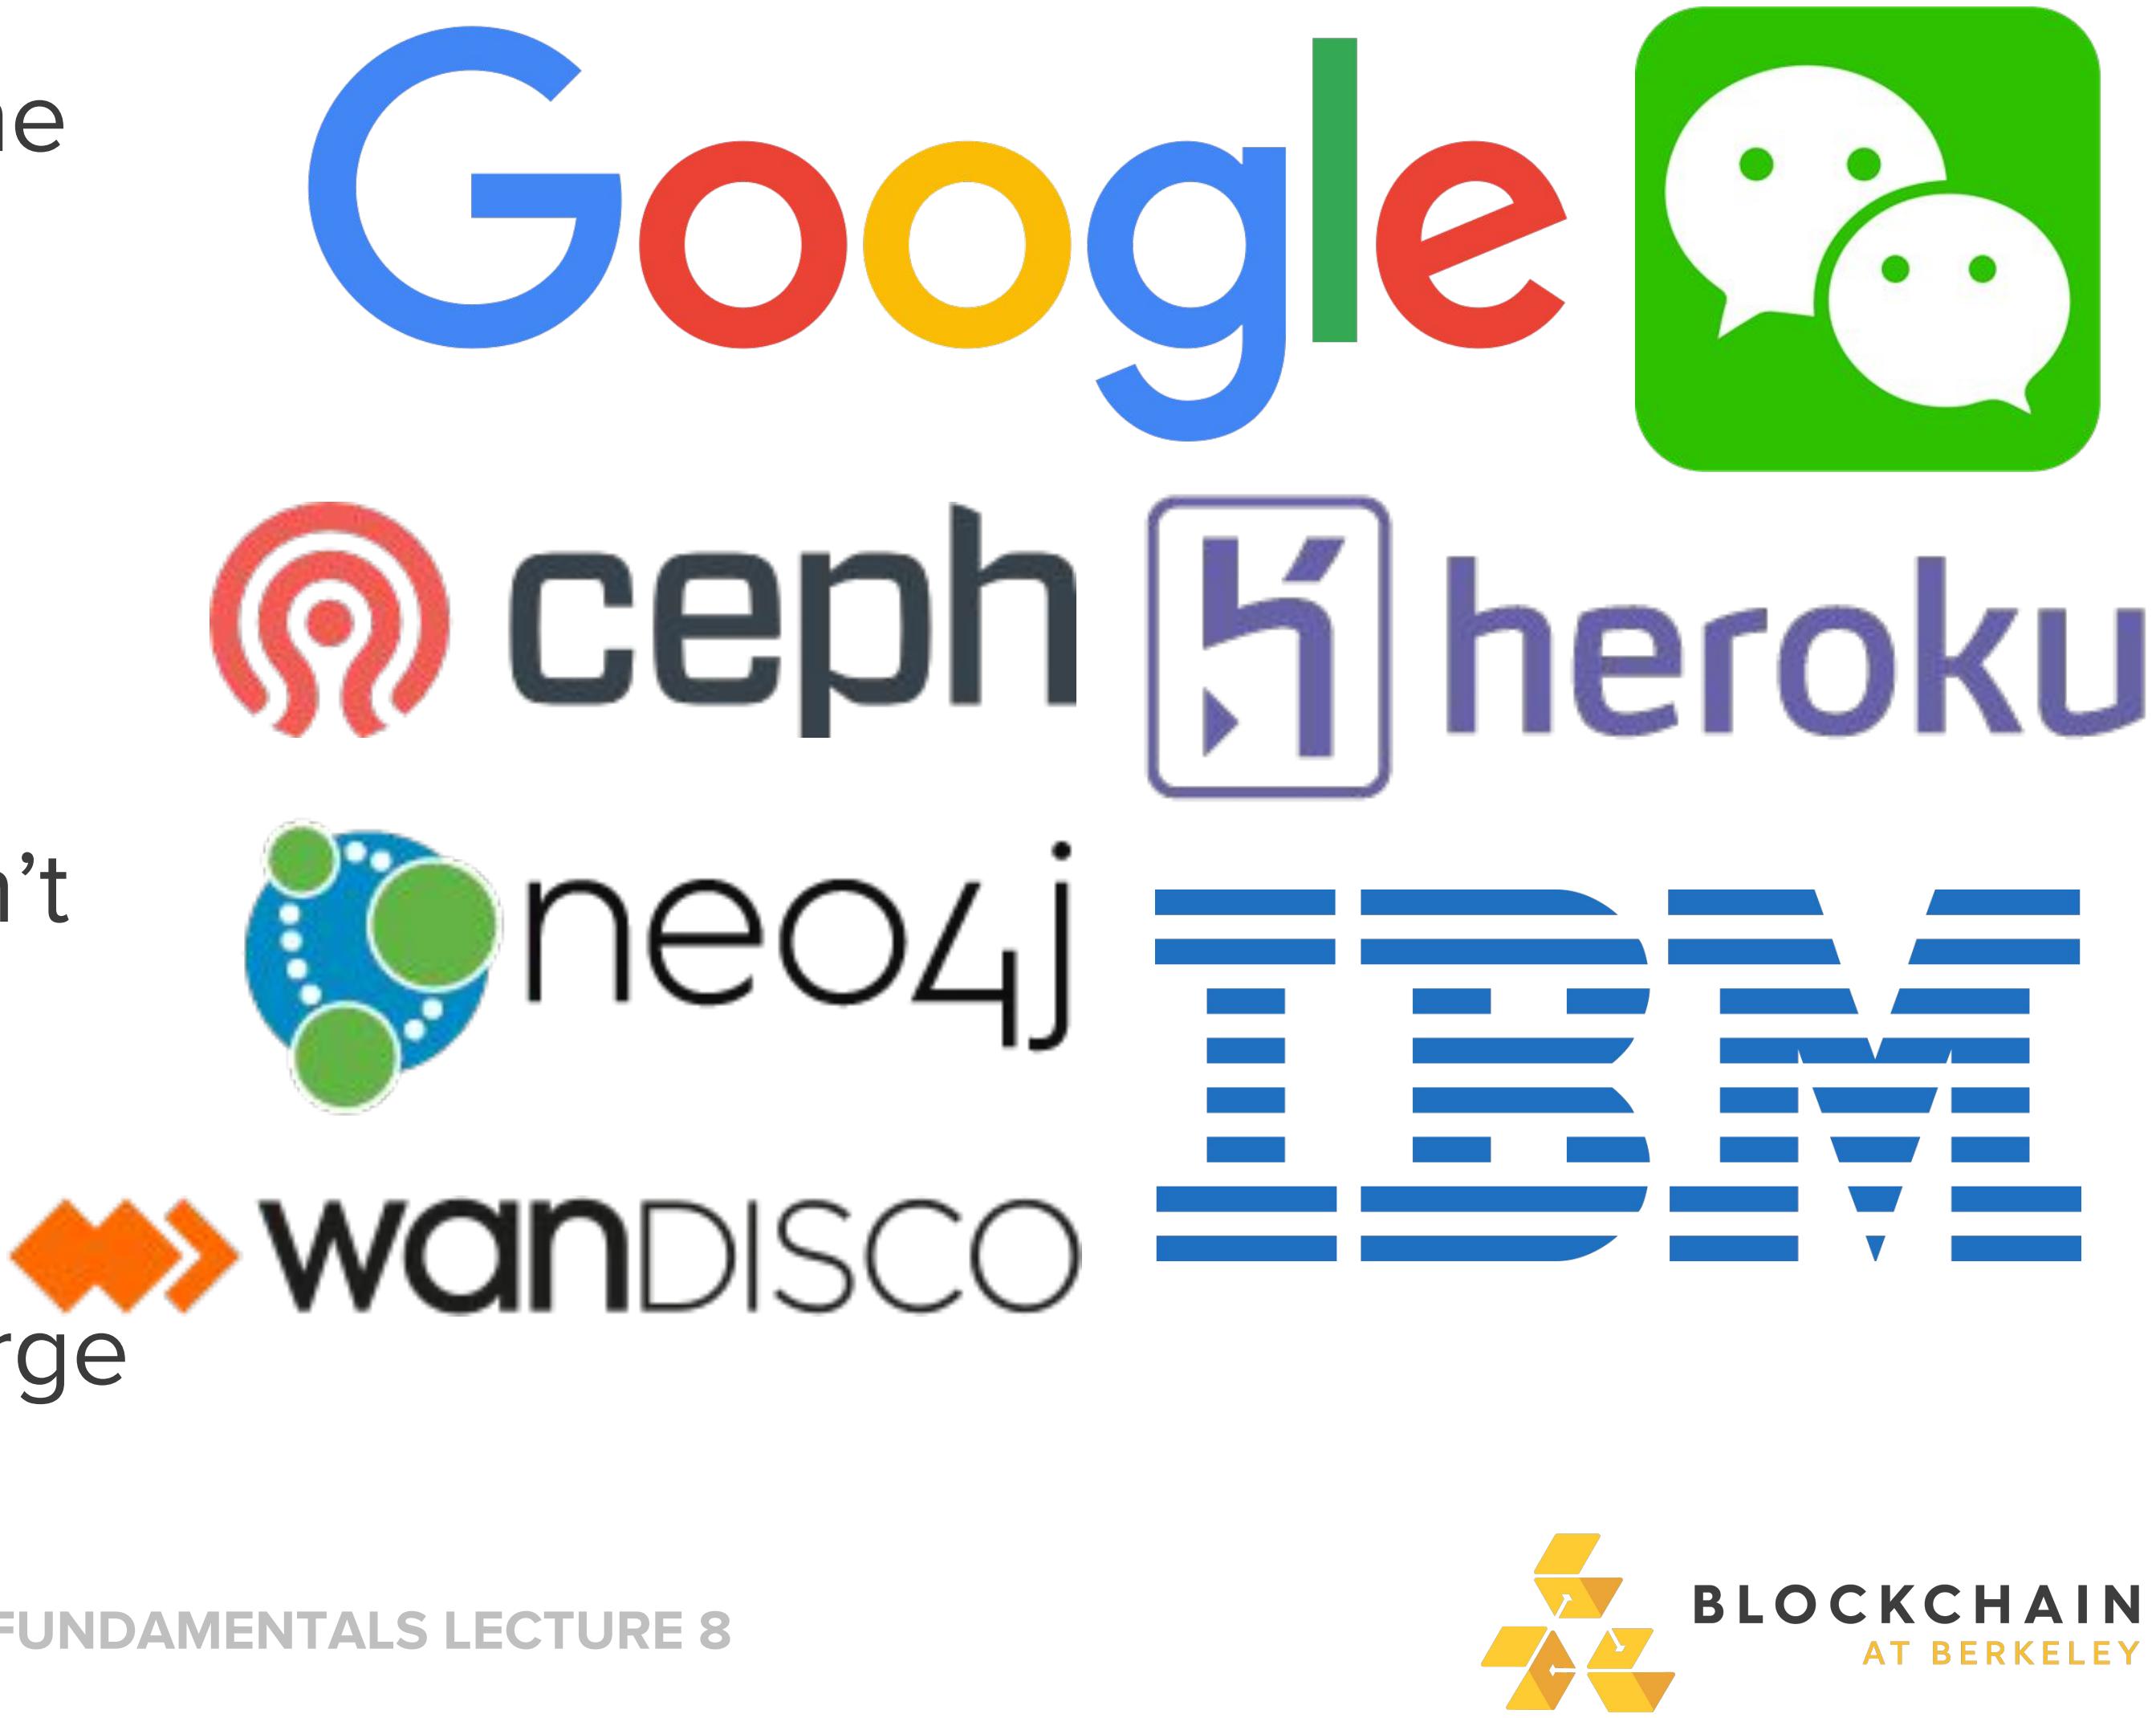
ASSUMPTIONS AND REAL WORLD USE

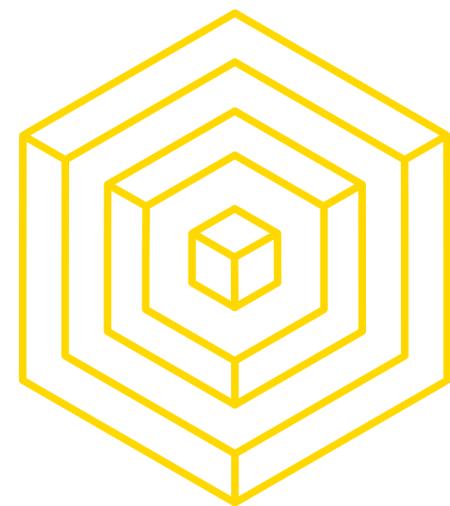
- Nodes do not try to subvert the protocol (messages delivered without corruption)
- Only works for fail-stop (no Byzantine failures) faults
- Failed nodes at any point won't affect the entire network
- Good performance (fast)
- Generally used to replicate large sets of data



AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8





RAFT

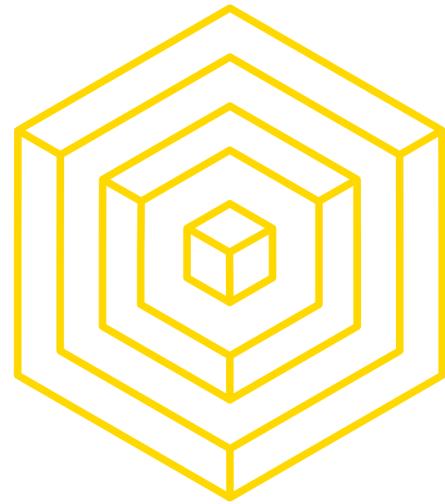
OVERVIEW

Raft is another consensus algorithm designed to be an alternative to Paxos

- Designed to be more understandable than Paxos
- Leader-based approach
- Easier to implement
- JP Morgan's Quorum: Raft-based consensus



AUTHOR: BRIAN HO



RAFT

HOW IT WORKS

A Raft cluster has one and only one elected **leader**

- Communicates with client directly
- Responsible for managing log replication on the other servers of the cluster
- Leads until it fails or disconnects, in which case a new leader is elected



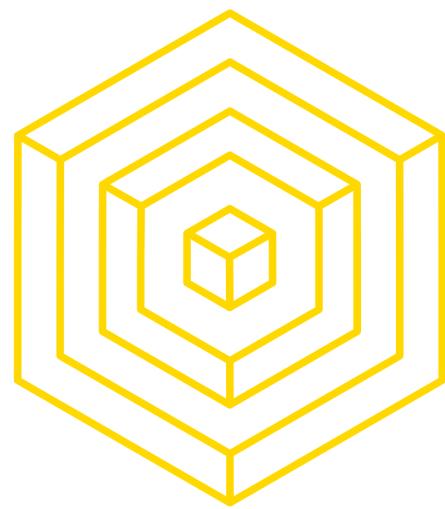
AUTHOR: BRIAN HO

Leader Election

1. Leader sends “heartbeats” to other nodes saying that it is online
2. If other nodes no longer receive “heartbeat,” they start an election cycle (and internal timer)
3. First candidate to timeout becomes new leader

Log Replication

1. Leader accepts client request
2. Leader ensures all other nodes have followed that request



PROOF-OF-STAKE PROBLEM

In Proof-of-Stake, we have the problem of “the rich get richer”

- Large stakeholders stake their balances
- Smaller stakeholders see fewer rewards and lower returns
- Centralizes staking rewards

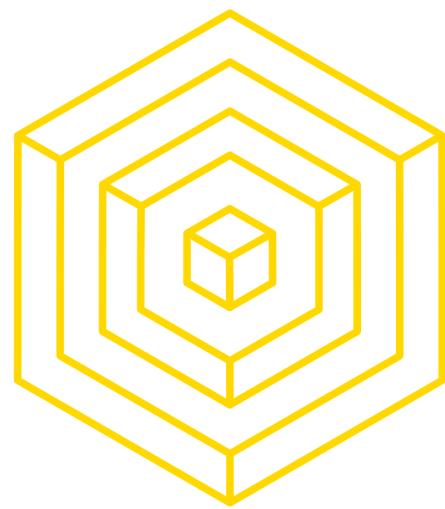


SOURCE:

<https://cointelegraph.com/news/the-inevitable-failure-of-proof-of-stake-blockchains-and-why-a-new-algorithm-is-needed>



AUTHOR: RUSTIE LIN



DELEGATED PROOF-OF-STAKE SOLUTION

Idea: Appoint **delegates** to secure the network on your behalf

- Everyone has a proportional voice in security of the network
- Delegates vote on “witnesses” to produce blocks
- If they you don’t like their behavior, “witnesses” can get voted out

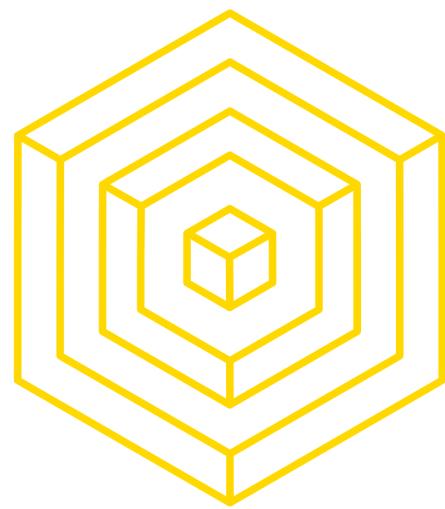


SOURCE:

<https://cointelegraph.com/news/the-inevitable-failure-of-proof-of-stake-blockchains-and-why-a-new-algorithm-is-needed>



AUTHOR: RUSTIE LIN



BITSHARES

OVERVIEW

Bitshares is a decentralized asset exchange started by Daniel Larimer in 2014

- First DPOS system
- Assets backed by core currency (BTS)

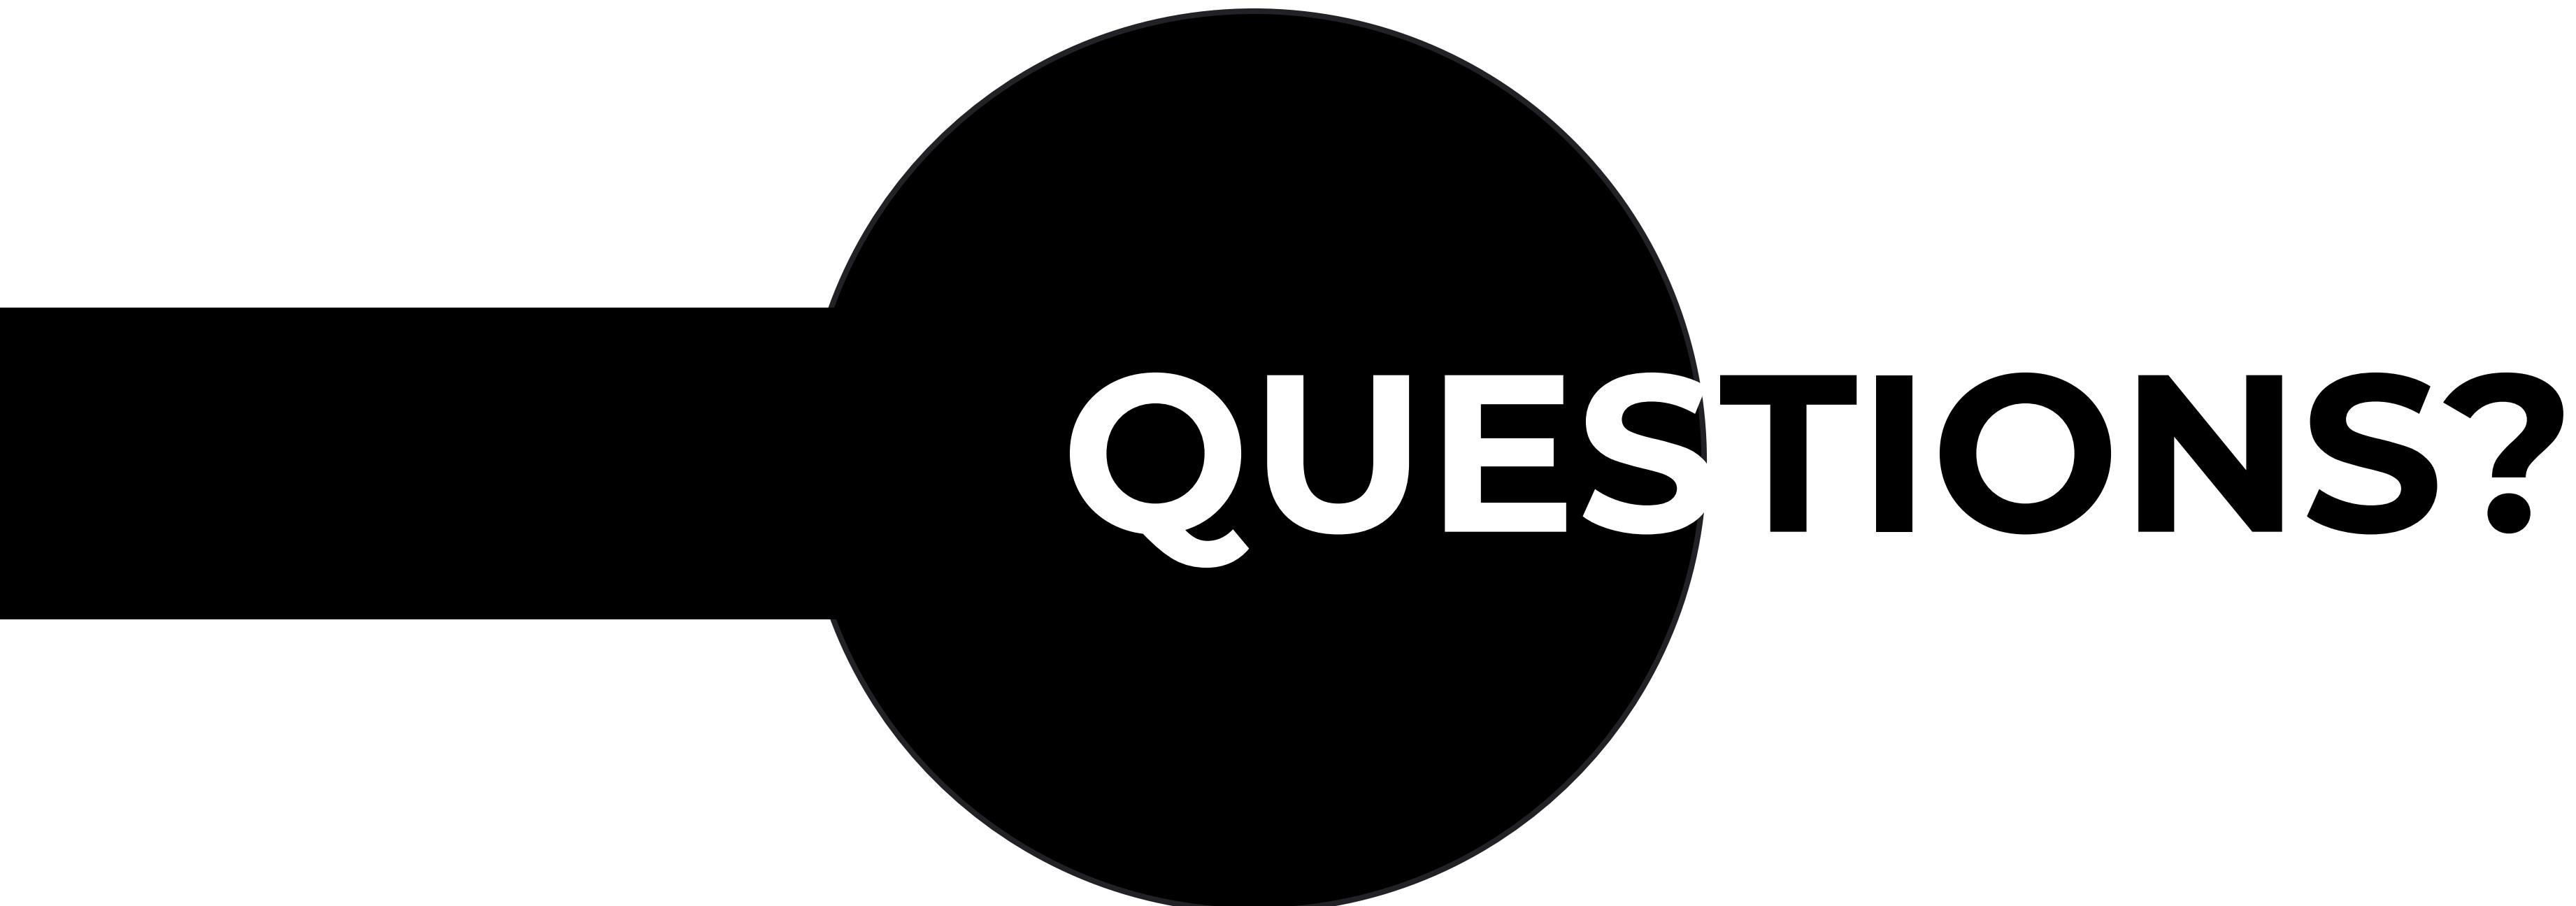


bitshares

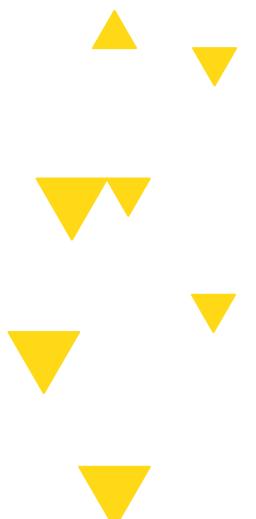


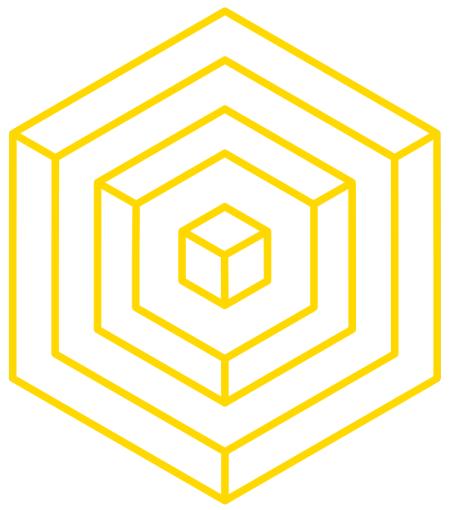
AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



QUESTIONS?

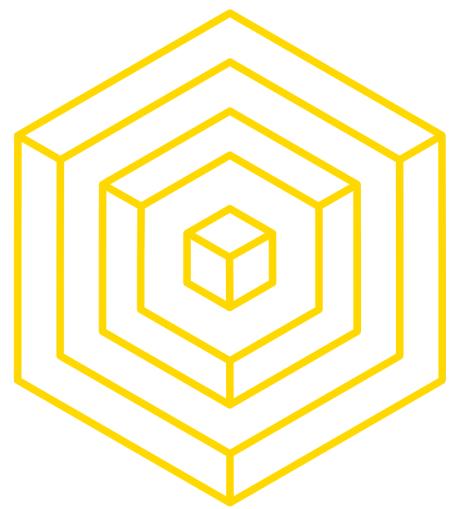




5

FEDERATED CONSENSUS

BLOCKCHAIN FUNDAMENTALS LECTURE 8



BYZANTINE AGREEMENT

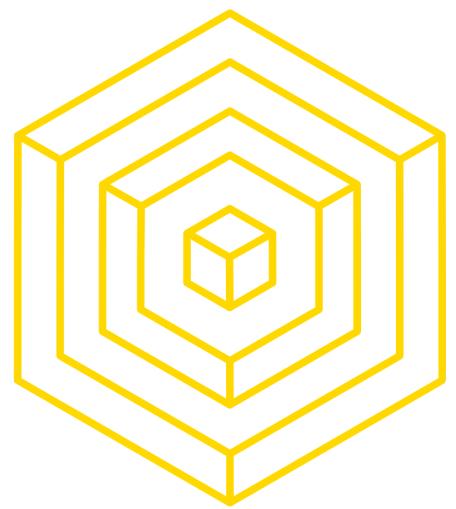
OVERVIEW

In a **distributed system**, a **quorum** is a set of nodes sufficient to reach agreement.



AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8



BYZANTINE AGREEMENT

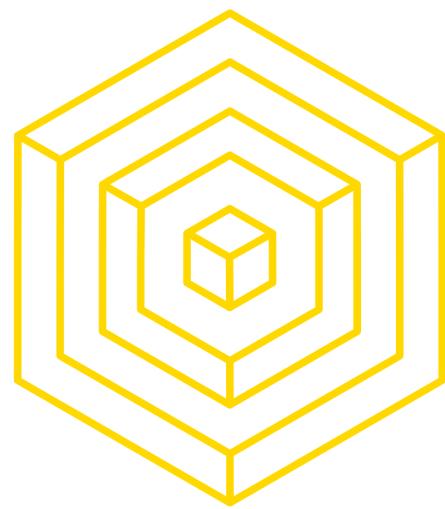
OVERVIEW

What if you don't necessarily trust certain nodes in the quorum? How can we still achieve consensus?



AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8



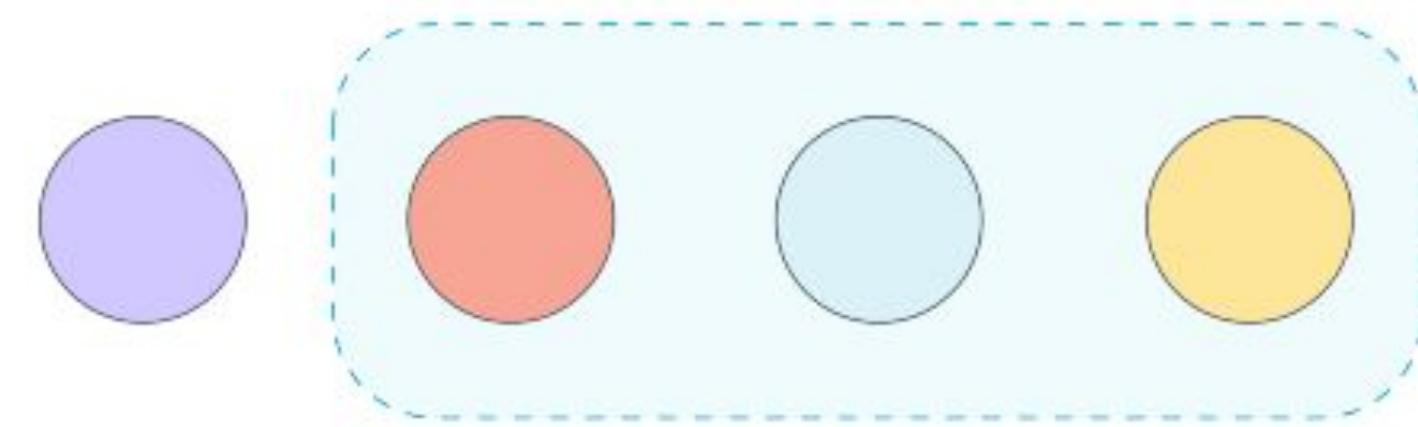
FEDERATED BYZANTINE AGREEMENT

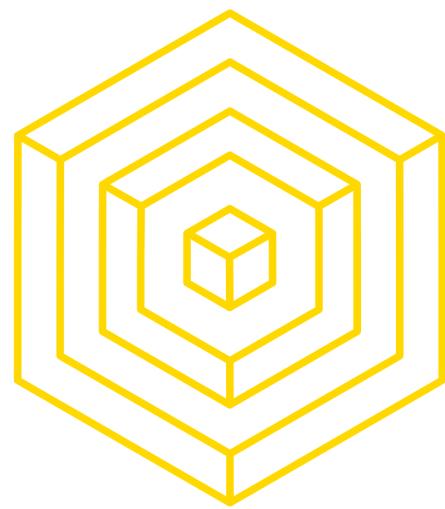
OVERVIEW

Problem: How do we determine **quorums** in a decentralized way?

Solution: introduce **quorum slices**

- Subset of a **quorum** that can convince one particular node of agreement
- Individual nodes decide on other participants they trust for information





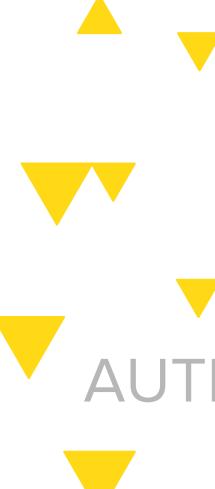
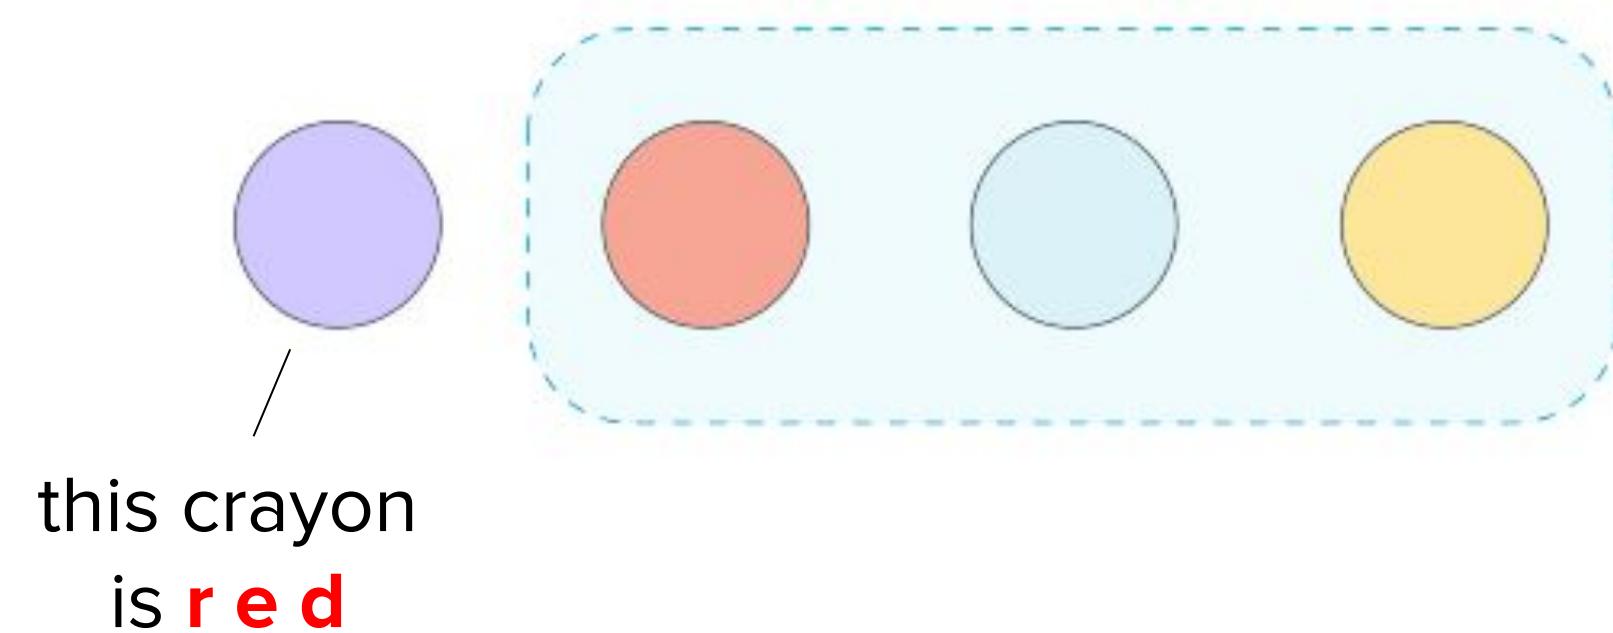
FEDERATED BYZANTINE AGREEMENT

OVERVIEW

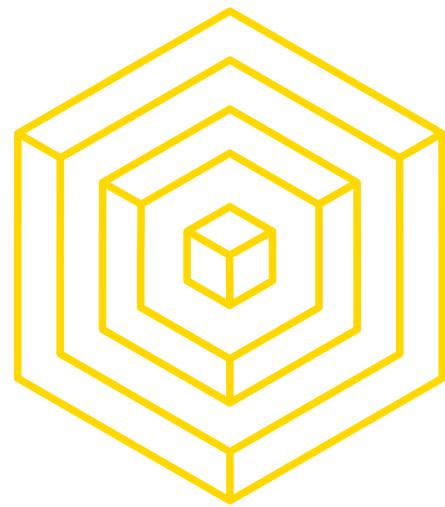
Problem: How do we determine **quorums** in a decentralized way?

Solution: introduce **quorum slices**

- Subset of a **quorum** that can convince one particular node of agreement
- Individual nodes decide on other participants they trust for information



AUTHOR: BRIAN HO



FEDERATED BYZANTINE AGREEMENT

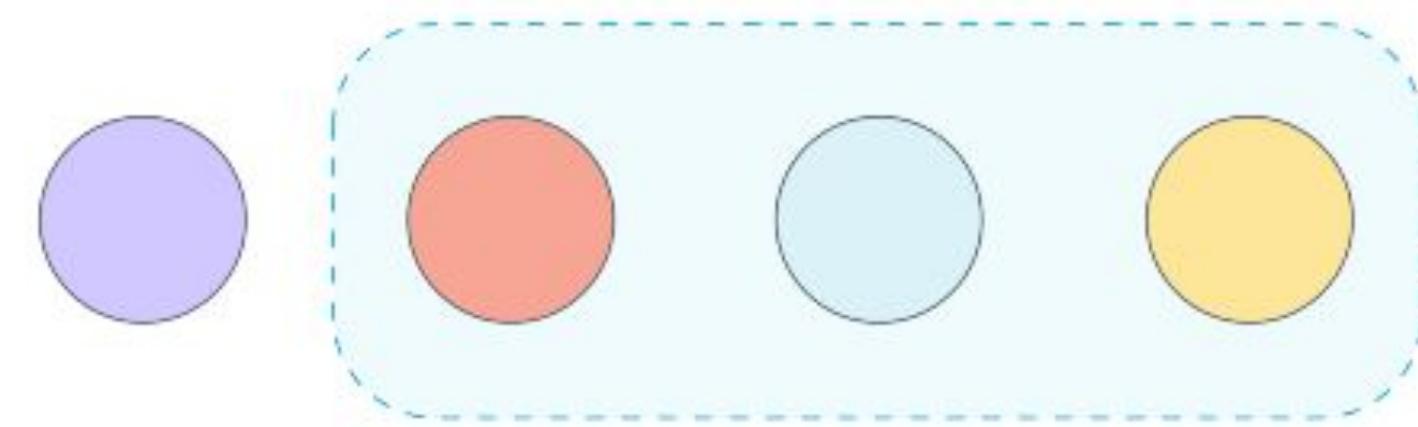
OVERVIEW

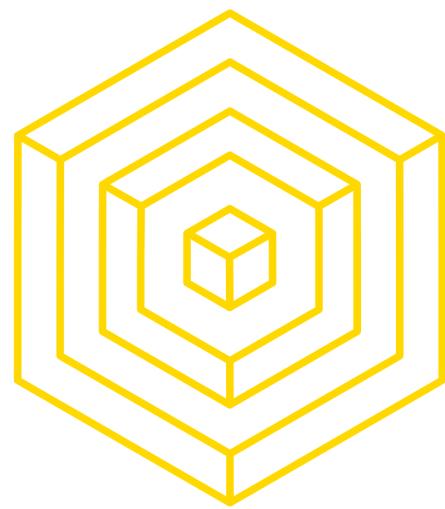
Problem: How do we determine **quorums** in a decentralized way?

yo it's actually
green

Solution: introduce **quorum slices**

- Subset of a **quorum** that can convince one particular node of agreement
- Individual nodes decide on other participants they trust for information





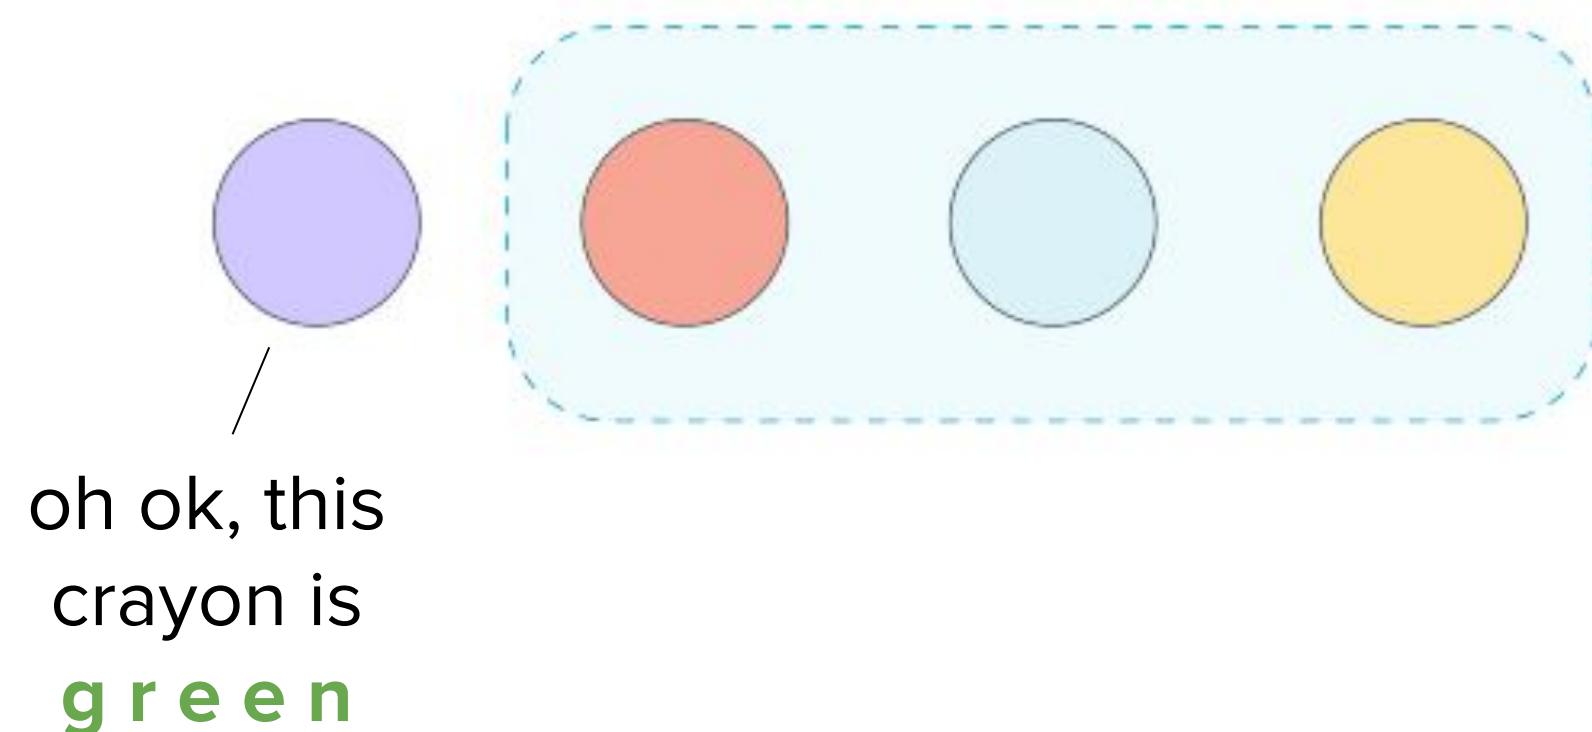
FEDERATED BYZANTINE AGREEMENT

OVERVIEW

Problem: How do we determine **quorums** in a decentralized way?

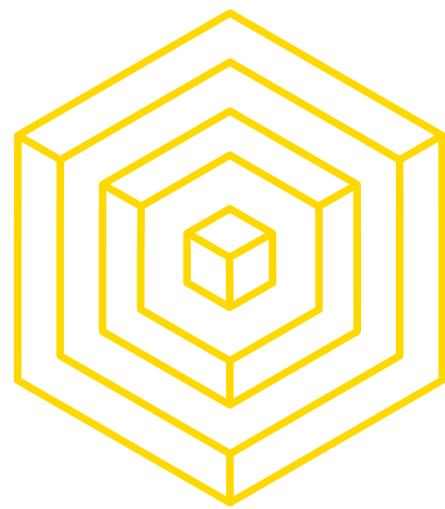
Solution: introduce **quorum slices**

- Subset of a **quorum** that can convince one particular node of agreement
- Individual nodes decide on other participants they trust for information



AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8



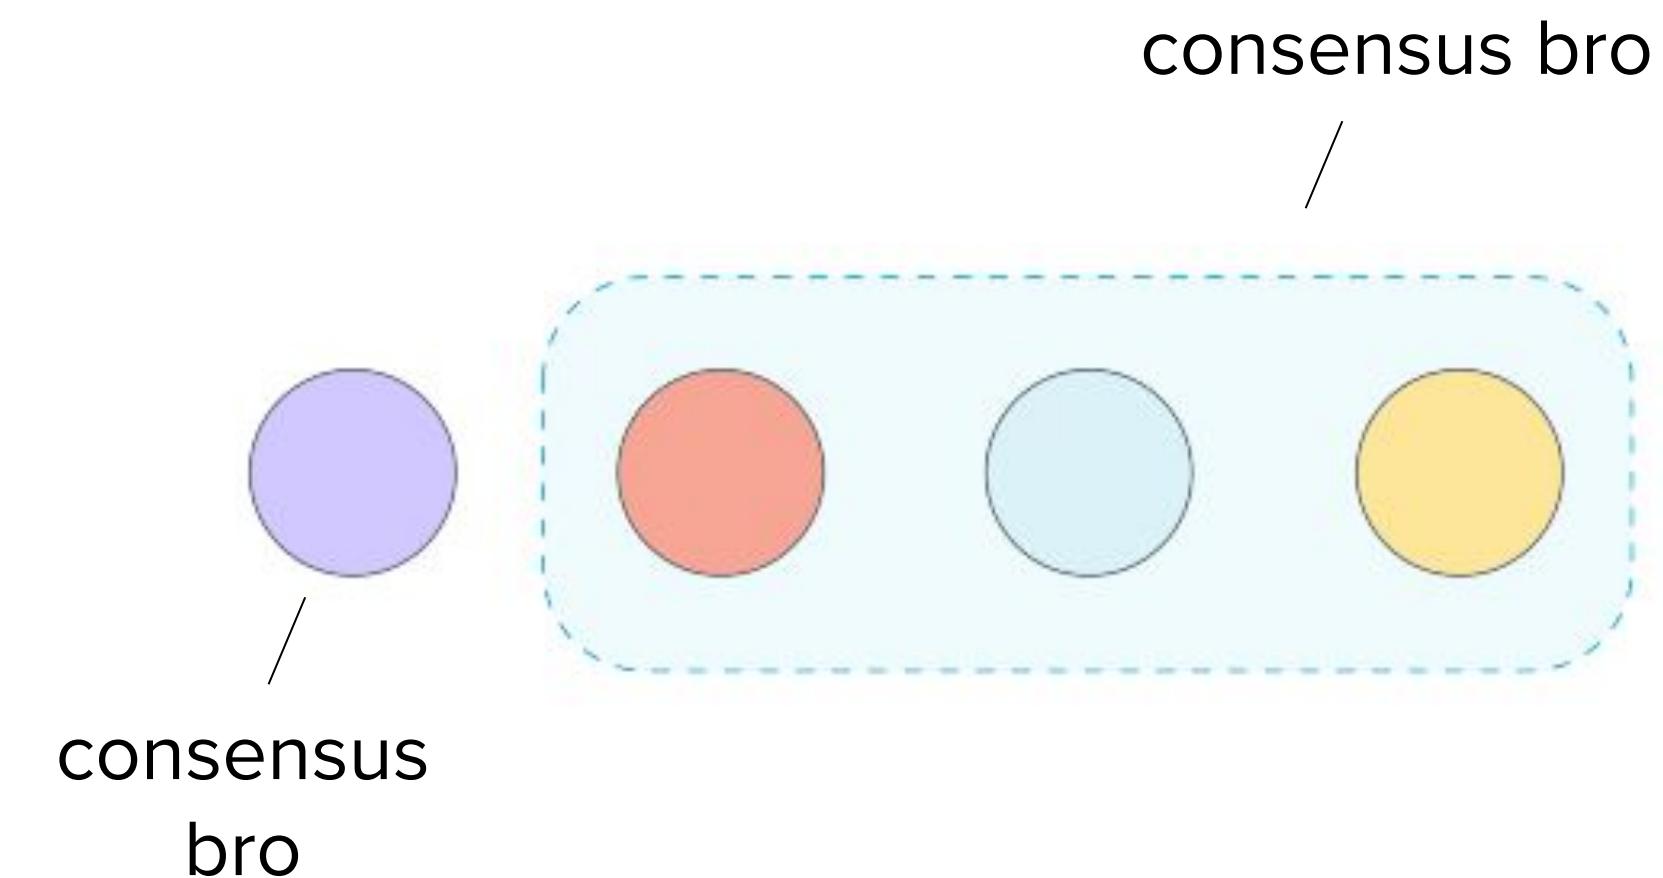
FEDERATED BYZANTINE AGREEMENT

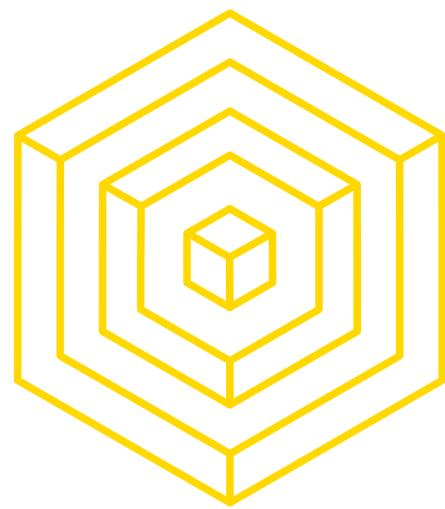
OVERVIEW

Problem: How do we determine **quorums** in a decentralized way?

Solution: introduce **quorum slices**

- Subset of a **quorum** that can convince one particular node of agreement
- Individual nodes decide on other participants they trust for information





FEDERATED BYZANTINE AGREEMENT

OVERVIEW

Idea: What happens when multiple quorum slices join together?

We get a **quorum intersection!**

- Quorum slices that come together will slowly convince other quorum slices and form a “larger” quorum

Otherwise we get **disjoint quorums** that agree on different things

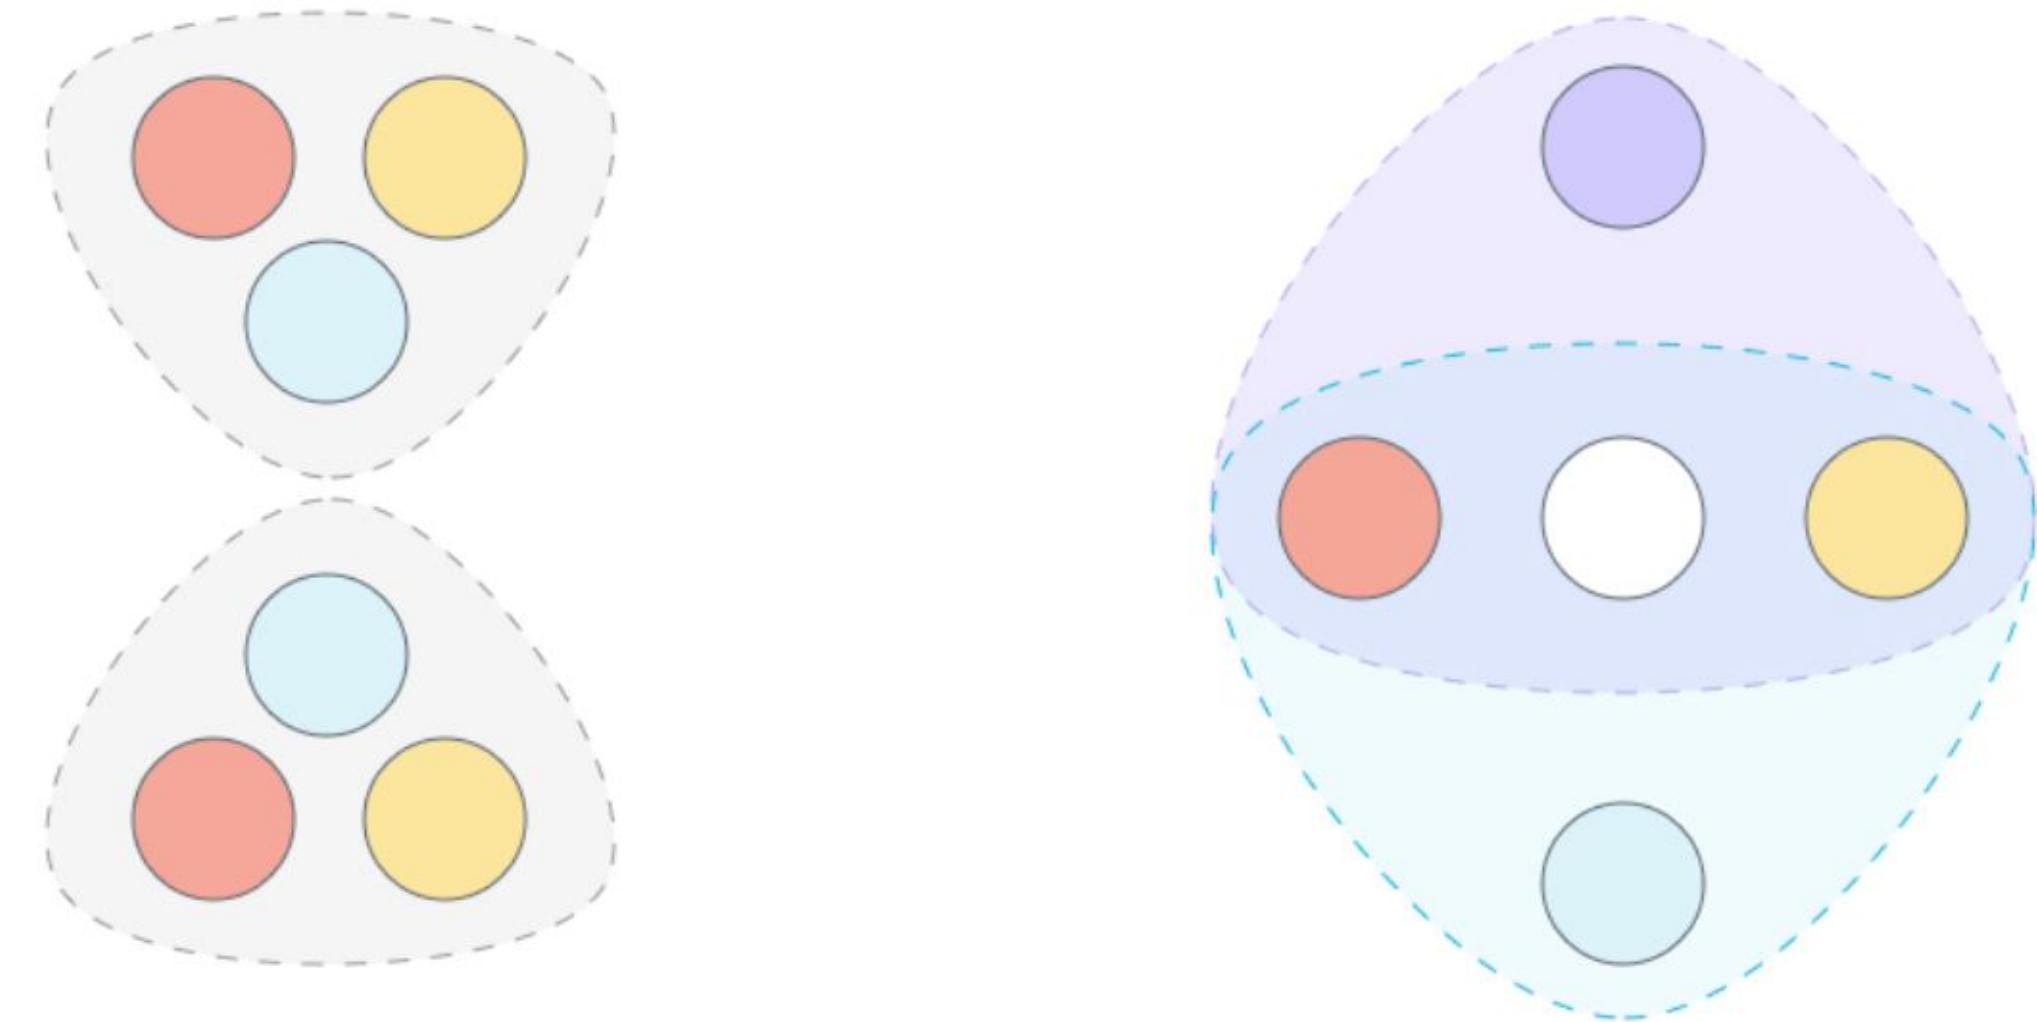
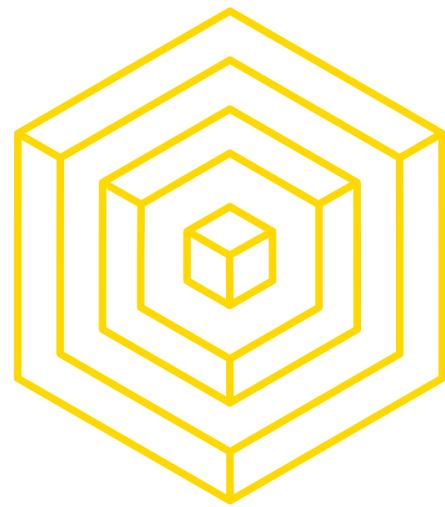


Image source: https://cdn-images-1.medium.com/max/1600/0*msL7MVVEy4p2VzhP

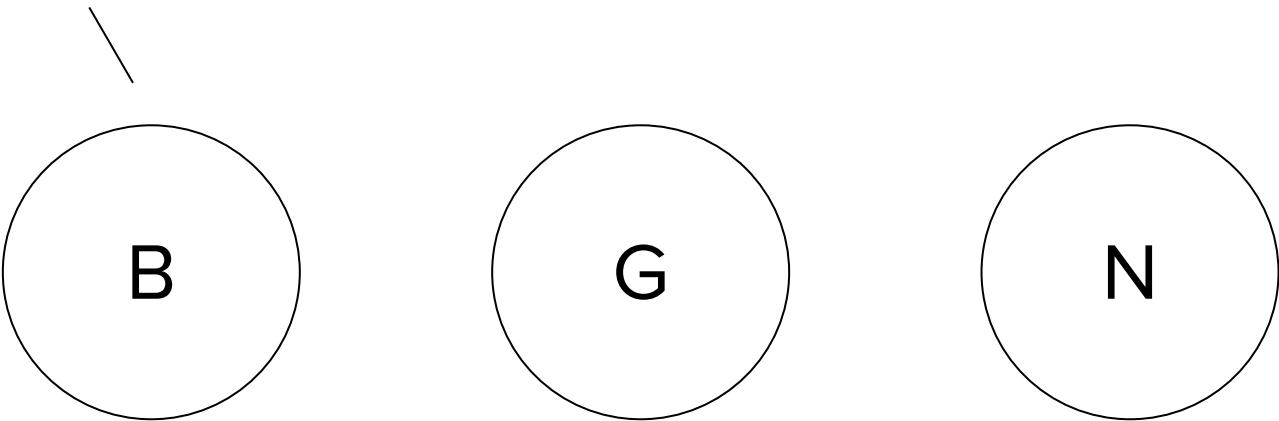


FEDERATED BYZANTINE AGREEMENT

EXAMPLE

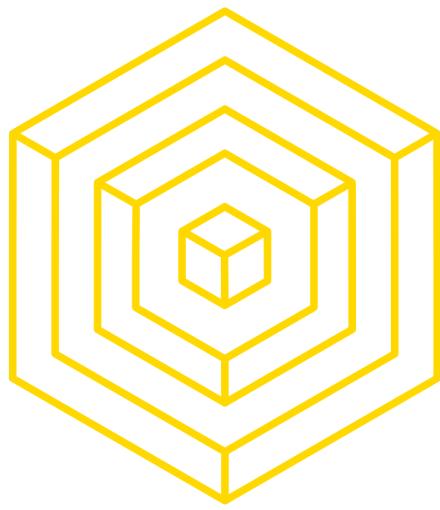
Lunchtime Consensus!

hey, let's get
lunch



AUTHOR: BRIAN HO

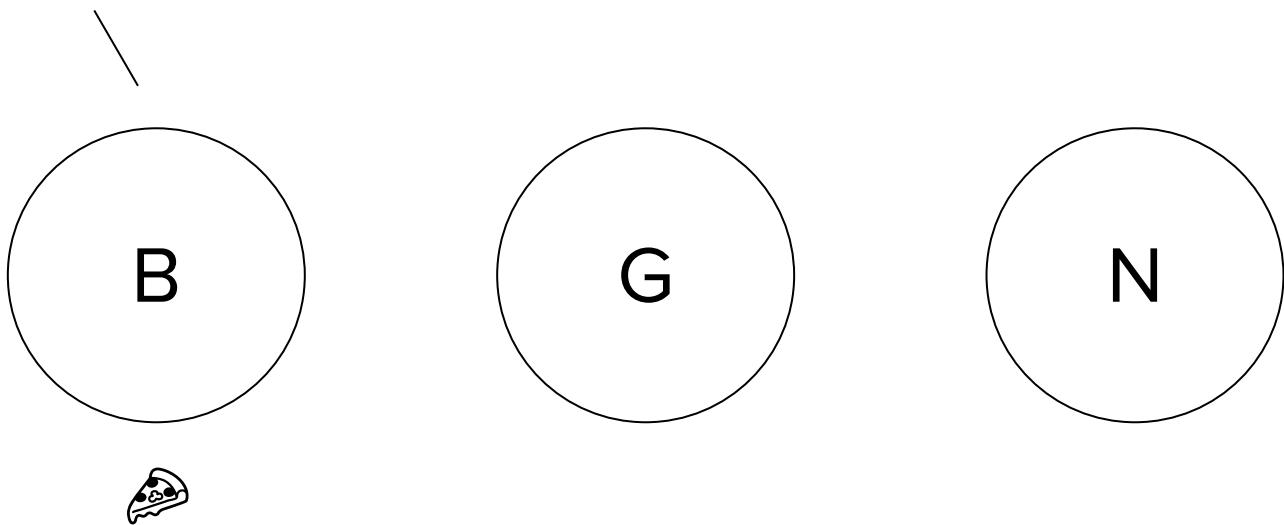
BLOCKCHAIN FUNDAMENTALS LECTURE 8



FEDERATED BYZANTINE AGREEMENT

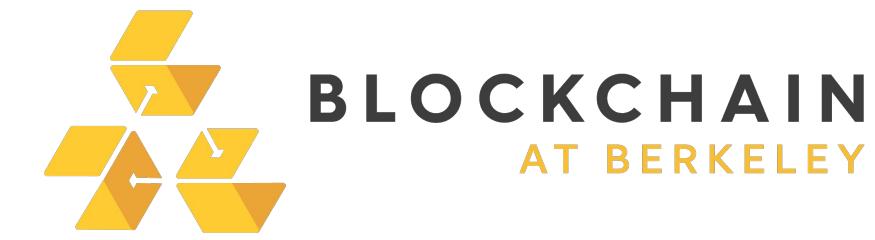
Lunchtime Consensus!

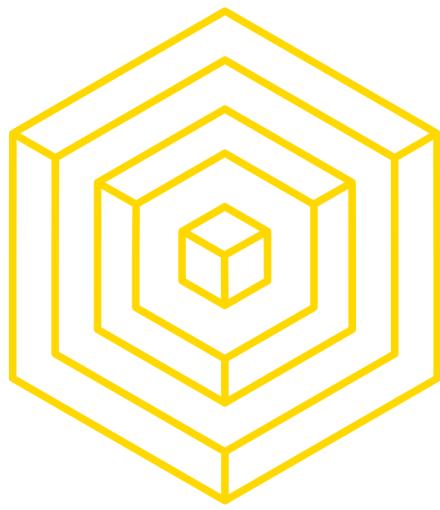
wanna get ?



AUTHOR: BRIAN HO

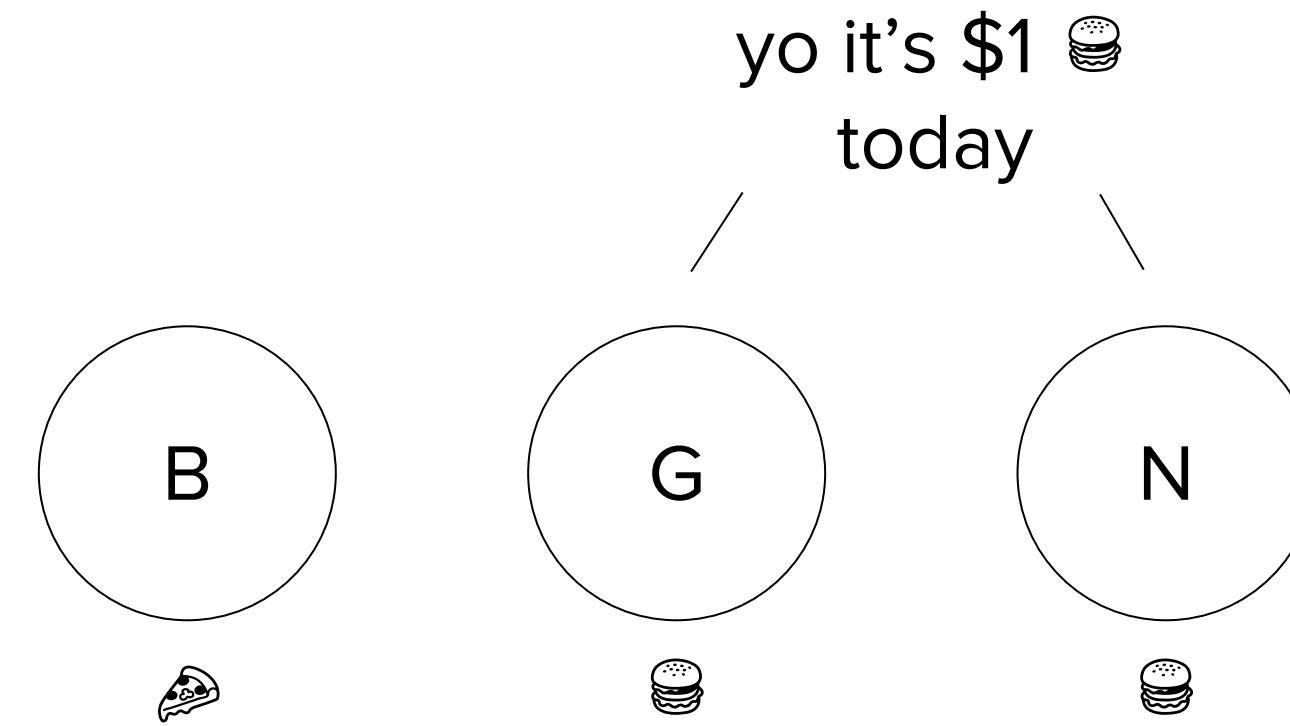
BLOCKCHAIN FUNDAMENTALS LECTURE 8





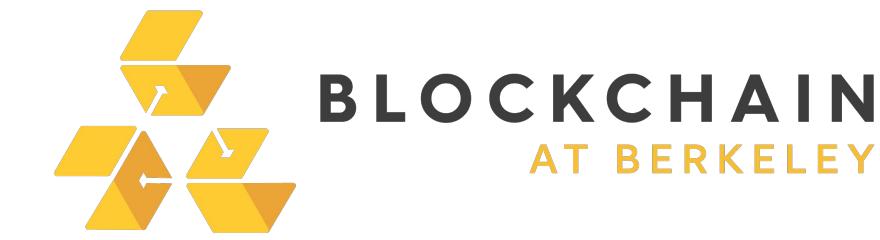
FEDERATED BYZANTINE AGREEMENT

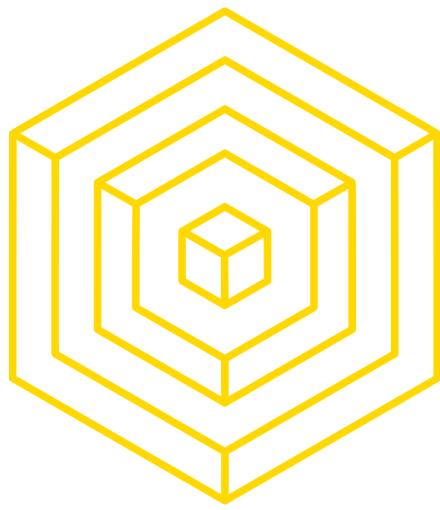
Lunchtime Consensus!



AUTHOR: BRIAN HO

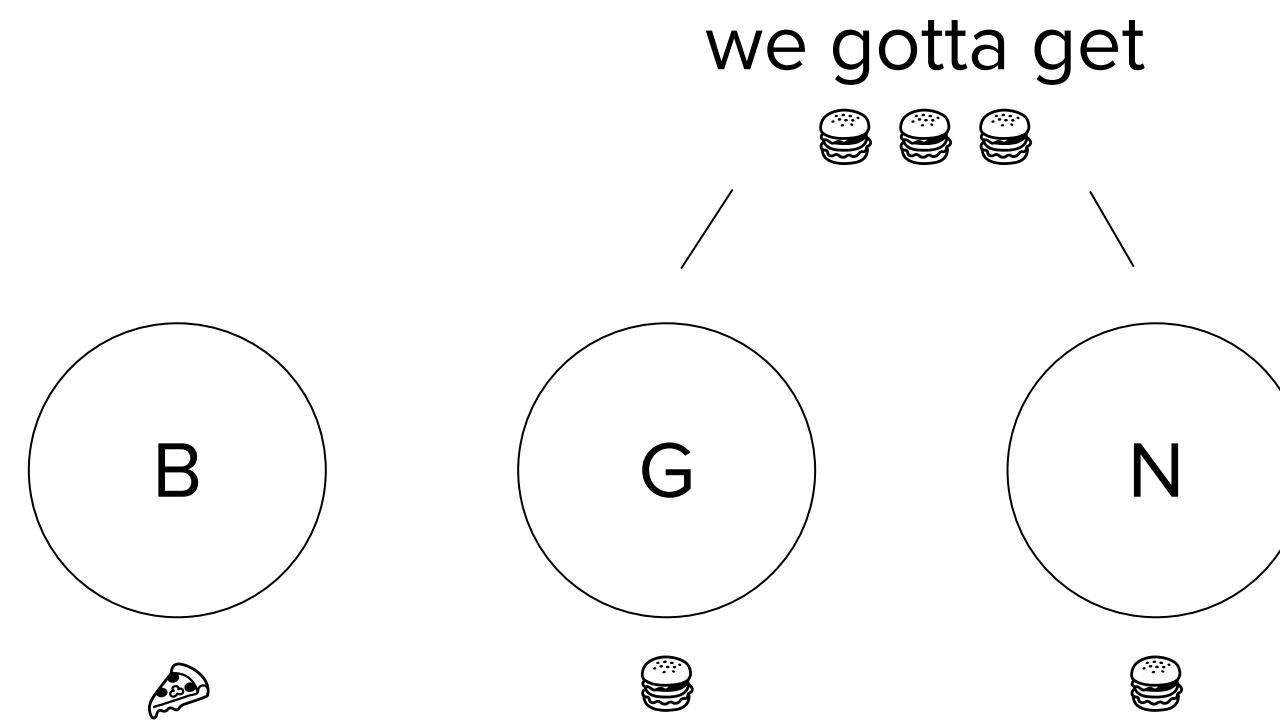
BLOCKCHAIN FUNDAMENTALS LECTURE 8



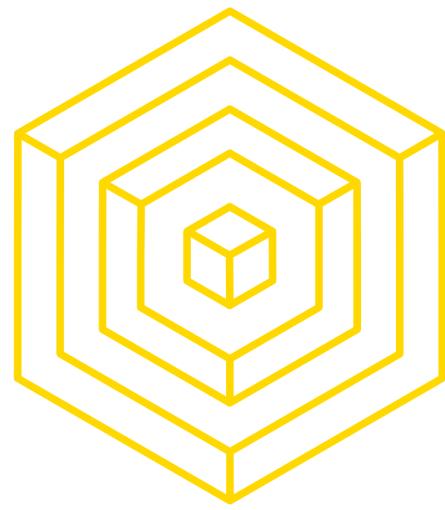


FEDERATED BYZANTINE AGREEMENT

Lunchtime Consensus!



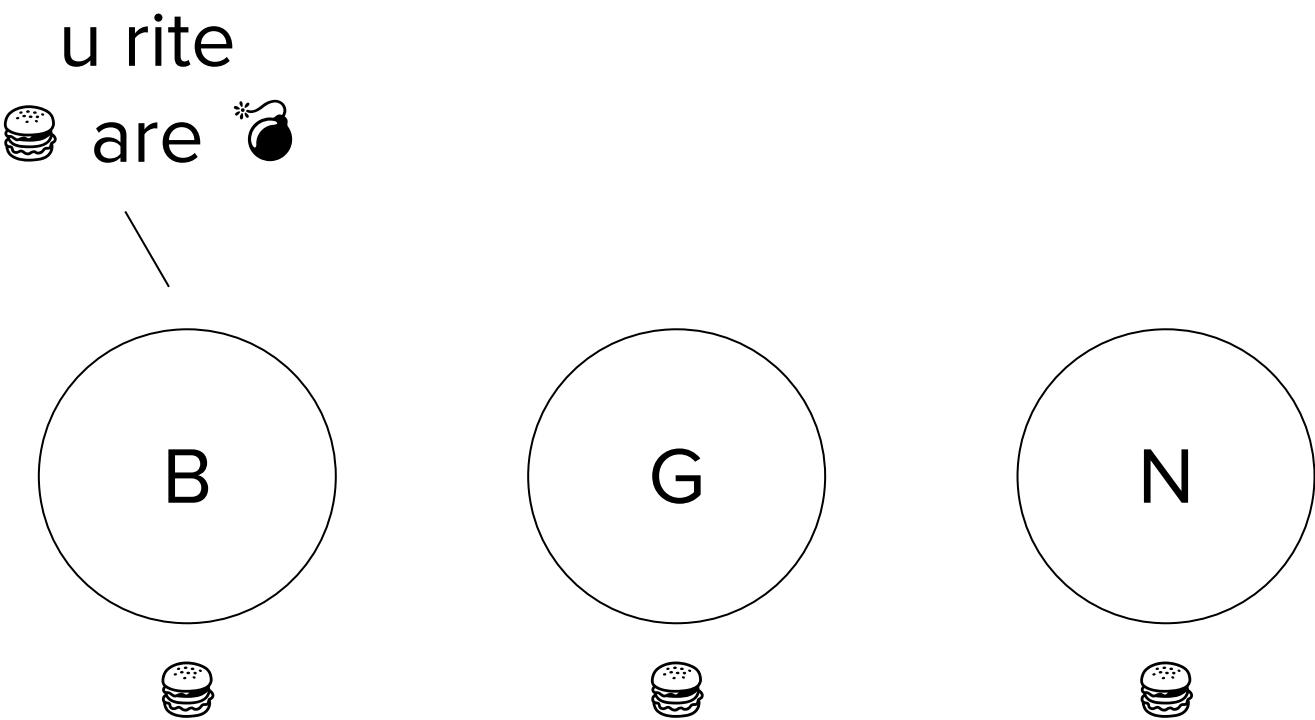
AUTHOR: BRIAN HO



FEDERATED BYZANTINE AGREEMENT

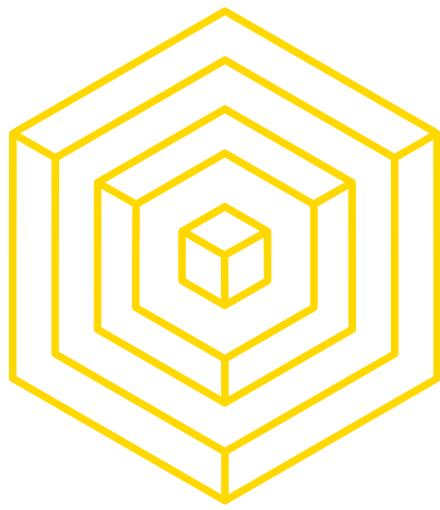
EXAMPLE

Lunchtime Consensus!



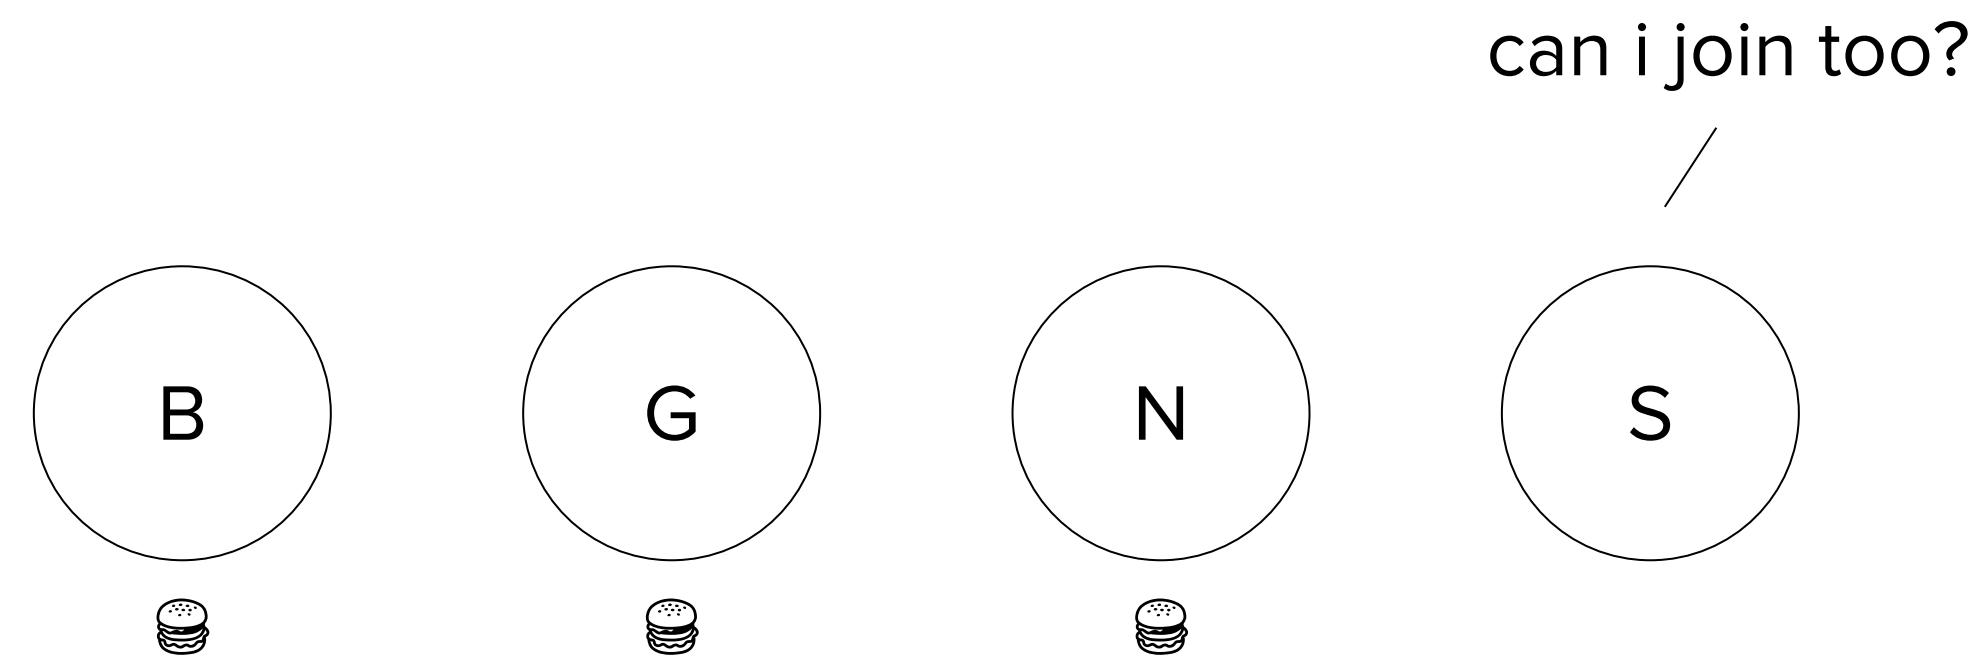
AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8

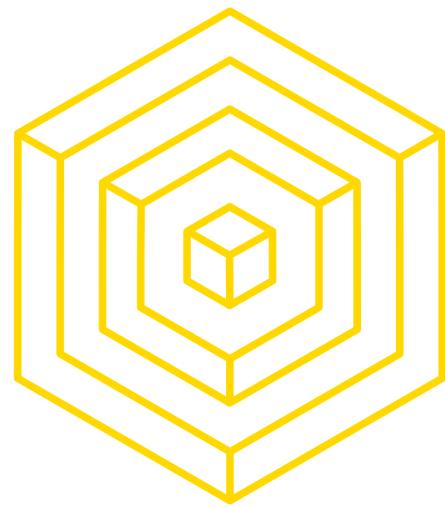


FEDERATED BYZANTINE AGREEMENT

Lunchtime Consensus!



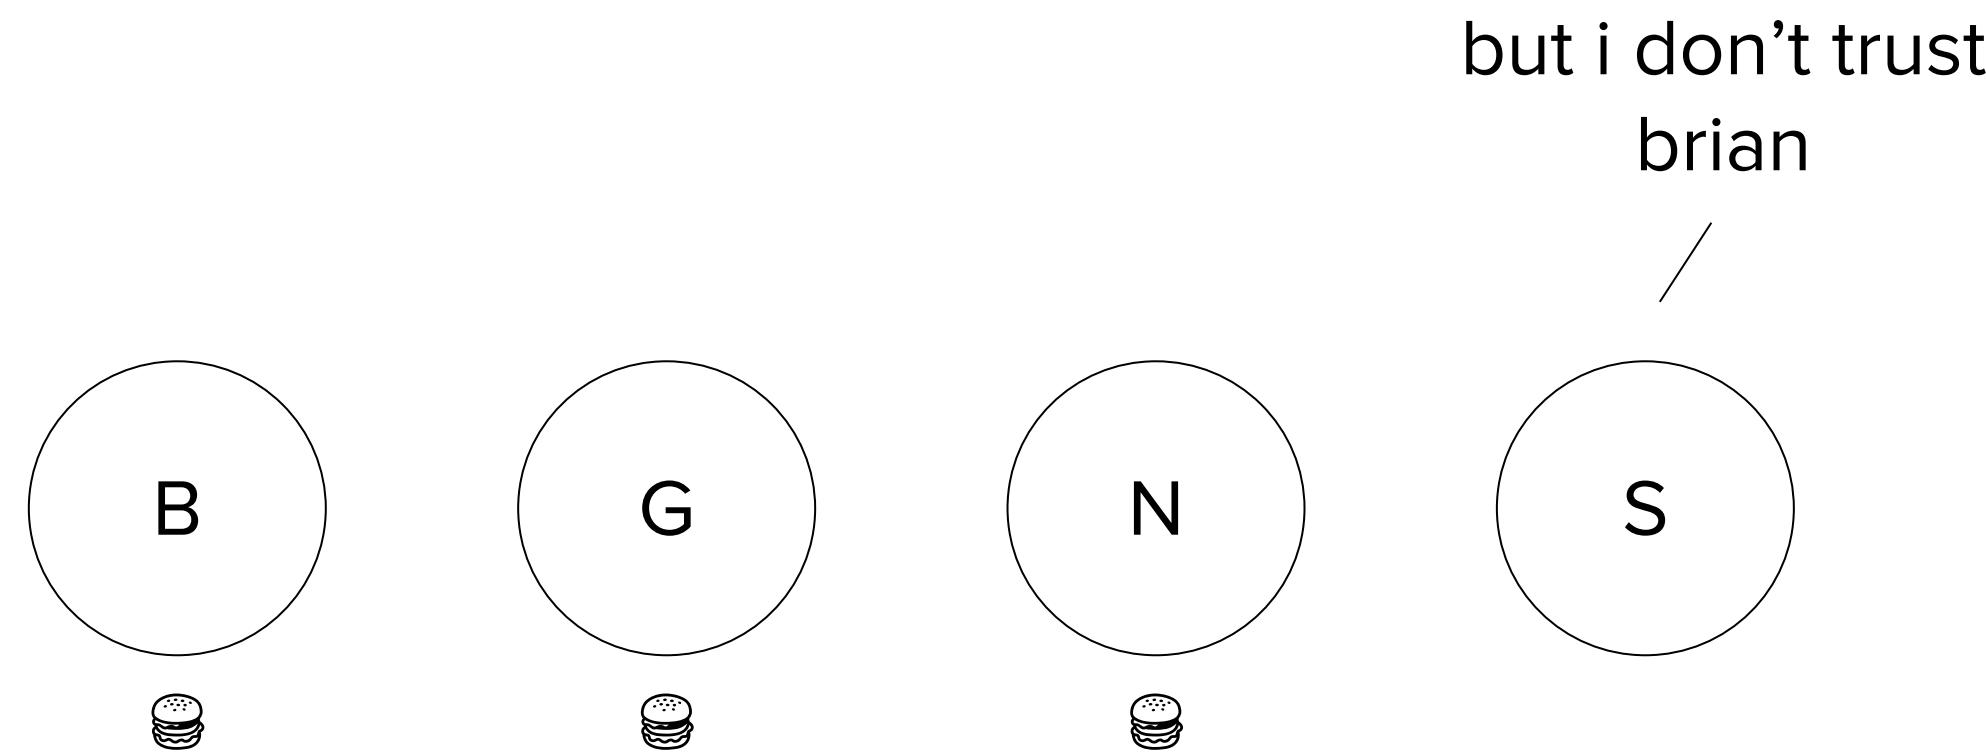
AUTHOR: BRIAN HO



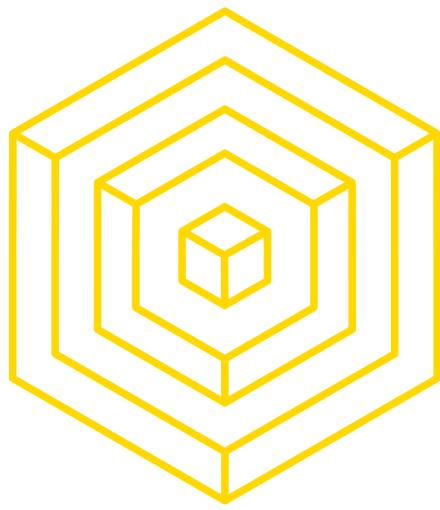
FEDERATED BYZANTINE AGREEMENT

EXAMPLE

Lunchtime Consensus!

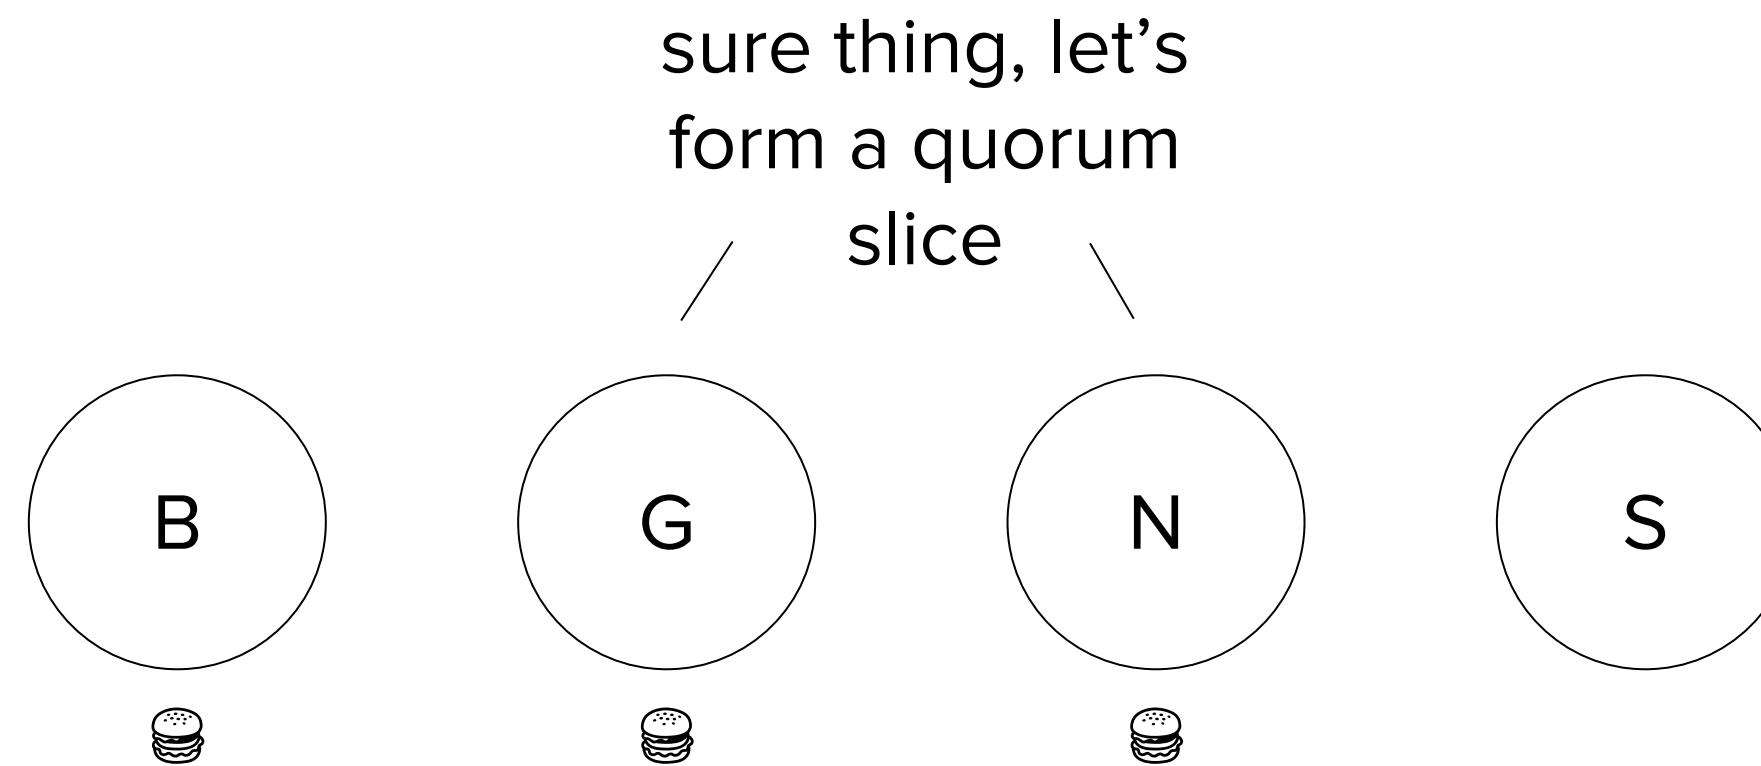


AUTHOR: BRIAN HO
▼ ▲ ▼ ▲ ▼ ▲

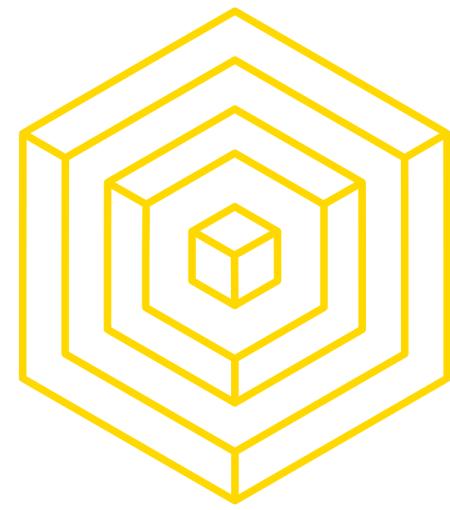


FEDERATED BYZANTINE AGREEMENT

Lunchtime Consensus!



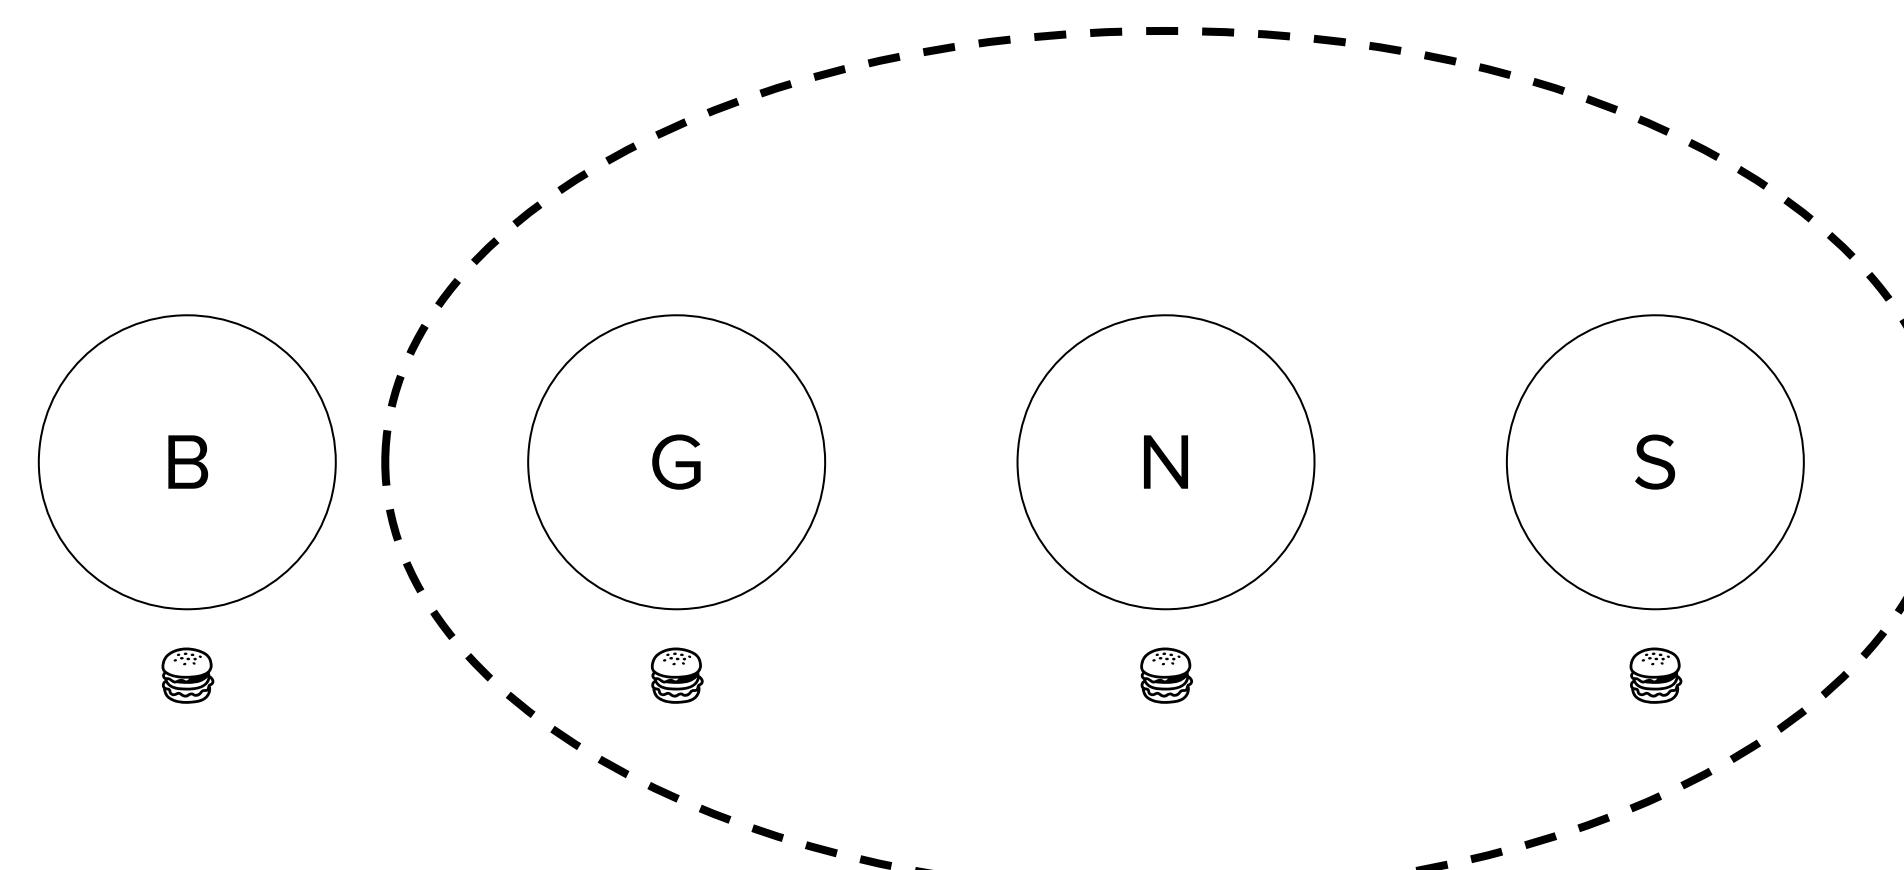
AUTHOR: BRIAN HO



FEDERATED BYZANTINE AGREEMENT

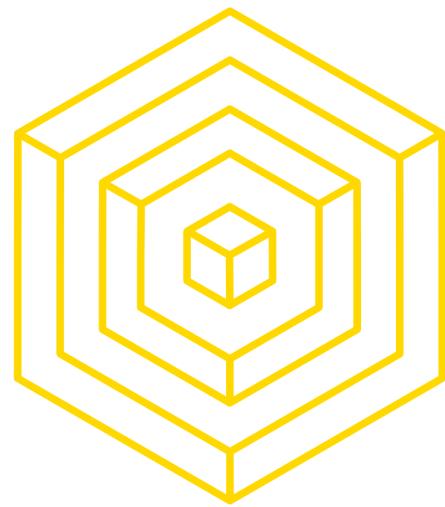
EXAMPLE

Lunchtime Consensus!



AUTHOR: BRIAN HO

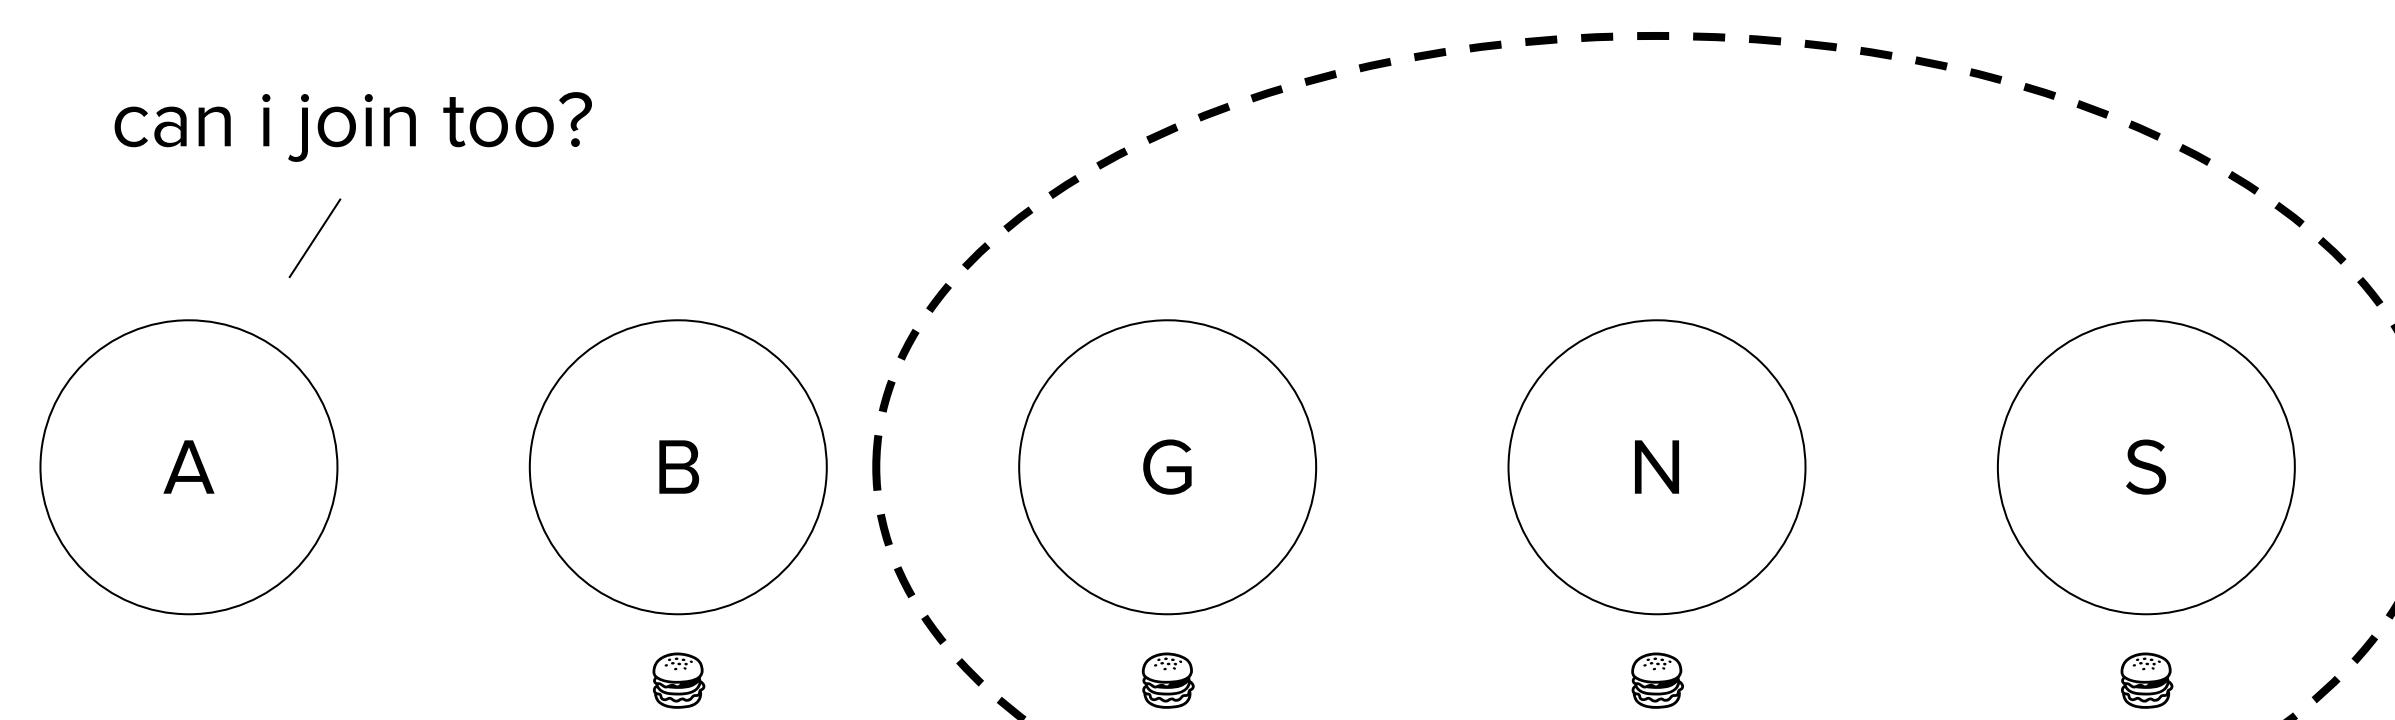
BLOCKCHAIN FUNDAMENTALS LECTURE 8



FEDERATED BYZANTINE AGREEMENT

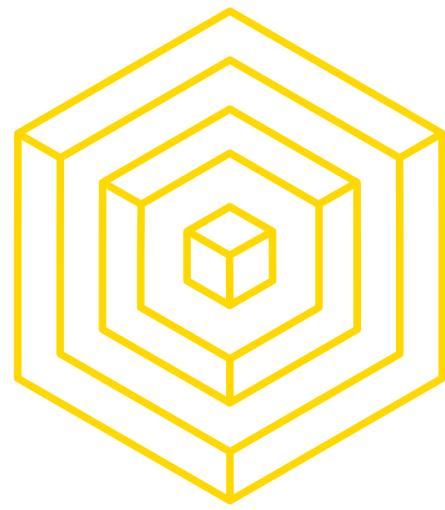
EXAMPLE

Lunchtime Consensus!



◀ ▶
▼ ▲
▼ AUTHOR: BRIAN HO
▼

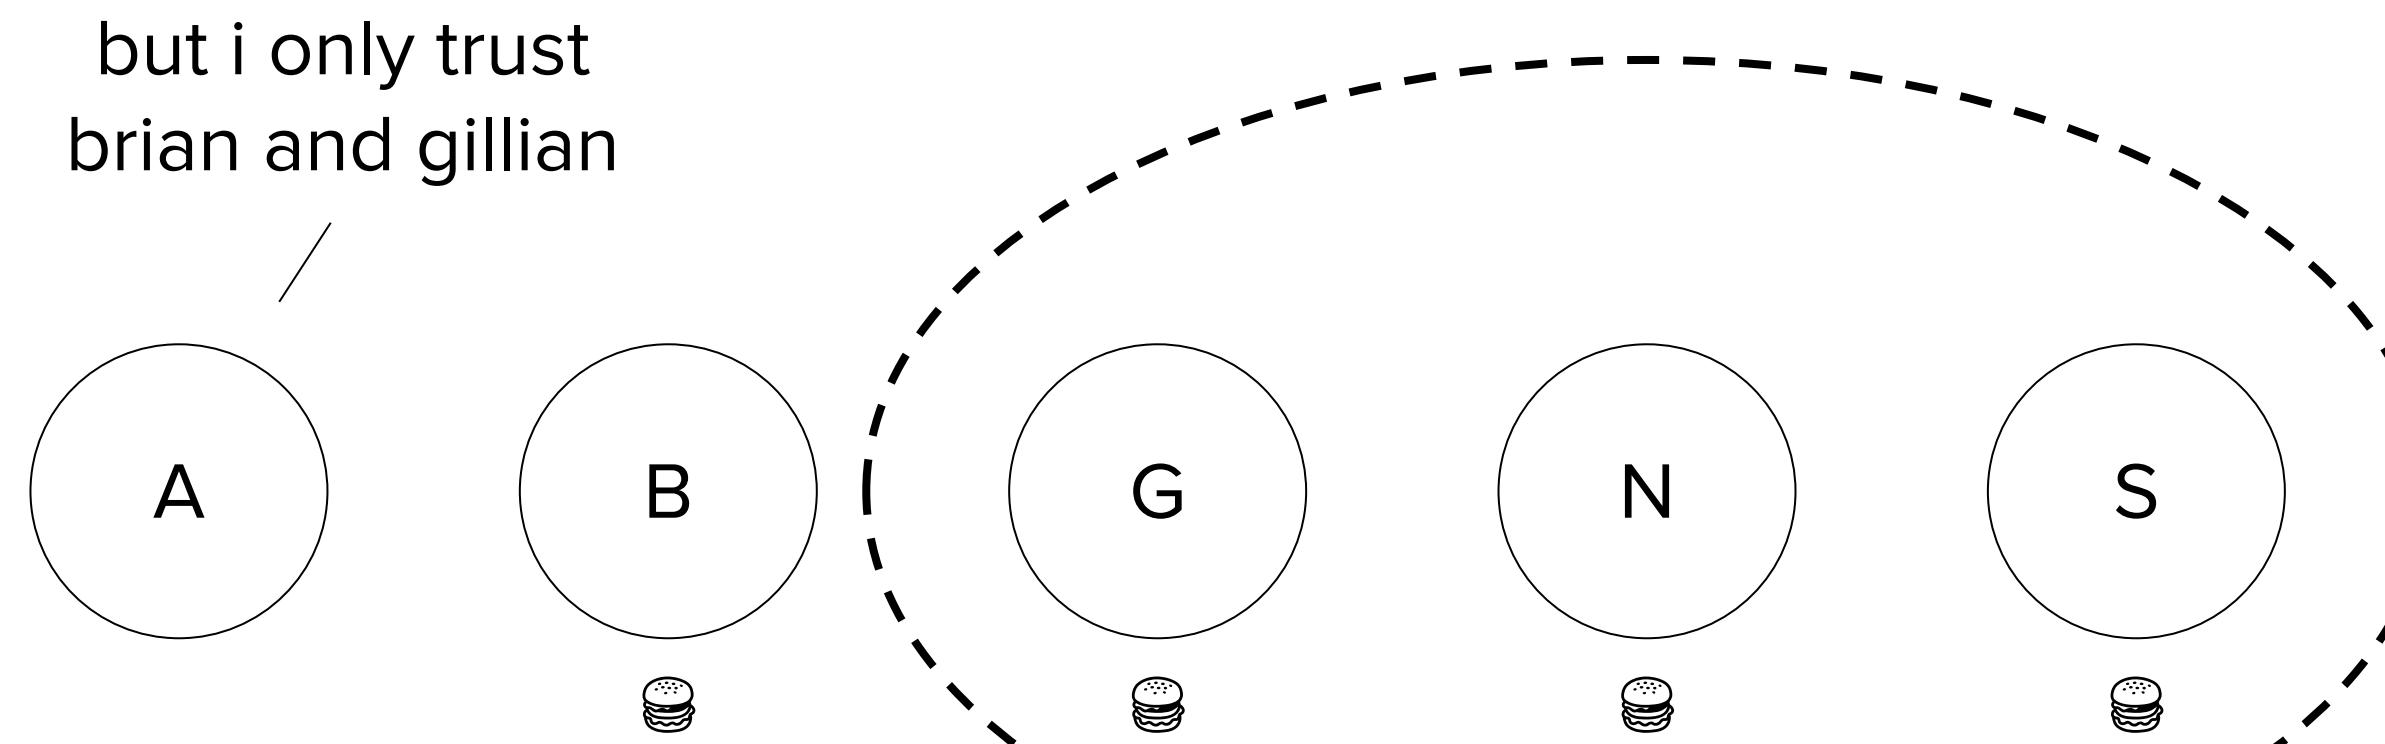
BLOCKCHAIN FUNDAMENTALS LECTURE 8



FEDERATED BYZANTINE AGREEMENT

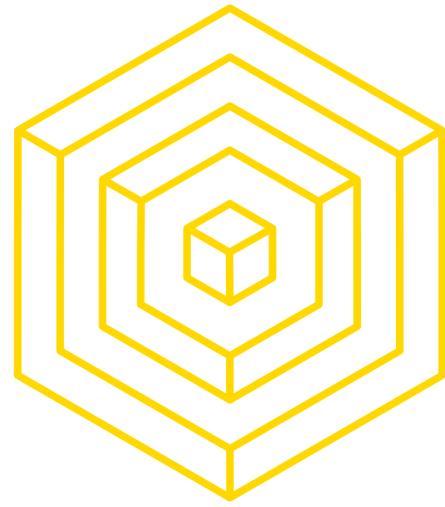
EXAMPLE

Lunchtime Consensus!



AUTHOR: BRIAN HO
▼ ▲ ▼ ▲ ▼

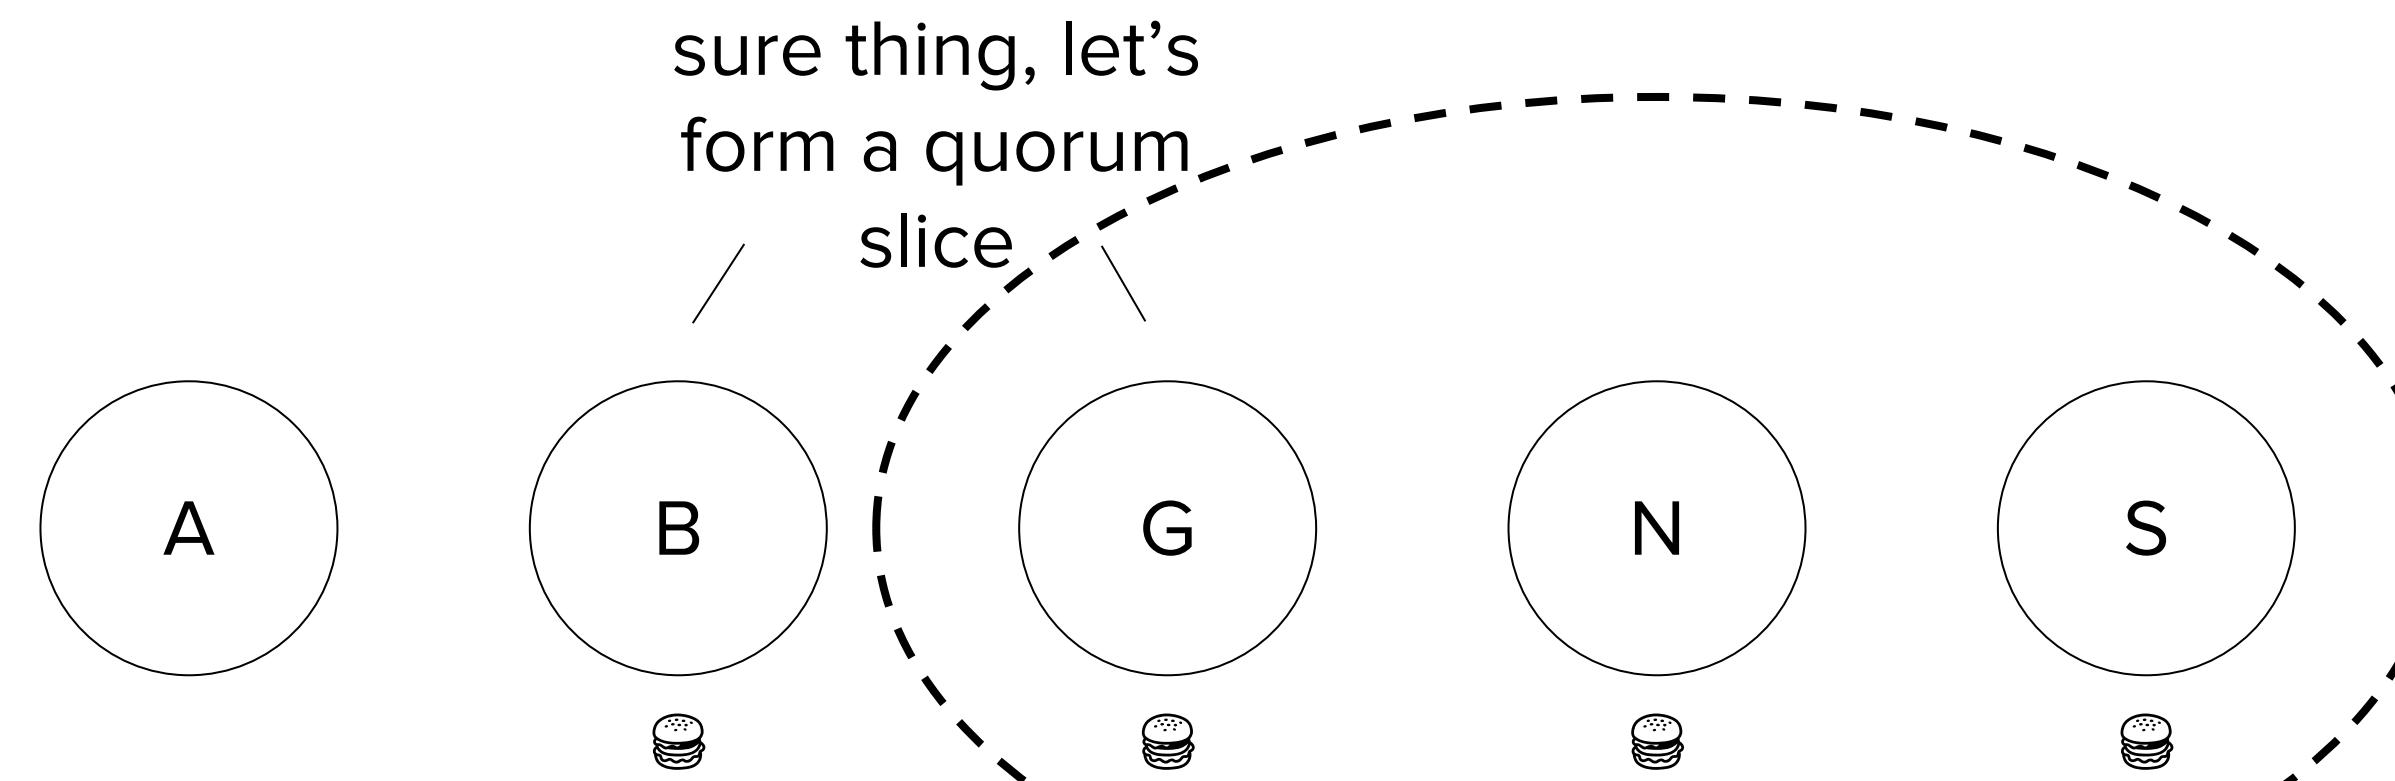
BLOCKCHAIN FUNDAMENTALS LECTURE 8



FEDERATED BYZANTINE AGREEMENT

EXAMPLE

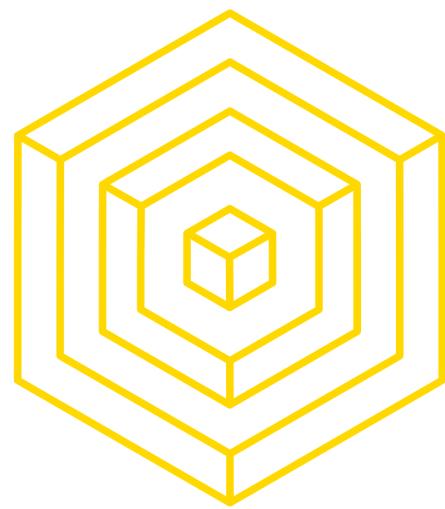
Lunchtime Consensus!



◀ ▶ ▲ ▼

AUTHOR: BRIAN HO

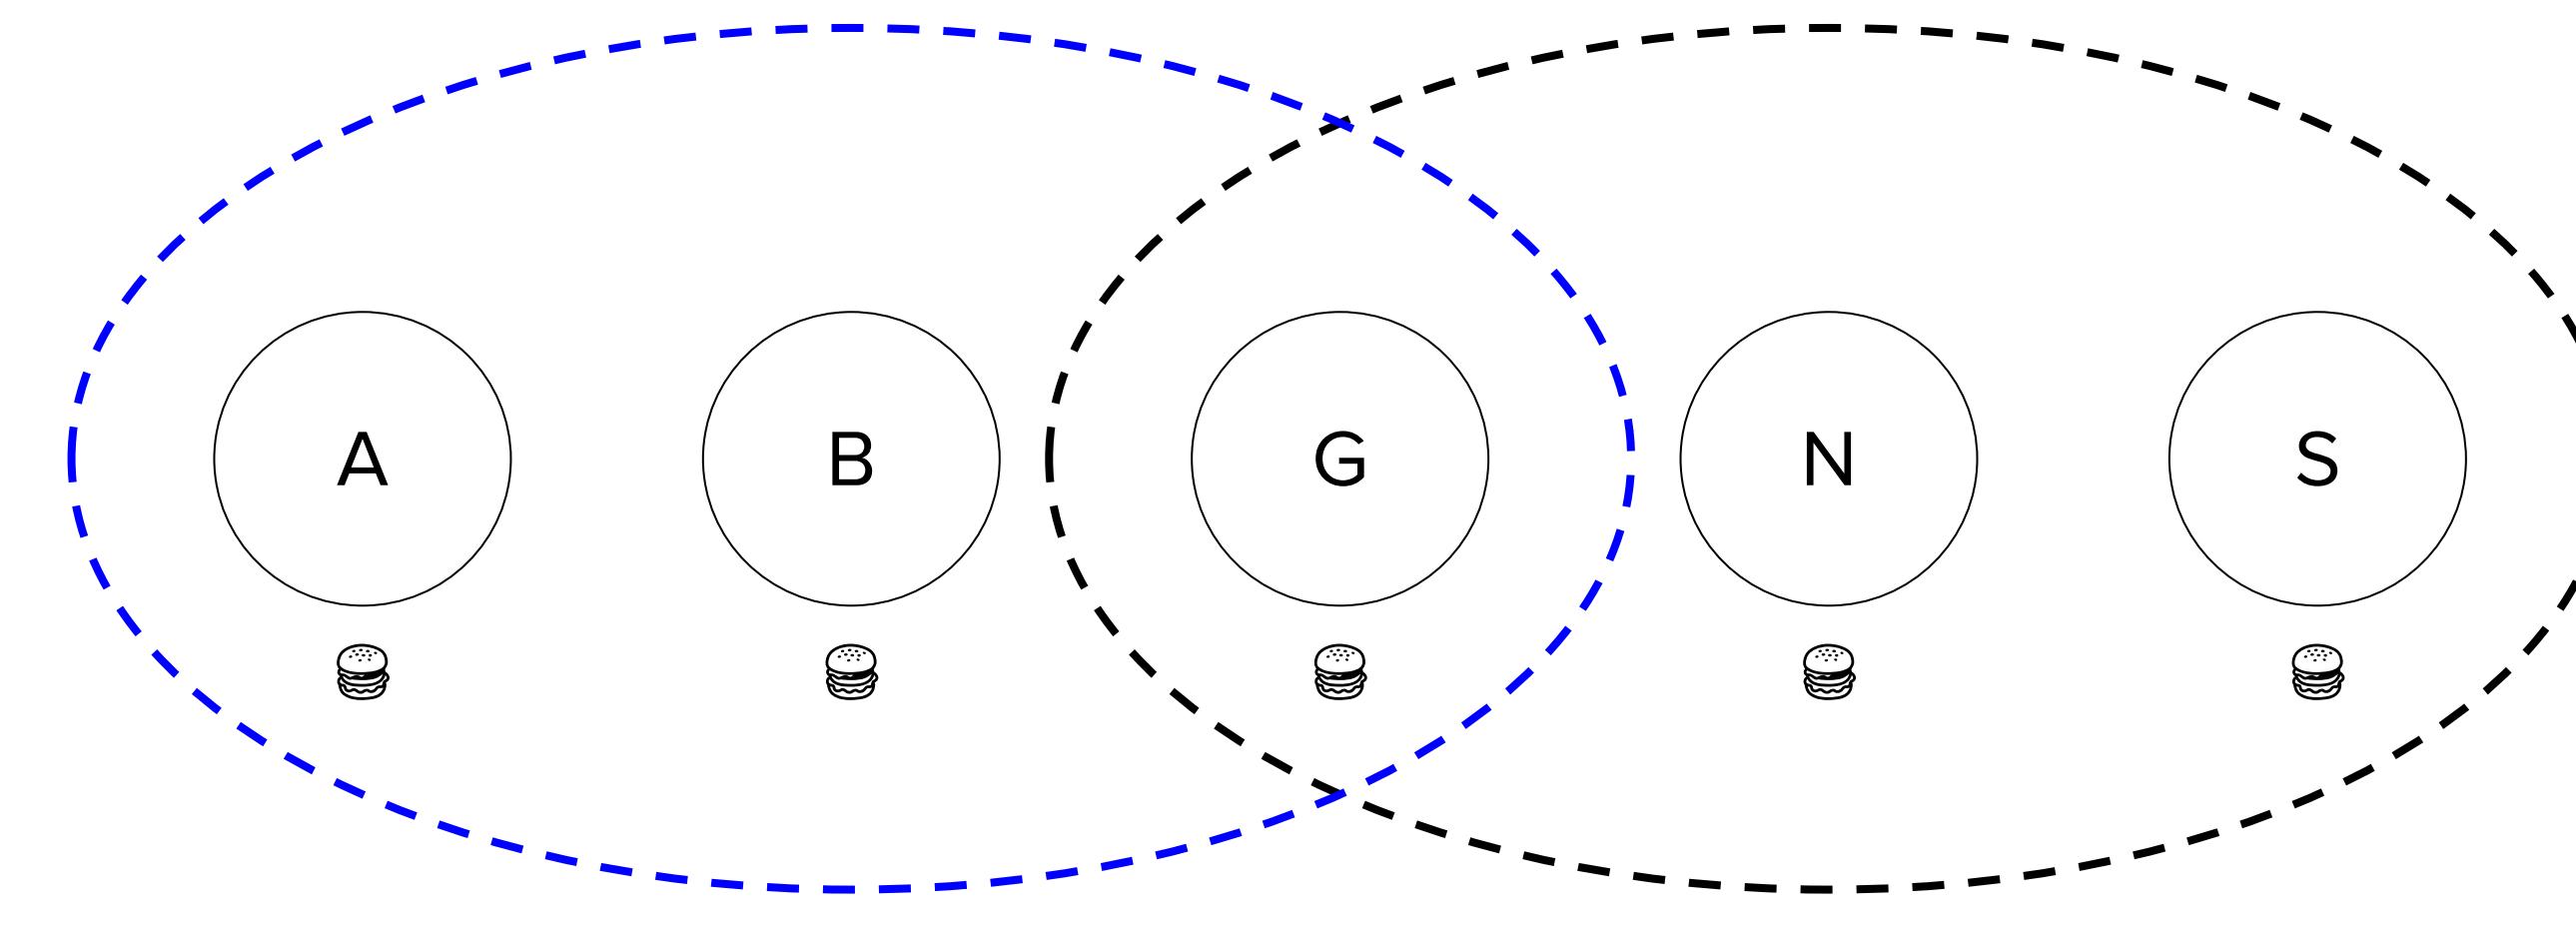
BLOCKCHAIN FUNDAMENTALS LECTURE 8



FEDERATED BYZANTINE AGREEMENT

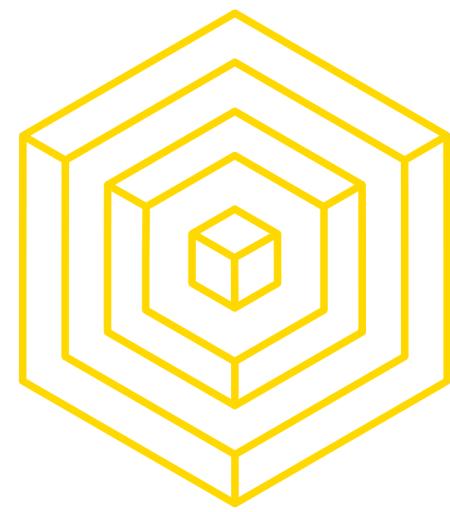
EXAMPLE

Lunchtime Consensus!



AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8



FEDERATED BYZANTINE AGREEMENT

OVERVIEW

Decentralized control: no central authority that authorizes consensus

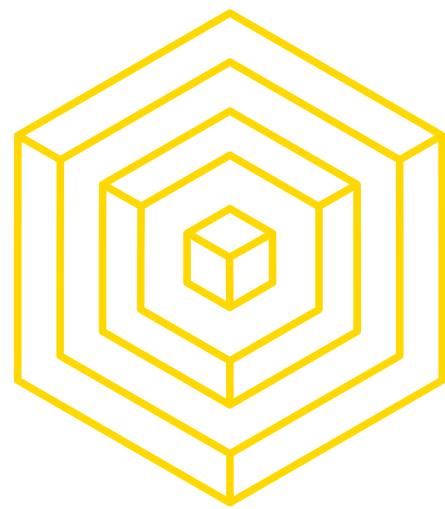
Low latency: consensus achieved in a few seconds

Flexible trust: nodes choose who they trust, don't have to trust the entire network



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



RIPPLE

OVERVIEW

Ripple is a real-time gross settlement system (RTGS), currency exchange, and remittance network

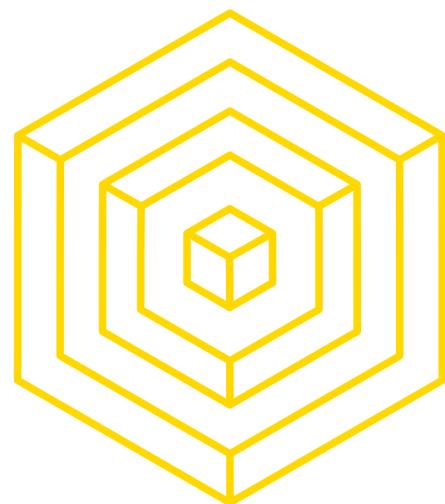
- Distributed ledger across financial institutions
- Allows banks to send money to overseas branches in real time with very low fees via Ripple's blockchain system
- Permissioned ledger



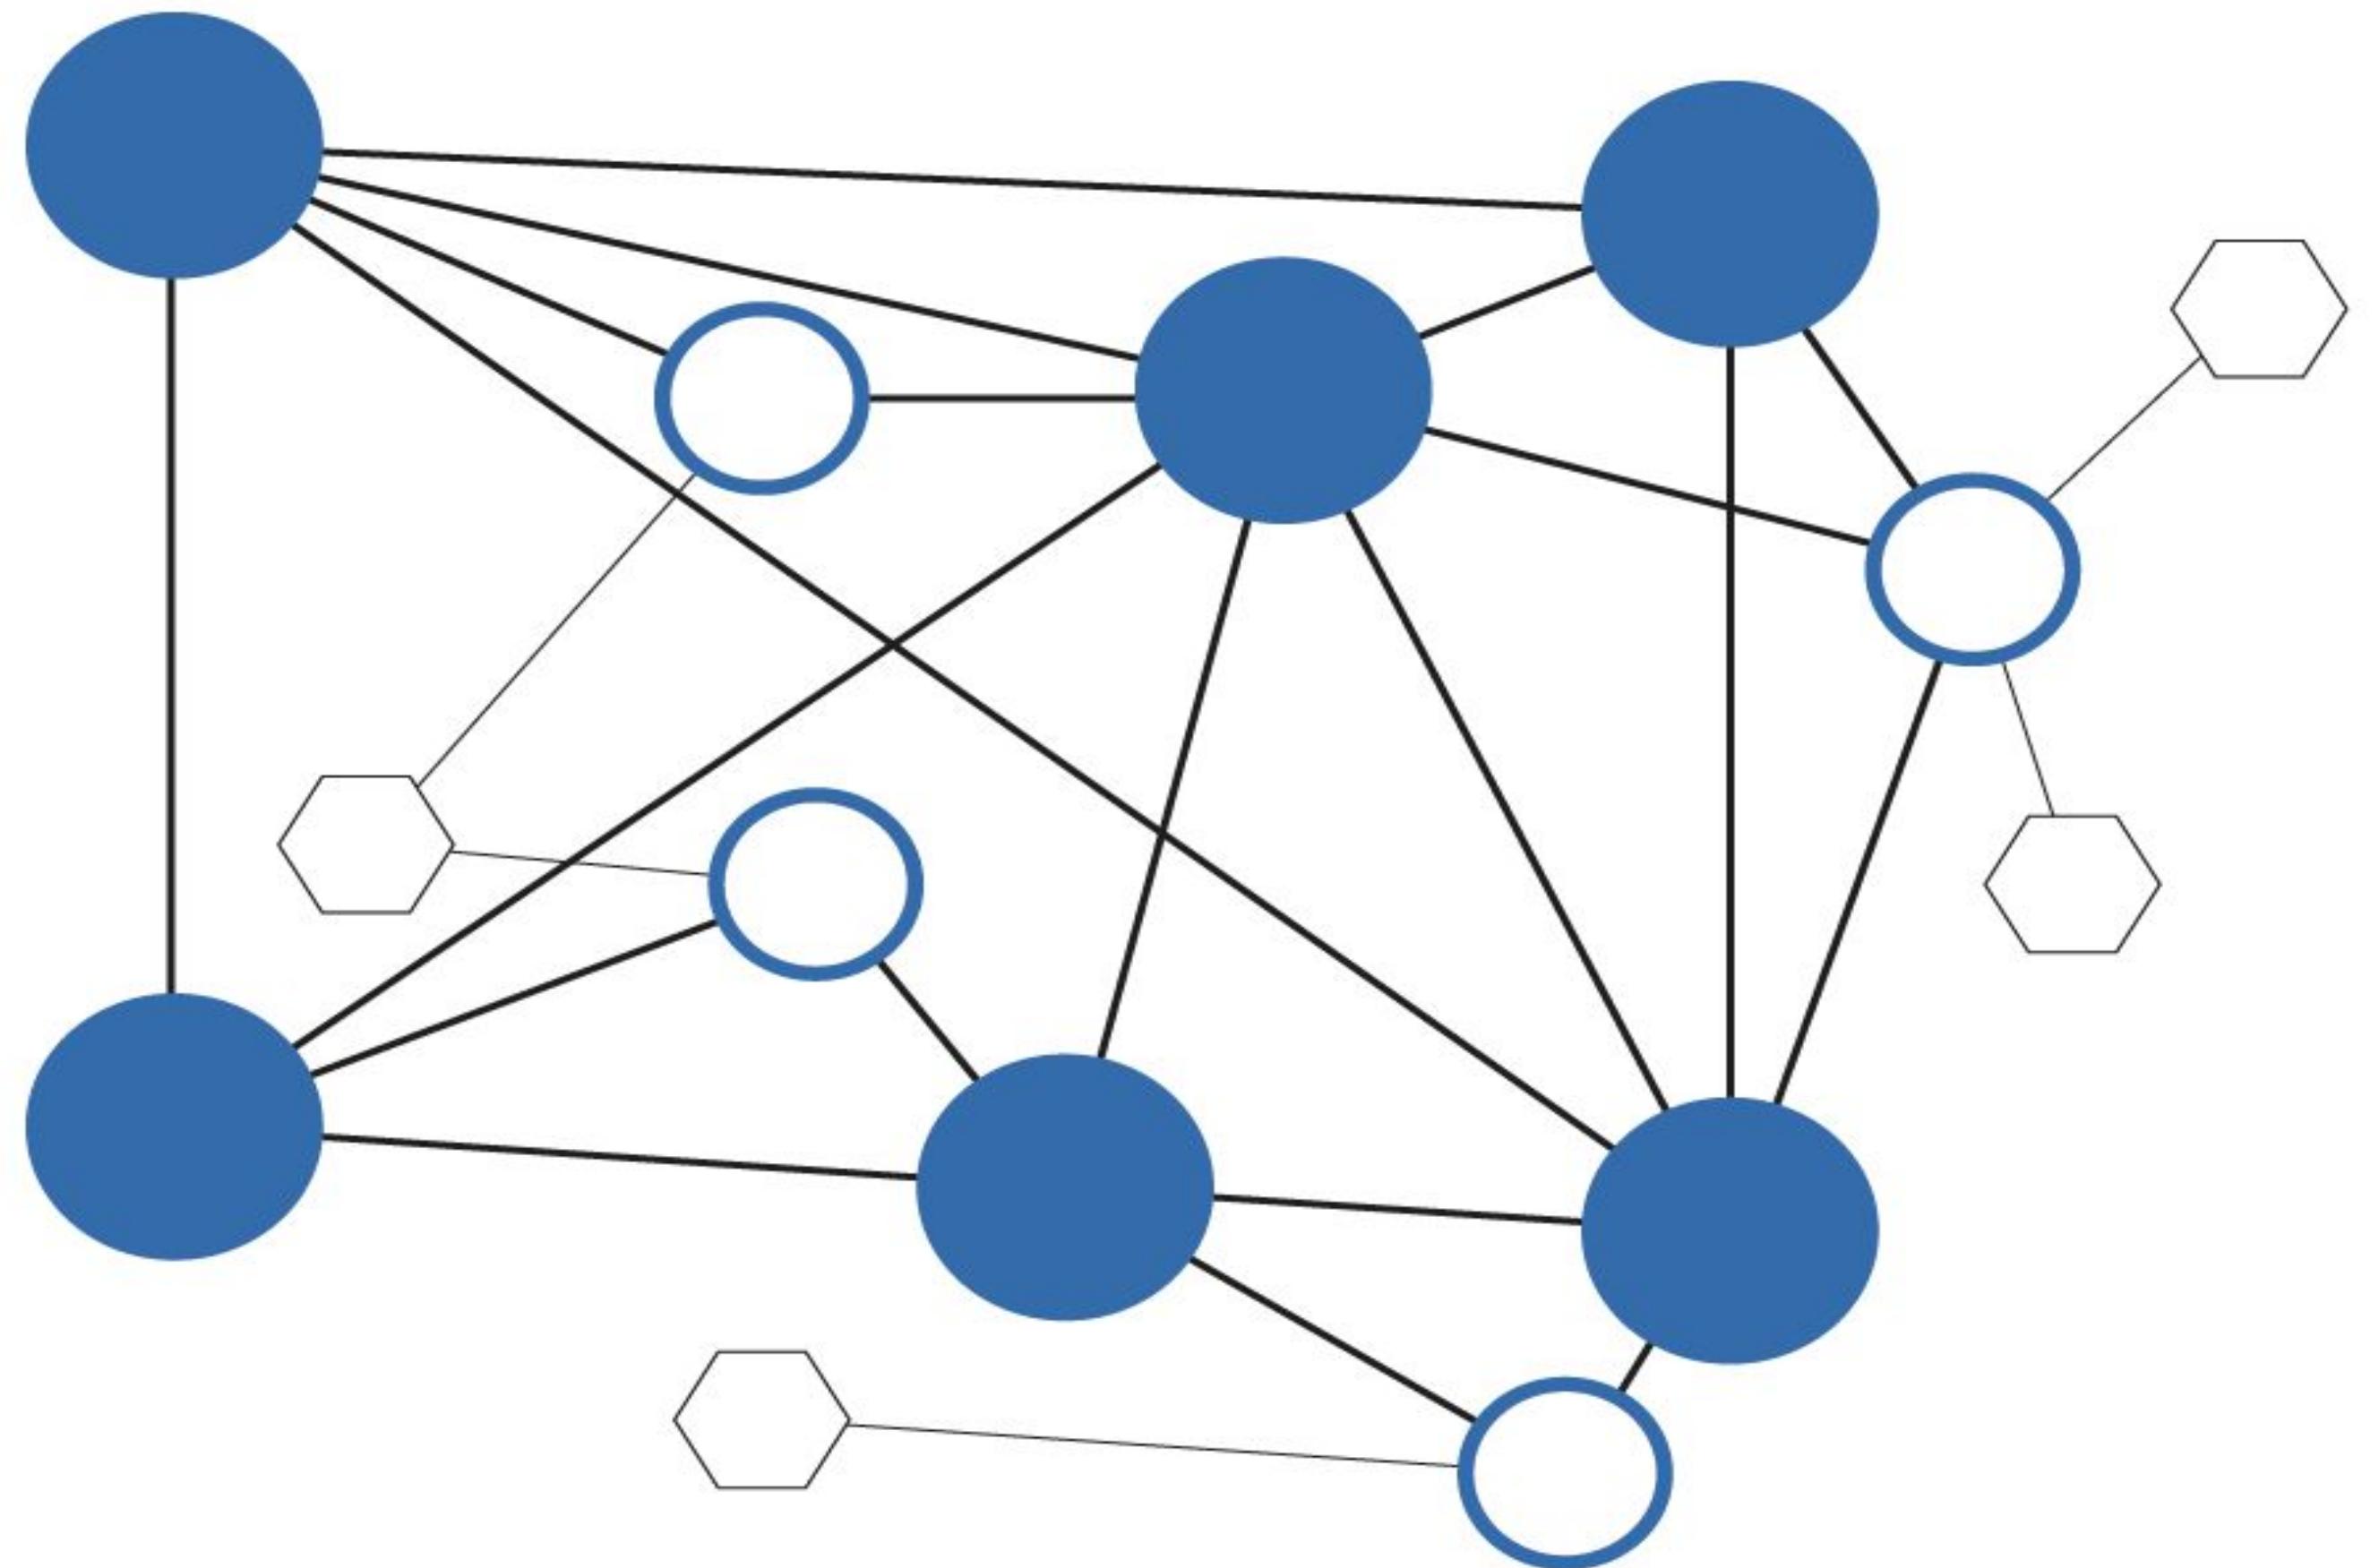
Focus: big financial institutions



AUTHOR: BRIAN HO



RIPPLE OVERVIEW

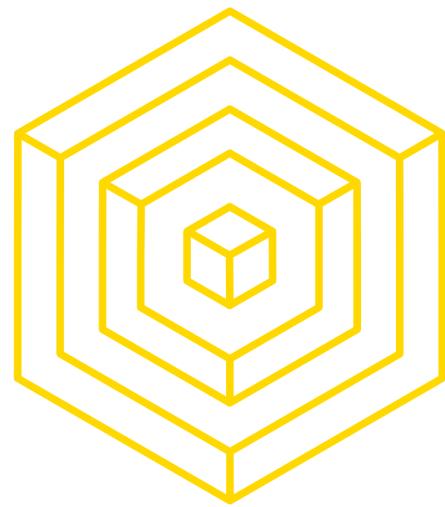


- submitted transactions
- network messages
- validating node
- tracking node
- application

SOURCE: <https://ripple.com/build/xrp-ledger-consensus-process/>

BLOCKCHAIN FUNDAMENTALS LECTURE 8

AUTHOR: RUSTIE LIN



STELLAR

OVERVIEW

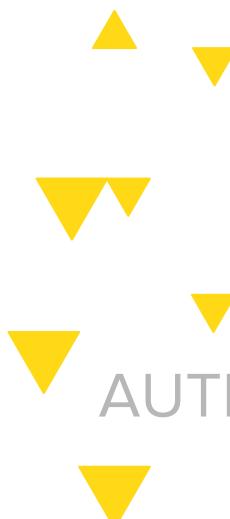
Stellar is an open-source protocol for exchanging money

- Platform that connects banks, payments systems, and people
- Open ledger

Focus: non-profit distribution of wealth

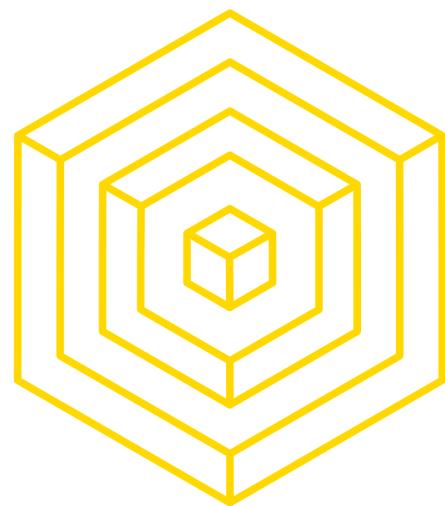


STELLAR

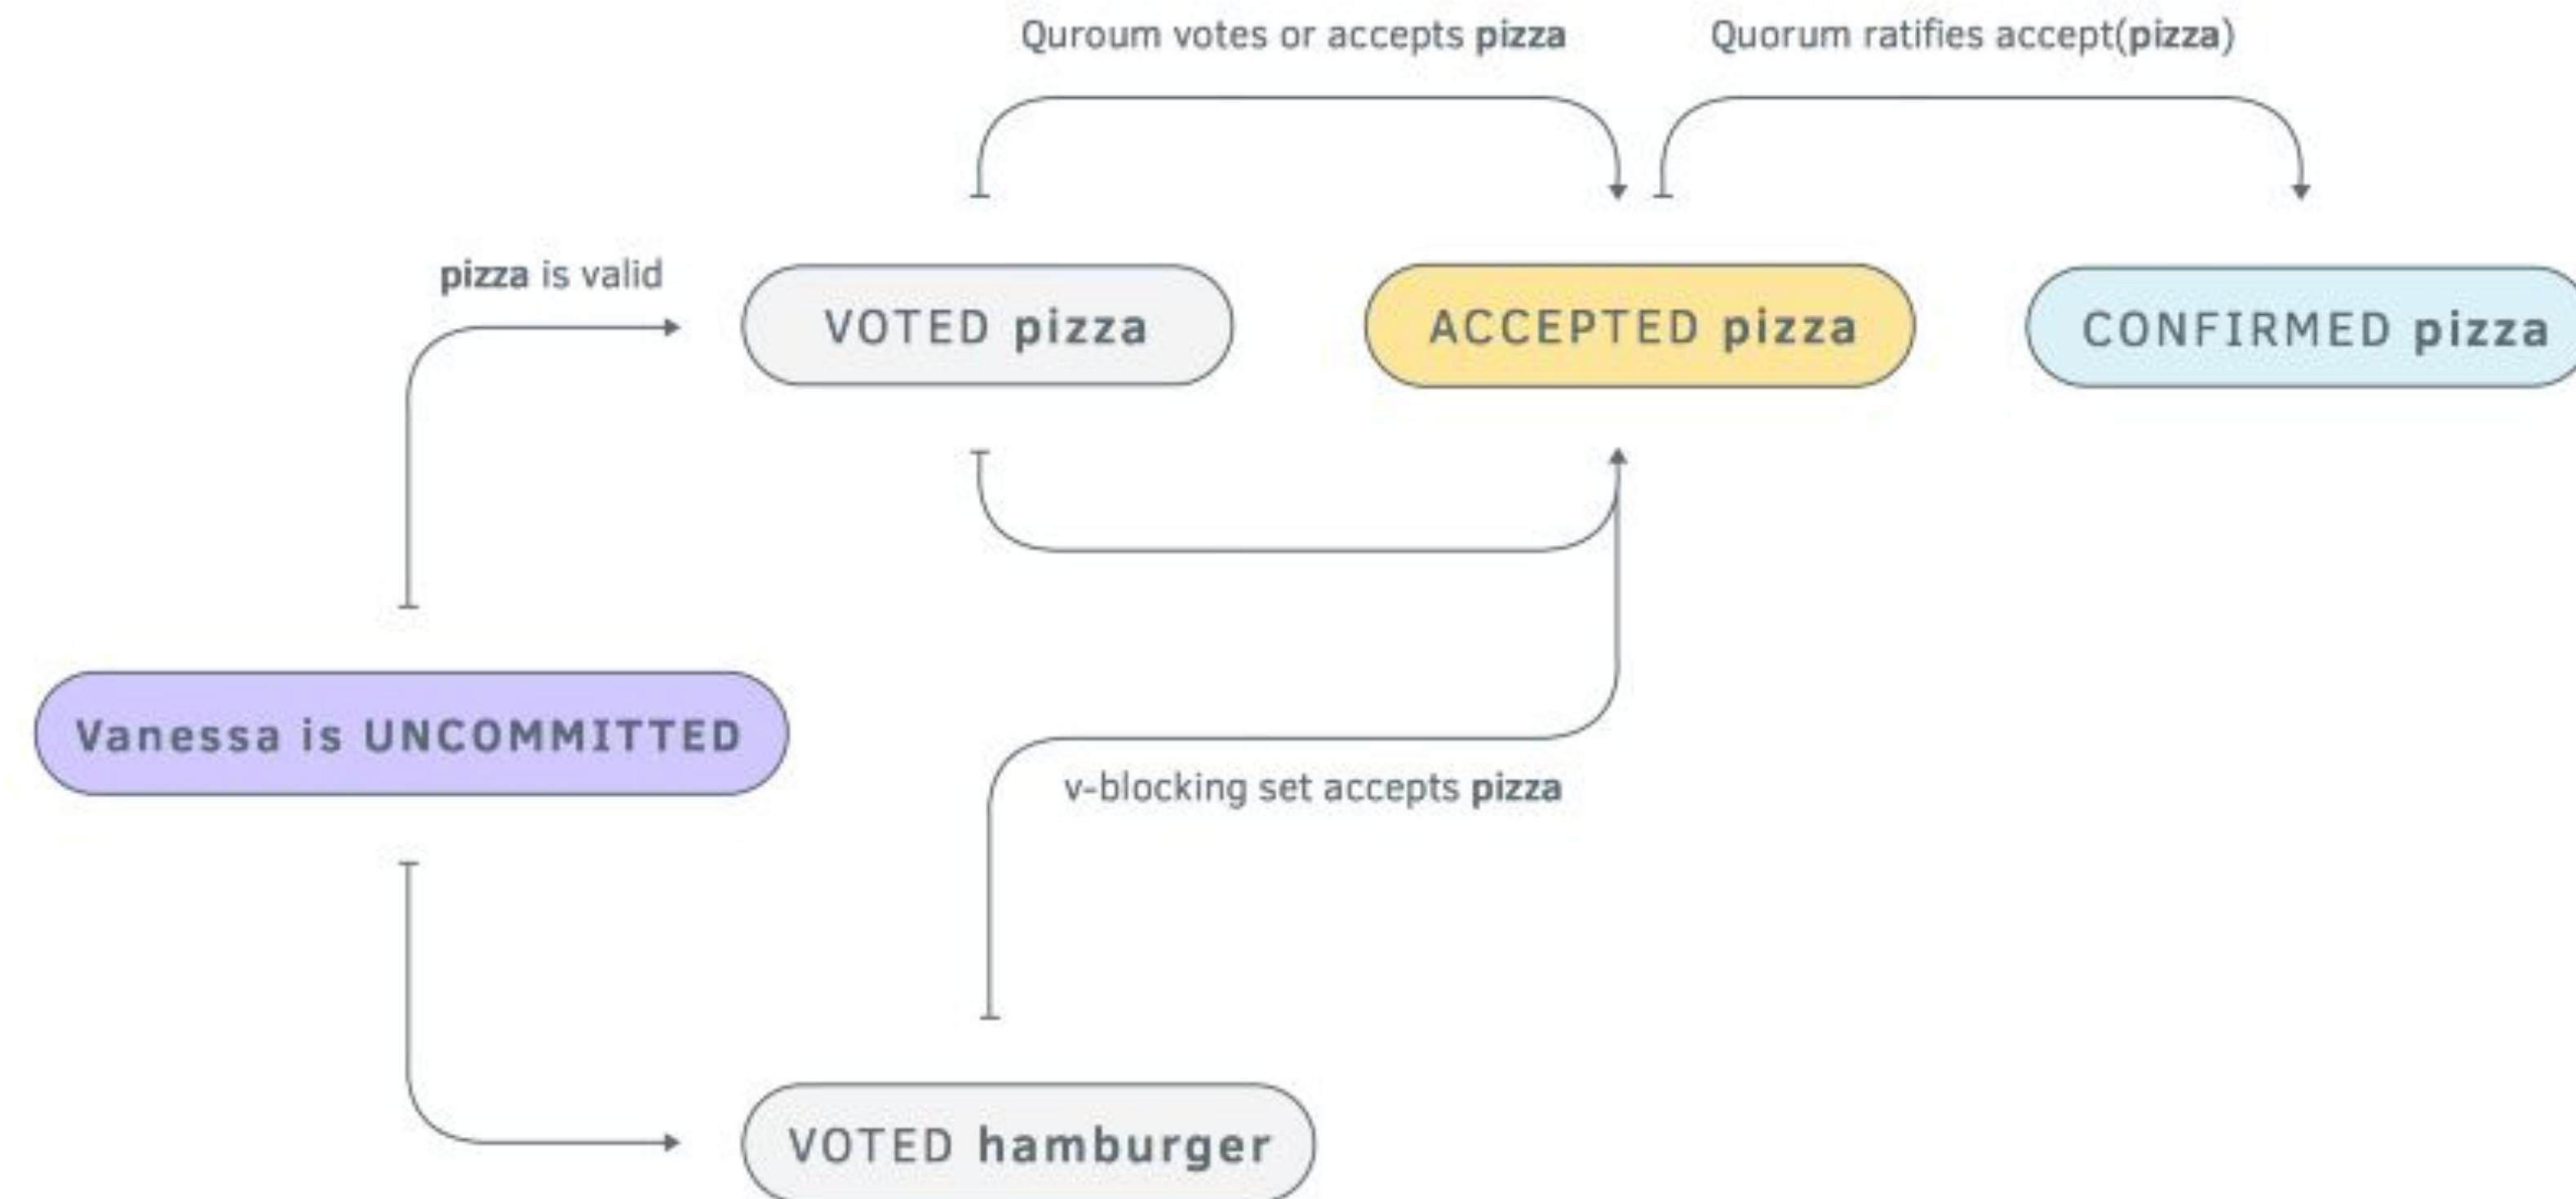


AUTHOR: BRIAN HO

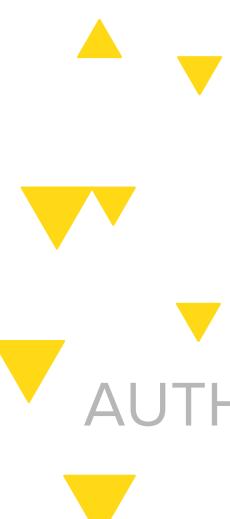
BLOCKCHAIN FUNDAMENTALS LECTURE 8



STELLAR OVERVIEW

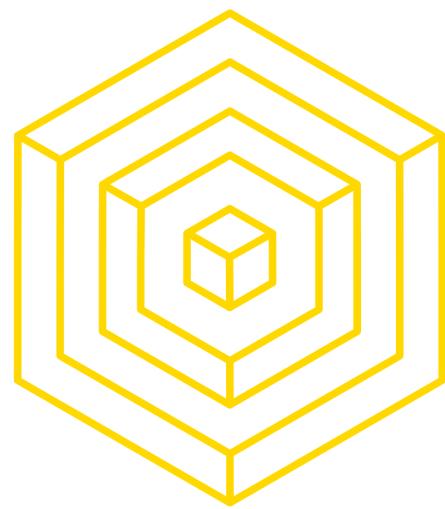


SOURCE: <https://medium.com/a-stellar-journey/on-worldwide-consensus-359e9eb3e949>



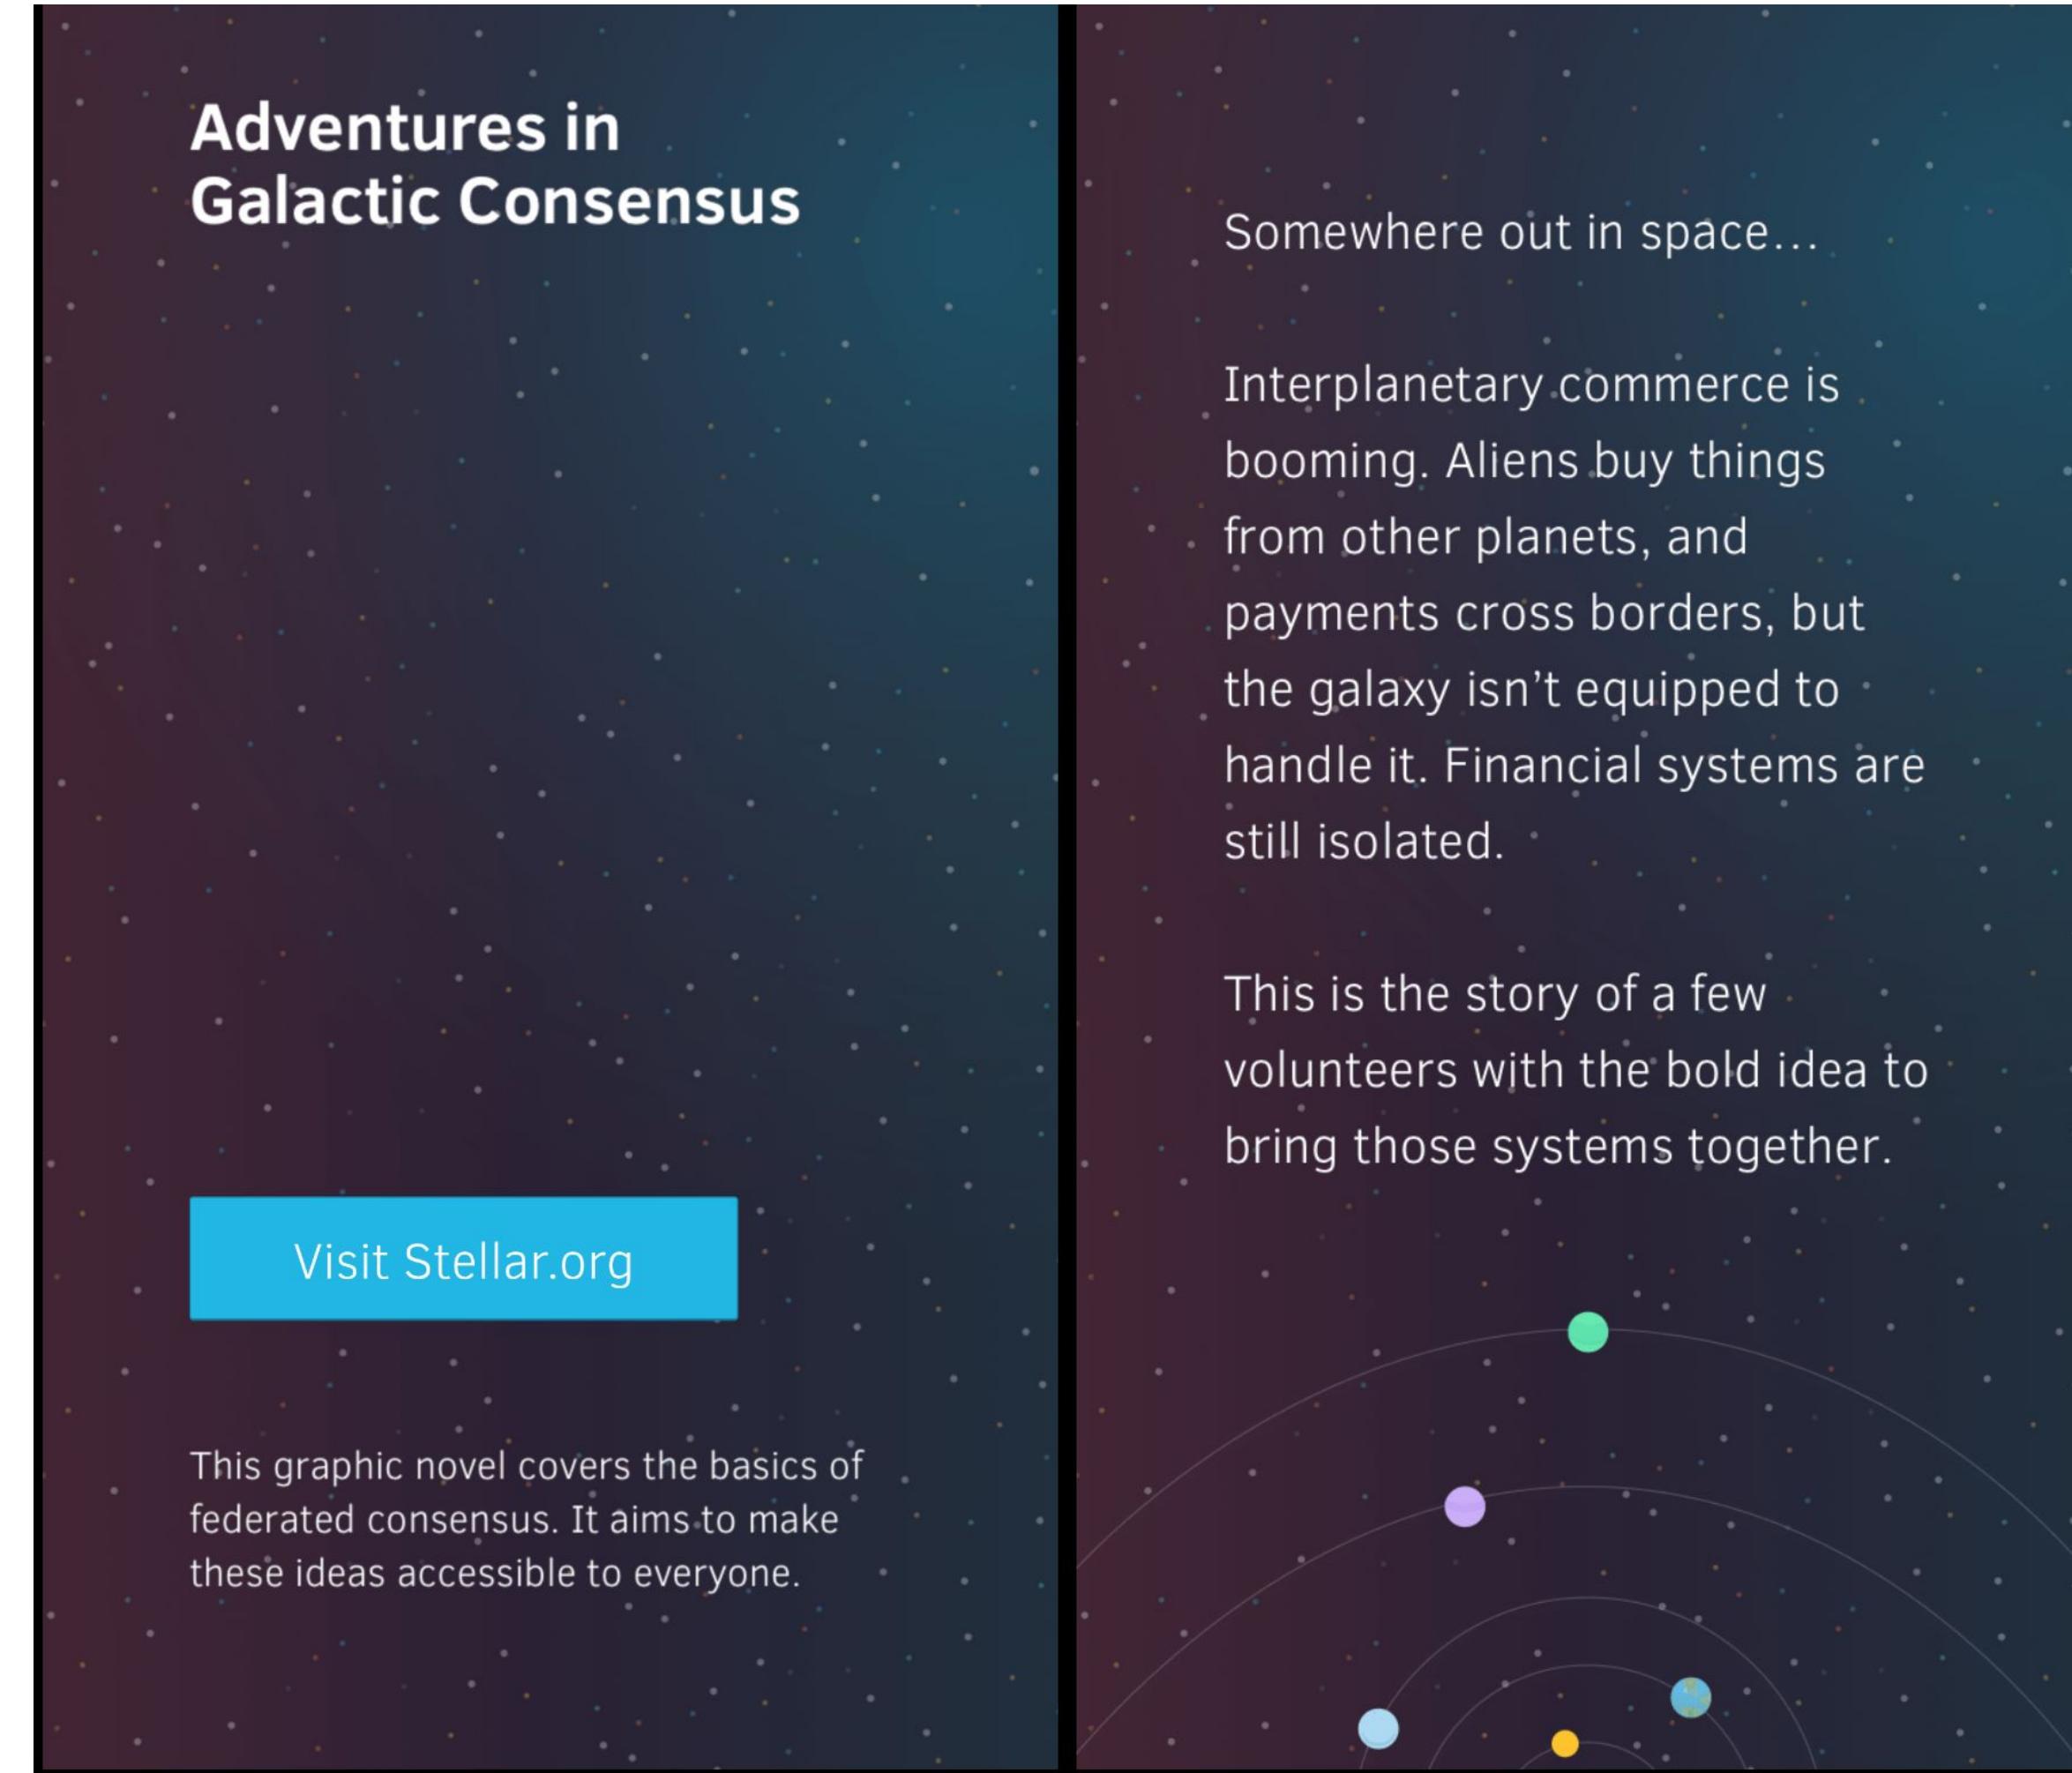
AUTHOR: BRIAN HO

BLOCKCHAIN FUNDAMENTALS LECTURE 8

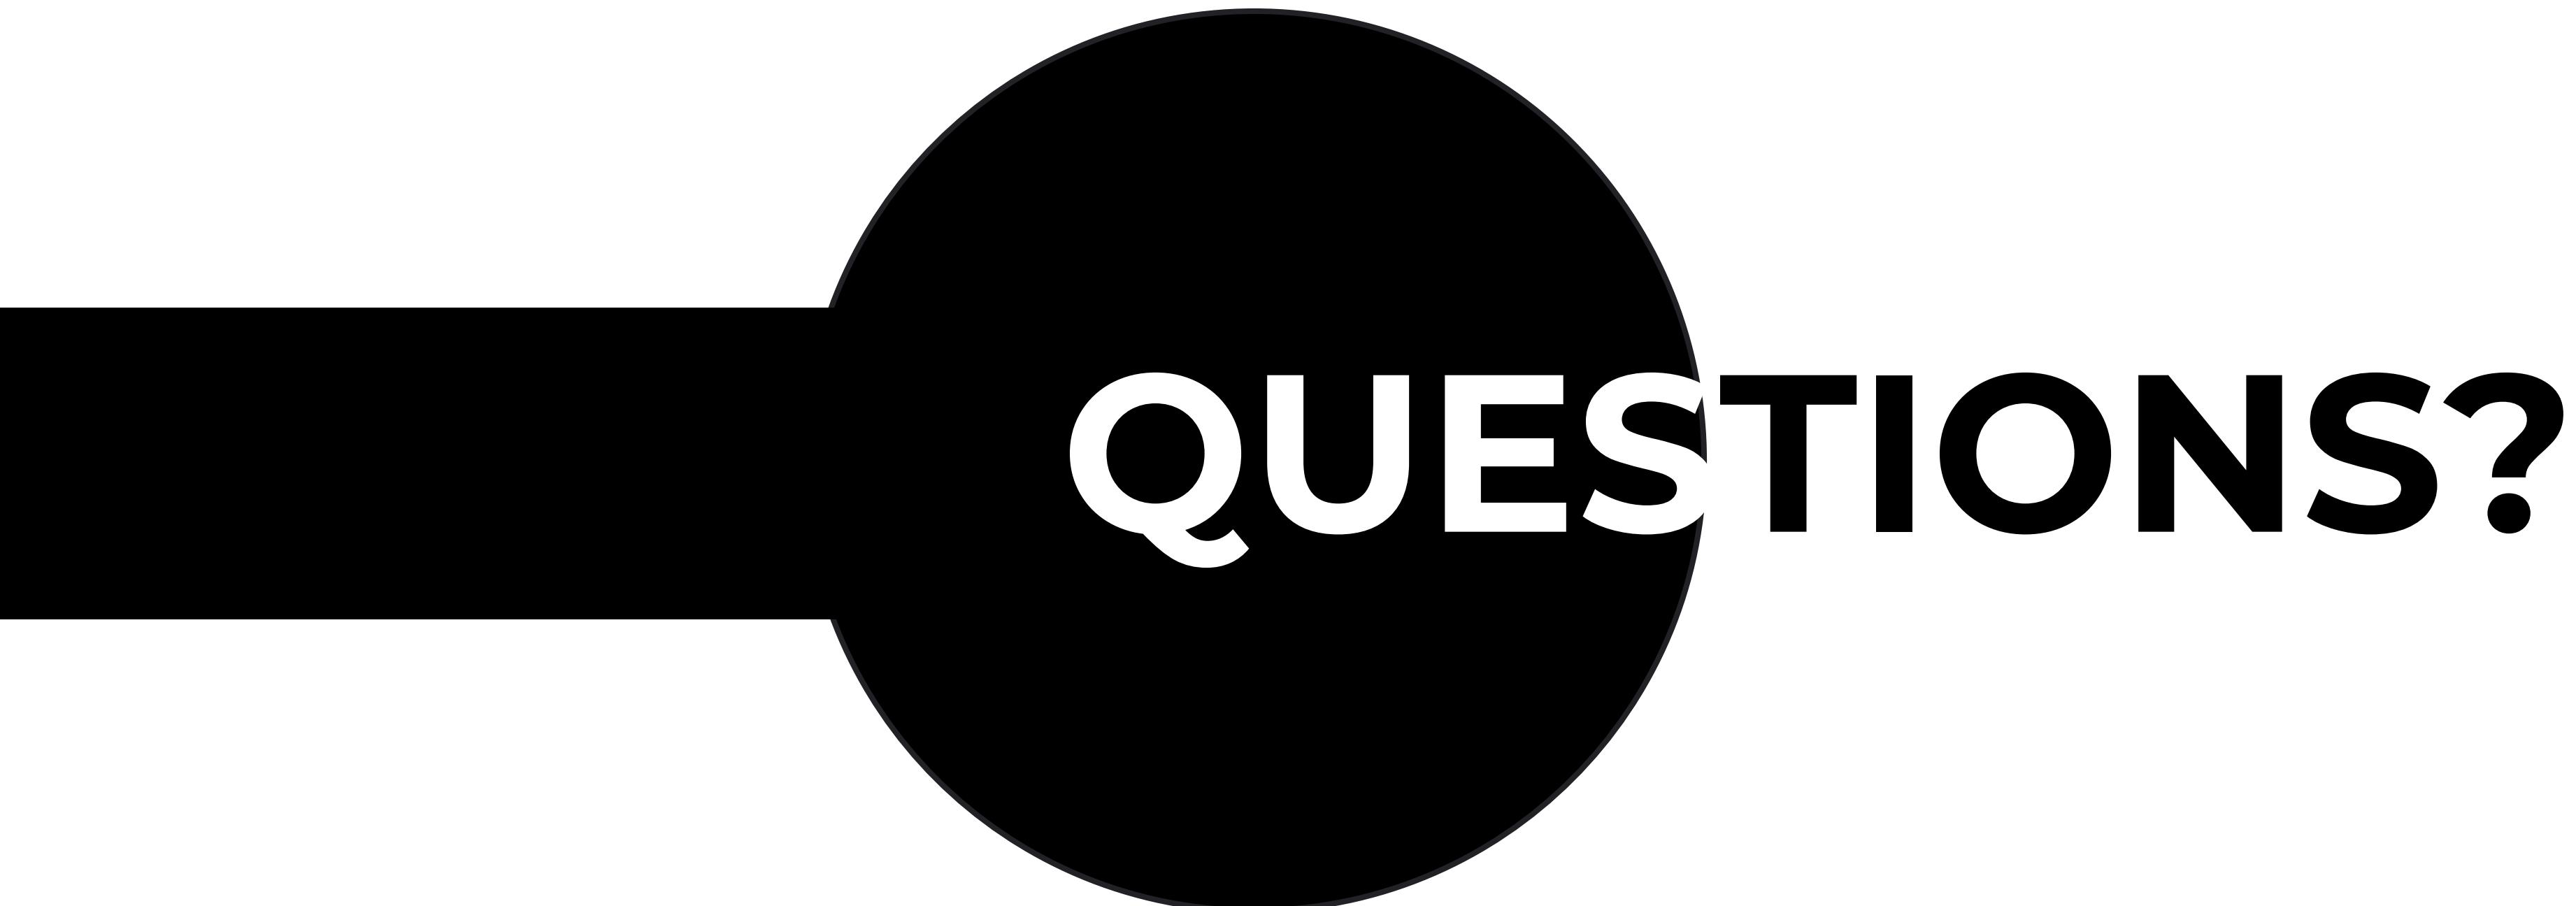


STELLAR

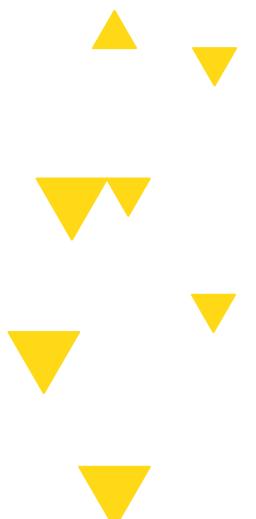
OVERVIEW

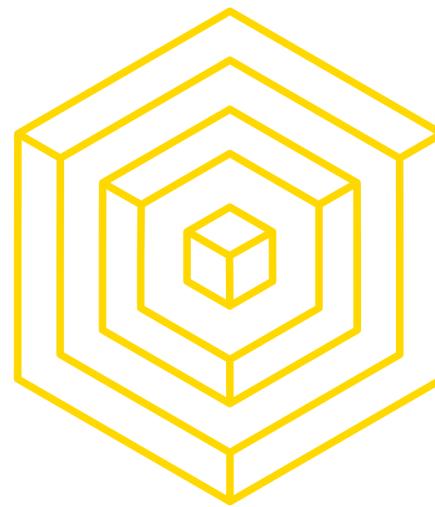


SOURCE: <https://www.stellar.org/stories/adventures-in-galactic-consensus-chapter-1/>



QUESTIONS?





READINGS

- Short overview of alternatives to PoW:

<https://www.coindesk.com/short-guide-blockchain-consensus-protocols/>

- Stellar Consensus Comic:

<https://www.stellar.org/stories/adventures-in-galactic-consensus-chapter-1/>

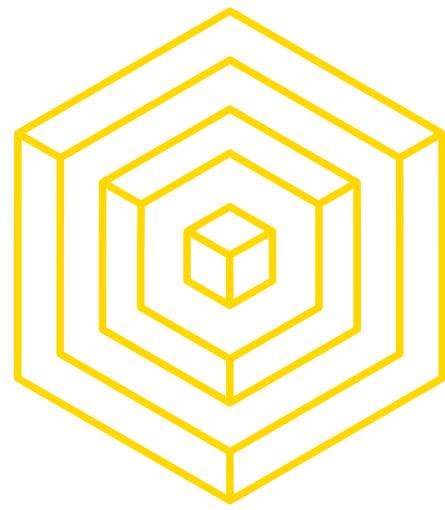
- Stellar Consensus Protocol Overview:

<https://medium.com/a-stellar-journey/on-worldwide-consensus-359e9eb3e949>

- CAP Theorem Overview: <https://www.youtube.com/watch?v=Jw1iFr4v58M>

- Raft Overview:

<http://container-solutions.com/raft-explained-part-1-the-consensus-problem/>



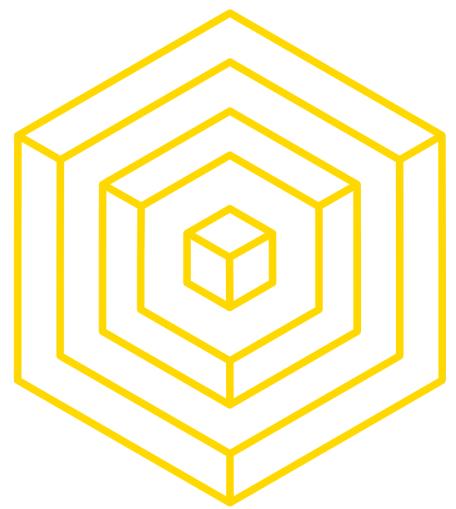
EXTRA: BYZANTINE GENERALS IRL

DIGITAL AVIONICS

- Super dependable computers pioneered by aircraft manufacturers
- Safety first, especially...
 - when you're flying 60-120 million USD aircraft
 - when there's 100s of passengers
- Fault Tolerant Avionics:
http://www.davi.ws/avionics/TheAvionicsHandbook_Cap_28.pdf



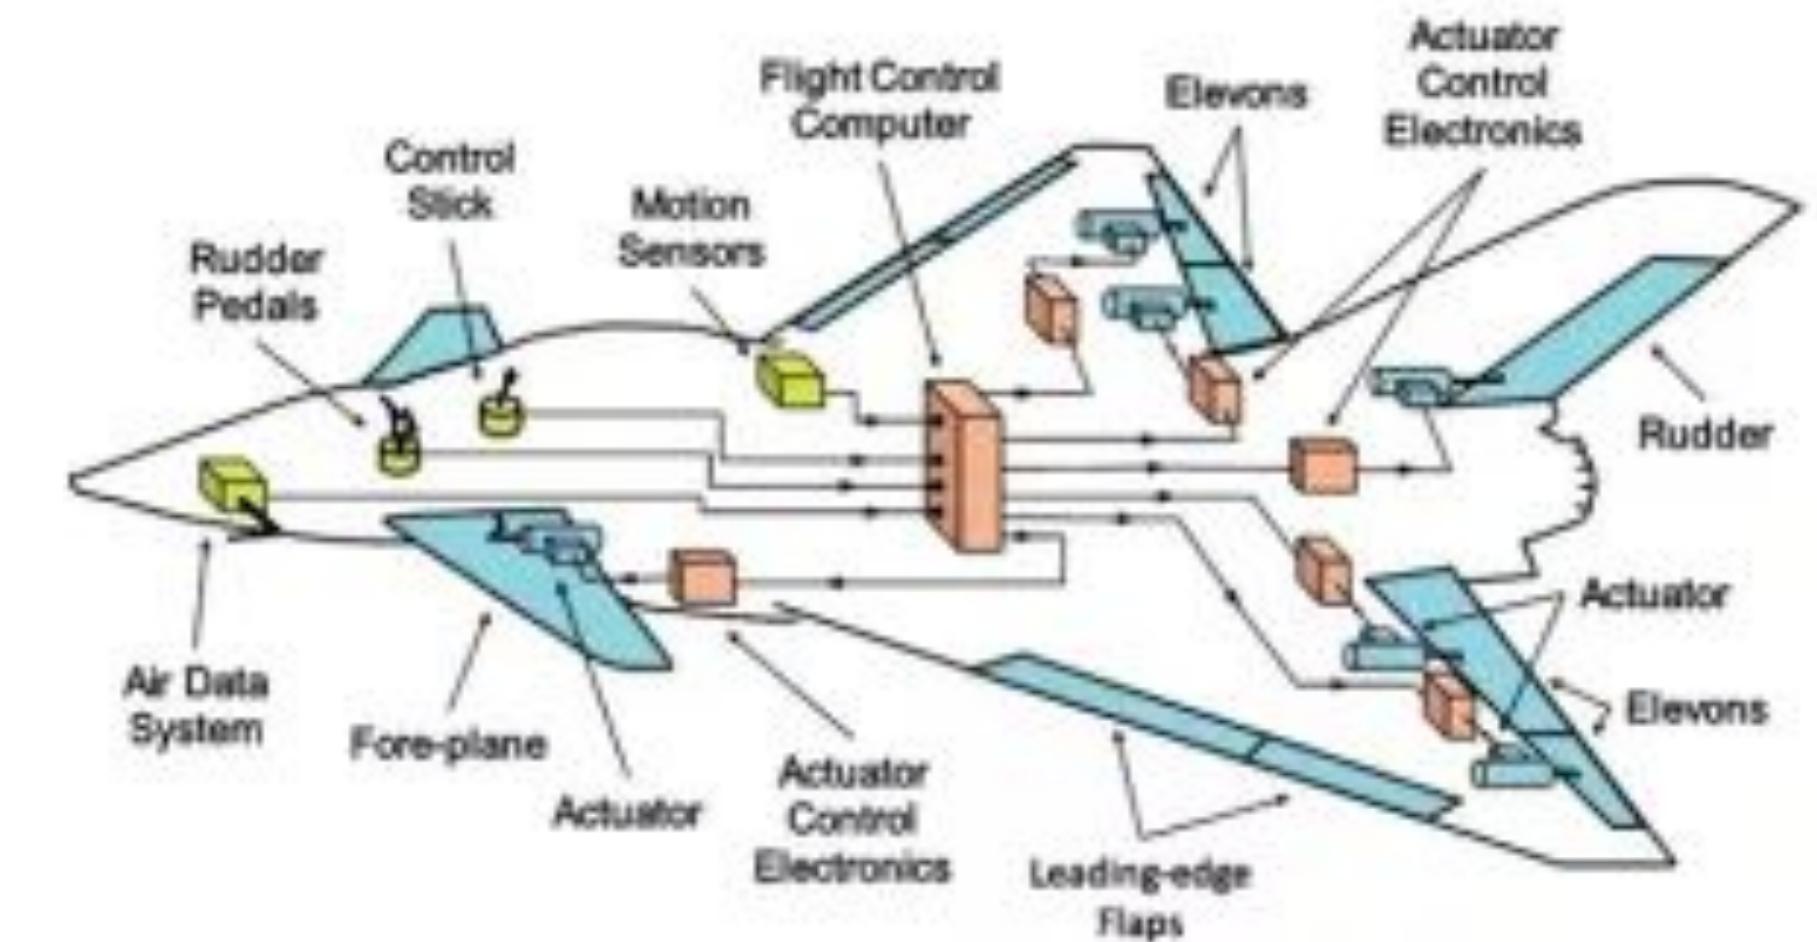
AUTHOR: RUSTIE LIN



EXTRA: BYZANTINE GENERALS IRL

DIGITAL AVIONICS

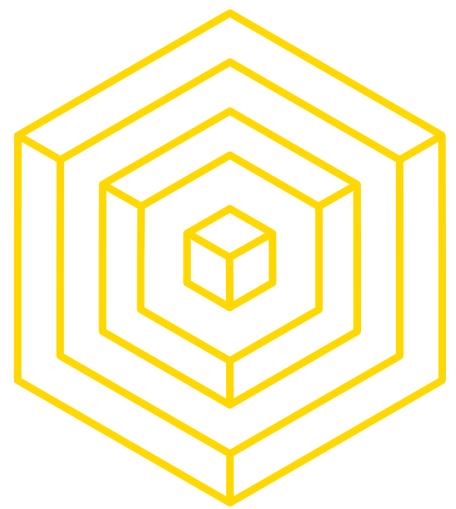
- e.g. 550 passenger Boeing 777
- ARINC 659 SAFEBus network
 - nodes: computers, sensors,
 - duplicate transmitters that control nodes
 - silence node if transmitters informed of too many errors
 - recipient nodes receive 4 copies of message
 - accept/record only if 4 identical messages



Basic elements of the FWB control system.



AUTHOR: RUSTIE LIN



EXTRA: BYZANTINE GENERALS IRL

DIGITAL AVIONICS

- e.g. SpaceX Dragon
- NASA requirements for ISS
- Triply redundant computers
- Radiation events changing memory or register values

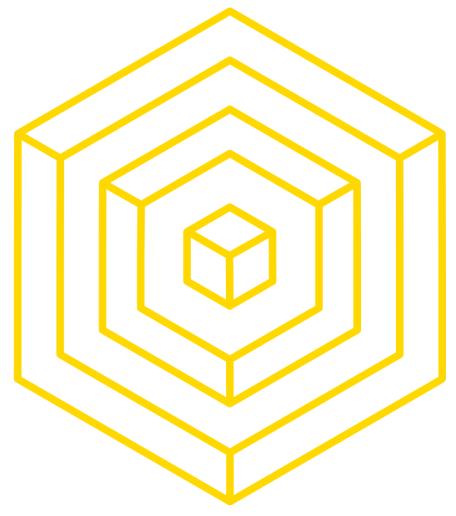


SOURCE:

<https://www.universetoday.com/112272/meet-spacexs-new-manned-dragon-cool-animation-shows-how-it-works/>



AUTHOR: RUSTIE LIN



EXTRA: BASIC PAXOS

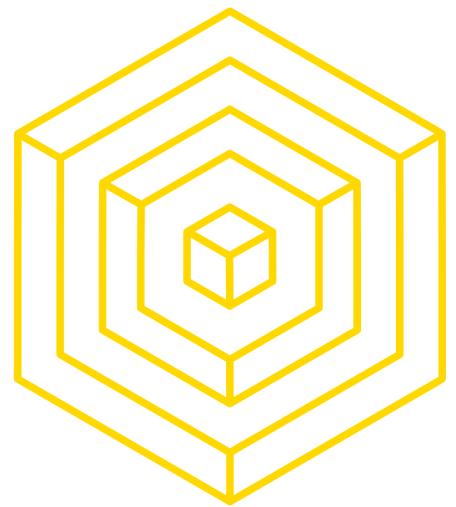
PHASE 1A: PREPARE

Proposer creates proposal N, sends to quorum of acceptors



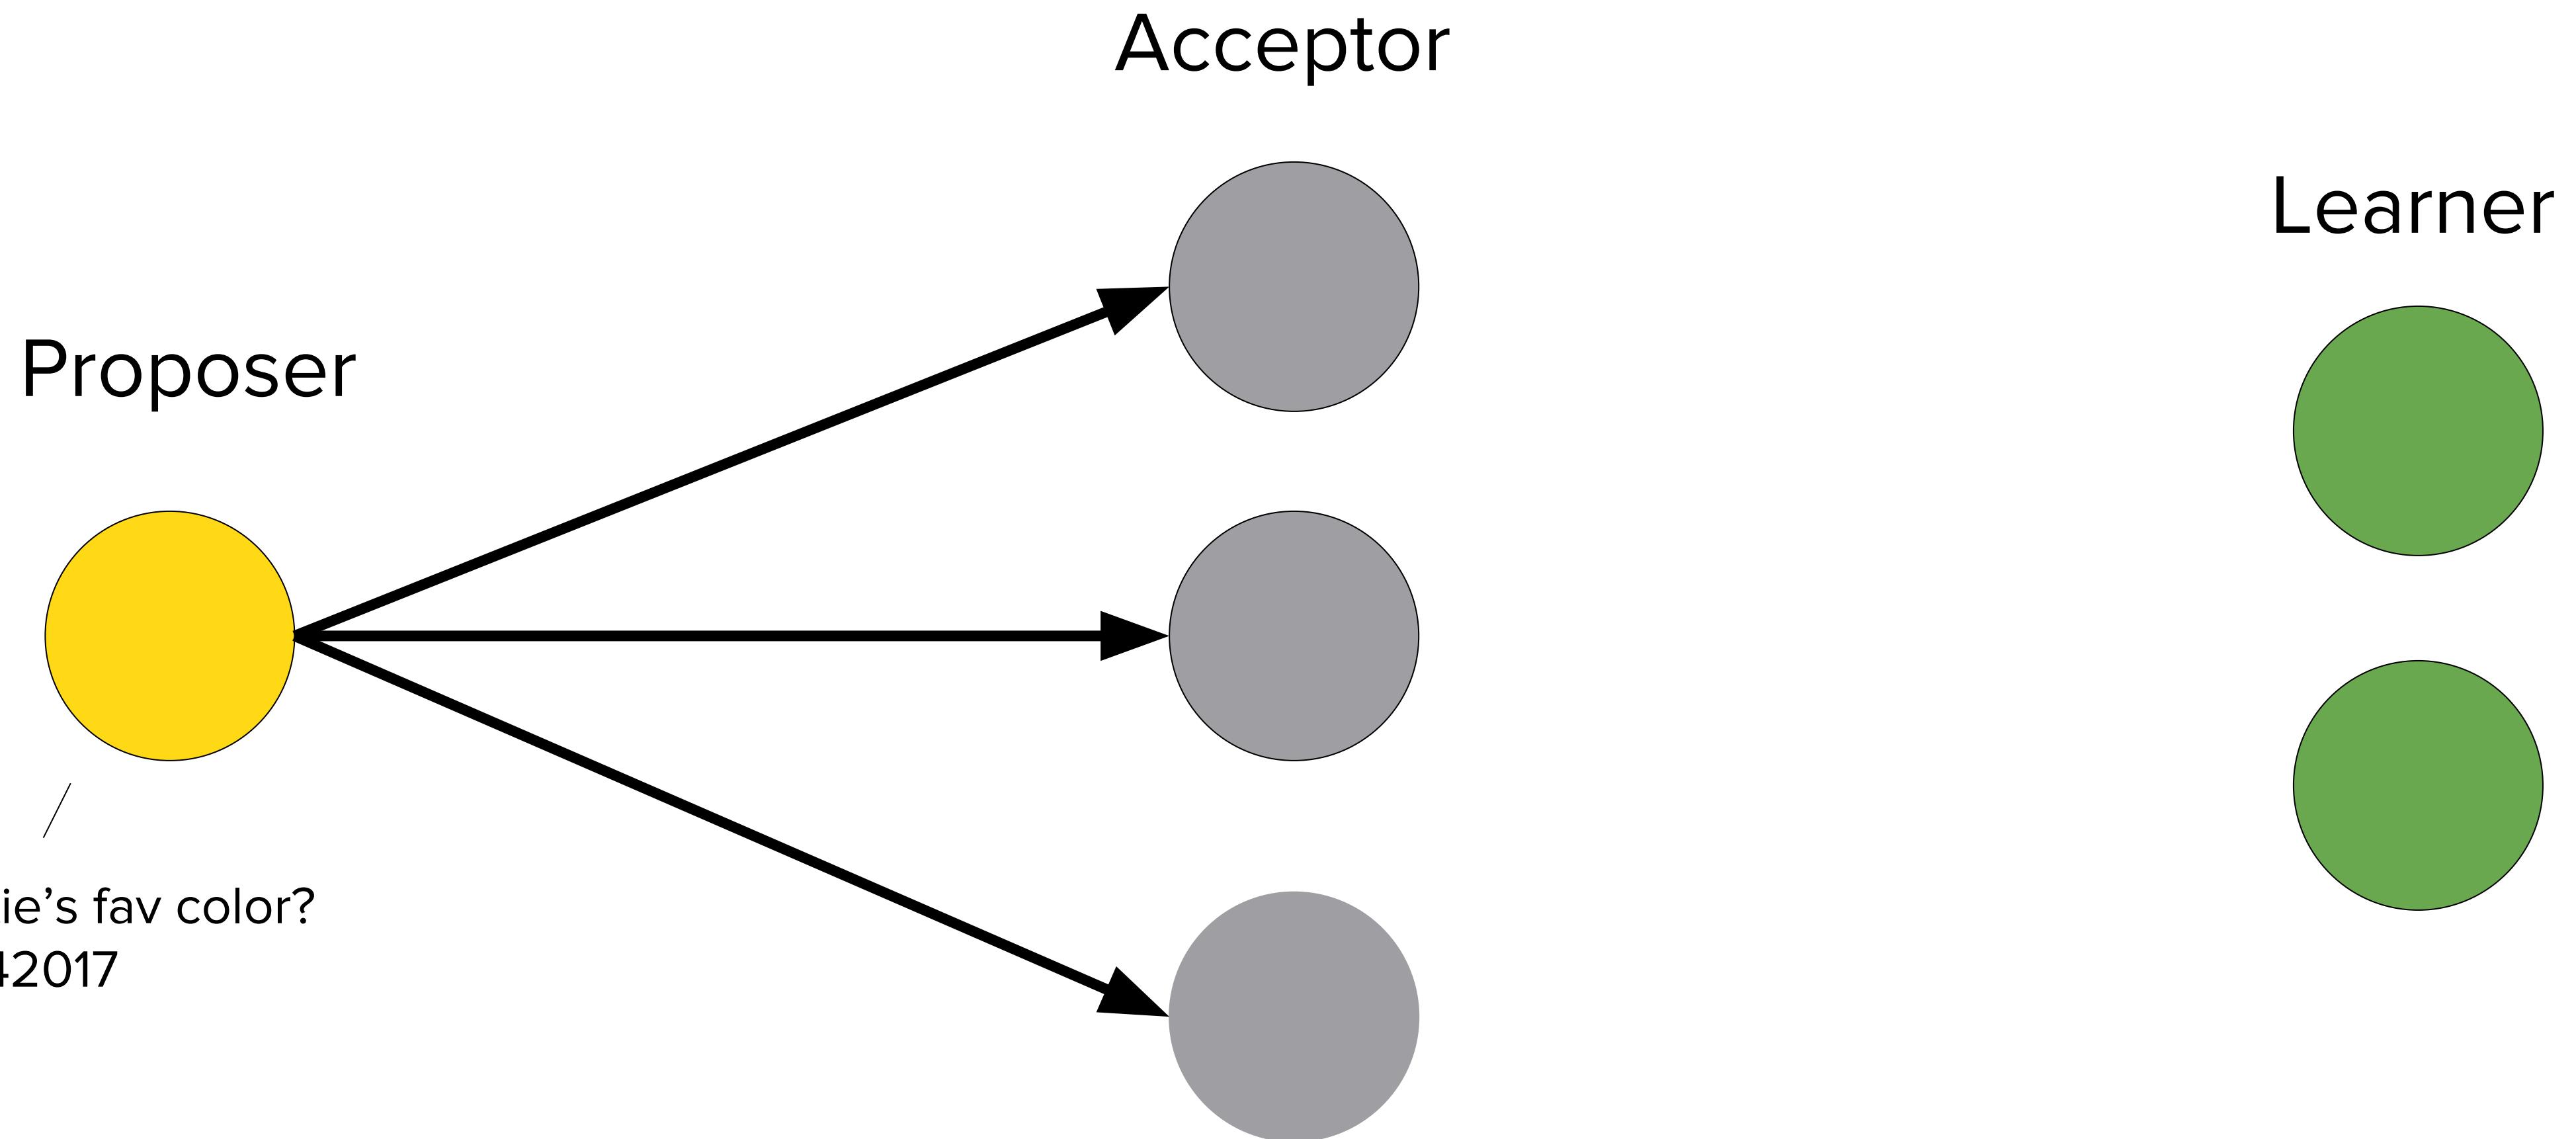
AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8

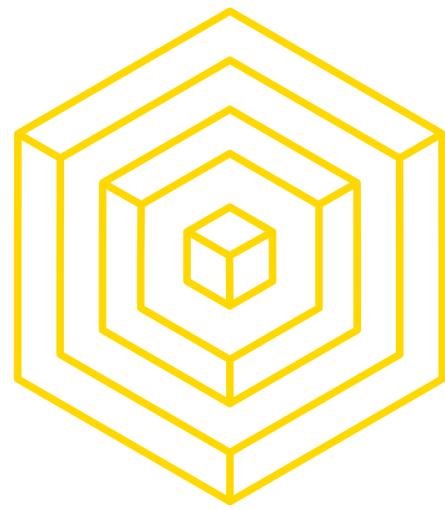


EXTRA: BASIC PAXOS

PHASE 1A: PREPARE



▲
▼
▼
▼
AUTHOR: RUSTIE LIN



EXTRA: BASIC PAXOS

PHASE 1B: PROMISE

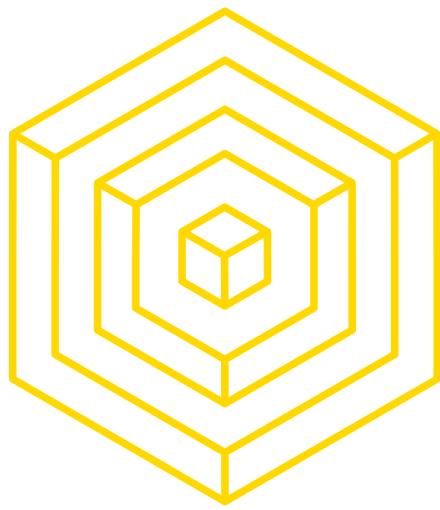
After getting proposal N, Acceptors promise to ignore proposals less than N, response to Proposer

If Acceptor had accepted proposal in the past, send previous proposal number and value back to Proposer, along with Promise



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8

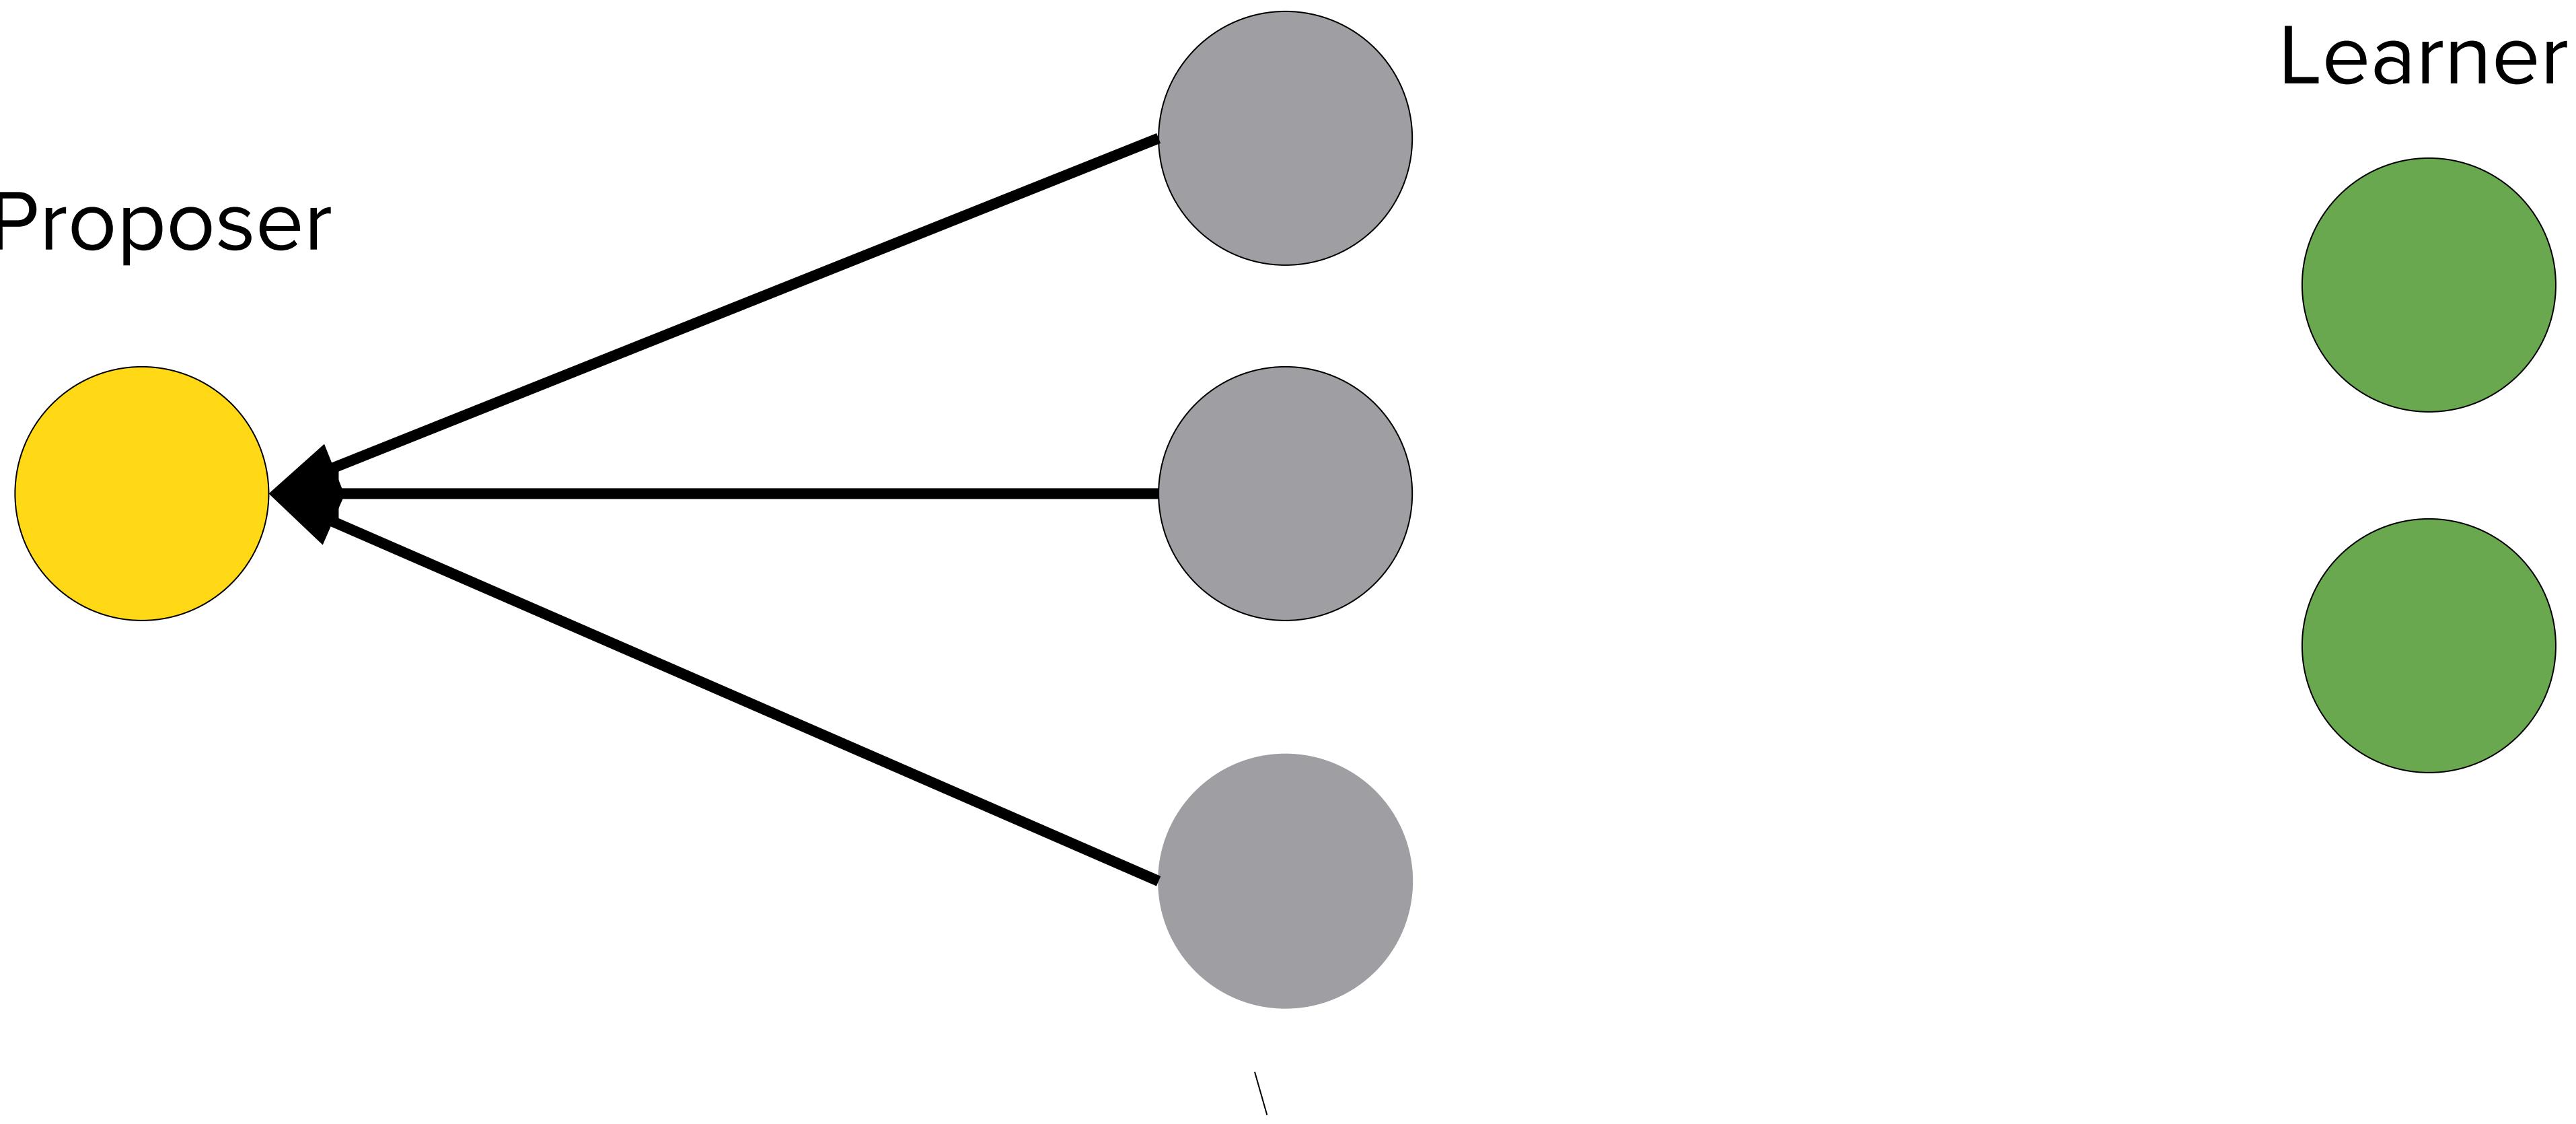


EXTRA: BASIC PAXOS

PHASE 1B: PROMISE

Acceptor

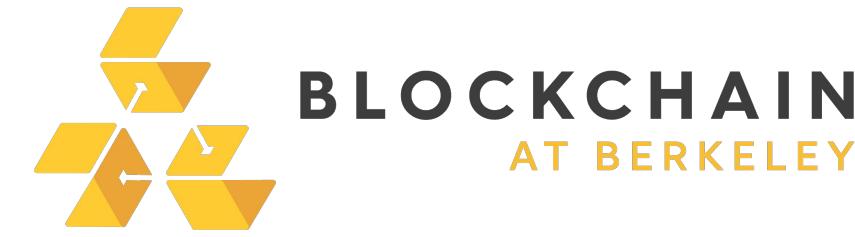
Proposer

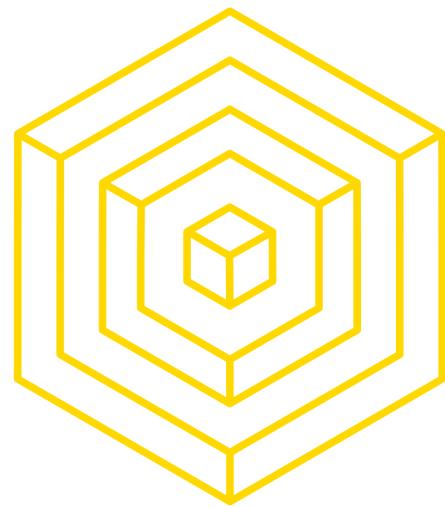


wait idk, lmk what it is pls

BLOCKCHAIN FUNDAMENTALS LECTURE 8

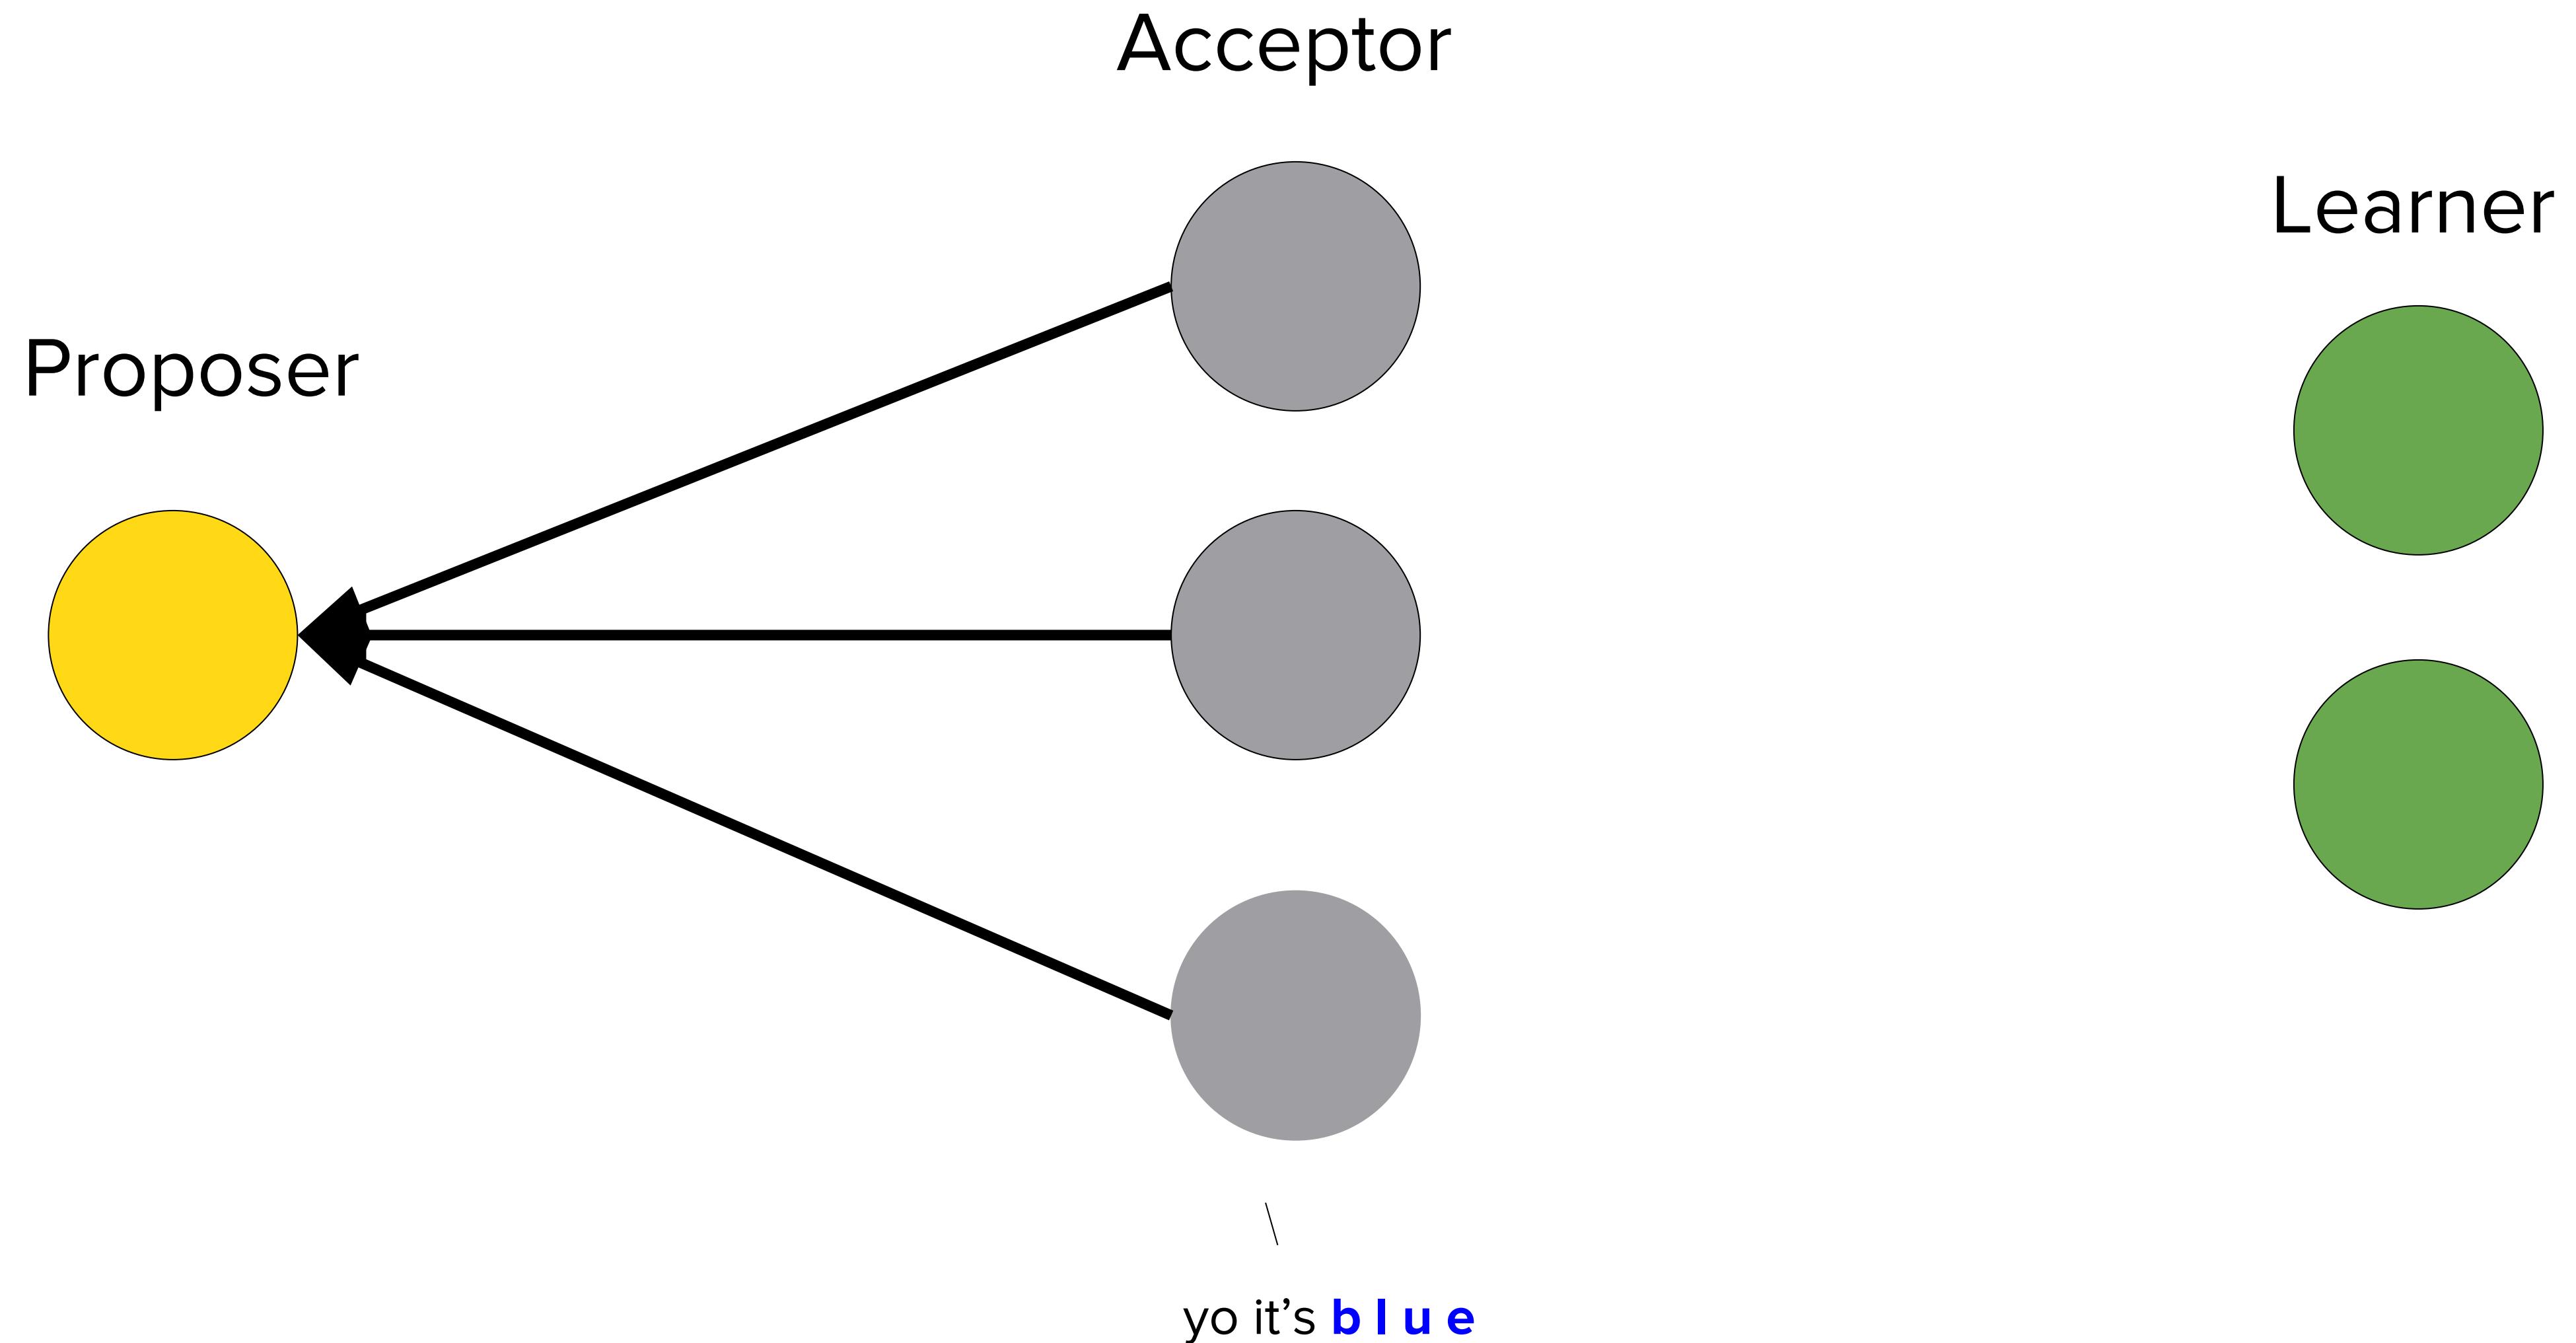
AUTHOR: RUSTIE LIN





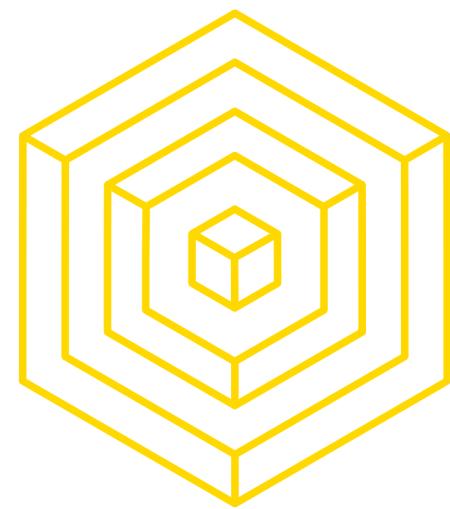
EXTRA: BASIC PAXOS

PHASE 1B: PROMISE



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



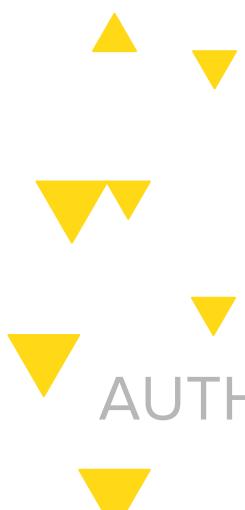
EXTRA: BASIC PAXOS

PHASE 2A: ACCEPT REQUEST

If Proposer gets enough promises from a Quorum of Acceptors, need to set a proposal value

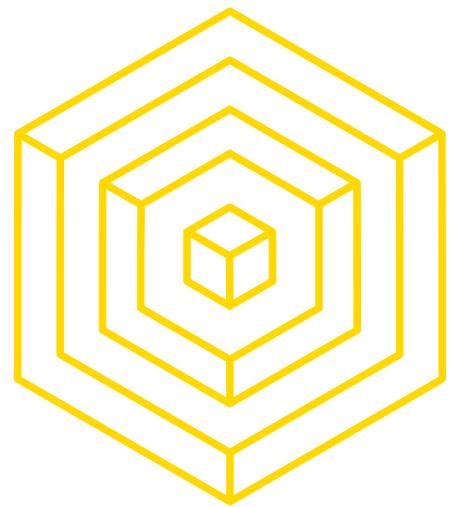
```
set proposal value:  
  if Acceptors had accept proposal before:  
    proposal value := value of proposal with the highest proposal number  
  else:  
    proposal value := any value determined by Proposer
```

Proposer then sends proposal value to Quorum of Acceptors



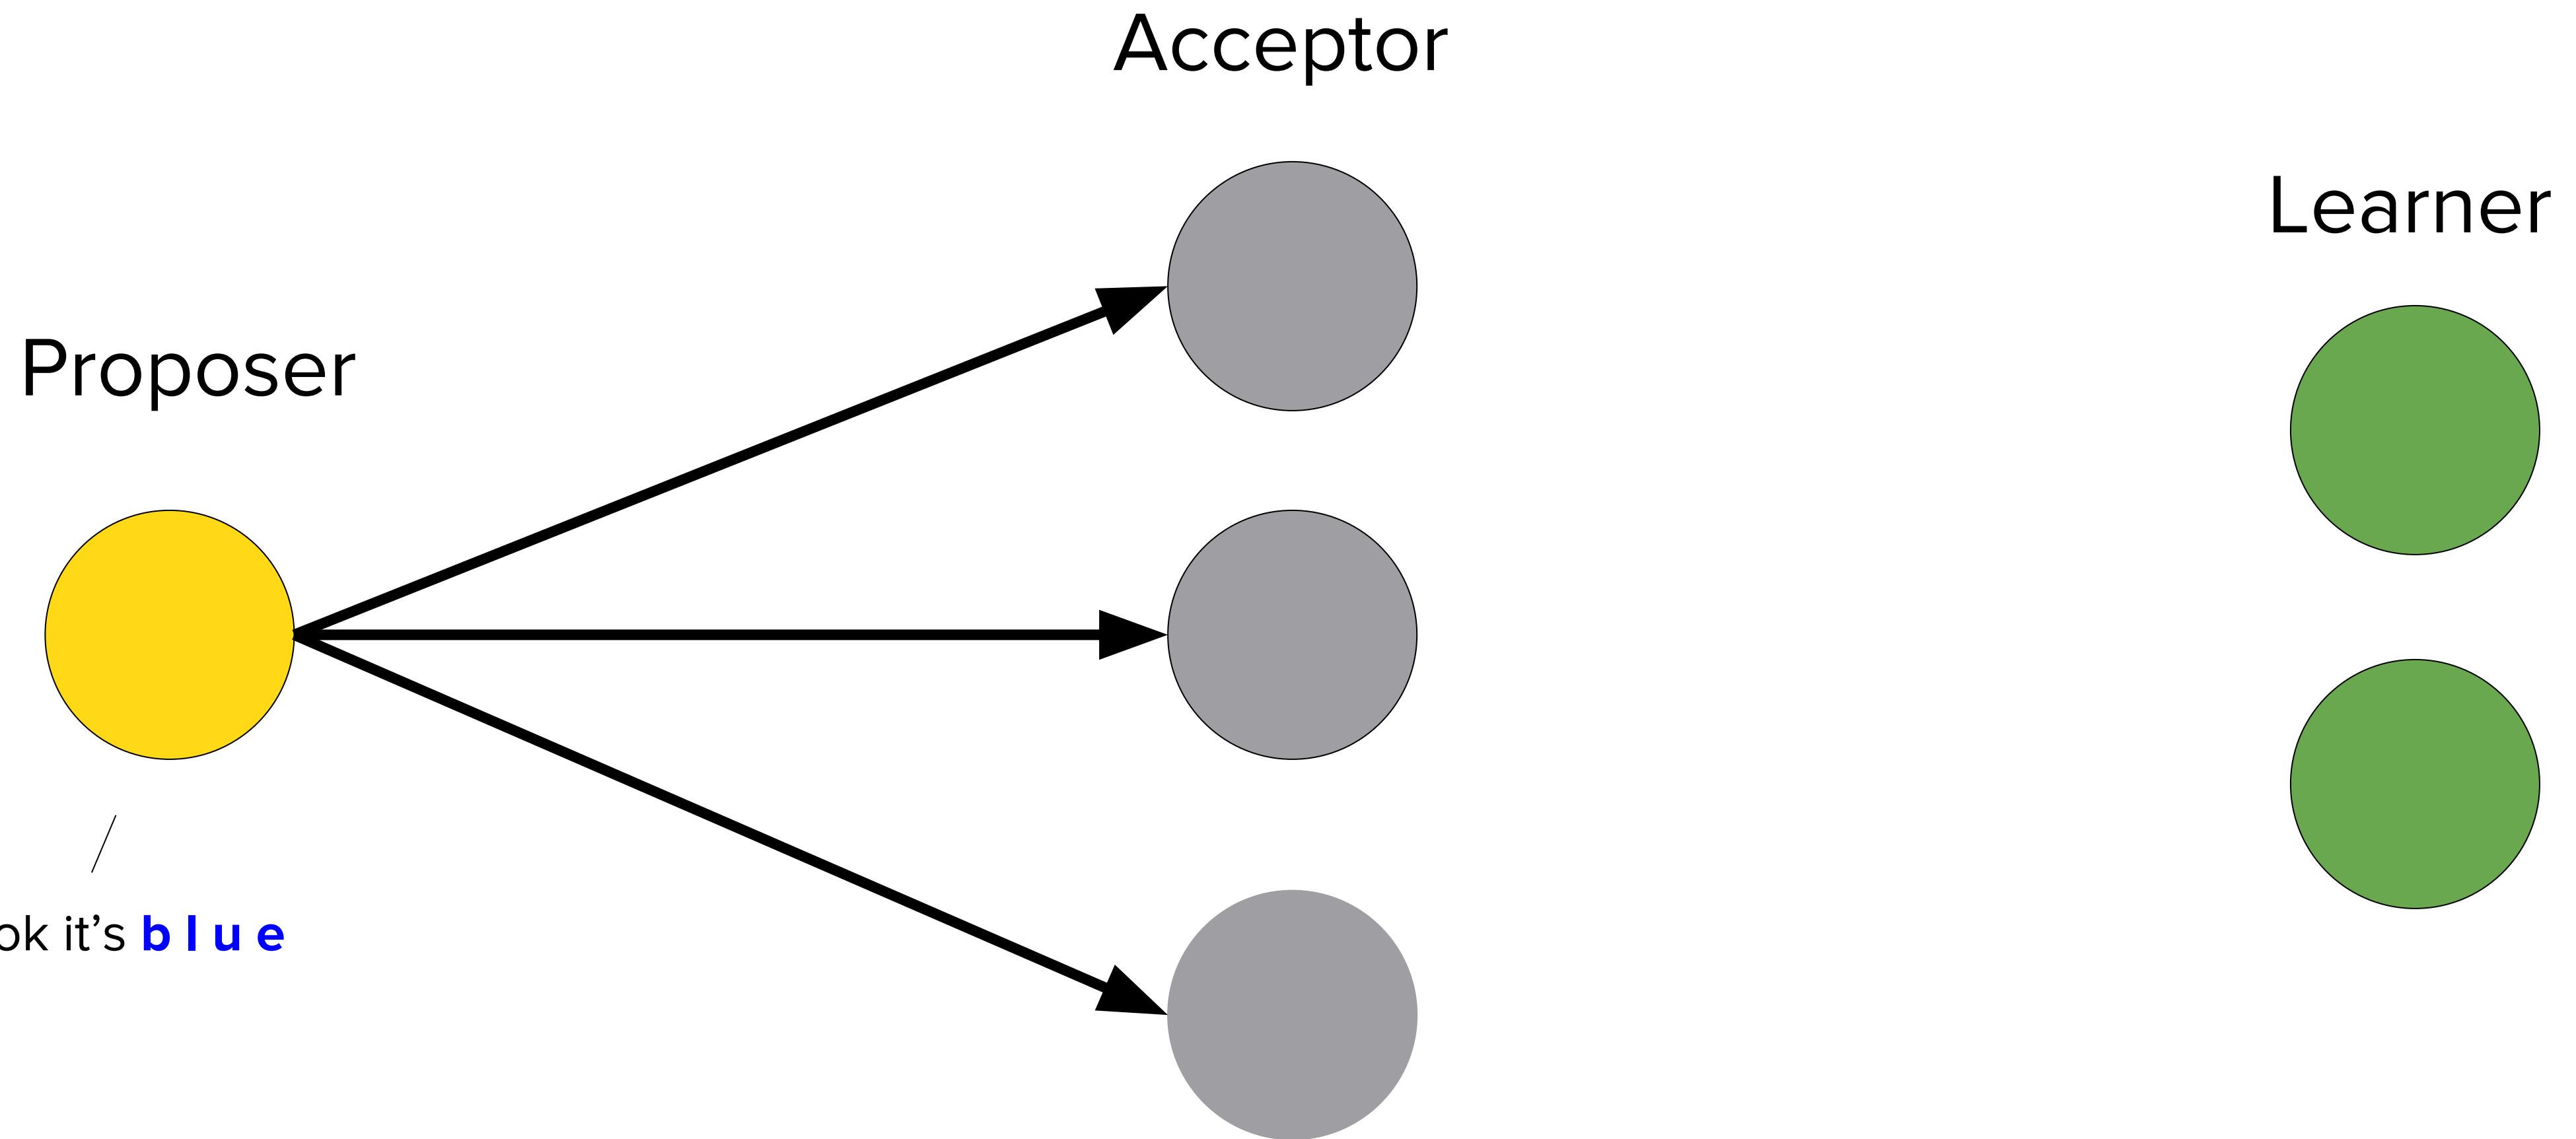
AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



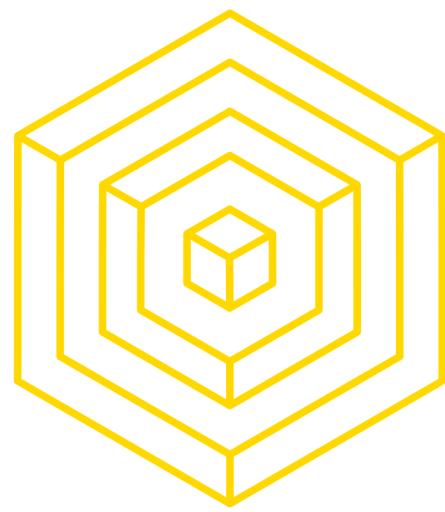
EXTRA: BASIC PAXOS

PHASE 2A: ACCEPT REQUEST



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



EXTRA: BASIC PAXOS

PHASE 2B: ACCEPTED

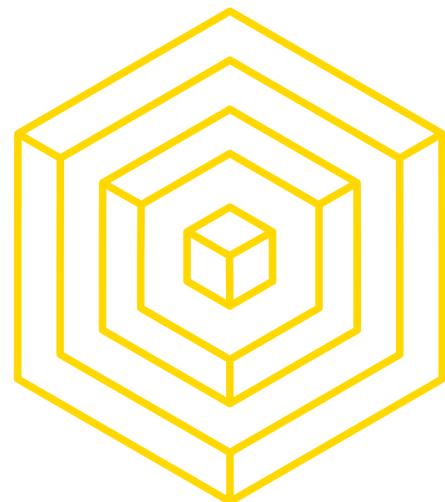
Acceptor accepts if and only if it has not already promised to only consider proposals greater than N

Register corresponding proposal value, and send message to Proposer and every Learner



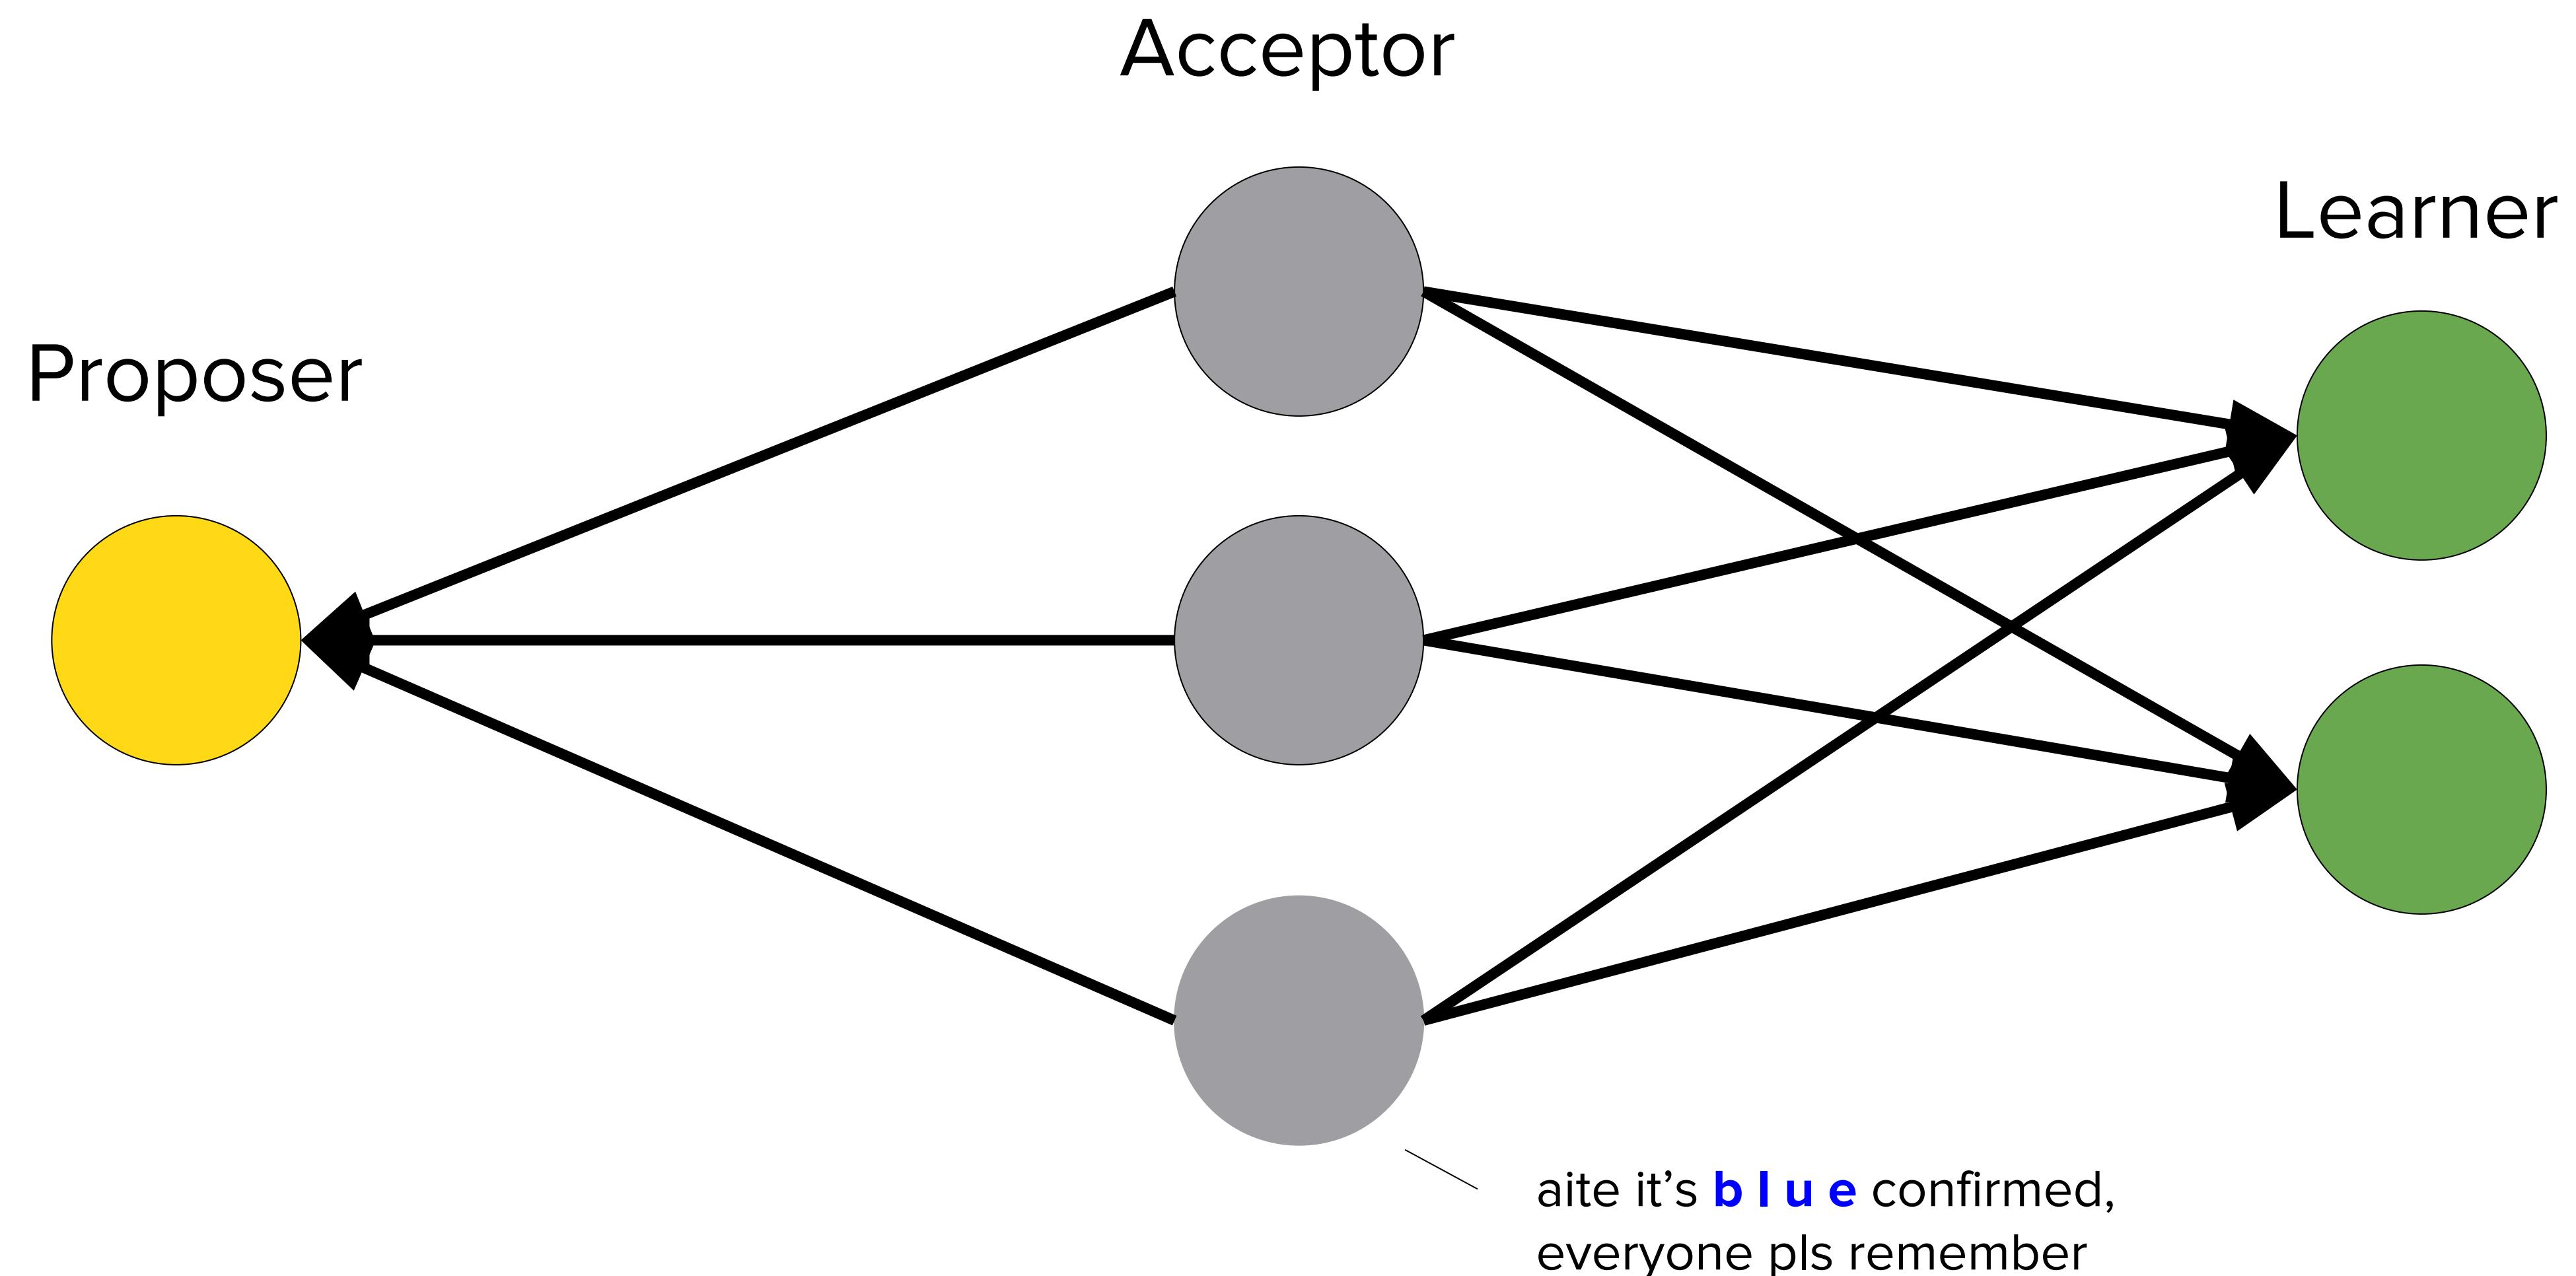
AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8

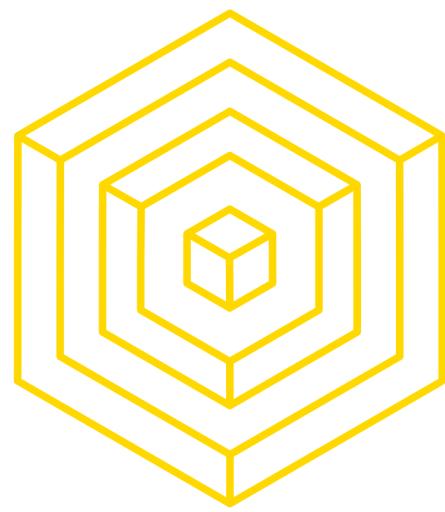


EXTRA: BASIC PAXOS

PHASE 1B: PROMISE



◀ ▶
▼
▼
▼ AUTHOR: RUSTIE LIN



EXTRA: PAXOS

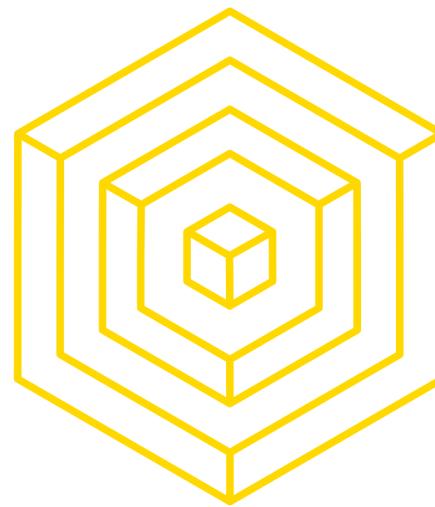
ASSUMPTIONS

- Nodes do not try to subvert the protocol
- Only works for fail-stop (no Byzantine failures)
- Messages delivered without corruption
- Failed nodes at any point won't affect the entire network



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8



EXTRA: PAXOS

SAFETY AND LIVENESS PROPERTIES

- Learners can only learn values that have been proposed
- If there's a proposal, a Learner will eventually learn some value
- At most one value can be learned (different learners cannot learn different values)



AUTHOR: RUSTIE LIN

BLOCKCHAIN FUNDAMENTALS LECTURE 8